

Cortex[™]-A Series

Version: 3.0

Programmer's Guide



Cortex-A Series Programmer's Guide

Copyright © 2011, 2012 ARM. All rights reserved.

Release Information

The following changes have been made to this book.

Change history

Date	Issue	Confidentiality	Change
25 March 2011	A	Non-Confidential	First release
10 August 2011	B	Non-Confidential	Second release. Virtualization chapter added. Updated to include Cortex-A15 processor, and LPAE. Corrected and revised throughout
25 June 2012	C	Non-Confidential	Updated for third release. Updated to include Cortex-A7 processor, and big.LITTLE. Index added. Corrected and revised throughout.

Proprietary Notice

This Cortex-A Series Programmer's Guide is protected by copyright and the practice or implementation of the information herein may be protected by one or more patents or pending applications. No part of this Cortex-A Series Programmer's Guide may be reproduced in any form by any means without the express prior written permission of ARM. **No license, express or implied, by estoppel or otherwise to any intellectual property rights is granted by this Cortex-A Series Programmer's Guide.**

Your access to the information in this Cortex-A Series Programmer's Guide is conditional upon your acceptance that you will not use or permit others to use the information for the purposes of determining whether implementations of the information herein infringe any third party patents.

This Cortex-A Series Programmer's Guide is provided "as is". ARM makes no representations or warranties, either express or implied, included but not limited to, warranties of merchantability, fitness for a particular purpose, or non-infringement, that the content of this Cortex-A Series Programmer's Guide is suitable for any particular purpose or that any practice or implementation of the contents of the Cortex-A Series Programmer's Guide will not infringe any third party patents, copyrights, trade secrets, or other rights.

This Cortex-A Series Programmer's Guide may include technical inaccuracies or typographical errors.

To the extent not prohibited by law, in no event will ARM be liable for any damages, including without limitation any direct loss, lost revenue, lost profits or data, special, indirect, consequential, incidental or punitive damages, however caused and regardless of the theory of liability, arising out of or related to any furnishing, practicing, modifying or any use of this Programmer's Guide, even if ARM has been advised of the possibility of such damages. The information provided herein is subject to U.S. export control laws, including the U.S. Export Administration Act and its associated regulations, and may be subject to export or import regulations in other countries. You agree to comply fully with all laws and regulations of the United States and other countries ("**Export Laws**") to assure that neither the information herein, nor any direct products thereof are; (i) exported, directly or indirectly, in violation of Export Laws, either to any countries that are subject to U.S. export restrictions or to any end user who has been prohibited from participating in the U.S. export transactions by any federal agency of the U.S. government; or (ii) intended to be used for any purpose prohibited by Export Laws, including, without limitation, nuclear, chemical, or biological weapons proliferation.

Words and logos marked with ® or ™ are registered trademarks or trademarks of ARM Limited, except as otherwise stated below in this proprietary notice. Other brands and names mentioned herein may be the trademarks of their respective owners.

Copyright © 2011, 2012 ARM Limited, 110 Fulbourn Road Cambridge, CB1 9NJ, England

This document is Non-Confidential but any disclosure by you is subject to you providing notice to and the acceptance by the recipient of, the conditions set out above.

In this document, where the term ARM is used to refer to the company it means "ARM or any of its subsidiaries as appropriate".

Web Address

<http://www.arm.com>

Contents

Cortex-A Series Programmer's Guide

	Preface	
	References	x
	Typographical conventions	xi
	Feedback on this book	xii
	Glossary	xiii
Chapter 1	Introduction	
	1.1 History	1-2
	1.2 System-on-Chip (SoC)	1-3
	1.3 Embedded systems	1-4
Chapter 2	ARM Architecture and Processors	
	2.1 Architecture versions	2-3
	2.2 Architecture history and extensions	2-4
	2.3 Key architectural points of ARM Cortex-A series processors	2-8
	2.4 Processors and pipelines	2-9
	2.5 The Cortex-A series processors	2-11
Chapter 3	Tools, Operating Systems and Boards	
	3.1 Linux distributions	3-2
	3.2 Useful tools	3-6
	3.3 Software toolchains for ARM processors	3-8
	3.4 ARM DS-5	3-11
	3.5 Example platforms	3-13
Chapter 4	ARM Registers, Modes and Instruction Sets	
	4.1 Instruction sets	4-2
	4.2 Modes	4-3

	4.3	Registers	4-4
	4.4	Instruction pipelines	4-7
	4.5	Branch prediction	4-10
Chapter 5		Introduction to Assembly Language	
	5.1	Comparison with other assembly languages	5-2
	5.2	Instruction sets	5-4
	5.3	Introduction to the GNU Assembler	5-5
	5.4	ARM tools assembly language	5-9
	5.5	Interworking	5-11
	5.6	Identifying assembly code	5-12
Chapter 6		ARM/Thumb Unified Assembly Language Instructions	
	6.1	Instruction set basics	6-2
	6.2	Data processing operations	6-6
	6.3	Multiplication operations	6-9
	6.4	Memory instructions	6-10
	6.5	Branches	6-13
	6.6	Integer SIMD instructions	6-14
	6.7	Saturating arithmetic	6-18
	6.8	Miscellaneous instructions	6-19
Chapter 7		Floating-Point	
	7.1	Floating-point basics and the IEEE-754 standard	7-2
	7.2	VFP support in GCC	7-8
	7.3	VFP support in the ARM Compiler	7-9
	7.4	VFP support in Linux	7-10
	7.5	Floating-point optimization	7-11
Chapter 8		Introducing NEON	
	8.1	SIMD	8-2
	8.2	NEON architecture overview	8-4
Chapter 9		Caches	
	9.1	Why do caches help?	9-3
	9.2	Cache drawbacks	9-4
	9.3	Memory hierarchy	9-5
	9.4	Cache architecture	9-6
	9.5	Cache policies	9-12
	9.6	Write and Fetch buffers	9-14
	9.7	Cache performance and hit rate	9-15
	9.8	Invalidating and cleaning cache memory	9-16
	9.9	Point of coherency and unification	9-17
	9.10	Level 2 cache controller	9-18
	9.11	Parity and ECC in caches	9-19
Chapter 10		Memory Management Unit	
	10.1	Virtual memory	10-3
	10.2	Level 1 page tables	10-4
	10.3	Level 2 page tables	10-7
	10.4	The Translation Lookaside Buffer	10-10
	10.5	TLB coherency	10-11
	10.6	Choice of page sizes	10-12
	10.7	Memory attributes	10-13
	10.8	Multi-tasking and OS usage of page tables	10-16
	10.9	Large Physical Address Extensions	10-19

Chapter 11	Memory Ordering	
	11.1 ARM memory ordering model	11-4
	11.2 Memory barriers	11-7
	11.3 Cache coherency implications	11-12
Chapter 12	Exception Handling	
	12.1 Types of exception	12-3
	12.2 Exception mode summary	12-5
	12.3 Entering an exception handler	12-7
	12.4 Exit from an exception handler	12-8
	12.5 Vector table	12-9
	12.6 Return instruction	12-10
Chapter 13	Interrupt Handling	
	13.1 External interrupt requests	13-2
	13.2 Generic Interrupt Controller	13-4
Chapter 14	Other Exception Handlers	
	14.1 Abort handler	14-2
	14.2 Undefined instruction handling	14-3
	14.3 SVC exception handling	14-4
	14.4 Linux exception program flow	14-5
Chapter 15	Boot Code	
	15.1 Booting a bare-metal system	15-2
	15.2 Configuration	15-6
	15.3 Booting Linux	15-7
Chapter 16	Porting	
	16.1 Endianness	16-2
	16.2 Alignment	16-5
	16.3 Miscellaneous C porting issues	16-7
	16.4 Porting ARM assembly code to ARMv7	16-10
	16.5 Porting ARM code to Thumb	16-11
Chapter 17	Application Binary Interfaces	
	17.1 Procedure Call Standard	17-2
	17.2 Mixing C and assembly code	17-7
Chapter 18	Profiling	
	18.1 Profiler output	18-3
Chapter 19	Optimizing Code to Run on ARM Processors	
	19.1 Compiler optimizations	19-3
	19.2 ARM memory system optimization	19-8
	19.3 Source code modifications	19-13
Chapter 20	Writing NEON Code	
	20.1 NEON C Compiler and assembler	20-2
	20.2 Optimizing NEON assembler code	20-7
	20.3 NEON power saving	20-9
Chapter 21	Introduction to Multi-processing	
	21.1 Multi-processing ARM systems	21-3
	21.2 Symmetric multi-processing	21-5
	21.3 Asymmetric multi-processing	21-7

Chapter 22	SMP Architectural Considerations	
	22.1 Cache coherency	22-2
	22.2 TLB and cache maintenance broadcast	22-5
	22.3 Handling interrupts in an SMP system	22-6
	22.4 Exclusive accesses	22-7
	22.5 Booting SMP systems	22-10
	22.6 Private memory region	22-12
Chapter 23	Parallelizing Software	
	23.1 Decomposition methods	23-2
	23.2 Threading models	23-4
	23.3 Threading libraries	23-5
	23.4 Synchronization mechanisms in the Linux kernel	23-8
Chapter 24	Issues with Parallelizing Software	
	24.1 Thread safety and reentrancy	24-2
	24.2 Performance issues	24-3
	24.3 Profiling in SMP systems	24-5
Chapter 25	Power Management	
	25.1 Power and clocking	25-2
Chapter 26	Security	
	26.1 TrustZone hardware architecture	26-2
Chapter 27	Virtualization	
	27.1 ARMv7-A Virtualization Extensions	27-3
	27.2 Hypervisor exception model	27-5
	27.3 Relationship between virtualization and ARM Security Extensions	27-6
Chapter 28	Introducing big.LITTLE	
	28.1 big.LITTLE configuration	28-2
	28.2 Structure of a big.LITTLE system	28-3
	28.3 Execution models in big.LITTLE	28-4
	28.4 big.LITTLE MP operation	28-9
Chapter 29	Debug	
	29.1 ARM debug hardware	29-2
	29.2 ARM trace hardware	29-4
	29.3 Debug monitor	29-7
	29.4 Debugging Linux applications	29-8
	29.5 DS-5 debug and trace	29-9
Appendix A	Instruction Summary	
	A.1 Instruction Summary	A-2
Appendix B	NEON and VFP Instruction Summary	
	B.1 NEON general data processing instructions	B-6
	B.2 NEON shift instructions	B-13
	B.3 NEON logical and compare operations	B-17
	B.4 NEON arithmetic instructions	B-23
	B.5 NEON multiply instructions	B-32
	B.6 NEON load and store element and structure instructions	B-35
	B.7 VFP instructions	B-41
	B.8 NEON and VFP pseudo-instructions	B-47

Appendix C	Building Linux for ARM Systems	
C.1	Building the Linux kernel	C-2
C.2	Creating the Linux filesystem	C-6
C.3	Putting it together	C-8

Preface

This book provides an introduction to ARM technology for programmers using ARM Cortex-A series processors that conform to the ARM ARMv7–A architecture. The v7 refers to version 7 of the architecture, while the *A* indicates the architecture profile that describes Application processors. This includes the Cortex-A5, Cortex-A7, Cortex-A8, Cortex-A9 and Cortex-A15 processors. The book complements rather than replaces other ARM documentation that is available for Cortex-A series processors, such as the ARM *Technical Reference Manuals* (TRMs) for the processors themselves, documentation for individual devices or boards and, most importantly, the *ARM Architecture Reference Manual* (or the “ARM ARM”).

The purpose of this book is to bring together information from a wide variety of sources to provide a single guide for programmers who want to develop applications for the latest Cortex-A series of processors. We will cover hardware concepts such as caches and Memory Management Units, but only where this is valuable to the application writer. The book is intended to provide information that will be useful to both assembly language and C programmers. We will look at how complex operating systems, such as Linux, make use of ARM features, and how to take full advantage of the many advanced capabilities of the ARM processor, in particular writing software for multi-processing and using the SIMD capabilities of the device.

Although much of the book is also applicable to other ARM processors, we do not explicitly cover processors that implement older versions of the Architecture. The Cortex-R series and M-series processors are mentioned but not described. Our intention is to provide an approachable introduction to the ARM architecture, covering the feature set in detail and providing practical advice on writing both C and assembly language programs to run efficiently on a Cortex-A series processor.

This is not an introductory level book. We assume knowledge of the C programming language and microprocessors, but not of any ARM-specific background. In the allotted space, we cannot hope to cover every topic in detail. In some chapters, we suggest further reading (referring either to books or websites) that can give a deeper level of background to the topic in hand, but in this book we will focus on the ARM-specific detail. We do not assume the use of any particular tool chain.

We will mention both GNU and ARM tools in the course of the book. We hope that the book will be well suited to programmers who have a desktop PC or x86 background and are taking their first steps into the ARM processor based world.

The first chapters of the book cover the basic features of the ARM Cortex-A series processors. An introduction to the fundamentals of the ARM architecture and some background on individual processors is provided in [Chapter 2](#). In [Chapter 3](#), we briefly consider some of the tools and platforms available to those getting started with ARM programming. Chapters 4, 5 and 6 provide a brisk introduction to ARM assembly language programming, covering the various registers, modes and assembly language instructions. We look at floating-point and the ARM Advanced SIMD extensions (NEON) in Chapters 7 and 8. These chapters are only an introduction to the relevant topics. It would take significantly longer to cover all of the powerful capabilities of NEON and how to apply these to common signal processing algorithms. We then switch our focus to the memory system and look at Caches, Memory Management and Memory Ordering in Chapters 9, 10 and 11. Dealing with interrupts and other exceptions is described in Chapters 12 to 14.

The remaining chapters of the book provide more advanced programming information. [Chapter 15](#) provides an overview of boot code. [Chapter 16](#) looks at issues with porting C and assembly code to the ARMv7 architecture, from other architectures and from older versions of the ARM architecture. [Chapter 17](#) covers the Application Binary Interface, knowledge of which is useful to both C and assembly language programmers. Profiling and optimizing of code is covered in Chapters 18 and 19. Many of the techniques presented are not specific to the ARM architecture, but we also provide some processor-specific hints.

Chapters 21-24 cover the area of multi-processing. We take a detailed look at how this is implemented by ARM and how you can write code to take advantage of it. Power management is an important part of ARM programming and is covered in Chapter 21. The final chapters of the book provide a brief coverage of the ARM Security Extensions (TrustZone®), the ARM Virtualization extensions ([Chapter 27](#)), big.LITTLE™ technology ([Chapter 28](#)), and the powerful hardware debug features available to programmers ([Chapter 29](#)). Appendices A and B give a summary of the available ARM, NEON and VFP instructions and Appendix C gives step-by-step instructions for configuring and building Linux for ARM systems.

References

Cohen, D. “*On Holy Wars and a Plea for Peace*”, USC/ISI IEN April, 1980, <http://www.ietf.org/rfc/ien/ien137.txt>.

Furber, Steve. “*ARM System-on-chip Architecture*”, 2nd edition, Addison-Wesley, 2000, ISBN: 9780201675191.

Hohl, William. “*ARM Assembly Language: Fundamentals and Techniques*” CRC Press, 2009. ISBN: 9781439806104.

Sloss, Andrew N.; Symes, Dominic.; Wright, Chris. “*ARM System Developer's Guide: Designing and Optimizing System Software*”, Morgan Kaufmann, 2004, ISBN: 9781558608740.

ANSI/IEEE Std 754-1985, “*IEEE Standard for Binary Floating-Point Arithmetic*”.

ANSI/IEEE Std 754-2008, “*IEEE Standard for Binary Floating-Point Arithmetic*”.

ANSI/IEEE Std 1003.1-1990, “*Standard for Information Technology - Portable Operating System Interface (POSIX) Base Specifications, Issue 7*”.

ANSI/IEEE Std 1149.1-2001, “*IEEE Standard Test Access Port and Boundary-Scan Architecture*”.

The *ARM Architecture Reference Manual* (known as the ARM ARM) is a must-read for any serious ARM programmer. It is available (after registration) from the ARM website. It fully describes the ARMv7 instruction set architecture, programmer’s model, system registers, debug features and memory model. It forms a detailed specification to which all implementations of ARM processors must adhere.

References to the *ARM Architecture Reference Manual* in this document are to:

ARM Architecture Reference Manual - ARMv7-A and ARMv7-R edition (ARM DDI 0406).

————— **Note** —————

In the event of a contradiction between this book and the ARM ARM, the ARM ARM is definitive and must take precedence. In most instances, however, the ARM ARM and the Cortex-A Programmer’s Guide cover two separate world views. The ARM ARM is for processor implementers, while this book is for processor users. The most likely scenario is that this book will describe something in a way that does not cover all architecturally permitted behaviors - or simply rewords an abstract concept in more practical terms.

ARM Generic Interrupt Controller Architecture Specification (ARM IHI 0048).

ARM Compiler Toolchain Assembler Reference (DUI 0489).

ARM C Language Extensions (IHI 0053).

The individual processor Technical Reference Manuals provide a detailed description of the processor behavior. They can be obtained from the ARM website documentation area, <http://infocenter.arm.com>.

Typographical conventions

This book uses the following typographical conventions:

<i>italic</i>	Highlights important notes, introduces special terminology, denotes internal cross-references, and citations.
bold	Used for terms in descriptive lists, where appropriate.
monospace	Denotes text that you can enter at the keyboard, such as commands, file and program names, instruction names, parameters and source code.
<i>monospace italic</i>	Denotes arguments to monospace text where the argument is to be replaced by a specific value.
< and >	Enclose replaceable terms for assembler syntax where they appear in code or code fragments. For example: MRC p15, 0, <Rd>, <CRn>, <CRm>, <Opcode_2>

Feedback on this book

We have tried to ensure that the Cortex-A Series Programmer's Guide is both easy to read and still covers the material in enough depth to provide the comprehensive introduction to using the processors that we originally intended.

If you have any comments on this book, don't understand our explanations, think something is missing, or think that it is incorrect, send an e-mail to errata@arm.com. Give:

- The title, The *Cortex-A Series Programmer's Guide*.
- The number, ARM DEN0013C.
- The page number(s) to which your comments apply.
- What you think needs to be changed.

ARM also welcomes general suggestions for additions and improvements.

Glossary

Abbreviations and terms used in this document are defined here.

AAPCS	ARM Architecture Procedure Call Standard.
ABI	Application Binary Interface.
ACP	Accelerator Coherency Port.
AHB	Advanced High-Performance Bus.
AMBA®	Advanced Microcontroller Bus Architecture.
AMP	Asymmetric Multi-Processing.
APB	Advanced Peripheral Bus.
ARM ARM	The ARM Architecture Reference Manual.
ASIC	Application Specific Integrated Circuit.
APSR	Application Program Status Register.
ASID	Address Space ID.
ATPCS	ARM Thumb® Procedure Call Standard.
AXI	Advanced eXtensible Interface.
BE8	Byte Invariant Big-Endian Mode.
BIU	Bus Interface Unit.
BSP	Board Support Package.
BTAC	Branch Target Address Cache.
BTB	Branch Target Buffer.
CISC	Complex Instruction Set Computer.
CP15	Coprocessor 15 - System control coprocessor.
CPSR	Current Program Status Register.
DAP	Debug Access Port.
DBX	Direct Bytecode Execution.
DDR	Double Data Rate (SDRAM).
DMA	Direct Memory Access.
DMB	Data Memory Barrier.
DPU	Data Processing Unit.
DS-5™	The ARM Development Studio.
DSB	Data Synchronization Barrier.
DSP	Digital Signal Processing.
DSTREAM®	An ARM debug and trace unit.

DVFS	Dynamic Voltage/Frequency Scaling.
EABI	Embedded ABI.
ECC	Error Correcting Code.
ECT	Embedded Cross Trigger.
EOF	End Of File.
ETB	Embedded Trace Buffer™.
ETM	Embedded Trace Macrocell™.
FDT	Flattened Device Tree.
FIQ	An interrupt type (formerly fast interrupt).
FPSCR	Floating-Point Status and Control Register.
GCC	GNU Compiler Collection.
GIC	Generic Interrupt Controller.
GIF	Graphics Interchange Format.
GPIO	General Purpose Input/Output.
Gprof	GNU profiler.
Harvard architecture	Architecture with physically separate storage and signal pathways for instructions and data.
HCR	Hyp Configuration Register.
ICU	Instruction Cache Unit.
IDE	Integrated development environment.
I/F	Interface (abbreviation used in some diagrams).
IPA	Intermediate Physical Address.
IRQ	Interrupt Request (normally external interrupts).
ISA	Instruction Set Architecture.
ISB	Instruction Synchronization Barrier.
ISR	Interrupt Service Routine.
Jazelle™	The ARM bytecode acceleration technology.
JIT	Just In Time.
L1/L2	Level 1/Level 2.
LPAAE	Large Physical Address Extension.
LSB	Least Significant Bit.
MESI	A cache coherency protocol with four states; Modified, Exclusive, Shared and Invalid.

MMU	Memory Management Unit.
MPU	Memory Protection Unit.
MSB	Most Significant Bit.
NEON™	The ARM Advanced SIMD Extensions.
NMI	Non-Maskable Interrupt.
Oprofile	A Linux system profiler.
QEMU	A processor emulator.
PCI	Peripheral Component Interconnect. A computer bus standard.
PCS	Procedure Call Standard.
PFU	Prefetch Unit.
PIPT	Physically Indexed, Physically Tagged.
PLE	Preload Engine.
PLI	Preload Instruction.
PMU	Performance Monitor Unit.
PoC	Point of Coherency.
PoU	Point of Unification.
PPI	Private Peripheral Input.
PSR	Program Status Register.
PTE	Page Table Entry.
RCT	Runtime Compiler Target.
RISC	Reduced Instruction Set Computer.
RVCT	RealView® Compilation Tools (the “ARM Compiler”).
SBZP	Should Be Preserved.
SCU	Snoop Control Unit.
Sgi	Software Generated Interrupt.
SIMD	Single Instruction, Multiple Data.
SiP	System in Package.
SMP	Symmetric Multi-Processing.
SoC	System on Chip.
SP	Stack Pointer.
SPI	Shared Peripheral Interrupt.
SPSR	Saved Program Status Register.
Streamline	A graphical performance analysis tool.

SVC	Supervisor Call instruction. (Previously SWI)
SWI	Software Interrupt instruction. (Replaced with SVC)
SYS	System Mode.
TAP	Test Access Port (JTAG Interface).
TCM	Tightly Coupled Memory.
TDMI[®]	Thumb, Debug, Multiplier, ICE.
TEX	Type Extension.
Thumb[®]	An instruction set extension to ARM.
Thumb-2	A technology extending the Thumb instruction set to support both 16-bit and 32-bit instructions.
TLB	Translation Lookaside Buffer.
TLS	Thread Local Storage.
TrustZone	The ARM security extension.
TTB	Translation Table Base.
TTBR	Translation Table Base Register.
UAL	Unified Assembly Language.
UART	Universal Asynchronous Receiver/Transmitter.
UEFI	Unified Extensible Firmware Interface.
U-Boot	A Linux Bootloader.
UNK	Unknown.
USR	User mode, a non-privileged processor mode.
VFP	The ARM floating-point instruction set. Before ARMv7, the VFP extension was called the Vector Floating-Point architecture, and was used for vector operations.
VIC	Vectored Interrupt Controller.
VIPT	Virtually Indexed, Physically Tagged.
VMID	Virtual Machine ID.
VMSA	Virtual Memory Systems Architecture.
XN	Execute Never.

Chapter 1

Introduction

ARM processors are everywhere. More than 10 billion ARM processor based devices had been manufactured by the end of 2008 and at the time of writing (early 2012), it is estimated that around one quarter of electronic products contain one or more ARM processors. At the end of 2011, over 25 billion ARM processors had been shipped. It is likely that readers of this book own products containing ARM processor based devices – a mobile phone, personal computer, television or car. It might come as a surprise to programmers more used to the personal computer to learn that the x86 architecture occupies a much smaller (but still highly lucrative) position in terms of total microprocessor shipments, with around three billion devices.

The ARM architecture has advanced significantly since the first ARM1 silicon in 1985. The ARM processor is not a single processor, but a whole family of processors which share common instruction sets and programmer's models and have some degree of backward compatibility.

Let's begin, however, with a brief look at the history of ARM.

1.1 History

The first ARM processor (ARM1) was designed within Acorn Computers Limited by a team led by Sophie Wilson and Steve Furber, with the first silicon (which worked first time!) produced in April 1985. The ARM1 was quickly replaced by the ARM2 (which added multiplier hardware), which was used in real systems, including Acorn's Archimedes personal computer.

ARM was formed in Cambridge, England in November 1990, as Advanced RISC Machines Ltd. It was a joint venture between Apple Computers, Acorn Computers and VLSI Technology and has outlived two of its parents. The original 12 employees came mainly from the team within Acorn Computers. One reason for spinning ARM off as a separate company was that the processor had been selected by Apple Computers for use in its Newton product.

The new company quickly decided that the best way forward for their technology was to license their *Intellectual Property* (IP). Instead of designing, manufacturing and selling the chips themselves, they would sell rights to their designs to semiconductor companies. These companies would design the ARM processor into their own products, in a partnership model. This IP licensing business is how ARM continues to operate today. ARM was quickly able to sign up licensees with Sharp, Texas Instruments and Samsung prominent names among early customers. In 1998, ARM floated on the London Stock Exchange and Nasdaq. At the time of writing, ARM has over 2000 employees and has expanded somewhat from its original remit of processor design. ARM also licenses Physical IP – libraries of cells (NAND gates, RAM and so forth), graphics and video accelerators and software development products such as compilers, debuggers and development boards.

1.2 System-on-Chip (SoC)

Chip designers today can produce chips with many millions of transistors. Designing and verifying such complex circuits has become an extremely difficult task. It is increasingly rare for all of the parts of such systems to be designed by a single company. In response to this, ARM and other semiconductor IP companies design and verify components (so-called IP blocks or processors). These are licensed by semiconductor companies who use these blocks in their own designs and include microprocessors, DSPs, 3D graphics and video controllers, along with many other functions.

The semiconductor companies take these blocks and integrate many other parts of a particular system onto the chip, to form a *System-on-Chip* (SoC). The architects of such devices must select the appropriate processor(s), memory controllers, on-chip memory, peripherals, bus interconnect and other logic (perhaps including analog or radio frequency components), in order to produce a system.

The term *Application Specific Integrated Circuit* (ASIC) is one that we will also use in the book. This is an Integrated Circuit design that is specific to a particular application. An individual ASIC might well contain an ARM processor, memory and other components. Clearly there is a large overlap with devices which can be termed SoCs. The term SoC usually refers to a device with a higher degree of integration, including many of the parts of the system in a single device, possibly including analog, mixed-signal or radio frequency circuits.

The large semiconductor companies investing tens of millions of dollars to create these devices will typically also make a large investment in software to run on their platform. It would be uncommon to produce a complex system with a powerful processor without at least having ported one or more operating systems to it and written device drivers for peripherals.

Of course, powerful operating systems like Linux require significant amounts of memory to run, more than is usually possible on a single silicon device. The term System-on-Chip is therefore not always named entirely accurately, as the device does not always contain the whole system. Apart from the issue of silicon area, it is also often the case that many useful parts of a system require specialist silicon manufacturing processes that preclude them from being placed on the same die. An extension of the SoC that addresses this to some extent is the concept of *System-in-Package* (SiP) that combines a number of individual chips within a single physical package. Also widely seen is package-on-package stacking. The package used for the SoC chip contains connections on both the bottom (for connection to a PCB) and top (for connection to a separate package that might contain a flash memory or a large SDRAM device).

1.3 Embedded systems

An embedded system is conventionally defined as a piece of computer hardware running software designed to perform a specific task. Examples of such systems might be TV set-top boxes, smartcards, routers, disk drives, printers, automobile engine management systems, MP3 players or photocopiers. These contrast with what is generally considered a computer system, that is, one that runs a wide range of general purpose software and possesses input and output devices such as a keyboard, and a graphical display of some kind.

This distinction is becoming increasingly blurred. Consider the cellular or mobile phone. A basic model might just perform the task of making phone calls, but modern smartphones can run a complex operating system to which many thousands of applications are available for download.

Embedded systems can contain very simple 8-bit microprocessors, such as an Intel 8051 or PIC micro-controllers, or some of the more complex 32 or 64-bit processors, such as the ARM family that form the subject matter for this book. They need some RAM and some form of non-volatile storage to hold the program(s) to be executed by the system. Systems will almost always have additional peripherals, relating to the actual function of the device – typically including UARTs, interrupt controllers, timers, GPIO controllers, but also potentially quite complex blocks such as DSPs, GPUs, or DMA controllers.

Software running on such systems is typically grouped into two separate parts: the *Operating System* (OS) and applications that run on top of the OS. A wide range of operating systems are in use, ranging from simple kernels, through complex *Real-Time Operating Systems* (RTOS), to full-featured complex operating systems, of the kind that might be found on a desktop computer. Microsoft Windows or Linux are familiar examples of the latter. In this book, we will concentrate mainly on examples from Linux. The source code for Linux is readily available for inspection by the reader and is likely to be familiar to many programmers. Nevertheless, lessons learned from Linux are equally applicable to other operating systems.

There are many constraints on embedded systems that can make programming them more of a challenge than writing an application for a general-purpose PC.

Memory Footprint

In many systems, to minimize cost (and power), memory size can be limited. The programmer could be forced to consider the size of the program and how to reduce memory usage while it runs.

Real-time behavior

A feature of certain systems is that there are deadlines to respond to external events. This might be a “hard” requirement (a car braking system *must* respond within a certain time) or “soft” requirement (audio processing must complete within a certain time-frame to avoid a poor user experience – but failure to do so under rare circumstances may not render the system worthless).

Power

In many embedded systems the power source is a battery, and programmers and hardware designers must take great care to minimize the total energy usage of the system. This can be done, for example, by slowing the clock, reducing supply voltage or switching off the processor when there is no work to be done.

Cost

Reducing the bill of materials can be a significant constraint on system design.

Time to market

In competitive markets, the time to develop a working product can significantly impact the success of that product.

Chapter 2

ARM Architecture and Processors

As described in [Chapter 1](#), ARM does not manufacture silicon devices. Instead, ARM creates microprocessor designs, which are licensed to semiconductor companies and OEMs, who integrate them into System-on-Chip devices.

To ensure compatibility between implementations, ARM defines architecture specifications which define how compliant products must behave. Processors implementing the ARM architecture conform to a particular version of the architecture. There might be multiple processors with different internal implementations and micro-architectures, different cycle timings and clock speeds which conform to the same version of the architecture.

The programmer must distinguish between behaviors which are specific to the following:

- | | |
|---------------------------|--|
| Architecture | This defines behavior common to a set, or family, of processor designs and is defined in the appropriate <i>ARM Architecture Reference Manual</i> (ARM ARM). It covers instruction sets, registers, exception handling and other programmer's model features. The architecture defines behavior that is visible to the programmer, for example, which registers are available, and what individual assembly language instructions actually do. |
| Micro-architecture | This defines how the visible behavior specified by the architecture is implemented. This could include the number of pipeline stages, for example. It can still have some programmer visible effects, such as how long a particular instruction takes to execute, or the number of stall cycles after which the result is available. |
| Processor | A processor is an individual implementation of a micro-architecture. A processor might be licensed and manufactured by many companies. It might therefore, have been integrated into a wide range of different devices and |

systems, with a correspondingly wide range of memory maps, peripherals, and other implementation specific features. Processors are documented in Technical Reference Manuals, available on the ARM website.

Core

We use this term to refer to a separate logical execution unit inside a multi-core processor.

SoC

A System-on-Chip contains one or more processors and typically also memory and peripherals. The device could be part of a system which contains one or more of additional processors, memory, and peripherals. Documentation is available, not from ARM, but from the supplier of the individual SoC or platform.

2.1 Architecture versions

Periodically, new versions of the architecture are announced by ARM. These add new features or make changes to existing behaviors. Such changes are typically backwards compatible, meaning that user code which ran on older versions of the architecture will continue to run correctly on new versions. Of course, code written to take advantage of new features will not run on older processors that lack these features.

In all versions of the architecture, some system features and behaviors are left as implementation-defined. For example, the architecture does not define cache sizes or cycle timings for individual instructions. These are determined by the individual processor and SoC.

Each architecture version can also define optional extensions. These may or may not be implemented in a particular implementation of a processor. For example, in the ARMv7 architecture, the Advanced SIMD (NEON) technology is available as an optional extension, and we describe this at length in [Chapter 8 Introducing NEON](#).

The ARMv7 architecture also has the concept of *profiles*. These are variants of the architecture describing processors targeting different markets and usages.

The profiles are as follows:

- A** The *Application* profile defines an architecture aimed at high performance processors, supporting a virtual memory system using a *Memory Management Unit* (MMU) and therefore capable of running complex operating systems. Support for the ARM and Thumb instruction sets is provided.
- R** The *Real-time* profile defines an architecture aimed at systems that need deterministic timing and low interrupt latency and which do not need support for a virtual memory system and MMU, but instead use a simpler memory protection unit (MPU).
- M** The *Microcontroller* profile defines an architecture aimed at low cost and low performance systems, where low-latency interrupt processing is vital. It uses a different exception handling model to the other profiles and supports only a variant of the Thumb instruction set.

Throughout this book, our focus will be on version 7 of the architecture (ARMv7), particularly ARMv7-A, the Application profile. This is the most recent version of the architecture. It is implemented by all Cortex-A series processors, and by processors from Marvell and Qualcomm. We will, where appropriate, point out differences between ARMv7 and older versions of the architecture.

2.2 Architecture history and extensions

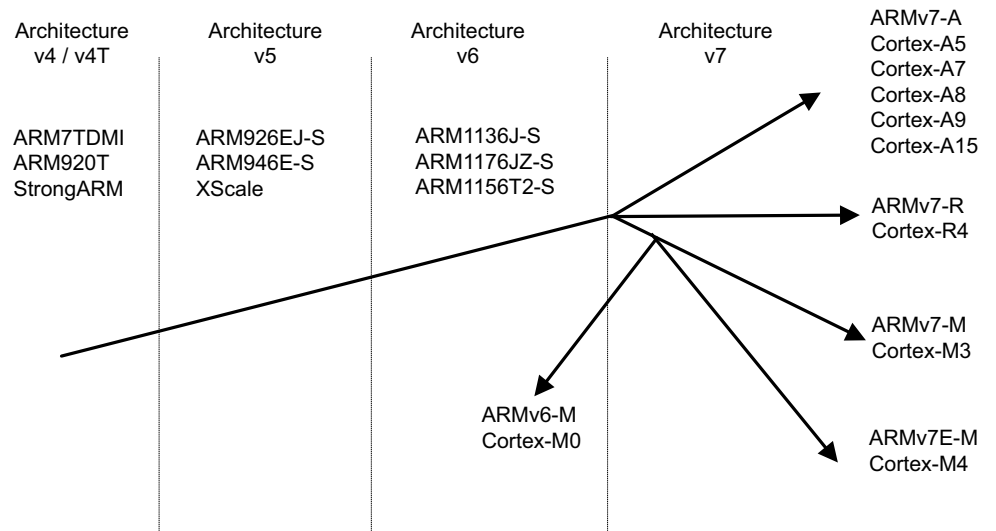
In this section, we look briefly at the development of the architecture through previous versions. Readers unfamiliar with the ARM architecture shouldn't worry if parts of this description use terms they don't know, as we'll describe all of these topics in the following chapters.

The ARM architecture changed relatively little between the first test silicon in the mid-1980s through to the first ARM6 and ARM7 devices of the early 1990s. In the first version of the architecture, the majority of the load, store and arithmetic operations along the exception model and register set was implemented by the ARM1. Version 2 added multiply and multiply-accumulate instructions and support for coprocessors, plus some further innovations. These early processors only supported 26-bits of address space. Version 3 of the architecture separated the program counter and program status registers and added several new modes, enabling support for 32-bits of address space. Version 4 adds support for halfword load and store operations and an additional kernel-level privilege mode.

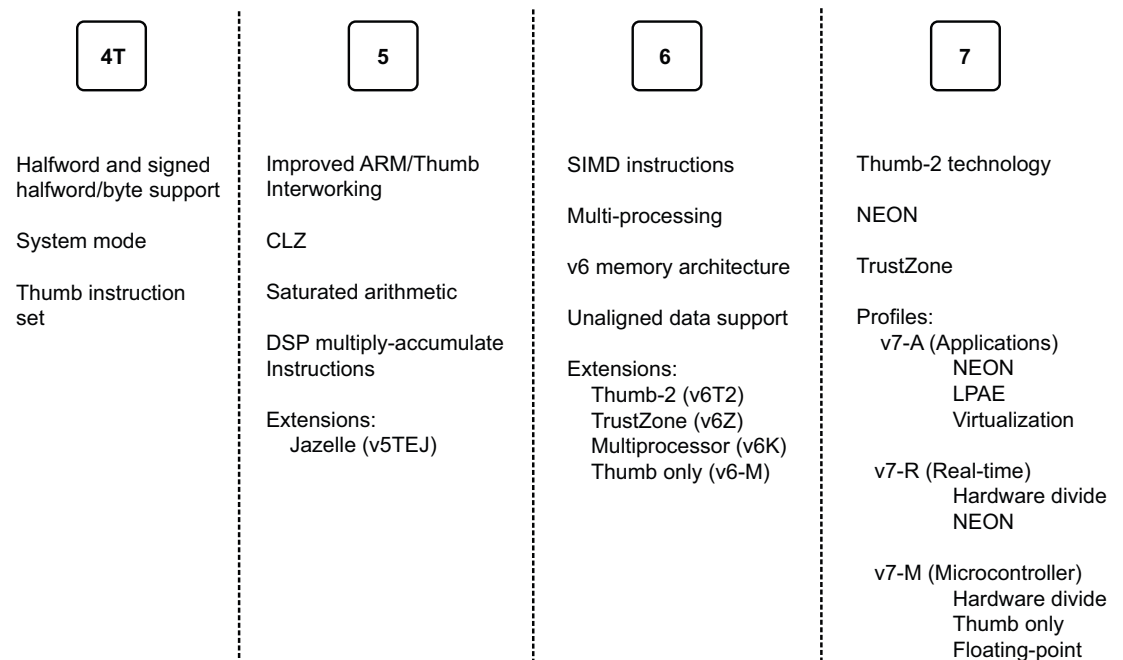
The ARMv4T architecture, which introduced the Thumb (16-bit) instruction set, was implemented by the ARM7TDMI® and ARM9TDMI® processors, products which have shipped in their billions. The ARMv5TE architecture added improvements for DSP-type operations and saturated arithmetic and to ARM/Thumb interworking. ARMv6 made a number of enhancements, including support for unaligned memory accesses, significant changes to the memory architecture and for multi-processor support, plus some support for SIMD operations operating on bytes or halfwords within the 32-bit registers. It also provided a number of optional extensions, notably Thumb-2 and Security Extensions (TrustZone). Thumb-2 extends Thumb to be a mixed length (16-bit and 32-bit) instruction set. The ARMv7-A architecture makes the Thumb-2 extensions mandatory and adds the Advanced SIMD extensions (NEON), described in [Chapter 8](#) and [Chapter 20](#).

For a number of years, ARM adopted a sequential numbering system for processors with ARM9 following ARM8, which came after ARM7. Various numbers and letters were appended to the base family to denote different variants. For example, the ARM7TDMI processor has T for Thumb, D for Debug, M for a fast multiplier and I for EmbeddedICE.

For the ARMv7 architecture, ARM Limited adopted the brand name *Cortex* for many of its processors, with a supplementary letter indicating which of the three profiles (A, R, or M) the processor supports. [Figure 2-1 on page 2-5](#) shows how different versions of the architecture correspond to different processor implementations. The figure is not comprehensive and does not include all architecture versions or processor implementations.

**Figure 2-1 Architecture and processors**

In [Figure 2-2](#), we show the development of the architecture over time, illustrating additions to the architecture at each new version. Almost all architecture changes are backwards-compatible, meaning software written for the ARMv4T architecture can still be used on ARMv7 processors.

**Figure 2-2 Architecture history**

Individual chapters of this book will cover these architecture topics in greater detail, but here we will briefly introduce a number of architecture elements.

2.2.1 DSP multiply-accumulate and saturated arithmetic instructions

These instructions, added in the ARMv5TE architecture, improve the capability for digital signal processing and multimedia software and are denoted by the letter E. The new instructions provide many variations of signed multiply-accumulate, saturated add and subtract, and count leading zeros and are present in all later versions of the architecture. In many cases, this made it possible to remove a simple separate DSP from the system.

2.2.2 Jazelle

Jazelle-DBX (Direct Bytecode eXecution) was introduced in ARMv5TEJ to accelerate Java performance while conserving power. A combination of increased memory availability and improvements in just-in-time (JIT) compilers have since reduced its value in application processors. As a result, many ARMv7-A processors do not implement this hardware acceleration.

Jazelle-DBX is best suited to providing high performance Java in systems with very limited memory (for example, feature phone or low-cost embedded use). In today's systems, it is mainly used for backwards compatibility.

2.2.3 Thumb Execution Environment (ThumbEE)

Introduced and required in ARMv7-A, ThumbEE is sometimes referred to as Jazelle-RCT (Runtime Compilation Target). It involves small changes to the Thumb instruction set that make it a better target for code generated at runtime in controlled environments (for example, by managed languages like Java, Dalvik, C#, Python or Perl).

ThumbEE is designed to be used by *Just-In-Time* (JIT) or *Ahead-Of-Time* (AOT) compilers, where it can reduce the code size of recompiled code. Compilation of managed code is outside the scope of this document.

2.2.4 Thumb-2

Thumb-2 technology was introduced in ARMv6T2, and is required in ARMv7. This technology extends the original 16-bit Thumb instruction set to include 32-bit instructions. The combined 16-bit and 32-bit Thumb instruction set achieves similar code density to the original Thumb instruction set, but with performance similar to the 32-bit ARM instruction set. The resulting Thumb instruction set provides virtually all the features of the ARM instruction set, plus some additional capabilities.

2.2.5 Security Extensions (TrustZone)

The optional Security Extensions referred to as TrustZone introduced with ARMv6K have been implemented in all ARM Cortex-A processors. TrustZone provides a separate Secure world to isolate sensitive code and data from the normal world that contains the operating system and applications. The software in the Secure world is therefore intended to provide security services to the Normal (non-secure) world applications. TrustZone is described in more detail in [Chapter 26 Security](#).

2.2.6 VFP

Before ARMv7, the VFP extension was called the Vector Floating-Point Architecture, and was used for vector operations. VFP is an extension which implements single-precision and optionally, double-precision floating-point arithmetic, compliant with the ANSI/IEEE Standard for Floating-Point Arithmetic.

2.2.7 Advanced SIMD (NEON)

The ARM NEON technology provides an implementation of the Advanced SIMD instruction set, with separate register files (shared with VFP). Some implementations have a separate NEON pipeline back-end. It supports 8, 16, 32 and 64-bit integer, and single-precision (32-bit) floating-point data, which can be operated on as vectors in 64-bit and 128-bit registers.

2.2.8 Large Physical Address Extension (LPAE)

LPAE is optional in the v7-A architecture and is presently supported by the Cortex-A7 and Cortex-A15 processors. It allows 32-bit processors that are normally limited to addressing a maximum of 4GB of RAM to access up to 1TB of memory by translating 32-bit virtual memory addresses into 40-bit physical memory addresses. See [Large Physical Address Extensions on page 10-19](#).

2.2.9 Virtualization

The ARM processor Virtualization Extensions are optional extensions to the ARMv7-A architecture profile. The extensions support the use of a virtual machine monitor, known as the hypervisor, to switch from one operating system to another. When implemented in a uniprocessor or in a multiprocessor system, the Virtualization Extensions support running multiple virtual machines on a single processor. See [Chapter 27 Virtualization](#).

2.2.10 big.LITTLE

big.LITTLE processing was introduced in ARMv7 to address one of industry's current challenges, namely how to create a System on Chip (SoC) that provides both high performance as well as power efficiency, to extend battery life. big.LITTLE uses a high performance Cortex-A15 processor, coupled with an energy efficient Cortex-A7 processor. The Cortex-A15 processor can be utilized for heavy workloads, while the Cortex-A7 processor can take over for the majority of mobile device workloads. big.LITTLE is described in more detail in [Chapter 28 Introducing big.LITTLE](#).

2.3 Key architectural points of ARM Cortex-A series processors

A number of key points are common to the Cortex-A family of devices:

- 32-bit RISC processor, with 16×32 -bit visible registers with mode-based register banking.
- Modified Harvard Architecture (separate, concurrent access to instructions and data).
- Load/Store Architecture.
- Thumb-2 technology as standard.
- VFP and NEON options which are expected to become standard features in general purpose applications processors.
- Backward compatibility with code from previous ARM processors.
- 4GB of virtual address space and a minimum of 4GB of physical address space.
- Hardware page table walking for virtual to physical address translation.
- Virtual page sizes of 4KB, 64KB, 1MB and 16MB. Cacheability attributes and access permissions can be set on a per-page basis.
- Big-endian and little-endian data access support.
- Unaligned access support for basic load/store instructions.
- SMP support on MPCore™ variants, with full data coherency from the L1 cache level. Automatic cache and Translation Lookaside Buffer (TLB) maintenance propagation provides high efficiency SMP operation.
- *Physically indexed, physically tagged* (PIPT) data caches.

2.4 Processors and pipelines

In this section, we briefly look at some ARM processors and identify which processor implements which architecture version. In [The Cortex-A series processors on page 2-11](#) we take a slightly more detailed look at some of the individual processors which implement architecture version v7-A, which forms the main focus of this book. Some terminology will be used in this chapter which may be unfamiliar to the first-time user of ARM processors and which will not be explained until later in the book.

[Table 2-1](#) lists the architecture version implemented by a number of older ARM processors.

Table 2-1 Older ARM processors and architectures

Architecture version	Applications processor	Embedded processor
v4T	ARM720T™ ARM920T™ ARM922T™	ARM7TDMI™
v5TE		ARM946E-S™ ARM966E-S™ ARM968E-S
v5TEJ	ARM926EJ-S™	
v6K	ARM1136J(F)-S™ ARM11™ MPCore™	
v6T2		ARM1156T2-S™
v6K + security extensions	ARM1176JZ(F)-S™	

[Table 2-2](#) lists the architecture version implemented by the Cortex family of processors.

Table 2-2 Cortex processors and architecture versions

v7-A (Applications)	v7-R (Real Time)	v6-M/v7-M (Microcontroller)
Cortex-A5 (Single/MP)	Cortex-R4	Cortex-M0+ (ARMv6-M)
Cortex-A7 (MP)	Cortex-R5	Cortex-M0 (ARMv6-M)
Cortex-A8 (Single)	Cortex-R7	Cortex-M1™ (ARMv6-M)
Cortex-A9 (Single/MP)		Cortex-M3™ (ARMv7-M)
Cortex-A15 (MP)		Cortex-M4(F) (ARMv7E-M)

[Table 2-3 on page 2-10](#) compares the properties of Cortex-A series processors. For additional information on the processor cache, see [Table 9-1 on page 9-11](#).

Table 2-3 Comparison of Cortex-A series processors

	Processor				
	Cortex-A5	Cortex-A7	Cortex-A8	Cortex-A9	Cortex-A15
Release date	Dec 2009	Oct 2011	July 2006	March 2008	April 2011
Clock speed	~1GHz	~1GHz on 28nm	~1GHz on 65nm	~2GHz on 40nm	~2.5GHz on 28nm
Execution order	In-order	In-order	In-order	Out of order	Out of order
Cores (Multicore support)	1 to 4	1 to 4	1	1 to 4	1 to 4
Peak throughput	1.6DMIPS/MHz	1.9DMIPS/MHz	2DMIPS/MHz	2.5DMIPS/MHz	3.5DMIPS/MHz
VFP/NEON architecture	VFPv4/NEON	VFPv4/NEON	VFPv3/NEON	VFPv3/NEON	VFPv4/NEON
Half precision extension	Yes	Yes	No	Yes	Yes
FP/NEON register renaming	No	No	No	No	Yes
GP register renaming	No	No	No	Yes	Yes
Hardware Divide	No	Yes	No	No	Yes
Fused Multiply Accumulate	Yes	Yes	No	No	Yes
Pipeline stages	8	8	13	9 to 12	15+
Decodes	1	Partial dual issue	2	2	3
Return stack entries	4	8	8	8	48
Floating Point Unit	Optional	Optional	Yes	Optional	Optional
AMBA interface	64-bit I/F AMBA 3	128-bit I/F AMBA 4	64 or 128-bit I/F AMBA 3	2× 64-bit I/F AMBA 3	128-bit I/F AMBA 4

2.5 The Cortex-A series processors

In this section, we take a closer look at each of the processors that implement the ARMv7-A architecture. Only a general description is given in each case, for more specific information on each processor, see [Table 2-3 on page 2-10](#) and [Table 9-1 on page 9-11](#).

2.5.1 The Cortex-A5 processor

The Cortex-A5 processor supports all ARMv7-A architectural features, including the TrustZone Security Extensions and the NEON Media Processing Engine. It is extremely area and power efficient, but has lower maximum performance than other Cortex-A series processors. Both single and multi-core versions of the Cortex-A5 processor are available.

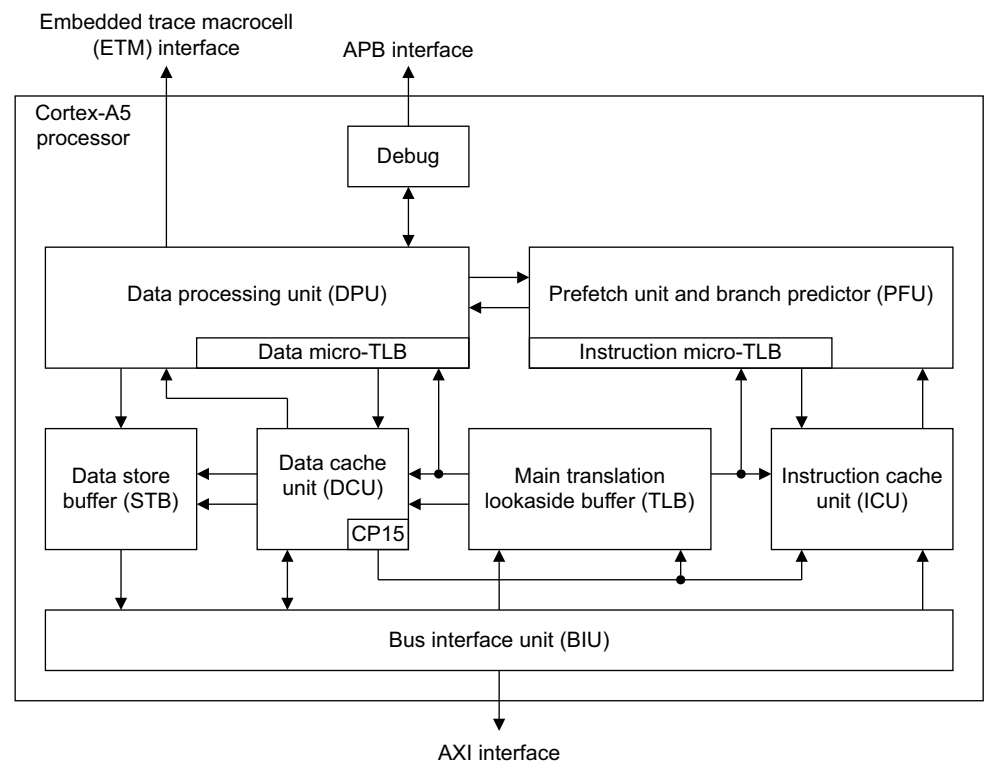


Figure 2-3 Block diagram of a single core Cortex-A5 processor

The Cortex-A5 processor shown in [Figure 2-3](#) can dual-issue a branch instruction with a preceding non-branch instruction in some circumstances and contains sophisticated branch prediction logic to reduce penalties associated with pipeline refills. Both NEON and floating-point hardware support are optional. It supports the ARM and Thumb instruction sets plus the Jazelle-DBX and Jazelle-RCT technology.

2.5.2 The Cortex-A7 processor

The Cortex-A7 MPCore processor is a high-performance, low-power processor fully compatible with other Cortex-A series processors described in this book.




Figure 2-4 Block diagram of single core Cortex-A7 processor

The Cortex-A7 processor incorporates all of the features of the high-performance Cortex-A15 processor including virtualization, *Large Physical Address Extensions* (LPAE), NEON, and AMBA 4 ACE coherency.

The Cortex-A7 MPCore processor has the following features:

- Improved memory management and bus interface.
- LPAE, addressing up to 1TB of memory.
- Multiple coherent multi-core processor clusters through AMBA4 technology.
- AMBA4 Cache Coherent Interconnect (CCI) allowing full cache coherency between multiple Cortex-A7 MPCore processors.

2.5.3 The Cortex-A8 processor

The Cortex-A8 processor was the first to implement the ARMv7-A architecture. It is available in a number of different devices, including the S5PC100 from Samsung, the OMAP3530 from Texas Instruments and the i.MX515 from Freescale. A wide range of device performances are available, with some giving clock speeds of more than 1GHz.

The Cortex-A8 processor has a considerably more complex micro-architecture compared with previous ARM processors.

Figure 2-5 is a block diagram showing the internal structure of the Cortex-A8 processor, including the pipelines.

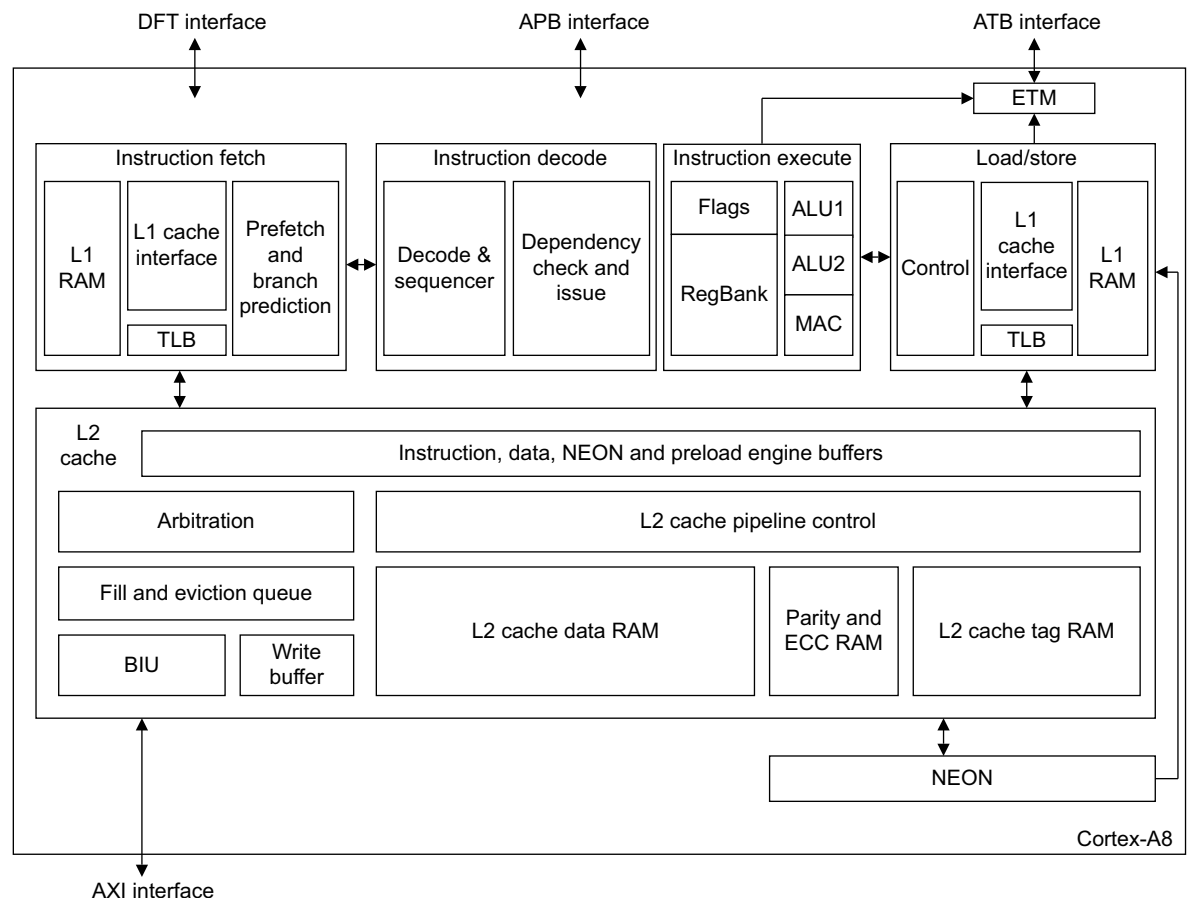


Figure 2-5 Block diagram of Cortex-A8 processor

The separate instruction and data level 1 caches are 16KB or 32KB in size. They are supplemented by an integrated, unified level 2 cache, which can be up to 1MB in size. The level 1 data cache and level 2 cache both have a 128-bit wide data interface to the processor. The level 1 data cache is virtually indexed, but physically tagged, while level 2 uses physical addresses for both index and tags. Data used by NEON is, by default, not allocated to L1 (although NEON can read and write data that is already in the L1 data cache).

2.5.4 The Cortex-A9 processor

The Cortex-A9 MPCore processor and the Cortex-A9 single processor provide higher performance than the Cortex-A5 or Cortex-A8 processor. ARM, Thumb, Thumb-2, TrustZone, both Jazelle RCT and DBX technologies are all supported.

The level 1 cache system provides hardware support for cache coherency for between one and four processors for multi-core software. ARM supplies an external level 2 cache controller (L2C-310, previously called the PL310) which supports caches of up to 8MB in size. The processor also contains an integrated interrupt controller, an implementation of the ARM *Generic Interrupt Controller* (GIC) architecture specification. This can be configured to provide support for up to 224 interrupt sources.

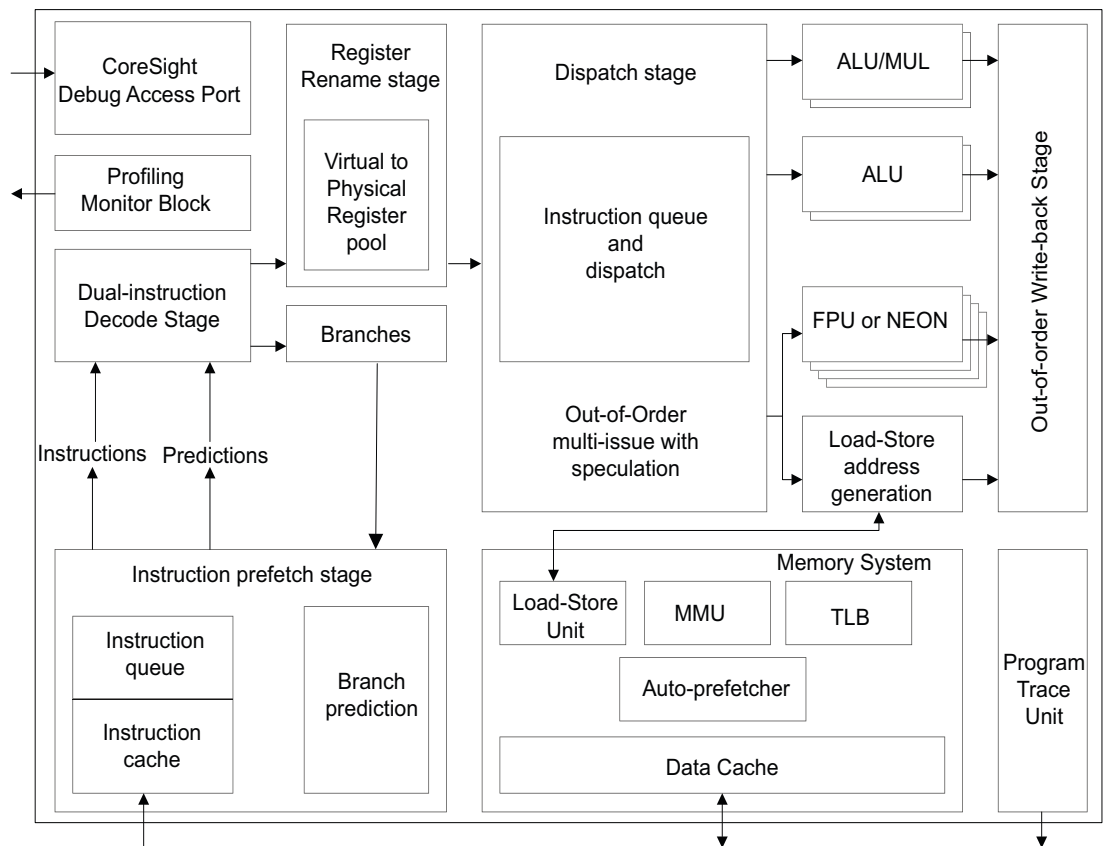


Figure 2-6 Block diagram of single core Cortex-A9 processor

Devices containing the Cortex-A9 processor include nVidia's dual-core Tegra-2, the SPEAr1300 from ST and TI's OMAP4 platform.

2.5.5 The Cortex-A15 processor

The Cortex-A15 MPCore processor is currently the highest performance ARM processor and is application compatible with the other ARM processors described in this book. The Cortex-A15 MPCore processor introduced some new capabilities, including support for full hardware virtualization and *Large Physical Address Extension* (LPAE), which enables addressing of up to 1TB of memory.

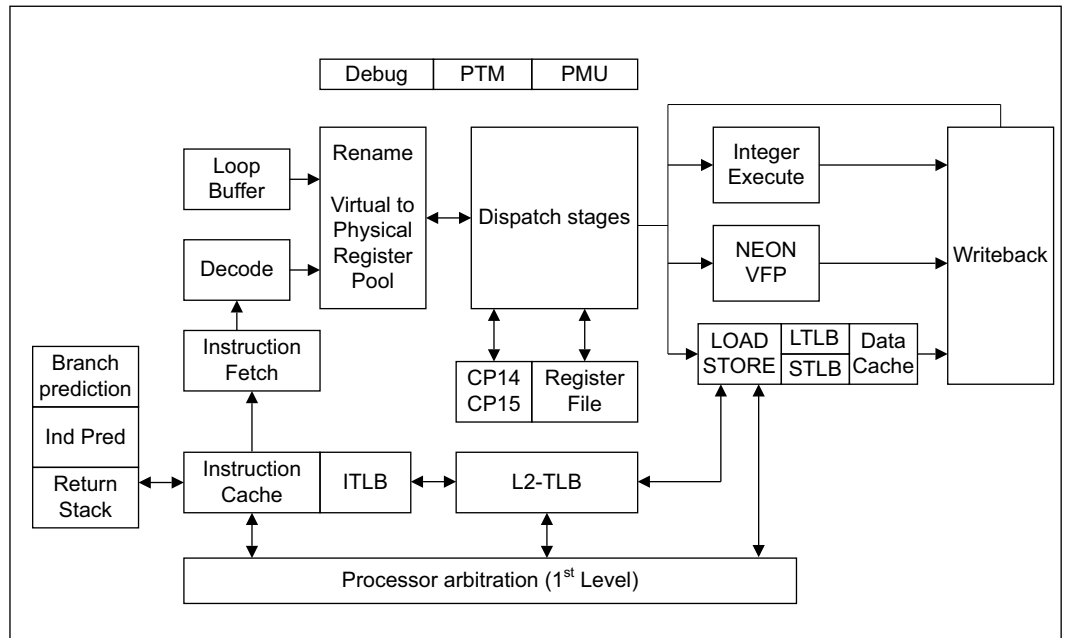


Figure 2-7 Block diagram of single core Cortex-A15 processor

The Cortex-A15 MPCore processor has the following features:

- An out-of-order superscalar pipeline.
- Tightly-coupled low-latency level-2 cache (up to 4MB in size).
- Improved floating-point and NEON media performance.
- Full hardware virtualization.
- Large Physical Address Extension (LPAE) addressing up to 1TB of memory.
- Error correction capability for fault-tolerance and soft-fault recovery.
- Multiple coherent multi-core processor clusters through AMBA4 technology.
- AMBA4 Cache Coherent Interconnect (CCI) allowing full cache coherency between multiple Cortex-A15 MPCore processors.

2.5.6 Qualcomm Scorpion

ARM is not the only company which designs processors compliant with the ARMv7-A instruction set architecture. In 2005, Qualcomm Inc. announced that it was creating its own implementation under license from ARM, with the name Scorpion. The Scorpion processor is available as part of Qualcomm's Snapdragon platform, which contains the features necessary to implement netbooks, smartphones or other mobile internet devices.

Relatively little information has been made publicly available by Qualcomm, although it has been mentioned that Scorpion has a number of similarities with the Cortex-A8 processor. It is an implementation of ARMv7-A, is superscalar and dual issue and has support for both VFP and NEON (called the VeNum media processing engine in Qualcomm press releases). There are a number of differences, however. Scorpion can process 128 bits of data in parallel in its NEON implementation. Scorpion has a 13-stage load/store pipeline and two integer pipelines. One of these is 10 stages long and can execute only simple arithmetic instructions (for example adds or

subtracts), while the other is 12 stages and can execute all data processing operations, including multiplies. Scorpion also has a 23-stage floating-point/SIMD pipeline, and VFPv3 operations are pipelined.

We will not specifically mention Scorpion again in this book. However, as the processor conforms to the ARMv7-A architecture specification, most of the information presented here will apply also to Scorpion.

Chapter 3

Tools, Operating Systems and Boards

ARM processors can be found in a very wide range of devices, running a correspondingly wide range of software. Many readers will have ready access to appropriate hardware, tools and operating systems, but before we proceed to look at the underlying architecture, it might be useful to some readers to present an overview of some of these readily available compilation tools, ARM processor based hardware and Linux operating system distributions.

In this chapter, we will provide a brief mention of a number of interesting commercially available development boards. We will provide some information about the Linux operating system and some useful associated tools. However, information about open source software and off-the-shelf boards is likely to change rapidly.

3.1 Linux distributions

Linux is a UNIX-like operating system kernel, originally developed by Linus Torvalds, who continues to maintain the official kernel. It is open source, distributed under the GNU Public License, widely-used and available on a large number of different processor architectures.

The Linux kernel is bundled with libraries and applications to make up a complete operating system in what is called a “Linux distribution”. A number of free Linux distributions exist for ARM processors, including Debian, Ubuntu, Fedora and Gentoo.

3.1.1 Linux for ARM systems

Support for the ARM architecture has been included in the standard Linux kernel for many years. Development of this port is ongoing, with significant input from ARM to provide kernel support for new processors and architecture versions.

You may wonder why a book about the Cortex-A series processors contains information about Linux. There are several reasons for this. Linux source code is available to all readers and represents a useful learning resource. In addition there are many useful resources with existing code and explanations. Many readers will be familiar with Linux, as it can be run on most processor architectures. By explaining how Linux features like virtual memory, multi-tasking, shared libraries and so forth are implemented, readers will be able to apply their understanding to other operating systems commonly used on ARM processors. The scalability of Linux is another factor – it can run on the most powerful ARM processors, and its derivative, uClinux, is also commonly used on much smaller processors, including the Cortex-M3 or ARM7TDMI processors. It can run on both the ARM and Thumb Instruction Set Architectures, using little-endian or big-endian data accesses and with or without a memory management unit.

One of the benefits of a modern operating system is that you do not need to know much detail about the underlying hardware in order to develop application software for it. This book focuses on the Cortex-A series of processors, and provides information about how to develop software that will run transparently on all of them.

In this book, we can merely scratch the surface of what there is to be said about Linux development. What we hope to do here is to show some ways in which programming for an ARM architecture based system differs from an x86 environment and to give some pointers to useful tools, which you might wish to investigate further.

3.1.2 Linux terminology

Here, we define some terms which we will use when describing how the Linux kernel interacts with the underlying ARM architecture:

- | | |
|----------------|---|
| Process | A process is the kernel's view of an executing unprivileged application. The same application (for example, <code>bin/bash</code>) can be running in several simultaneous instances in the system – and each of these instances will be a separate process. The process has resources associated with it, such as a memory map and file descriptors. A process can consist of one or more <i>threads</i> . |
| Thread | A thread is a context of software execution within a process. It is the entity which is scheduled by the kernel, and actually executes the instructions that make up the application. A process can consist of multiple threads, each executing with their own program counter, stack pointer and register set – all existing within the same memory map and operating on the file descriptors held by the process as a whole. In a multi-processor system, threads inside the same process can execute concurrently on separate processors. Different threads within the same process can be configured to have different scheduling priorities. |

There are also threads executing inside the kernel, to manage various tasks asynchronously, such as file cache management, or watchdog tickling (which is not as exciting as it sounds).

- Scheduler** This is a vital part of the kernel which has a list of all the current threads. It knows which threads are ready to be run and which are currently not able to run. It dynamically calculates priority levels for each thread and schedules the highest priority thread to be run next. It is called after an interrupt has been handled. The scheduler is also explicitly called by the kernel via the `schedule()` function, for example, when an application executing a *system call* needs to sleep. The system will have a timer based interrupt which results in the scheduler being called at regular intervals. This enables the OS to implement time-division multiplexing, where many threads share the processor, each running for a certain amount of time, giving the user the illusion that many applications are running simultaneously.
- System calls** Linux applications run in User (unprivileged) mode. Many parts of the system are not directly accessible in User mode. For example, the kernel might prevent User mode programs from accessing peripherals, kernel memory space and the memory space of other User mode programs. Access to some features of the system control coprocessor (CP15) is not permitted in User mode. The kernel provides an interface (via the SVC instruction) which permits an application to call kernel services, this is what forms a system call. Execution is transferred to the kernel through the SVC exception handler, which returns to the user application when the system call is complete.
- Libraries** Linux applications are, with very few exceptions, not loaded as complete pre-built binaries. Instead, the application relies on external support code linked from files called shared libraries. This has the advantage of saving memory space, in that the library only needs to be loaded into RAM once and is more likely to be in the cache as it can be used by other applications. Also, updates to the library do not require every application to be rebuilt. However, this dynamic loading means that the library code must not rely on being in a particular location in memory.
- Files** These are essentially blocks of data which are referred to using a pathname attached to them. Device nodes have pathnames like files, but instead of being linked to blocks of data, they are linked to device drivers which handle real I/O devices like an LCD display, disk drive or mouse. When an application opens, reads from or writes to a device, control is passed to specific routines in the kernel that handle that device.

3.1.3 Embedded Linux

Linux-based systems cover the range from servers, via the desktop, through mobile devices, right down to high-performance micro-controllers in the form of uClinux for processors lacking an MMU. However, while the kernel source code base is the same, different priorities and constraints mean that there can be some fundamental differences between the Linux running on your desktop and the one running in your set-top box, as well as between the development methodologies used.

In a desktop system, a form of bootloader executes from ROM – be it BIOS or UEFI. This has support for mass-storage devices and can then load a second-stage loader (for example GRUB) from a CD, a hard drive or even a USB memory stick. From this point on, everything is loaded from a general-purpose mass storage device.

In an embedded device, the initial bootloader is likely to load a kernel directly from on-board flash into RAM and execute it. In severely memory constrained systems, it might have a kernel built to *execute in place* (XiP), where all of the read-only portions of the kernel remain in ROM, and only the writable portions use RAM. Unless the system has a hard drive (or for fault tolerance reasons), the root filesystem on the device is likely to be located in flash. This can be a read-only filesystem, with portions that need to be writable overlaid by tmpfs mounts, or it can be a read-write filesystem. In both cases, the storage space available is likely to be significantly less than in a typical desktop computer. For this reason, they might use software components such as uClibc and BusyBox to reduce the overall storage space required for the base system. A general desktop Linux distribution is usually supplied preinstalled with a lot of software that you might find useful at some point. In a system with limited storage space, this is not really optimal. Instead, you want to be able to select exactly the components you need to achieve what you want with your system. Various specific embedded Linux distributions exist to make this easier.

In addition, embedded systems often have lower performance than general purpose computers. In this situation, speed of development can be significantly increased by compiling software for the target device on a faster desktop computer and then moving it across to the target device. This process is called *cross-compiling*.

3.1.4 Board Support Package

Getting Linux to run on a particular platform requires a *Board Support Package* (BSP). We can divide the platform-specific code into a number of areas:

- Architecture-specific code. This is found in the `arch/arm/` directory of the Linux kernel source code and forms part of the kernel porting effort carried out by the ARM Linux maintainers.
- Processor-specific code. This is found in the `arch/arm/mm/` and `arch/arm/include/asm/` directories of the Linux kernel source code. This takes care of MMU and cache functions (for example, page table setup, Translation Lookaside Buffer and cache invalidation and memory barriers). On SMP processors, spinlock code will be enabled.
- Generic device drivers are found under `drivers/`.
- Platform-specific code will be placed in the `arch/arm/mach-*/` directory of the Linux kernel source code. This is code which is most likely to be altered by people porting to a new board containing a processor with existing Linux support. The code will define the physical memory map, interrupt numbers, location of devices and any initialization code specific to that board.

3.1.5 Linaro

Linaro is a non-profit organization which works on a range of open source software running on ARM processors, including kernel related tools and software and middleware. It is a collaborative effort between a number of technology companies to provide engineering help and resources to the open source community.

Linaro does not produce a Linux distribution, nor is it tied to any particular distribution or board. Instead, Linaro works on generic ARM technology to provide a common software platform for use by board support package developers. Its focus is on tools to helping developers write and debug code, on low-level software which interacts with the underlying hardware and on key pieces of middleware. Linaro's members have a reduced time to market and can deliver unique open source based products using ARM technology.

Linaro engineers work on the kernel and tools, graphics and multimedia and power management. Linaro provides patches to upstream projects and makes monthly source tree tarballs available, with an occasional integrated build every six months to consolidate the work.

In this way, code can easily be transferred to the mainline Linux kernel and other open source projects. Evaluation builds of Android and Ubuntu, plus generic Linux, toolchain and other downloads for ARMv7 processors are available from <http://www.linaro.org/downloads/>.

See <http://www.linaro.org/> for more information about Linaro.

3.2 Useful tools

Let's take a brief look at some available tools which can be useful to developers of ARM architecture based Linux systems. These are all extensively documented elsewhere. In this section, we merely point out that these tools can be useful, and provide short descriptions of their purpose and function.

3.2.1 QEMU

QEMU is a fast, open source machine emulator. It was originally developed by Fabrice Bellard and is available for a number of architectures, including ARM. It can run operating systems and applications made for one machine (for example, an ARM processor) on a different machine, such as a PC or Mac. It uses dynamic translation of instructions and can achieve useful levels of performance, enabling it to boot complex operating systems like Linux, without the need for any target hardware.

3.2.2 BusyBox

BusyBox is a piece of open source software which provides many standard Unix tools, in a very small executable, which is ideal for many embedded systems and could be considered to be a *de facto* standard. It includes most of the Unix tools which can be found in the GNU Core Utilities, and many other useful tools including `init`, `dhclient`, `wget` and `tftp`. Less commonly used command switches are removed.

BusyBox calls itself the “Swiss Army Knife of Embedded Linux” – a reference to the large number of tools packed into a small package. BusyBox is a single binary executable which combines many applications. This reduces the overheads introduced by the executable file format and enables code to be shared between multiple applications without needing to be part of a library.

3.2.3 Scratchbox

The general principle of cross-compiling is to use one system (the host) to compile software which runs on some other system (the target).

The target is a different architecture to the host and so the host cannot natively run the resulting image. For example, you might have a powerful desktop x86 machine and want to develop code for a small battery-powered ARM processor based device which has no keyboard. Using the desktop machine will make code development simpler and compilation faster. There are some difficulties with this process. Some build environments will try to run programs on the target machine during compilation, and of course this is not possible. In addition, tools which during the build process try to discover information about the machine (for software portability reasons), do not work correctly when cross-compiling.

Scratchbox is a cross-compilation toolkit which solves these problems and gives the necessary tools to cross-compile a complete Linux distribution. It can use either QEMU or a target board to execute the cross-compiled binaries it produces.

3.2.4 U-Boot

Das U-Boot (Universal Bootloader) is a universal bootloader that can easily be ported to new hardware processors or boards. It provides serial console output which makes it easy to debug and is designed to be small and reliable. In an x86 system, we have BIOS code which initializes the processor and system and then loads an intermediate loader such as GRUB or syslinux, which then in turn loads and starts the kernel. U-Boot essentially covers both functions.

3.2.5 UEFI and Tianocore

The Unified Extensible Firmware Interface (UEFI) is the specification of a programmable software interface that sits on top a computer's hardware and firmware. Rather than all of the boot code being stored in the motherboard's BIOS, UEFI sits in non-volatile memory. A computer boots into UEFI, a set of actions are carried out, before loading an operating system, such as Windows or Linux. While BIOS is limited to 16-bit processes and 1MB of memory addressing, UEFI can function in 32-bit and 64-bit modes, allowing much more RAM to be addressed by more complex processes. It also can be architecture independent and provide drivers for components that are also independent of what kind of processor you have. The UEFI forum is a non-profit collaborative trade organization formed to promote and manage the UEFI standard.

UEFI is processor architecture independent and the Tianocore EFI Development Kit 2 (EDK2) is available under a BSD license. It contains UEFI support for ARM platforms, including ARM Versatile Express boards and the BeagleBoard (see *BeagleBoard* on page 3-13).

See <http://www.uefi.org> and <http://sourceforge.net/apps/mediawiki/tianocore> for more information.

3.3 Software toolchains for ARM processors

There are a wide variety of compilation and debug tools available for ARM processors. In this section, we will focus on two toolchains (a collection of programming tools), the GNU toolchain which includes the GNU Compiler (gcc), and the ARM Compiler toolchain which includes the armcc compiler.

Figure 3-1 shows how the various components of a software toolchain interact to produce an executable image.

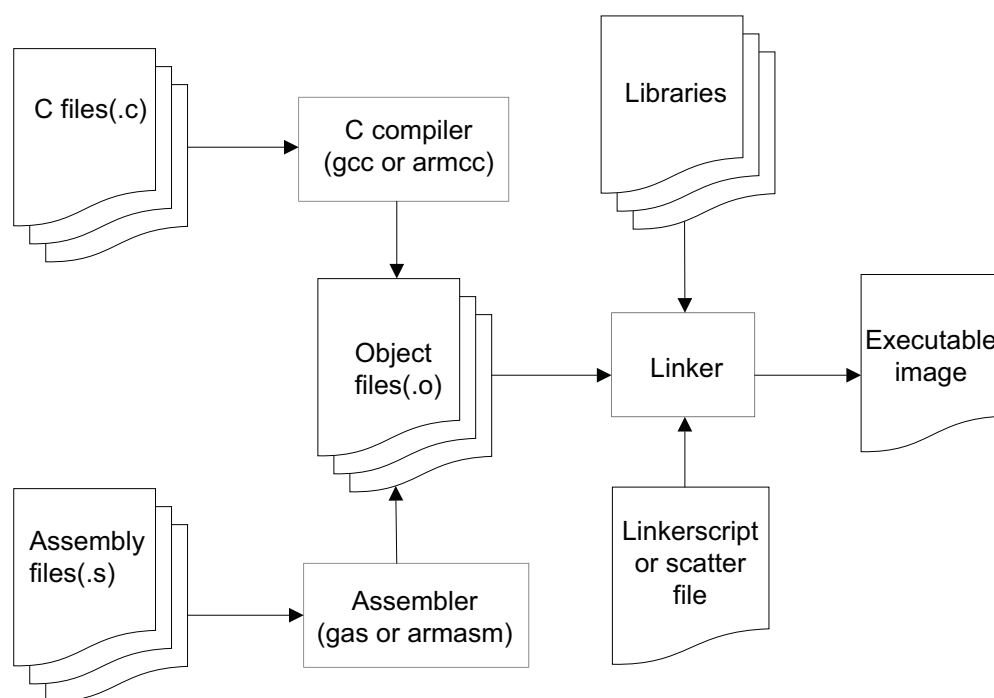


Figure 3-1 Using a software toolchain to produce an image

3.3.1 GNU toolchain

The GNU toolchain is used both to develop the Linux kernel and to develop applications (and indeed other operating systems). Like Linux, the GNU tools are available on a large number of processor architectures and are actively developed to make use of the latest features incorporated in ARM processors.

The toolchain includes the following components:

- GNU make.
- GNU Compiler Collection (GCC).
- GNU binutils linker, assembler and other object/library manipulation tools.
- GNU Debugger (GDB).
- GNU build system (autotools).
- GNU C library (glibc or eglibc).

glibc is available on all GNU Linux host systems and provides portability, wide compliance with standards, and is performance optimized. However, it is quite large for some embedded systems (approaching 2MB in size) so other libraries may be preferred in smaller systems. For example, uClibc provides most features and is around 400KB in size, and produces significantly smaller application binaries. Android does not use glibc, but instead has its own BSD-derived system C library called Bionic.

Prebuilt versions of GNU toolchains

If you are using a complete Linux distribution on your target platform, and you are not cross-compiling, you can install the toolchain packages using the standard package manager. For example, on a Debian-based distribution such as Ubuntu you can use the command:

```
sudo apt-get install gcc g++ gcc-doc
```

Additional required packages such as `binutils` will also be pulled in by this command, or you can add them explicitly on the command line. In fact, if `g++` is specified this way, `gcc` is automatically pulled in. This toolchain will then be accessible in the way you would expect in any Linux system, by just calling `gcc`, `g++`, `as`, or similar.

If you are cross-compiling, you will need to install a suitable cross-compilation toolchain. The cross-compilation toolchain consists of the GNU Compiler Collection (GCC) but also the GNU C library (glibc) which is necessary for building applications (but not the kernel).

Ubuntu distributions from Maverick (10.10) onwards include specific packages for this. These can be run using the command:

```
sudo apt-get install gcc-arm-linux-gnueabi
```

The resulting toolchain will be able to build Linux kernels, applications and libraries for the same Ubuntu version that is used on the build platform. It will however, have a prefix added to all of the individual tool commands in order to avoid problems distinguishing it from the native tools for the workstation. For example, the cross-compiling `gcc` will be accessible as `arm-linux-gnueabi-gcc`.

If your workstation uses an older Ubuntu distribution or an alternative Linux distribution, another toolchain must be used.

Linaro provide up-to-date source packages for ARM toolchains from <http://www.linaro.org/downloads/>. These can be used for generating both cross and native toolchains.

3.3.2 ARM Compiler toolchain

The ARM Compiler toolchain can be used to build programs from C, C++, or ARM assembly language source. It generates optimized code for the 32-bit ARM and mixed length (16-bit and 32-bit) Thumb instruction sets, and supports full ISO standard C and C++. It also supports the NEON SIMD instruction set with the vectorizing (multiple operations simultaneously) NEON compiler.

The ARM Compiler toolchain comprises the following components:

armcc The ARM and Thumb compiler. This compiles your C and C++ code. It supports inline and embedded assembly code, and also includes the NEON vectorizing compiler, invoked using the command:

```
armcc --vectorize
```

armasm The ARM and Thumb assembler. This assembles ARM and Thumb assembly language sources.

- armlink** The linker. This combines the contents of one or more object files with selected parts of one or more object libraries to produce an executable program.
- armar** The librarian. This enables sets of ELF format object files to be collected together and maintained in libraries. You can pass such a library to the linker in place of several ELF files. You can also use the library for distribution to a third party for further application development.
- fromelf** The image conversion utility. This can also generate textual information about the input image, such as disassembly and its code and data size.
- C libraries** The ARM C libraries provide:
- an implementation of the library features as defined in the C and C++ standards
 - extensions specific to the ARM Compiler
 - GNU extensions
 - common nonstandard extensions to many C libraries
 - POSIX extended functionality
 - functions standardized by POSIX (See [Threading libraries on page 23-5](#)).

C++ libraries

The ARM C++ libraries provide:

- helper functions when compiling C++
- additional C++ functions not supported by the Rogue Wave library.

Rogue Wave C++ libraries

The Rogue Wave library provides an implementation of the standard C++ library.

3.4 ARM DS-5

ARM DS-5 is a professional software development solution for Linux, Android and bare-metal embedded systems based on ARM processor based hardware platforms. DS-5 covers all the stages in development, from boot code and kernel porting to application debug. See <http://www.arm.com/products/tools/software-tools/ds-5/index.php>.

ARM DS-5 features an application and kernel space graphical debugger with trace, system-wide performance analyzer, real-time system simulator, and compiler. These features are included in an Eclipse-based IDE.




Figure 3-2 DS-5 Debugger

A full list of the hardware platforms that are supported by DS-5 is available from <http://www.arm.com/products/tools/software-tools/ds-5/supported-platforms.php>.

ARM DS-5 includes the following components:

- Eclipse-based IDE combines software development with the compilation technology of the DS-5 tools. Tools include a powerful C/C++ editor, project manager and integrated productivity utilities such as the Remote System Explorer (RSE), SSH and Telnet terminals.
- DS-5 Compilation Tools. Both GCC and the ARM Compiler are provided. See [ARM Compiler toolchain on page 3-9](#) for more information about the ARM Compiler.
- Real-time simulation models (RTSMs) of complete ARM Cortex-A8 and Cortex-A9 MPCore processor-based devices. Typical simulation speeds are above 250 MHz.

- The DS-5 Debugger, shown in [Figure 3-2 on page 3-11](#), together with a supported debug target, enables debugging of bare-metal applications, Linux kernel, Linux applications and Linux kernel modules. It gives complete control over the flow of program execution to quickly isolate and correct errors. It provides comprehensive and intuitive views, including synchronized source and disassembly, call stack, memory, registers, expressions, variables, threads, breakpoints, and trace and a number of Example projects, including bare-metal startup code examples for the range of ARM processors, and Linux applications example projects that can run models or (via JTAG-based debug hardware) on real hardware.
- DS-5 Streamline, a system-wide software profiling and performance analysis tool for ARM processor based Linux and Android platforms. DS-5 Streamline supports SMP configurations, native Android applications and libraries.

Streamline only requires a standard TCP/IP network connection to the target in order to acquire and analyze system-wide performance data from Linux and Android systems, therefore making it an affordable solution to make software optimization possible from the early stages of the development cycle.

See [DS-5 Streamline on page 18-4](#) for more information.

3.5 Example platforms

In this section we'll mention a few widely available, off-the-shelf ARM processor based platforms which are suitable for use by students or hobbyists for Linux development. This list is likely to become outdated quickly, as newer and better boards are frequently announced. Of course, for mobile application development, your nearest smartphone is a good development platform.

3.5.1 BeagleBoard

The BeagleBoard is a readily available, inexpensive board which provides performance levels similar to that of a laptop from a single fan-less board, powered through a USB connection. It contains an OMAP device from Texas Instruments, which includes a Cortex-A8 processor with a 256KB level 2 cache. It is intended for use by the Open Source community and not to form a part of any commercial product. It has been released in different versions with different amounts of RAM, and clocked at different speeds. Since the design was also freely published, several other cloned or derivative boards have also been produced by third parties.

3.5.2 Pandora

The Pandora device also uses OMAP3530 (a Cortex-A8 processor clocked at 600MHz). It has controls typically found on a gaming console and in fact, looks like a typical handheld gaming device, with an 800x480 LCD.

3.5.3 ST Ericsson Snowball

The Snowball developer board was introduced to drive the software innovation for Android, Linux and Linaro based embedded devices on the NovaTM A9500 platform. Snowball combines ST-Ericsson's Nova A9500 featuring a dual-core Cortex-A9 processor at 1 GHz and Mali-400 MP GPU with GPS and advanced connectivity plus a complete set of STMicroelectronics sensors, including a 3D gyroscope, accelerometer, magnetometer, barometer.

3.5.4 Gumstix

This derives its name from the fact that the board is the same size as a stick of chewing gum. The Gumstix Overo uses the OMAP3503 device from TI, containing a Cortex-A8 processor clocked at 600MHz and runs Linux 2.6 with the BusyBox utilities and OpenEmbedded build environment.

3.5.5 PandaBoard

PandaBoard is a single-board computer based on the Texas Instruments OMAP4430 device, including a dual-core 1GHz ARM Cortex-A9 processor, a 3D Accelerator video processor and 1GB of DDR2 RAM. Its features include Ethernet and Bluetooth plus DVI and HDMI interfaces.

Chapter 4

ARM Registers, Modes and Instruction Sets

In this chapter, we introduce the fundamental features of ARM processors, including details of registers, modes and instruction sets. We will also touch on some details of processor implementation features including instruction pipelines and branch prediction.

The ARMv7 architecture is a 32-bit processor architecture. It is a load/store architecture, meaning that data-processing instructions operate on values in general purpose registers. Only load and store instructions access memory. General purpose registers are also 32 bits. Throughout the book, when we refer to a word, we mean 32 bits. A doubleword is therefore 64 bits and a halfword is 16 bits wide.

Though the ARMv7 architecture is a 32-bit architecture, individual processor implementations do not necessarily have 32-bit width for all blocks and interconnections. For example, it is possible to have 64-bit, or wider paths for instruction fetches or data accesses.

Processors that implement the ARMv7-A architecture do not have a memory map which is fixed by the architecture. The processor has access to a 4GB virtual address space addressed as bytes, and both memory and peripherals can be mapped freely within that space. We will describe memory management further, in [Chapter 9](#) and [Chapter 10](#), where we look at the caches and the Memory Management Unit (MMU).

4.1 Instruction sets

Historically, most ARM processors support more than one instruction set:

- ARM – a 32-bit instruction set.
- Thumb – a 16-bit instruction set, with better code density (but reduced performance compared with ARM code).

The processor can switch back and forth between these two instruction sets, under program control.

All Cortex-A series processors, implement Thumb-2 technology, which extends the Thumb instruction set. This gives a mixture of 32-bit and 16-bit instructions with approximately the code density of the original Thumb instruction set and the performance of the ARM instruction set. Since all Cortex-A series processors support this extension, software targeting them is often compiled to the Thumb instruction set.

4.2 Modes

The ARM architecture has nine processor modes, summarized in [Table 4-1](#). There are eight privileged modes and a non-privileged User mode. *Privilege* is the ability to perform certain tasks that cannot be done from User (*Unprivileged*) mode. In User mode, there are limitations on certain operations, such as MMU access, for example, changing the operating mode is a privileged operation. Note that modes are associated with exception events, which are described further in [Chapter 12 Exception Handling](#).

Table 4-1 ARM processor modes

Mode	Encoding	Function
User (USR)	10000	Unprivileged mode in which most applications run
FIQ	10001	Entered on an FIQ interrupt exception
IRQ	10010	Entered on an IRQ interrupt exception
Supervisor (SVC)	10011	Entered on reset or when a Supervisor Call instruction (SVC) is executed
Monitor (MON)	10110	With Security Extensions (Secure only). See Chapter 26
Abort (ABT)	10111	Entered on a memory access exception
Hyp (HYP)	11010	With Virtualization Extensions (Non-secure only). See Chapter 27
Undef (UND)	11011	Entered when an undefined instruction executed
System (SYS)	11111	Privileged mode, sharing the register view with User mode

4.3 Registers

The ARM architecture has a number of general purpose registers, providing data storage within the body of the processor. In addition to general data storage, there is R15, the program counter, and a Program Status Register which contains ALU flags and other execution state information. Many of these registers are banked, that is, not accessible to the processor except in specific processor modes. These banked-out registers are automatically switched in and out when a different processor mode is entered. Each exception mode (which does not include System mode) additionally has a Saved Program Status Register containing a copy of the Program Status Register at the time the triggering exception was taken.

R0	User mode R0-R7, R15 and CPSR	User mode R0-R12, R15 and CPSR	User mode R0-R12, R15 and CPSR	User mode R0-R12, R15 and CPSR	User mode R0-R12, R15 and CPSR	User mode R0-R12, R15 and CPSR	User mode R0-R12, R15 and CPSR
R1							
R2							
R3							
R4							
R5							
R6							
R7							
R8	R8	R13 (SP) R14 (LR)	R13 (SP) R14 (LR)	R13 (SP) R14 (LR)	R13 (SP) R14 (LR)	R13 (SP) R14 (LR)	R13 (SP) R14 (LR)
R9	R9						
R10	R10						
R11	R11						
R12	R12						
R13 (sp)	R13 (SP)						
R14 (lr)	R14 (LR)						
R15 (pc)							
(A/C)PSR							
	SPSR	SPSR	SPSR	SPSR	SPSR	SPSR	SPSR
User	FIQ	IRQ	ABT	SVC	UND	MON	HYP

Figure 4-1 The ARM register set

So, for example, in [Figure 4-1](#), if the processor is in IRQ mode, we can see R0, R1 ... R12 (the same registers we can see in User mode), plus SP_IRQ and LR_IRQ (registers accessible only while we are in IRQ mode) and R15 (the program counter, PC). SP_USR and LR_USR are not directly accessible. We do not normally need to specify the mode in the register name in the way we have just done. If we (for example) refer to R13 in a line of code, the processor will access the SP register of the mode we are currently in.

At any given moment, the programmer has access to 16 registers (R0–R15) and the *Current Program Status Register* (CPSR), explained further in Section 4.3.1. In User mode, a restricted form of the CPSR called the *Application Program Status Register* (APSR) is accessed instead. R15 is the program counter and holds the current program address (actually, it always points eight bytes ahead of the current instruction in ARM state and four bytes ahead of the current instruction in Thumb state, a legacy of the three stage pipeline of the original ARM1 processor). For more information on the CPSR, see [Program Status Registers on page 4-5](#).

Implementations which support the Virtualization Extensions have additional registers available in Hypervisor (Hyp) mode. Hyp mode has access to its own versions of R13 (SP) and SPSR. It uses the User mode link register as well as a dedicated new register (ELR). We'll discuss this in [Chapter 27 Virtualization](#).

We can write to the PC to change the flow of the program. LR is the link register, which holds a return address for a function or exception (although it can occasionally be used as a general purpose register when not holding either of these values). R13, by convention is used as a stack pointer. R0-R12 are general purpose registers. Some 16-bit Thumb instructions have limitations on which registers they can access – the accessible subset is called the low registers and comprises R0-R7. [Figure 4-2](#) shows the subset of registers visible to general data processing instructions.

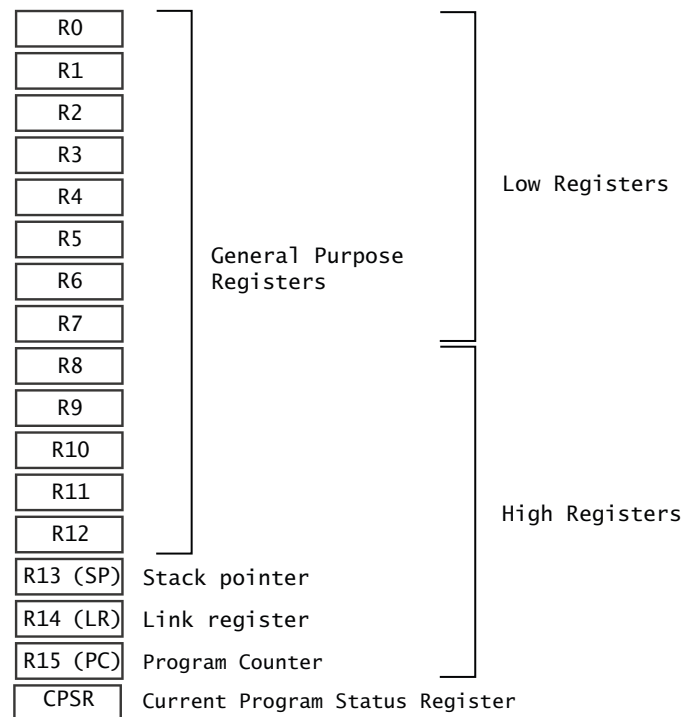


Figure 4-2 Programmer visible registers for user code

The reset values of R0-R14 are unpredictable. SP, the stack pointer, must be initialized (for each mode) by boot code before making use of the stack. The AAPCS or AEABI (see [Chapter 17 Application Binary Interfaces](#)) specifies how software should use the general purpose registers in order to interoperate between different toolchains or programming languages.

4.3.1 Program Status Registers

The program status registers form an additional set of banked registers. Each exception mode has its own *Saved Program Status Register* (SPSR) where a copy of the pre-exception CPSR is stored automatically when an exception occurs. These are not accessible from User modes. While it could be argued that you only have indirect access to banked registers even from the privileged modes (including System), the “indirect” access uses the same mnemonic (though a different encoding) as the “direct” access.

The *ARM Architecture Reference Manual* describes how program status is reported in the 32-bit *Application Program Status Register* (APSR), with other status and control bits (system level information) remaining in the CPSR. In the ARMv7-A architecture covered in this book, the

APSR is in fact the same register as the CPSR, despite the fact that they have two separate names. The APSR must be used only to access the N, Z, C, V, Q, and GE[3:0] bits. These bits are not normally accessed directly, but instead set by condition code setting instructions and tested by instructions which are executed conditionally. The renaming is therefore an attempt to clean-up the mixed access CPSR of the older ARM architectures. [Figure 4-3](#) shows the make-up of the CPSR.




Figure 4-3 CPSR bits

The individual bits represent the following:

- *N* – Negative result from ALU.
- *Z* – Zero result from ALU.
- *C* – ALU operation Carry out.
- *V* – ALU operation oVerflowed.
- *Q* – cumulative saturation (also described as *sticky*).
- *J* – indicates whether the processor is in Jazelle state.
- *GE*[3:0] – used by some SIMD instructions.
- *IT* [7:2] – If-Then conditional execution of Thumb-2 instruction groups.
- *E* bit controls load/store endianness.
- *A* bit disables asynchronous aborts.
- *I* bit disables IRQ.
- *F* bit disables FIQ.
- *T* bit – indicates whether the processor is in Thumb state.
- *M*[4:0] – specifies the processor mode (FIQ, IRQ, etc. as described in [Table 4-1 on page 4-3](#)).

The processor can change between modes using instructions which directly write to the CPSR mode bits (not possible when in User mode). More commonly, the processor changes mode as a result of exception events.

We will consider these bits in more detail in [Chapter 6](#) and [Chapter 12](#).

4.4 Instruction pipelines

The instruction pipeline is a technique used in the design of processors to increase the instruction throughput. Processing of information consists of performing a succession of different operations on input data. To optimize performance, processor cores are designed as blocks of cascaded elementary functional units, each one fed by the output of the previous unit: at every clock tick unit N processes the data output of unit N-1. The main reason for this choice is that simple processing blocks are able to run faster than complex ones, so pipelined structures allow the design of digital circuits working at GHz clock frequencies. Each instruction moves from one step to another, over a number of clock cycles. Each pipeline stage handles a part of the process of executing an instruction, so that on any given clock cycle, a number of different instructions can be in different stages of the pipeline. The total time to execute an individual instruction does not change much compared with a non-pipelined implementation, but the overall throughput is significantly raised. The overall speed of the processor is then governed by the speed of the slowest step, which is significantly less than the time needed to perform all steps. A non-pipelined architecture is less efficient because some blocks within the processor will be idle most of the time during the instruction execution.




Figure 4-4 Pipeline instruction flow

The classic pipeline comprises three stages – Fetch, Decode and Execute as shown in [Figure 4-4](#). More generally, an instruction pipeline might be divided into the following broad definitions:

- Instruction prefetch (deciding from which locations in memory instructions are to be fetched, and performing associated bus accesses).
- Instruction fetch (reading instructions to be executed from the memory system).
- Instruction decode (working out what instruction is to be executed and generating appropriate control signals for the datapaths).
- Register fetch (providing the correct register values to act upon).
- Issue (issuing the instruction to the appropriate execution unit).
- Execute (the actual ALU or multiplier operation, for example).
- Memory access (performing data loads or stores).
- Register write-back (updating processor registers with the results).

In individual processor implementations, some of these steps can be combined into a single pipeline stage, or some steps can be spread over several cycles. A longer pipeline means fewer logic gates in the critical path between each pipeline stage which results in faster execution. However, there are typically many dependencies between instructions. If an instruction depends

on the result of a previous instruction, the control logic might need to insert a stall (or bubble) into the pipeline until the dependency is resolved. Additional logic is needed to detect and resolve such dependencies (for example, forwarding logic, which feeds the output of a pipeline stage back to earlier pipeline stages).

In general, the ARM architecture tries to hide pipeline effects from the programmer. This means that the programmer can determine the pipeline structure only by reading the processor manual. Some pipeline artifacts are still present, however. For example, the program counter register (R15) points two instructions ahead of the instruction that is currently executing when in ARM state, a legacy of the three stage pipeline of the original ARM1 processor.

A further drawback of a long pipeline is that sometimes the sequential execution of instructions from memory will be interrupted. This can happen as a result of execution of a branch instruction, or by an exception event (such as an interrupt). When this happens, the processor cannot determine the correct location from which the next instruction should be fetched until the branch is resolved. In typical code, many branch instructions are conditional as a result of loops or `if` statements. Therefore, whether or not the branch will be taken cannot be determined at the time the instruction is fetched. If we fetch instructions which follow a branch, and the branch is taken, the pipeline must be flushed and a new set of instructions from the branch destination must be fetched from memory instead. As pipelines get longer, the cost of this *branch penalty* becomes higher.

Cortex-A series processors have branch prediction logic which aims to reduce the effect of the branch penalty. In essence, the processor predicts whether a branch will be taken or not and fetches instructions either from the instructions immediately following the branch (if the prediction is that the conditional branch will *not* be taken), or from the target instruction of the branch (if the prediction is that the branch *will* be taken). If the prediction is correct, the branch does not flush the pipeline. If the prediction is wrong, the pipeline must be flushed and instructions from the correct location fetched to refill it. We will look at this in more detail in [Branch prediction on page 4-10](#).

4.4.1 Multi-issue pipelines

A refinement of the processor pipeline is that we can have more than one hardware unit to handle a pipeline stage. In the ARM11 processor family, for example, there are three parallel back-end pipelines – an ALU pipeline, a load/store pipeline and a multiply pipeline. Instructions can be issued into any of these pipelines. A logical development of this idea is to have multiple instances of the execute hardware – for example two ALU pipelines. We can then issue more than one instruction per cycle into these parallel pipelines – an example of instruction level parallelism. Such a processor is said to be *superscalar*. The Cortex-A8, Cortex-A9, and Cortex-A15 processors are superscalar processors – they can potentially decode and issue more than one instruction in a single clock cycle. The Cortex-A5 and Cortex-A7 processors are more limited and can only dual-issue certain combinations of instructions – for example, a branch and a data-processing instruction can be issued in the same cycle. The instructions are still issued from a sequential stream of instructions in memory. Extra hardware logic is required to check for dependencies between instructions, as, for example, in the case where one instruction must wait for the result of the other.

Out-of-order execution provides scope for increasing pipeline efficiency. If instructions are processed sequentially, one instruction is completely retired before the next is dealt with. In out-of-order processing, multiple memory accesses can be outstanding at once, and can complete in a different order from their original program order.

Often, an instruction must be stalled due to a dependency (for example, the need to use a result from a previous instruction). We can execute following instructions which do not share this dependency, provided that logical hazards between instructions are rigorously respected. The Cortex-A9 and Cortex-A15 processors achieve very high levels of efficiency and instruction

throughput using this technique. They can be considered to have a pipeline of variable length, as the pipeline length depends upon which back-end execution pipeline an instruction uses. They can execute instructions speculatively and can decode two instructions per clock, but have the ability to issue up to four instructions on an individual clock cycle. This can improve performance if the pipeline has become unblocked having previously been stalled for some reason.

4.4.2 Register renaming

The Cortex-A9 processor has an interesting micro-architectural implementation which makes use of a register renaming scheme. The set of registers which form a standard part of the ARM architecture are visible to the programmer, but the hardware implementation of the processor actually has a much larger pool of physical registers, with logic to dynamically map the programmer visible registers to the physical ones. Figure 4-5 shows the separate pools of architectural and physical registers.

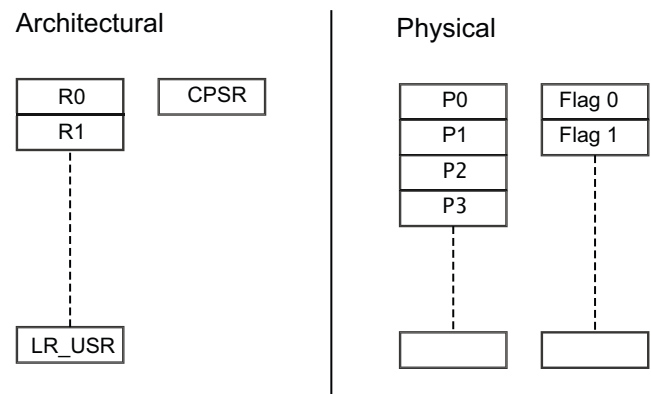


Figure 4-5 Register renaming

Consider the case where code writes the value of a register to external memory and shortly thereafter reads the value of a different memory location into the same register. This might cause a pipeline stall in previous processors, even though in this particular case, there is no actual data dependency. Register renaming avoids this problem by ensuring that the two instances of R0 are renamed to different physical registers, removing the dependency. This permits a compiler or assembler programmer to reuse registers without the need to consider architectural penalties for reusing registers when there are no inter-instruction dependencies. Importantly, it also allows out-of-order execution of write-after-write and write-after-read sequences. (A write-after-write hazard could occur when we write values to the same register in two separate instructions. The processor must ensure that an instruction which comes after the two writes sees the result of the later instruction.)

To avoid dependencies between instructions related to flag setting and comparisons, the APSR flags also use a similar technique.

4.5 Branch prediction

As we have seen, branch prediction logic is an important factor in achieving high throughput in Cortex-A series processors. With no branch prediction, we would have to wait until a conditional branch executes before we could determine where to fetch the next instruction from.

The first time that a conditional jump instruction is fetched, there is little information on which to base a prediction about the address of the next instruction. Older ARM processors used static branch prediction. This is the simplest branch prediction method as it needs no prior information about the branch. We speculate that backward branches will be taken, and forward branches will not. A backward branch has a target address that is lower than its own address. We can therefore look at a single opcode bit to determine the branch direction. This technique can give reasonable prediction accuracy owing to the prevalence in code of loops, which almost always contain backward-pointing branches and are taken more often than not taken. Due to the pipeline length of Cortex-A series processors, we get better performance by using more complex branch prediction schemes, which give better prediction accuracy.

Dynamic prediction hardware can further reduce the average branch penalty by making use of historical information about whether conditional branches were taken or not taken on previous execution. A *Branch Target Address Cache* (BTAC), also called *Branch Target Buffer* (BTB) in the Cortex-A8 processor, is a cache which holds information about previous branch instruction execution. It enables the hardware to speculate on whether a conditional branch will or will not be taken.

The processor must still evaluate the condition code attached to a branch instruction. If the branch prediction hardware predicts correctly, the pipeline does not need to be stalled. If the branch prediction hardware speculation was wrong, the processor will flush the pipeline and refill it.

4.5.1 Return stack

Readers who are not at all familiar with ARM assembly language may want to skip this section until they have read [Chapter 5](#) and [Chapter 6](#).

The description in [Branch prediction](#) looked at strategies the processor can use to predict whether branches are taken or not. For most branch instructions, the target address is fixed and encoded in the instruction. However, there is a class of branches where the branch target destination cannot be determined by looking at the instruction. For example, if we perform a data processing operation which modifies the PC (for example, MOV, ADD or SUB) we must wait for the ALU to evaluate the result before we can know the branch target. Similarly if we load the PC from memory, using an LDR, LDM or POP instruction, we cannot know the target address until the load completes.

Such branches (often called *indirect branches*) cannot, in general, be predicted in hardware. There is, however, one common case that can usefully be optimized, using a last-in-first-out stack in the pre-fetch hardware (the *return stack*). Whenever a function call (BL or BLX) instruction is executed, we enter the address of the following instruction into this stack. Whenever we encounter an instruction which can be recognized as being a function return instructions (BX LR, or a stack pop which contains the PC in its register list), we can speculatively pop an entry from the stack and start fetching instructions from that address. When the return instruction actually executes, the hardware compares the address generated by the instruction with that predicted by the stack. If there is a mismatch, the pipeline is flushed and we restart from the correct location.

The return stack is of a fixed size (eight entries in the Cortex-A8 or Cortex-A9 processors, for example). If a particular code sequence contains a large number of nested function calls, an eight entry return stack can predict only the first eight function returns.

4.5.2 Programmer's view

For the majority of application level programmers, branch prediction is a part of the hardware implementation which can safely be ignored. However, knowledge of the processor behavior with branches can be useful when writing highly optimized code. The hardware performance monitor counters can generate information about the numbers of branches correctly or incorrectly predicted. This hardware is described further in [Chapter 19](#).

Branch prediction logic is disabled at reset. Part of the boot code sequence will typically be to set the Z bit in the CP15:SCTLR, System Control Register, which enables branch prediction. There is one other situation where the programmer might need to take care. When moving or modifying code at an address from which code has already been executed in the system, it might be necessary (and is always prudent) to remove stale entries from the branch history logic by using the CP15 instruction which invalidates all entries.

Chapter 5

Introduction to Assembly Language

Assembly language is a human-readable representation of machine code. There is in general a one-to-one relationship between assembly language instructions (mnemonics) and the actual binary opcode executed by the processor. The purpose of this chapter is not to teach assembly language programming. We describe the ARM and Thumb instruction sets, highlighting features and idiosyncrasies that differentiate it from instruction sets used by other microprocessor families.

Many programmers writing at the application level will have little need to code in assembly language. However, knowledge of assembly code can be useful in cases where highly optimized code is required, when writing compilers, or where low level use of features not directly available in C is needed. It might be required for portions of boot code, device drivers or when performing OS development. Finally, it can be useful to be able to read assembly code when debugging C, and particularly, to understand the mapping between assembly instructions and C statements.

Programmers seeking a more detailed description of ARM assembly language should also refer to the *ARM Compiler Toolchain Assembler Reference*.

The ARM architecture supports implementations across a very wide range of performance points. Its simplicity leads to very small implementations, and this enables very low power consumption. Implementation size, performance, and very low power consumption are key attributes of the ARM architecture.

5.1 Comparison with other assembly languages

All processors have basic data processing instructions which permit them to perform arithmetic operations (such as ADD) and logical bit manipulation (for example AND). They also need to transfer program execution from one part of the program to another, in order to support loops and conditional statements. Processors always have instructions to read and write external memory, too.

The ARM instruction set is generally considered to be simple, logical and efficient. It has features not found in other processors, while at the same time lacking operations found in some other processors. For example, it cannot perform data processing operations directly on memory. To increment a value in a memory location, the value must be loaded to an ARM register, the register incremented and a third instruction is required to write the updated value back to memory. The *Instruction Set Architecture* (ISA) includes instructions that combine a shift with an arithmetic or logical operation, auto-increment and auto-decrement addressing modes for optimized program loops, Load, and Store Multiple instructions which allow efficient stack and heap operations, plus block copying capability and conditional execution of almost all instructions.

As many readers will already be familiar with one or more assembly languages, it might be useful to compare some code sequences, showing the x86, 68K and ARM instructions to perform equivalent tasks.

Like the x86 (but unlike the 68K), ARM instructions typically have a two or three operand format, with the first operand in most cases specifying the destination for the result (LDM and store instructions, for example, being an exception to this rule). The 68K, by contrast, places the destination as the last operand. For ARM instructions, there are generally no restrictions on which registers can be used as operands. [Example 5-1](#) and [Example 5-2](#) give a flavor of the differences between the different assembly languages.

Example 5-1 Instructions to add 100 to a value in a register

x86:	add	eax, #100
68K:	ADD	#100, D0
ARM:	add	r0, r0, 100

Example 5-2 Load a register with a 32-bit value from a register pointer

x86:	mov	eax, DWORD PTR [ebx]
68K:	MOVE.L	(A0), D0
ARM:	ldr	r0, [r1]

An ARM processor is a *Reduced Instruction Set Computer* (RISC) processor. *Complex Instruction Set Computer* (CISC) processors, like the x86, have a rich instruction set capable of doing complex things with a single instruction. Such processors often have significant amounts of internal logic which decode machine instructions to sequences of internal operations (microcode). RISC architectures, in contrast, have a smaller number of more general purpose instructions, which might be executed with significantly fewer transistors, making the silicon cheaper and more power efficient. Like other RISC architectures, ARM processors have a large

number of general-purpose registers and many instructions execute in a single cycle. It has simple addressing modes, where all load/store addresses can be determined from just register contents and instruction fields.

5.2 Instruction sets

As described in [Chapter 4](#), many ARM processors are able to execute two or even three different instruction sets, while some (for example, the Cortex-M3 processor) do not in fact execute the original ARM instruction set. There are at least two instruction sets that ARM processors can use.

ARM (32-bit instructions)

This is the original ARM instruction set.

Thumb The Thumb instruction set was first added in the ARM7TDMI processor and contained only 16-bit instructions, which gave much smaller programs (memory footprint can be a major concern in smaller embedded systems) at the cost of some performance. Recent processors, including those in the Cortex-A series, support Thumb-2 technology, which extends the Thumb instruction set to provide a mix of 16-bit and 32-bit instructions. This gives the best of both worlds, performance similar to that of ARM, with code size similar to that of Thumb. Due to its size and performance advantages, it is increasingly common for all code to be compiled or assembled to take advantage of Thumb-2 technology.

The currently used instruction set is indicated by the CPSR T bit and the processor is said to be in *ARM state* or *Thumb state*. Code has to be explicitly compiled or assembled to one state or the other. An explicit instruction is used to change between instruction sets. Calling functions which are compiled for a different state is known as interworking. We'll take a more detailed look at this in [Interworking on page 5-11](#).

For Thumb assembly code, there is often a choice of 16-bit and 32-bit instruction encodings, with the 16-bit versions being generated by default. The .W (32-bit) and .N (16-bit) width specifiers can be used to force a particular encoding (if such an encoding exists).

5.3 Introduction to the GNU Assembler

The GNU Assembler, part of the GNU tools, is used to convert assembly language source code into binary object files. The assembler is extensively documented in the GNU Assembler Manual, which can be found online at <http://sourceware.org/binutils/docs/as/index.html> or (if you have GNU tools installed on your system) in the `gnutools/doc` sub-directory. GNU Assembler documentation is also available in the `/gcc-doc/` package on Ubuntu.

What follows is a brief description, intended to highlight differences in syntax between the GNU Assembler and standard ARM assembly language, and to provide enough information to allow programmers to get started with the tools.

The names of GNU tool components will have prefixes indicating the target options selected, including operating system. An example would be `arm-none-eabi-gcc`, which might be used for bare metal systems using the ARM EABI (described in [Chapter 20 Writing NEON Code](#)).

5.3.1 Invoking the GNU Assembler

You can assemble the contents of an ARM assembly language source file by running the `arm-none-eabi-as` program.

```
arm-none-eabi-as -g -o filename.o filename.s
```

The option `-g` requests the assembler to include debug information in the output file.

When all of your source files have been assembled into binary object files (with the extension `.o`), you use the GNU Linker to create the final executable in ELF format.

This is done by executing:

```
arm-none-eabi-ld -o filename.elf filename.o
```

For more complex programs, where there are many separate source files, it is more common to use a utility like `make` to control the build process.

You can use the debugger provided by either `arm-none-eabi-gdb` or `arm-none-eabi-insight` to run the executable files on your host machine, as an alternative to a real target processor.

5.3.2 GNU Assembler syntax

The GNU Assembler can target many different processor architectures and is not ARM-specific. This means that its syntax is somewhat different from other ARM assemblers, such as the ARM toolchain. The GNU Assembler uses the same syntax for all of the many processor architectures that it supports.

Assembly language source files consist of a sequence of statements, one per line.

Each statement has three optional parts, ordered as follows:

```
label: instruction @ comment
```

A *label* lets you identify the address of this instruction. This can then be used as a target for branch instructions or for load and store instructions. A label can be a letter followed (optionally) by a sequence of alphanumeric characters, followed by a colon.

The *instruction* can be either an ARM assembly instruction, or an assembler directive. These are pseudo-instructions that tell the assembler itself to do something. These are required, amongst other things, to control sections and alignment, or create data.

Everything on the line after the `@` symbol is treated as a comment and ignored (unless it is inside a string). C style comment delimiters `/*` and `*/` can also be used.

At link time an entry point can be specified on the command line if one has not been explicitly provided in the source code.

5.3.3 Sections

An executable program with code will have at least one section, which by convention will be called `.text`. Data can be included in a `.data` section.

Directives with the same names enable you to specify which of the two sections should hold what follows in the source file. Executable code should appear in a `.text` section and read or write data in the `.data` section. Also read-only constants can appear in a `.rodata` section. Zero initialized data will appear in `.bss`. The Block Started by Symbol (bss) segment defines the space for uninitialized static data.

5.3.4 Assembler directives

This is a key area of difference between GNU tools and other assemblers.

All assembler directives begin with a period “.” A full list of these is described in the GNU documentation. Here, we give a subset of commonly used directives.

- `.align` This causes the assembler to pad the binary with bytes of zero value, in data sections, or NOP instructions in code, ensuring the next location will be on a word boundary. `.align n` gives 2^n alignment on ARM processors.
- `.ascii “string”`
 Insert the string literal into the object file exactly as specified, without a NUL character to terminate. Multiple strings can be specified using commas as separators.
- `.asciiz` Does the same as `.ascii`, but this time additionally followed by a NUL character (a byte with the value 0 (zero)).
- `.byte expression, .hword expression, .word expression`
 Inserts a byte, halfword, or word value into the object file. Multiple values can be specified using commas as separators. The synonyms `.2byte` and `.4byte` can also be used.
- `.data` Causes the following statements to be placed in the data section of the final executable.
- `.end` Marks the end of this source code file. The assembler does not process anything in the file after this point.
- `.equ symbol, expression`
 Sets the value of *symbol* to *expression*. The “=” symbol and `.set` have the same effect.
- `.extern symbol`
 Indicates that *symbol* is defined in another source code file.
- `.global symbol`
 Tells the assembler that *symbol* is to be made globally visible to other source files and to the linker.

`.include "filename"`

Inserts the contents of *filename* into the current source file and is typically used to include header files containing shared definitions.

`.text`

This switches the destination of following statements into the text section of the final output object file. Assembly instructions must always be in the text section.

For reference, [Table 5-1](#) shows common assembler directives alongside GNU and ARM tools. Not all directives are listed, and in some cases there is not a 100% correspondence between them.

Table 5-1 Comparison of syntax

GNU Assembler	armasm	Description
@	;	Comment
#&	#0x	An immediate hex value
.if	IFDEF, IF	Conditional (not 100% equivalent)
.else	ELSE	
.elseif	ELSEIF	
.endif	ENDIF	
.ltorg	LTORG	
	:OR:	OR
&	:AND:	AND
<<	:SHL:	Shift Left
>>	:SHR:	Shift Right
.macro	MACRO	Start macro definition
.endm	ENDM	End macro definition
.include	INCLUDE	GNU Assembler needs "filename"
.word	DCD	A data word
.short	DCW	
.long	DCD	
.byte	DCB	
.req	RN	
.global	IMPORT, EXPORT	
.equ	EQU	

5.3.5 Expressions

Assembly instructions and assembler directives often require an integer operand. In the assembler, this is represented as an expression to be evaluated. Typically, this will be an integer number specified in decimal, hexadecimal (with a *0x* prefix) or binary (with a *0b* prefix) or as an ASCII character surrounded by single quotes.

In addition, standard mathematical and logical expressions can be evaluated by the assembler to generate a constant value. These can utilize labels and other pre-defined values. These expressions produce either *absolute* or *relative* values. Absolute values are position-independent and constant. Relative values are specified relative to some linker-defined address, determined when the executable image is produced – such as target addresses for branches.

5.3.6 GNU tools naming conventions

Registers are named in GCC as follows:

- General registers: R0 - R15.
- Stack pointer register: SP (R13).
- Frame pointer register: FP (R11).
- Link register: LR (R14).
- Program counter: PC (R15).
- Program Status Register flags: *xPSR*, *xPSR_all*, *xPSR_f*, *xPSR_x*, *xPSR_ctl*, *xPSR_fs*, *xPSR_fx*, *xPSR_f*, *xPSR_cs*, *xPSR_cf*, *xPSR_cx* (where *x* = C current or S saved). See [Program Status Registers](#) on page 4-5.

Note

In [Chapter 17 Application Binary Interfaces](#) we will see how all of the registers are assigned a role within the procedure call standard and that the GNU Assembler lets us refer to the registers using their *Procedure Call Standard* (PCS) names. See [Table 17-1 on page 17-2](#).

5.4 ARM tools assembly language

The *Unified Assembly Language* (UAL) format now used by ARM tools enables the same canonical syntax to be used for both ARM and Thumb instruction sets. The assembler syntax of ARM tools is not identical to that used by the GNU Assembler, particularly for preprocessing and pseudo-instructions which do not map directly to opcodes. In the next chapter, we will look at the individual assembly language instructions in a little more detail. Before doing that, we take a look at the basic syntax used to specify instructions and registers. Assembly language examples in this book use both UAL and GNU Assembly syntax.

UAL gives the ability to write assembler code which can be assembled to run on all ARM processors. In the past, it was necessary to write code explicitly for ARM or Thumb state. Using UAL the same code can be assembled for different instruction sets at the time of assembly, not at the time the code is written. This can be either through the use of command line switches or inline directives. Legacy code will still assemble correctly. It is worth noting that GNU Assembler now supports UAL through use of the `.syntax` directive, though it might not be identical syntax to the ARM tools assembler.

5.4.1 ARM assembler syntax

ARM assembler source files consist of a sequence of statements, one per line.

Each statement has three optional parts, ordered as follows:

```
label instruction ; comment
```

A *label* lets you identify the address of this instruction. This can then be used as a target for branch instructions or for load and store instructions.

The *instruction* can be either an assembly instruction, or an assembler directive. These are pseudo-instructions that tell the assembler itself to do something. These are required, amongst other things, to control sections and alignment, or create data.

Everything on the line after the `;` symbol is treated as a comment and ignored (unless it is inside a string). C style comment delimiters `“/*”` and `“*/”` can also be used.

5.4.2 Label

A label is required to start in the first character of a line. If the line does not have a label, a space or tab delimiter is needed to start the line. If there is a label, the assembler makes the label equal to the address in the object file of the corresponding instruction. Labels can then be used as the target for branches or for loads and stores.

Example 5-3 A simple example showing use of a label

```
Loop      MUL R5, R5, R1
          SUBS R1, R1, #1
          BNE Loop
```

In [Example 5-3](#), `Loop` is a label and the conditional branch instruction (`BNE Loop`) will be assembled in a way which makes the offset encoded in the branch instruction point to the address of the `MUL` instruction which is associated with the label `Loop`.

5.4.3 Directives

Most lines will normally have an assembly language instruction, to be converted by the tool into its binary equivalent, or a directive which tells the assembler to do something. It can also be a pseudo-instruction (one which will be converted into one or more real instructions by the assembler). We'll look at the actual instructions available in hardware in [Chapter 6](#) and focus mainly on the assembler directives here. These perform a wide range of tasks. They can be used to place code or data at a particular address in memory, create references to other programs and so forth.

The Define Constant (DCD, DCB, DCW) directives let us place data into a piece of code. This can be expressed numerically (in decimal, hex, binary) or as ASCII characters. It can be a single item or a comma separated list. DCB is for byte sized data, DCD can be used for word sized data, and DCW for half-word sized data items.

For example, we might have:

```
MESSAGE DCB "Hello World!",0
```

This will produce a series of bytes corresponding to the ASCII characters in the string, with a 0 termination. MESSAGE is a label which we can use to get the address of this data. Similarly, we might have data items expressed in hex:

```
Masks DCD 0x100, 0x80, 0x40, 0x20, 0x10
```

The EQU directive lets us assign names to address or data values. For example:

```
CtrID EQU 4
TUBE EQU 0x30000000
```

We can then use these labels in other instructions, as parts of expressions to be evaluated. EQU does not actually cause anything to be placed in the program executable – it merely sets a name to a value, for use in other instructions, in the symbol table for the assembler. It is convenient to use such names to make code easier to read, but also so that if we change the address or value of something in a piece of code, we need only modify the original definition, rather than having to change all of the references to it individually. It is usual to group EQU definitions together, often at the start of a program or function, or in separate include files.

The AREA pseudo-instruction is used to tell the assembler about how to group together code or data into logical sections for later placement by the linker. For example, exception vectors might need to be placed at a fixed address. The assembler keeps track of where each instruction or piece of data is located in memory. The AREA directive can be used to modify that location.

The ALIGN directive lets you align the current location to a specified boundary. It usually does this by padding (where necessary) with zeros or NOP instructions, although it is also possible to specify a pad value with the directive. The default behavior is to set the current location to the next word (four byte) boundary, but larger boundary sizes and offsets from that boundary can also be specified. This can be required to meet alignment requirements of certain instructions (for example LDRD and STRD doubleword memory transfers), or to align with cache boundaries. As with the .align directive on GNU Assembler, the ALIGN n directive gives 2^n alignment on ARM processors.

END is used to denote the end of the assembly language source program. Failure to use the END directive will result in an error being returned. INCLUDE tells the assembler to include the contents of another file into the current file. Include files can be used as an easy mechanism for sharing definitions between related files.

5.5 Interworking

When the processor executes ARM instructions, it is said to be operating in *ARM state*. When it is operating in *Thumb state*, it is executing Thumb instructions. A processor in a particular state can only sensibly execute instructions from that instruction set. We must make sure that the processor does not receive instructions of the wrong instruction set.

Each instruction set includes instructions to change processor state. ARM and Thumb code can be mixed, if the code conforms to the requirements of the ARM and Thumb Procedure Call Standards (described in [Chapter 17](#)). Compiler generated code will always do so, but assembly language programmers must take care to follow the specified rules.

Selection of processor state is controlled by the T bit in the current program status register (CPSR). See [Figure 4-3 on page 4-6](#). When T is 1, the processor is in Thumb state. When T is 0, the processor is in ARM state. However, when the T bit is modified, it is also necessary to flush the instruction pipeline (to avoid problems with instructions being decoded in one state and then executed in another). Special instructions are used to accomplish this. These are BX (Branch with eXchange) and BLX (Branch and Link with eXchange). LDR of PC and POP/LDM of PC also have this behavior. In addition to changing the processor state with these instructions, assembly programmers must also use the appropriate directive to tell the assembler to generate code for the appropriate state.

The BX or BLX instruction branches to an address contained in the specified register, or an offset specified in the opcode. The value of bit [0] of the branch target address determines whether execution continues in ARM state or Thumb state. Both ARM (aligned to a word boundary) and Thumb (aligned to a halfword boundary) instructions do not use bit [0] to form an address. This bit can therefore safely be used to provide the additional information about whether the BX or BLX instruction should change the state to ARM (address bit [0] = 0) or Thumb (address bit [0] = 1). The BL *label* will be turned into a BLX *label* as appropriate at link time if the instruction set of the caller is different from the instruction set of *label*, assuming that it is unconditional.

A typical use of these instructions is when a call from one function to another is made using the BL or BLX instruction, and a return from that function is made using the BX LR instruction. Alternatively, we can have a non-leaf function, which pushes the link register onto the stack on entry and pops the stored link register from the stack into the program counter, on exit. Here, instead of using the BX LR instruction to return, we instead have a memory load. Memory load instructions which modify the PC might also change the processor state depending upon the value of bit [0] of the loaded address.

5.6 Identifying assembly code

When faced with a piece of assembly language source code, it can be useful to be able to quickly determine which instruction set will be used and which kind of assembler it is targeted at.

Older ARM Assembly language code can have three (or even four) operand instructions present (for example, `ADD R0, R1, R2`) or conditional execution of non-branch instructions (for example, `ADDNE R0, R0, #1`). Filename extensions will typically be `.s` or `.S`.

Code targeted for the newer UAL, will contain the directive `.syntax unified` but will otherwise appear similar to traditional ARM Assembly language. The pound (or hash) symbol `#` can be omitted in front of immediate operands. Conditional instruction sequences must be preceded immediately by the `IT` instruction (described in [Chapter 6](#)). Such code assembles either to fixed-size 32-bit (ARM) instructions, or mixed-size (16-bit and 32-bit) Thumb instructions, depending on the presence of the directives `.thumb` or `.arm`.

You can, on occasion, encounter code written in 16-bit Thumb assembly language. This can contain directives like `.code 16`, `.thumb` or `.thumb_func` but will not specify `.syntax unified`. It uses two operands for most instructions, although `ADD` and `SUB` can sometimes have three. Only branches can be executed conditionally.

All GCC inline assembler (`.c`, `.h`, `.cpp`, `.cxx`, `.c++` and so on) code can be built for Thumb or ARM, depending on GCC configuration and command-line switches (`-marm` or `-mthumb`).

Chapter 6

ARM/Thumb Unified Assembly Language Instructions

This chapter is a general introduction to ARM/Thumb assembly language. We do not aim to provide detailed coverage of every instruction. Instructions can broadly be placed in one of a number of classes:

- *Data operations* (ALU operations such as ADD).
- *Memory operations* (load and stores to memory).
- *Branches* (for loops, goto, conditional code and other program flow control).
- *DSP* (operations on packed data, saturated mathematics and other special instructions, targeting codecs).
- *Miscellaneous* (coprocessor, debug, mode changes and so forth).

We'll take a brief look at each of those in turn. Before we do that, let us examine capabilities which are common to different instruction classes.

6.1 Instruction set basics

There are a number of features common to all parts of the instruction set.

6.1.1 Constant values

ARM or Thumb assembly language instructions have a length of only 16 or 32 bits. This presents something of a problem. It means that we cannot encode an arbitrary 32-bit value within the opcode.

Constant values encoded in an instruction can be one of the following in Thumb:

- a constant that can be produced by rotating an 8-bit value by any even number of bits within a 32-bit word
- a constant of the form `0x00XY00XY`
- a constant of the form `0xXY00XY00`
- a constant of the form `0xXYXYXYXY`.

where `XY` is a hexadecimal number in the range `0x00` to `0xFF`.

In the ARM instruction set, as opcode bits are used to specify condition codes, the instruction itself and the registers to be used, only 12 bits are available to specify an immediate value. We have to be somewhat creative in how these 12 bits are used. Rather than enabling a constant of size -2048 to $+2047$ to be specified, instead the 12 bits are divided into an 8-bit constant and 4-bit rotate value. The rotate value enables the 8-bit constant value to be rotated right by a number of places from 0 to 30 in steps of 2 (that is, 0, 2, 4, 6, 8 and so on)

So, we can have immediate values like `0x23` or `0xFF`. We can produce other useful immediate values (for example, addresses of peripherals or blocks of memory), for example, `0x23000000` can be produced by expressing it as `0x23 ROR 8` (see [ROR](#) on page A-31). But many other constants, like `0x3FF`, cannot be produced within a single instruction. For these values, you must either construct them in multiple instructions, or load them from memory. Programmers do not typically concern themselves with this, except where the assembler gives an error complaining about an invalid constant. Instead, we can use assembly language pseudo-instructions to do whatever is necessary to generate the required constant.

The `MOVW` instruction (move wide), will move a 16-bit constant into a register, while zeroing the top 16 bits of the target register. `MOVT` (move top) will move a 16-bit constant into the top half of a given register, without changing the bottom 16 bits. This permits a `MOV32` pseudo-instruction which is able to construct any 32-bit constant. The assembler provides some further help here. The prefixes `:upper16:` and `:lower16:` allow you to extract the corresponding half from a 32-bit constant:

```
MOVW R0, #:lower16:label
MOVT R0, #:upper16:label
```

Although this needs two instructions, it does not require any extra space to store the constant, and there is no need to read a data item from memory.

We can also use pseudo-instructions `LDR Rn, =<constant>` or `LDR Rn, =label`. (This was the only option for older processors which lacked `MOVW` and `MOVT`.) The assembler will then use the best sequence to generate the constant in the specified register (one of `MOV`, `MVN` or an `LDR` from a *literal pool*). A literal pool is an area of constant data held within the code section, typically after the end of a function and before the start of another. If it is necessary to manually control literal pool placement, this can be done with an assembler directive – `LTORG` for `armasm`, or `.ltorg` when using GNU tools. The register loaded could be the program counter, which would cause a

branch. This can be useful for absolute addressing or for references outside the current section; obviously this will result in position-dependent code. The value of the constant can be determined either by the assembler, or by the linker.

ARM tools also provides the related pseudo-instruction `ADR Rn, =label`. This uses a PC-relative `ADD` or `SUB`, to place the address of the label into the specified register, using a single instruction. If the address is too far away to be generated this way, the `ADRL` pseudo-instruction is used. This requires two instructions, which gives a better range. This can be used to generate addresses for position-independent code, but only within the same code section.

6.1.2 Conditional execution

A feature of the ARM instruction set is that nearly all instructions are conditional. On most other architectures, only branches or jumps can be executed conditionally. This can be useful in avoiding conditional branches in small `if/then/else` constructs or for compound comparisons.

As an example of this, consider code to find the smaller of two values, in registers `R0` and `R1` and place the result in `R2`. This is shown in [Example 6-1](#). The suffix `LT` indicates that the instruction should be executed only if the most recent flag-setting instruction returned *less than*; `GE` means *greater than or equal*.

Example 6-1 Example code showing branches (GNU)

```
@ Code using branches
CMP    R0, R1
BLT    .Lsmaller @ if R0<R1 jump over
MOV    R2, R1    @ R1 is smaller
B      .Lend     @ finish
.Lsmaller:
MOV    R2, R0    @ R0 is smaller
.Lend:
```

Now look at the same code written using conditional `MOV` instructions, rather than branches, in [Example 6-2](#)

Example 6-2 Same example using conditional execution

```
CMP    R0, R1
MOVGE   R2, R1 @ R1 is smaller
MOVLT   R2, R0 @ R0 is smaller
```

The latter piece of code is both smaller and, on older ARM processors, is faster to execute. However, as we shall see when we look further at optimization in [Chapter 19](#), this code can actually be slower on processors like the Cortex-A9, where inter-instruction dependencies could cause longer stalls than a branch.

As a reminder, this style of programming relies on the fact that status flags can be set optionally on some instructions. If the `MOVGE` instruction above automatically set the flags, the program might not work correctly. Load and Store instructions never set the flags. For data processing operations, however, the programmer has a choice. By default, flags will be preserved during such instructions. If the instruction is suffixed with an `S` (for example, `MOVVS` rather than `MOV`), the instruction will set the flags. The `S` suffix is not required, or permitted, for the explicit comparison instructions. The flags can also be set manually, by using the dedicated `PSR`

manipulation instruction (MSR). Some instructions set the Carry flag (C) based on the carry from the ALU and others based on the barrel shifter carry (which shifts a data word by a specified number of bits in one clock cycle).

Thumb-2 technology also introduced an If-Then (IT) instruction, providing conditional execution for up to four consecutive instructions. The conditions may all be identical, or some may be the inverse of the others. Instructions within an IT block must also specify the condition code to be applied.

IT is a 16-bit instruction that enables nearly all Thumb instructions to be conditionally executed, depending on the value of the ALU flags, using the condition code suffix (see [IT on page A-12](#)). The syntax of the instruction is `IT{x}{y}{z}}` where *x*, *y* and *z* specify the condition switch for the optional instructions in the IT block, either Then (T) or Else (E), for example, ITTET.

```
ITT    EQ
SUBEQ  r1, r1, #1
ADDEQ  r0, r0, #60
```

Typically, IT instructions are auto-generated by the assembler, rather than being hand-coded. 16-bit instructions which normally change the condition code flags, will not do so inside an IT block, except for CMP, CMN and TST whose only action is to set flags. There are some restrictions on which instructions can be used within an IT block. Exceptions can occur within IT blocks, the current if-then status is stored in the CPSR and so is copied into the SPSR upon exception entry, so that when the exception returns, the execution of the IT block resumes correctly.

Certain instructions always set the flags and have no other effect. These are CMP, CMN, TST and TEQ, which are analogous to SUBS, ADDS, ANDS and EORS but with the result of the ALU calculation being used only to update the flags and not being placed in a register.

[Table 6-1](#) lists the 15 condition codes that can be attached to most instructions.

Table 6-1 Conditional execution suffixes

Suffix	Flags	Description
EQ	Z = 1	zero (EQual to 0)
NE	Z = 0	not zero (Not Equal to 0)
CS/HS	C = 1	Carry Set / unsigned Higher or Same
CC/LO	C = 0	Carry Clear / unsigned LOwer
MI	N = 1	negative (MInus)
PL	N = 0	positive or zero (PLus)
VS	V = 1	sign overflow (oVerflow Set)
VC	V = 0	no sign overflow (oVerflow Clear)
HI	C = 1 AND Z = 0	unsigned HIgher
LS	C = 0 OR Z = 1	unsigned LOwer or Same
GE	N = V	signed GEater or Equal
LT	N != V	signed LEss Than

Table 6-1 Conditional execution suffixes (continued)

Suffix	Flags	Description
GT	Z = 0 AND N = V	signed Greater Than
LE	Z = 1 OR N != V	signed Less or Equal
AL	-	ALways (default)

Thumb code has a somewhat different mechanism for conditional execution. Branches can be executed conditionally. Instructions can also be conditionally executed by using the Compare and Branch on Zero (CBZ) and Compare and Branch on Non-Zero (CBNZ) instructions. These compare the value of a register against zero and branch on the result. In addition, instructions can be conditionally executed using the If-Then construct.

6.1.3 Status flags and condition codes

When we looked at the register set in [Program Status Registers on page 4-5](#), we stated that the ARM processor has a *Current Program Status Register* (CPSR) which contains four status flags, (Z)ero, (N)egative, (C)arry and o(V)erflow. [Table 6-2](#) indicates the value of these bits for flag setting operations.

Table 6-2 Summary of PSR flag bits

Flag	Bit	Name	Description
N	31	Negative	Set to the same value as bit[31] of the result. For a 32-bit signed integer, bit[31] being set indicates that the value is negative.
Z	30	Zero	Set to 1 if the result is zero, otherwise it is set to 0.
C	29	Carry	Set to the carry-out value from result, or to the value of the last bit shifted out from a shift operation.
V	28	Overflow	Set to 1 if signed overflow or underflow occurred, otherwise it is set to 0.

The operation of the negative and zero flags should be easily understood. The C flag will be set if the result of an unsigned operation overflows the 32-bit result register. This bit might be used to implement 64-bit (or longer) arithmetic using 32-bit operations, for example.

The V flag operates in the same way as the C flag, but for signed operations. 0x7FFFFFFF is the largest signed positive integer that can be represented in 32 bits. If, for example, we add 2 to this value, we will produce 0x80000001, a large negative number. The V bit is set to indicate the overflow or underflow, from bit [30] to bit [31].

6.2 Data processing operations

These are essentially the fundamental arithmetic and logical operations of the processor. Multiplies can be considered a special case of these – they typically have slightly different format and rules and are executed in a dedicated unit of the processor.

The ARM processors can only perform data processing on registers, never directly on memory. Data processing instructions (for the most part) use one destination register and two source operands. The basic format can be considered to be the opcode, optionally followed by a condition code, optionally followed by S (set flags), as follows:

Operation{cond}{s} Rd, Rn, Operand2

[Table 6-3](#) summarizes the data processing assembly language instructions, giving their mnemonic opcode, operands and a brief description of their function. [Appendix A](#) gives a fuller description of all of the available instructions.

Table 6-3 Summary of data processing operations in assembly language

Opcode	Operands	Description	Function
Arithmetic operations			
ADC	Rd, Rn, Op2	Add with carry	$Rd = Rn + Op2 + C$
ADD	Rd, Rn, Op2	Add	$Rd = Rn + Op2$
MOV	Rd, Op2	Move	$Rd = Op2$
MVN	Rd, Op2	Move NOT	$Rd = \sim Op2$
RSB	Rd, Rn, Op2	Reverse Subtract	$Rd = Op2 - Rn$
RSC	Rd, Rn, Op2	Reverse Subtract with Carry	$Rd = Op2 - Rn - !C$
SBC	Rd, Rn, Op2	Subtract with carry	$Rd = Rn - Op2 - !C$
SUB	Rd, Rn, Op2	Subtract	$Rd = Rn - Op2$
Logical operations			
AND	Rd, Rn, Op2	AND	$Rd = Rn \& Op2$
BIC	Rd, Rn, Op2	Bit Clear	$Rd = Rn \& \sim Op2$
EOR	Rd, Rn, Op2	Exclusive OR	$Rd = Rn \wedge Op2$
ORR	Rd, Rn, Op2	OR	$Rd = Rn Op2$ (OR NOT) $Rd = Rn \sim Op2$
Flag setting instructions			
CMP	Rn, Op2	Compare	$Rn - Op2$
CMN	Rn, Op2	Compare Negative	$Rn + Op2$
TEQ	Rn, Op2	Test EQuivalence	$Rn \& Op2$
TST	Rn, Op2	Test	$Rn \wedge Op2$

The purpose and function of many of these instructions will be readily apparent to most programmers, but some require additional explanation.

In the arithmetic operations, notice that the move operations MOV and MVN require only one operand (and this is treated as an operand 2 for maximum flexibility, as we shall see). RSB does a reverse subtract – that is to say it subtracts the first operand from the second operand. This instruction is needed because the first operand is inflexible – it can only be a register value. So to write $R0 = 100 - R1$, we must do `RSB R0, R1, #100`, as we cannot write `SUB R0, #100, R1`. The operations ADC and SBC perform additions and subtractions with carry. This lets the programmer synthesize arithmetic operations on values larger than 32 bits.

The logical operations are essentially the same as the corresponding C operators. Notice the use of ORR rather than OR (this is because the original ARM instruction set had three letter acronyms for all data-processing operations). The BIC instruction does an AND of a register with the inverted value of operand 2. If, for example, we wish to clear bit [11] of register R0, we can do this with the instruction `BIC R0, R0, #0x800`.

The second operand `0x800` has only bit [11] set to one, with all other bits at zero. The BIC instruction inverts this operand, setting all bits except bit [11] to logical one. ANDing this value with the value in R0 has the effect of clearing bit [11] and this result is then written back into R0.

The compare and test instructions modify the CPSR (and have no other effect).

6.2.1 Operand 2 and the barrel shifter

The first operand for all data processing operations must always be a register. The second operand is much more flexible and can be either an immediate value (`#x`), a register (`Rm`), or a register shifted by an immediate value or register “`Rm, shift #x`” or “`Rm, shift Rs`”. There are five shift operations: left shift (LSL), logical right-shift (LSR), arithmetic right-shift (ASR), rotate-right (ROR) and rotate-right extended (RRX).

A right shift creates empty positions at the top of the register. In that case, we must differentiate between a logical shift, which inserts 0 into the most significant bit(s) and an arithmetic shift, which fills vacant bits with the sign bit, from bit [31] of the register. So an ASR operation might be used on a signed value, with LSR used on an unsigned value. No such distinction is required on left-shifts, which always insert 0 to the least significant position.

So, unlike many assembly languages, ARM assembly language does not require explicit shift instructions. Instead, the MOV instruction can be used for shifts and rotates. $R0 = R1 \gg 2$ is done as `MOV R0, R1, LSR #2`. Equally, it is common to combine shifts with ADD, SUB or other instructions. For example, to multiply R0 by 5, we might write:

```
ADD R0, R0, R0, LSL #2
```

A left shift of n places is effectively a multiply by 2 to the power of n , so this effectively makes $R0 = R0 + (4 \times R0)$. A right shift provides the corresponding divide operation, although ASR rounds negative values differently than would division in C.

Apart from multiply and divide, another common use for shifted operands is array index look-up. Consider the case where R1 points to the base element of an array of `int` (32-bit) values and R2 is the index which points to the n th element in that array. We can obtain the array value with a single load instruction which uses the calculation $R1 + (R2 \times 4)$ to get the appropriate address. [Example 6-3 on page 6-8](#) provides examples of differing operand 2 types used in ARM instructions.

Example 6-3 ARM instructions showing a variety of operand 2 types

add	R0, R1, #1	$R0 = R2 + 1$
add	R0, R1, R2	$R0 = R1 + R2$
add	R0, R1, R2, LSL #4	$R0 = R1 + R2 \ll \#4$
add	R0, R1, R2, LSL R3	$R0 = R1 + R2 \ll R3$

6.3 Multiplication operations

The multiply operations are readily understandable. A key limitation to note is that there is no scope to multiply by an immediate value. Multiplies operate only on values in registers. Multiplication by a constant may need that constant to be loaded into a register first. Later versions of the ARM processor add significantly more multiply instructions, giving a range of possibilities for 8, 16 and 32-bit data. We will consider these in [Integer SIMD instructions on page 6-14](#) when looking at the DSP instructions.

[Table 6-4](#) summarizes the multiplication assembly language instructions, giving their mnemonic opcode, operands and a brief description of their function.

Table 6-4 Summary of multiplication operations in assembly language

Opcode	Operands	Description	Function
Multiplies			
MLA	Rd, Rn, Rm, Ra	Multiply accumulate (MAC)	$Rd = Ra + (Rn \times Rm)$
MLS	Rd, Rn, Rm, Ra	Multiply and Subtract	$Rd = Ra - (Rn \times Rm)$
MUL	Rd, Rn, Rm	Multiply	$Rd = Rn \times Rm$
SMLAL	RdLo, RdHi, Rn, Rm	Signed 32-bit multiply with a 64-bit accumulate	$RdHiLo += Rn \times Rm$
SMULL	RdLo, RdHi, Rn, Rm	Signed 64-bit multiply	$RdHiLo = Rn \times Rm$
UMLAL	RdLo, RdHi, Rn, Rm	Unsigned 64-bit MAC	$RdHiLo += Rn \times Rm$
UMULL	RdLo, RdHi, Rn, Rm	Unsigned 64-bit multiply	$RdHiLo = Rn \times Rm$

6.3.1 Additional multiplies

We saw in the data-processing instructions that we have the ability to multiply one 32-bit register with another, to produce either a 32-bit result or a 64-bit signed or unsigned result. In all cases, there is the option to accumulate further a 32-bit or 64-bit value into the result. Additional multiply instructions have been added. There are signed most-significant word multiplies, SMMUL, SMMLA and SMMLS. These perform a 32×32 -bit multiply in which the result is the top 32 bits of the product, with the bottom 32 bits discarded. The result may be rounded by applying an R suffix, otherwise it is truncated. The UMAAL (Unsigned Multiply Accumulate Accumulate Long) instruction performs a 32×32 -bit multiply and adds in the contents of two 32-bit registers.

6.4 Memory instructions

ARM processors perform ALU operations only on registers. The only supported memory operations are the *load* (which reads data from memory into registers) or *store* (which writes data from registers to memory). A LDR and STR can be conditionally executed, in the same fashion as other instructions.

The LDR (Load Register) pseudo-instruction has the following general format:

LDR Rd, [Rn, Op2]

Rn + op2 provides the memory address read (or just Rn, in some addressing modes). Rd is the register to write to. Rn is known as the base register with Op2 providing an offset in bytes from that base.

A STR (Store Register) instruction has the following general format:

STR Rd, [Rn, Op2]

Again, Rn + Op2 gives the address (or just Rn, in some addressing modes) and Rd the register to be stored out to memory.

As we shall see in the next section, this general format has a number of variants, which allows op2 to be specified in a number of different ways, providing a flexible set of memory access instructions.

We can specify the size of the Load or Store transfer by appending a B for Byte, H for Halfword, or D for doubleword (64 bits) to the instruction, for example, LDRB. For loads only, an extra S can be used to indicate a signed byte or halfword (SB for Signed Byte or SH for Signed Halfword). See [STR on page A-47](#) for examples of this. This approach can be useful, because if we load an 8-bit or 16-bit quantity into a 32-bit register we must decide what to do with the most significant bits of the register. For an unsigned number, we zero-extend (that is, we write the most significant 16 or 24 bits of the register to zero), but for a signed number, it is necessary to copy the sign bit (bit [7] for a byte, or bit [15] for a halfword) into the top 16 (or 24) bits of the register.

6.4.1 Addressing modes

There are multiple addressing modes (shown in [Example 6-4 on page 6-11](#)) which can be used for loads and stores:

- *Register addressing*— the address is in a register (1).
- *Pre-indexed addressing*— an offset to the base register is added before the memory access. The base form of this is LDR Rd, [Rn, Op2]. The offset can be positive or negative and can be an immediate value or another register with an optional shift applied.(2),(3).
- *Pre-indexed with write-back*— this is indicated with an exclamation mark (!) added after the instruction. After the memory access has occurred, this updates the base register by adding the offset value (4).
- *Post-index with write-back*— here, the offset value is written after the square bracket. The address from the base register only is used for the memory access, with the offset value added to the base register after the memory access (5).

Examples of each of these are shown in [Example 6-4](#):

Example 6-4 Examples of addressing modes

-
- | | | |
|-----|-----|---|
| (1) | LDR | R0, [R1] @ address pointed to by R1 |
| (2) | LDR | R0, [R1, R2] @ address pointed to by R1 + R2 |
| (3) | LDR | R0, [R1, R2, LSL #2] @ address is R1 + (R2*4) |
| (4) | LDR | R0, [R1, #32]! @ address pointed to by R1 + 32, then R1:=R1 + 32 |
| (5) | LDR | R0, [R1], #32 @ read R0 from address pointed to by R1, then R1:=R1 + 32 |
-

6.4.2 Multiple transfers

Load and Store Multiple instructions allow successive words to be read from or written to memory. These are extremely useful for stack operations and for memory copying. Only word values can be operated on in this way and a word aligned address should be used. The ABI, for example, requires an 8-byte stack alignment at any external boundaries.

The operands are a base register (with an optional ! denoting write-back of the base register) with a list of registers between braces. The register list is comma separated, with hyphens used to indicate ranges. The order in which the registers are loaded or stored has nothing to do with the order specified in this list. Instead, the operation proceeds in a fixed fashion, with the lowest numbered register always mapped to the lowest address.

For example:

```
LDMIA R10!, { R0-R3, R12 }
```

This instruction reads five registers from the addresses pointed to by register (R10) and because write-back is specified, increments R10 by 20 (5×4 bytes) at the end.

The instruction must also specify how to proceed from the base register Rd. The four possibilities are: IA/IB (Increment After/Before) and DA/DB (Decrement After/Before). These may also be specified using aliases (FD, FA, ED and EA) which work from a stack point of view and specify whether the stack pointer points to a full or empty top of the stack, and whether the stack ascends or descends in memory.

By convention, only the FD option is used for stacks in ARM processor based systems. This means that the stack pointer points to the last filled location in stack memory and will decrement with each new item of data pushed to the stack.

[Figure 6-1 on page 6-12](#) shows a push of two registers to the stack. Before the STMFD (PUSH) instruction is executed, the stack pointer points to the last occupied word of the stack. After the instruction is completed, the stack pointer has been decremented by 8 (two words) and the contents of the two registers have been written to memory, with the lowest numbered register being written to the lowest memory address.




Figure 6-1 Stack push operation

6.5 Branches

The instruction set provides a number of different kinds of branch instruction. For simple relative branches (those to an offset from the current address), the B instruction is used. Calls to subroutines, where it is necessary for the return address to be stored in the link register, use the BL instruction.

If we wish to change instruction set (from ARM to Thumb or vice-versa), we use BX, or BLX. These latter cases may also use a value in a register, enabling a jump to anywhere within the 32-bit address space.

We can also specify the PC as the destination register for the result of normal data processing operations such as ADD or SUB, but this is generally deprecated and is unsupported in Thumb. An additional type of branch instruction can be implemented using either a load (LDR) with the PC as the target, load multiple (LDM), or stack-pop (POP) instruction with PC in the list of registers to be loaded.

Thumb has the compare and branch instruction, which fuses a CMP instruction and a conditional branch, but does not change the CPSR condition code flags. There are two opcodes, CBZ (compare and branch to label if Rn is zero) and CBNZ (compare and branch to label if Rn is not zero). These instructions can only branch forward between 4 and 130 bytes. Thumb also has the TBB (Table Branch Byte) and TBH (Table Branch Halfword) instructions. These instructions read a value from a table of offsets (either byte or halfword size) and perform a forward PC-relative branch of twice the value of the byte or the halfword returned from the table. These instructions require the base address of the table to be specified in one register, and the index in another.

6.6 Integer SIMD instructions

This section describes the SIMD (*Single Instruction, Multiple Data*) operations added in v6 of the ARM architecture. These should not be confused with the significantly more powerful Advanced SIMD (NEON) operations which were introduced in the ARM-v7 architecture and are covered in detail in [Chapter 8](#), [Chapter 20](#) and [Appendix B](#). SIMD is one of four classifications of computer architectures defined by Michael J. Flynn in 1966 based upon the number of instruction and data streams available in the architecture.

These instructions were first added in the ARMv6 architecture and provide the ability to pack, extract and unpack 8-bit and 16-bit quantities within 32-bit registers and to perform multiple arithmetic operations such as add, subtract, compare or multiply to such packed data, with a single instruction.

6.6.1 Integer register SIMD instructions

The v6 SIMD operations make use of the GE (greater than or equal) flags within the CPSR. These are distinct from the normal condition flags. There is a flag corresponding to each of the four byte positions within a word. Normal data processing operations produce one value and set the N, Z, C and V flags (as seen in [Figure 4-3 on page 4-6](#)). The SIMD operations produce up to four outputs and set only the GE flags, to indicate overflow. The MSR and MRS instructions can be used to write or read these flags directly.

The general form of the SIMD instructions are that subword quantities in each register are operated on in parallel (for example, four ADDs on four bytes can be performed) and the GE flags are set or cleared according to the results of the instruction. Different types of add and subtract can be specified using appropriate prefixes. For example, QADD16 performs saturating addition on halfwords within a register. SADD/UADD8 and SSUB/USUB8 set the GE bits individually while SADD/UADD16 and SSUB/USUB16 set GE bits [3:2] together based on the top halfword result, and [1:0] together on the bottom halfword result.

Also available are the ASX and SAX class of instructions, which reverse halfwords of one operand and add/subtract or subtract/add parallel pairs. Like the previously described ADD and Subtract instructions, these exist as unsigned (UASX/USAX), signed (SASX/SSAX) and saturated (QASX/QSAX) versions.




Figure 6-2 ADD v6 SIMD example

The SADD16 instruction shown in [Figure 6-2](#) shows how two separate addition operations are performed by a single instruction. The top halfwords of registers R3 and R0 are added, with the result going into the top halfword of register R1 and the bottom halfwords of registers R3 and

R0 are added, with the result going into the bottom halfword of register R1. GE[3:2] bits in the CPSR are set based on the top halfword result and GE[1:0] based on the bottom halfword result. In each case the carry flag is duplicated in the specified pair of bits.

6.6.2 Integer register SIMD multiplies

Like the other SIMD operations, these operate in parallel, on subword quantities within registers. The instruction can also include an accumulate option, with add or subtract being able to be specified. The instructions are SMUAD (SIMD multiply and add with no accumulate), SMUSD (SIMD multiply and subtract with no accumulate), SMLAD (multiply and add with accumulate) and SMLSD (multiply and subtract with accumulate).

Adding an L (long) before D indicates 64-bit accumulation.

Using the X (eXchange) suffix indicates halfwords in R_m are swapped before calculation.

The Q flag is set if accumulation overflows.

The SMUSD instruction shown in [Figure 6-3](#) performs two signed 16-bit multiplies (top \times top and bottom \times bottom) and then subtracts the two results. This kind of operation is useful when performing operations on complex numbers (with a real and imaginary component), a common task for filter algorithms.




Figure 6-3 v6 SIMD signed dual multiply subtract example

6.6.3 Sum of absolute differences

Calculating the sum of absolute differences is a key operation in the motion vector estimation component of common video codecs and is carried out over arrays of pixel data. The USADA8 Rd, Rn, Rm, Ra instruction is illustrated in [Figure 6-4 on page 6-16](#). It calculates the sum of absolute differences of the bytes within a word in registers Rn and Rm, adds in the value stored in Ra and places the result in Rd.




Figure 6-4 Sum of absolute differences

6.6.4 Data packing and unpacking

Packed data is common in many video and audio codecs (video data is usually expressed as packed arrays of 8-bit pixel data, audio data may use packed 16-bit samples), and also in network protocols. Before additional instructions were added in architecture v6, this data had to be either loaded with LDRH and LDRB instructions or loaded as words and then unpacked using Shift and Bit Clear operations; both are relatively inefficient. Pack (PKHBT, PKHTB) instructions allow 16-bit or 8-bit values to be extracted from any position in a register and packed into another register. Unpack instructions (UXTH, UXTB, plus many variants, including signed, with addition) can extract 8-bit or 16-bit values from any bit position within a register.

This enables sequences of packed data in memory to be loaded efficiently using word or doubleword loads, unpacked into separate register values, operated upon and then packed back into registers for efficient writing out to memory.




Figure 6-5 Packing and unpacking of 16-bit data in 32-bit registers

In the simple example shown in Figure 6-5, R0 contains two separate 16-bit values, denoted A and B. We can use the UXTH instruction to unpack the two halfwords into registers for further processing and we can then use the PKHBT instruction to pack halfword data from two registers

into one. It would be possible to replace the unpack instruction in each case with a MOV and either LSL or LSR instructions, but in this case we use a single instruction intended to work on parts of registers.

6.6.5 Byte selection

The SEL instruction enables us to select each byte of the result from the corresponding byte in either the first or the second operand, based on the value of the GE[3:0] bits in the CPSR. The packed data arithmetic operations set these bits as a result of add or subtract operations, and SEL can be used after these to extract parts of the data – for example, to find the smaller of the two bytes in each position.

6.7 Saturating arithmetic

Saturated arithmetic is commonly used in audio and video codecs. Calculations which return a value higher (or lower) than the largest positive (or negative) number which can be represented do not overflow. Instead the result is set to the largest positive or negative value (saturated). The ARM instruction set includes a number of instructions which allows easy implementation of such algorithms.

6.7.1 Saturated math instructions

The ARM saturated math instructions operate on either word or halfword sized values. The 8 of the QADD8 and QSUB8 instructions indicate that they operate on byte sized values. The result of the operation will be saturated to the largest possible positive or negative number. If the result would have overflowed and has been saturated, the overflow flag (CPSR Q bit) is set. This flag is said to be *sticky*. When set it will remain set until explicitly cleared by a write to the CPSR.

The instruction set provides special instructions with this behavior, QSUB and QADD. Additionally, we have QDSUB and QDADD which are provided in support of Q15 or Q31 fixed point arithmetic. These instructions double and saturate their second operand before performing the specified add or subtract.

The Count Leading Zeros (CLZ) instruction returns the number of 0 bits before the most significant bit that is set. This can be useful for normalization and for certain division algorithms. To saturate a value to a specific bit position (effectively saturate to a power of two), we can use the USAT or SSAT (unsigned or signed) saturate operations. USAT16 and SSAT16 allow saturation of two halfword values packed within a register.

6.8 Miscellaneous instructions

The remaining instructions cover coprocessor, supervisor call, PSR modification, byte reversal, cache preload, bit manipulation and a few others.

6.8.1 Coprocessor instructions

Coprocessor instructions occupy part of the ARM instruction set. Up to 16 coprocessors can be implemented, numbered 0 to 15 (CP0, CP1 ... CP15). These can either be internal (built-in to the processor) or connected externally, through a dedicated interface. Use of external coprocessors is uncommon in older processors and is not supported at all in the Cortex-A series.

- Coprocessor 15 is a built-in coprocessor which provides control over many processor features, including cache and MMU.
- Coprocessor 14 is a built-in coprocessor which controls the hardware debug facilities of the processor, such as breakpoint units (described in [Chapter 29](#)).
- Coprocessors 10 and 11 give access to the floating-point and NEON hardware in the system (described in [Chapter 7](#) and [Chapter 8](#)).

If a coprocessor instruction is executed, but the appropriate coprocessor is not present in the system, an undefined instruction exception occurs.

There are five classes of coprocessor instruction

- CDP – initiate a coprocessor data processing operation.
- MRC – move to ARM register from coprocessor register.
- MCR – move to coprocessor register from ARM register.
- LDC – load coprocessor register from memory.
- STC – store from coprocessor register to memory.

Multiple register and other variants of these instructions are also available, for example, MRRC, MCCR, LDCL, STCL.

6.8.2 Coprocessor 15

CP15, the System Control coprocessor, is (despite the name coprocessor) an integral part of the processor and provides control of many features. It can contain up to 16 primary registers, each of size 32 bits. Note, however, that access to CP15 is privilege controlled and not all registers are available in User mode. The CP15 register access instructions specify the required primary register, with the other fields in the instruction used to define the access more precisely and increase the number of physical 32-bit registers in CP15. The 16 primary registers in CP15 are named c0 to c15, but are often referred to by a series of letters. For example, the CP15 System

Control Register is named CP15.SCTLR. [Table 6-5](#) summarizes the function of some of the more relevant registers used in Cortex-A series processors. We will consider some of these in more detail when we look at the operation of the cache and MMU.

Table 6-5 CP15 Register summary

Register	Description
Main ID Register (MIDR)	Gives identification information for the processor (including part number and revision).
Multiprocessor Affinity Register (MPIDR)	Provides a way to uniquely identify individual processors within a multi-processor system.
CP15 c1 System Control registers	
System Control Register (SCTLR)	The main processor control register (see Chapter 10 Memory Management Unit).
Auxiliary Control Register (ACTLR)	IMPLEMENTATION DEFINED. Implementation specific additional control and configuration options.
Coprocessor Access Control Register (CPACR)	Controls access to all coprocessors except CP14 and CP15.
Secure Configuration Register (SCR)	Used by TrustZone (Chapter 26 Security).
CP15 c2 and c3, memory protection and control registers	
Translation Table Base Register 0 (TTBR0)	Base address of level 1 page table (see Chapter 10 Memory Management Unit).
Translation Table Base Register 1 (TTBR1)	Base address of level 1 page table (see Chapter 10 Memory Management Unit).
Translation Table Base Control Register (TTBCR)	Controls use of TTBR0 and TTBR1 (see Chapter 10 Memory Management Unit).
CP15 c5 and c6, memory system fault registers	
Data Fault Status Register (DFSR)	Gives status information about the last data fault (see Chapter 14 Other Exception Handlers).
Instruction Fault Status Register (IFSR)	Gives status information about the last instruction fault (see Chapter 14 Other Exception Handlers).
Data Fault Address Register (DFAR)	Gives the virtual address of the access that caused the most recent precise data abort (see Chapter 14 Other Exception Handlers).
Instruction Fault Address Register (IFAR)	Gives the virtual address of the access that caused the most recent precise prefetch abort (see Chapter 14 Other Exception Handlers).
CP15 c7, cache maintenance and other functions	
Cache and branch predictor maintenance functions	See Chapter 9 Caches .

Table 6-5 CP15 Register summary (continued)

Register	Description
Data and instruction barrier operations	See Chapter 11 Memory Ordering .
CP15 c8, TLB maintenance operations	
CP15 c9, performance monitors	
CP15 c12, Security Extensions registers	
Vector Base Address Register (VBAR)	Provides the exception base address for exceptions that are not handled in Monitor mode.
Monitor Vector Base Address Register (MVBAR)	Holds the exception base address for all exceptions that are taken to Monitor mode.
CP15 c13, process, context and thread ID registers	
Context ID Register (CONTEXTIDR)	See description of ASID Chapter 10 Memory Management Unit .
Software thread ID registers	See description of ASID Chapter 10 Memory Management Unit .
CP15 c15, IMPLEMENTATION DEFINED registers	
Configuration Base Address Register (CBAR)	Provides a base address for the GIC and Local timer type peripherals.

All system architecture functions are controlled by reading or writing a general purpose processor register (Rt) from or to a set of registers (CRn) located within Coprocessor 15. The Op1, Op2, and CRm fields of the instruction can also be used to select registers or operations. The format is shown in [Example 6-5](#).

Example 6-5 CP15 Instruction syntax

MRC p15, Op1, Rt, CRn, CRm, Op2 ; read a CP15 register into an ARM register

MCR p15, Op1, Rt, CRn, CRm, Op2 ; write a CP15 register from an ARM register

We will not go through each of the various CP15 registers in detail, as this would duplicate reference information which can readily be obtained from the *ARM Architecture Reference Manual* or product documentation. We will look at one example – the read-only Main ID Register (MIDR), the format of which is shown in [Figure 6-6](#).




Figure 6-6 Main ID Register format

In a privileged mode, we can read this register, using the instruction

MRC p15, 0, <Rt>, c0, c0, 0

The result, placed in register Rt, tells software which processor it is running on. For an ARM Cortex processor the interpretation of the results is as follows:

- Bits [31:24] – implementer, will be 0x41 for an ARM designed processor.
- Bits [23:20] – variant, shows the revision number of the processor.
- Bits [19:16] – architecture, will be 0xF for ARM architecture v7.
- Bits [15:4] – part number (for example 0xC08 for the Cortex-A8 processor).
- Bits [3:0] – revision, shows the patch revision number of the processor.

6.8.3 SVC

The SVC (supervisor call) instruction, when executed, causes a supervisor call exception. This is described further in [Chapter 12 Exception Handling](#). The instruction includes a 24-bit (ARM) or 8-bit (Thumb) number value, which can be examined by the SVC handler code. Through the SVC mechanism, an operating system can specify a set of privileged operations which applications running in User mode can request. This instruction was originally called SWI (Software Interrupt).

6.8.4 PSR modification

Several instructions allow the PSR to be written to, or read from:

- MRS transfers the CPSR or SPSR value to a general purpose register. MSR transfers a general purpose register to the CPSR or SPSR. Either the whole status register, or just part of it can be updated. In User mode, all bits can be read, but only the condition flags (_f) are permitted to be modified.
- In a privileged mode the Change Processor State (CPS) instruction can be used to directly modify the mode and interrupt-enable or disable (I and F) bits in the CPSR. See [Figure 4-3 on page 4-6](#).
- SETEND modifies a single CPSR bit, the E (Endian) bit. This can be used in systems with mixed endian data to temporarily switch between little- and big-endian data access.

6.8.5 Bit manipulation

There are instructions which allow bit manipulation of values in registers:

- The Bit Field Insert (BFI) instruction allows a series of adjacent bits from the bottom of one register (specified by supplying a width value and LSB position) to be placed into any position in the destination register.
- The Bit Field Clear (BFC) instruction allows adjacent bits within a register to be cleared.
- The SBFX and UBFX instructions (Signed and Unsigned Bit Field Extract) copy adjacent bits from one register to the least significant bits of a second register, and sign extend or zero extend the value to 32 bits.
- The RBIT instruction reverses the order of all bits within a register.

6.8.6 Cache preload

Cache preloading is described further in [Chapter 19 Optimizing Code to Run on ARM Processors](#). Two instructions are provided, PLD (data cache preload) and PLI (instruction cache preload). Both instructions act as hints to the memory system that an access to the specified

address is likely to occur soon. Implementations that do not support these operations will treat a preload as a NOP, but all of the Cortex-A family processors described in this book are able to preload the cache. Any illegal address specified as a parameter to the PLD instruction will not result in a data abort exception.

6.8.7 Byte reversal

Instructions to reverse byte order can be useful for dealing with quantities of the opposite endianness or other data re-ordering operations.

- The REV instruction reverses the bytes in a word
- The REV16 reverses the bytes in each halfword of a register
- The REVSH reverses the bottom two bytes, and sign extends the result to 32 bits.

Figure 6-7 illustrates the operation of the REV instruction, showing how four bytes within a register have their ordering within a word reversed.




Figure 6-7 Operation of the REV instruction

6.8.8 Other instructions

A few other instructions are available:

- The breakpoint instruction (BKPT) will either cause a prefetch abort (see [Types of exception on page 12-3](#)) or cause the processor to enter debug state (depending on whether the processor is configured for monitor or halt mode debug). This instruction is used by debuggers. See [Debug events on page 29-2](#).
- Wait For Interrupt (WFI) puts the processor into standby mode, which is described further in [Chapter 25 Power Management](#). The processor stops execution until woken by an interrupt or debug event. Note that if WFI is executed with interrupts disabled, an interrupt will still wake the processor, but no interrupt exception is taken. The processor proceeds to the instruction after the WFI. In older ARM processors, WFI was implemented as a CP15 operation. WFE (Wait for Event) will also be described in [Chapter 25](#).
- A NOP instruction (no-operation) does nothing. It may or may not take time to execute, so the NOP instruction should not be used to insert timing delays into code. It is intended to be used as padding.

Chapter 7

Floating-Point

All computer programs deal with numbers. Floating-point numbers, however, can sometimes appear counter-intuitive to programmers who are not familiar with their detailed implementation. For example, numbers which look equivalent to the human eye can return `false` using a C language `==` comparison. Before looking at floating-point implementation on ARM processors, a short overview of floating-point fundamentals is included. Programmers with prior floating-point experience may wish to skip the following section.

7.1 Floating-point basics and the IEEE-754 standard

The IEEE-754 standard is the reference for almost all modern computer floating-point mathematics implementations, including ARM floating-point systems. The original IEEE-754-1985 standard was updated with the publication of IEEE-754-2008. The standard defines precisely what result will be produced by each of the fundamental floating-point operations over all of the possible input values. It describes what a compliant implementation should do with respect to rounding of results which cannot be expressed precisely. A simple example of such a calculation would be $1.0 \div 3.0$, which would require an infinite number of digits to express precisely in decimal or binary notation. IEEE-754 provides a number of different rounding options to cope with this (round towards positive infinity, round towards negative infinity, round toward zero, and two forms of round to nearest, see [Rounding algorithms on page 7-4](#)). IEEE-754 also specifies the outcome when an *exceptional* operation occurs. This means a calculation which potentially represents a problem. These conditions can be tested, either by querying the FPSCR (on ARM processors) or by setting up trap handlers (on some systems). Examples of exceptional operations are as follows:

Overflow A result which is too large to represent.

Underflow A result which is so small that precision is lost.

Inexact A result which cannot be represented without some loss of precision. It is clear that many floating-point calculations will fall into this category.

Invalid For example, attempting to calculate the square root of a negative number.

Division by zero

The specification also describes what action should be taken when one of the above exceptional operations is detected. Possible outcomes include the generation of a *NaN* (Not a Number) result (for invalid operations), positive or negative infinity (for overflow or division by zero) or *denormalized* numbers in the case of underflow. The standard further defines what results should be produced if subsequent floating-point calculations operate on NaN or infinities.

One of the things that IEEE-754 defines is how floating-point numbers are represented within the hardware. Floating-point numbers are typically represented using either single precision (32-bit) or double-precision (64-bit). VFP (see [VFP on page 2-6](#)) supports single-precision (32-bit) and double-precision (64-bit) formats in hardware. In addition, VFPv3 can have half-precision extensions to allow 16-bit values to be used for storage.

Floating-point formats use the available space to store three pieces of information about a floating-point number:

- A sign bit (S) which shows whether the number is positive (0) or negative (1).
- An exponent giving its order of magnitude.
- A mantissa giving the fractional binary digits of the number.

For a single precision float, for example, bit [31] of the word is the sign bit [S], bits [30: 23] give the exponent and bits [22:0] give the mantissa. See [Figure 7-1 on page 7-3](#).

The value of the number is then $\pm m \times 2^{\text{exp}}$, where “m” is derived from the mantissa and “exp” is derived from the exponent.

Because a 32-bit floating-point number has a 23-bit mantissa there are many values of a 32-bit int that if converted to 32-bit float cannot be represented exactly. This is referred to as *loss of precision*. If you convert one of these values to float and back to int you will get a different, nearby value. In the case of double-precision floating-point numbers, the exponent field has 11 bits (giving an exponent range from -1022 to $+1023$) and a mantissa field with 52 bits.

7.1.1 Rounding algorithms

The IEEE 754-1985 standard defines four different ways in which results can be rounded, as follows:

- *Round to nearest* (ties to even). This mode causes rounding to the nearest value. If a number is exactly midway between two possible values, it is rounded to the nearest value with a zero least significant bit.
- *Round toward 0*. This causes numbers to always be rounded towards zero (this can be also be viewed as truncation).
- *Round toward $+\infty$* . This selects rounding towards positive infinity.
- *Round toward $-\infty$* . This selects rounding towards negative infinity.

The IEEE 754-2008 standard adds an additional rounding mode. In the case of round to nearest, it is now also possible to round numbers which are exactly halfway between two values, away from zero (in other words, upwards for positive numbers and downwards for negative numbers). This is in addition to the option to round to the nearest value with a zero least significant bit. At present VFP does not support this rounding mode.

7.1.2 ARM VFP

VFP is an optional (but rarely omitted) extension to the instruction sets in the ARMv7-A architecture conforming to the IEEE 754 standard. It can be implemented with either thirty-two, or sixteen double-word registers. The terms VFPv3-D32 and VFPv3-D16 are used to distinguish between these two options. If the Advanced SIMD (NEON) extension is implemented together with VFPv3, VFPv3-D32 is always present. VFPv3 can also be optionally extended by the half-precision extensions that provide conversion functions in both directions between half-precision floating-point (16-bit) and single-precision floating-point (32-bit). These operations only allow half-precision floats to be converted to and from other formats.

VFPv4 adds both the half-precision extensions and the Fused Multiply-Add instructions to the features of VFPv3. In a Fused Multiply-Add operation, only a single rounding occurs at the end. This is one of the new facets of the IEEE 754-2008 specification. Fused operations can improve the accuracy of calculations which repeatedly accumulate products, such as matrix multiplication or dot product calculation. The VFP version supported by Cortex-A series processors is given in [Table 2-3 on page 2-10](#).

In addition to the registers described above, there are a number of other VFP registers. These are listed below.

Floating-Point System ID Register (FPSID)

This can be read by system software to determine which floating-point features are supported in hardware.

Floating-Point Status and Control register (FPSCR)

This holds comparison results and flags for exceptions. Control bits select rounding options and enable floating-point exception trapping.

Floating-Point Exception Register (FPEXC)

The FPEXC register contains bits which allow system software that handles exceptions to determine what has happened.

Media and VFP Feature registers 0 and 1 (MVFR0 and MVFR1)

These registers allow system software to determine which Advanced SIMD and floating-point features are provided on the processor implementation.

User mode code can only access the FPSCR. One implication of this is that applications cannot read the FPSID to determine which features are supported unless the host OS provides this information. Linux provides this via `/proc/cpuinfo`, for example, but the information is not nearly as detailed as that provided by the VFP hardware registers.

Unlike ARM integer processor instructions, no VFP operations will affect the flags in the APSR directly. The flags are stored in the FPSCR. Before the result of a floating-point comparison can be used by the integer processor, the flags set by a floating-point comparison must be transferred to the APSR, using the VMRS instruction. This includes use of the flags for conditional execution, even of other VFP instructions. [Example 7-1](#) shows a simple piece of code to illustrate this. The VCMPI instruction performs a comparison the values in VFP registers d0 and d1 and sets FPSCR flags as a result. These flags must then be transferred to the integer processor APSR, using the VMRS instruction. We can then conditionally execute instructions based on this.

Example 7-1 Example code illustrating usage of floating-point flags

```
VCMPI d0, d1
VMRS APSR_nzcv, FPSCR
BNE label
```

Flag meanings

The integer comparison flags support comparisons which are not applicable to floating-point numbers. For example, floating-point values are always signed, so there is no need for unsigned comparisons. On the other hand, floating-point comparisons can result in the unordered result (meaning that one or both operands was NaN, or *Not a Number*). IEEE-754 defines four testable relationships between two floating-point values, which map onto the ARM condition codes as follows:

Table 7-2 ARM APSR flags

IEEE-754 relationship	ARM APSR flags			
	N	Z	C	V
Equal	0	1	1	0
Less Than (LT)	1	0	0	0
Greater Than (GT)	0	0	1	0
Unordered (At least one argument was NaN)	0	0	1	1

Compare with zero

Unlike the integer instructions, most VFP (and NEON) instructions can operate only on registers, and cannot accept immediate values encoded in the instruction stream. The VCMPI instruction is a notable exception in that it has a special-case variant that enables quick and easy comparison with zero.

Interpreting the flags

When the flags are in the APSR, they can be used almost as if an integer comparison had set the flags. However, floating-point comparisons support different relationships, so the integer condition codes do not always make sense. Table 7-3 describes floating-point comparisons rather than integer comparisons:

Table 7-3 Interpreting the flags

Code	Meaning (when set by vcmp)	Meaning (when set by cmp)	Flags tested
EQ	Equal to	Equal to	Z = 1
NE	Unordered, or not equal to	Not equal to.	Z = 0
CS or HS	Greater than, equal to, or unordered	Greater than or equal to (unsigned).	C = 1
CC or LO	Less than.	Less than (unsigned).	C = 0
MI	Less than	Negative.	N = 1
PL	Greater than, equal to, or unordered	Positive or zero.	N = 0
VS	Unordered. (At least one argument was NaN.)	Signed overflow.	V = 1
VC	Not unordered. (No argument was NaN.)	No signed overflow.	V = 0
HI	Greater than or unordered	Greater than (unsigned).	(C = 1) && (Z = 0)
LS	Less than or equal to	Less than or equal to (unsigned).	(C = 0) (Z = 1)
GE	Greater than or equal to.	Greater than or equal to (signed).	N==V
LT	Less than or unordered.	Less than (signed).	N!=V
GT	Greater than.	Greater than (signed).	(Z==0) && (N==V)
LE	Less than, equal to or unordered.	Less than or equal to (signed).	(Z==1) (N!=V)
AL (or omitted)	Always executed.	Always executed.	None tested.

It should be obvious that the condition code is attached to the instruction reading the flags, and the source of the flags makes no difference to the flags that are tested. It is the meaning of the flags that differs when you perform a vcmp rather than a cmp. Similarly, it is clear that the opposite conditions still hold. (For example, HS is still the opposite of LO.)

The flags when set by CMP generally have analogous meanings when set by VCMPI. For example, GT still means *greater than*. However, the unordered condition and the removal of the signed conditions can confuse matters. Often, for example, it is desirable to use LO, normally an unsigned *less than* check, in place of LT, because it does not match in the unordered case.

7.1.3 Instructions

VFP instructions are provided which perform arithmetic and data processing, load and stores to memory, and register transfers (between VFP registers and to or from ARM registers). These instructions are encoded with ARM coprocessor instructions, but are typically viewed as part of the main instruction set, rather than as coprocessor operations. VFP offers all the common arithmetic operations, format conversions, a few complex arithmetic operations (for example, Multiply accumulate, VMLA, and square root, VSQRT), along with memory access instructions. [Appendix B](#) of this book provides a full list of the supported instructions.

7.1.4 Enabling VFP

If an ARMv7 processor includes VFP hardware, it must be explicitly enabled before applications can make use of it. Several steps are required to do this:

- The EN bit in the FPEXC register must be set.
- If access to VFP is required in the Normal world, access to CP10 and CP11 must be enabled in the *Non-Secure Access Control Register* (CP15.NSACR). This would normally be done inside the Secure bootloader.
- Access to CP10 and CP11 must be enabled in the Coprocessor Access Control Register (CP15.CACR). This can be done on demand by the operating system.

7.2 VFP support in GCC

Use of VFP is fully supported by GCC (although some builds can be configured to default to assume no VFP support, in which case floating-point calculations will use library code).

The main option to use for VFP support is:

- `-mfpu=vfp` specifies that the target has VFP hardware. (As does specifying the option `-mfpu=neon`.)

Other options can be used to specify support for a specific VFP implementation on an ARM Cortex-A series processor:

- `-mfpu=vfpv3` or `-mfpu=vfpv3-d16` (for Cortex-A8 and Cortex-A9 processors).
- `-mfpu=vfpv4` or `-mfpu=vfpv4-d16` (for Cortex-A5 and Cortex-A15 processors).

These options can be used for code which will run only on these VFP implementations, and do not require backward compatibility with older VFP implementations.

- `-mfloat-abi=softfp` (or `hard`) specify which ABI to use to enable the use of VFP.

`softfp` uses a Procedure Call Standard compatible with software floating-point, and so provides binary compatibility with legacy code. This permits running older `soft float` code with new libraries which support hardware floating-point, but still makes use of hardware floating-point registers between function calls. `hard` has floating-point values passed in floating-point registers. This is more efficient but is not backward compatible with the `softfp` ABI variant. Particular care is needed with libraries, including the C platform library. See [VFP and NEON register usage on page 17-4](#) for more information on efficient parameter passing.

C programmers should note that there can be a significant function call overhead when using `-mfloat-abi=softfp`, if many floating-point values are being passed.

7.3 VFP support in the ARM Compiler

Use of VFP is fully supported by the ARM Compiler (although some builds can be configured by default to assume no VFP support, in which case floating-point calculations will use library code).

The main option to use with the ARM Compiler for VFP support is:

- `--fpu=name` which lets you specify the target floating-point hardware.

The options used to specify support for a specific VFP implementation on an ARM Cortex-A series processor are:

- `--fpu=vfpv3` or `--fpu=vfpv3_d16` (for the Cortex-A8 and Cortex-A9 processors).
- `--fpu=vfpv4` or `--fpu=vfpv4_d16` (for the Cortex-A5, Cortex-A7 and Cortex-A15 processors).

These options can be used for code which will run only on these VFP implementations, and do not require backward compatibility with older VFP implementations. Use `--fpu=list` to see the full list of FPUs supported.

The following options can be used for linkage support:

- `--apcs=/hardfp` generates code for hardware floating-point linkage.
- `--apcs=/softfp` generates code for software floating-point linkage.

Hardware floating-point linkage uses the FPU registers to pass the arguments and return values. Software floating-point linkage means that the parameters and return value for a function are passed using the ARM integer registers R0 to R3 and the stack. `--apcs=/hardfp` and `--apcs=/softfp` interact with or override explicit or implicit use of `--fpu`.

7.4 VFP support in Linux

An application which uses VFP (or calls into a library which uses VFP) places some additional requirements on the Linux kernel. For the application to run correctly, the kernel must save and restore the VFP registers during context switches. The kernel may also need to decode and emulate VFP instructions where the VFP hardware is not present.

7.4.1 Context switching

In addition to saving and restoring integer registers, the kernel may also need to perform saving and restoring of VFP registers on a context switch. To avoid wasting cycles, this is done only when an application actually used VFP. As the VFP initialization code leaves VFP disabled, the first time a thread actually tries to access the floating-point hardware, an undefined exception occurs. The kernel function which handles this, sees that VFP is disabled and that a new thread wishes to use VFP. It saves the current VFP state and restores the state for the new thread.

On systems with multiple processors, where threads can migrate to a different processor, this simple system will no longer work correctly. Instead, the kernel saves the state if the VFP was used by the previous thread.

7.5 Floating-point optimization

This section contains some suggestions for developers writing FP assembly code. Some caution is needed when applying these points, as recommendations can be specific to a particular piece of hardware. A code sequence which is optimal for one processor can be sub-optimal on different hardware.

- Avoid mixing of VFP and NEON instructions on the Cortex-A9 processor, as there is a significant overhead in switching between data engines.
 - VLDR or VSTM can be executed by both VFP and NEON, but VLD1 and VLD2 can only be executed by NEON. VFP Double-Precision instructions which are not NEON can only be executed by the VFP.
- Moves to and from VFP system control registers, such as FPSCR are not typically present in high-performance code, and may not be optimized. These should therefore not be placed in time-critical loops, if possible. For example, accesses to control registers on the Cortex-A9 processor are serializing, and will have a significant performance impact if used in tight loops or performance-critical code.
- Register transfer between the integer processor register bank and the floating-point register bank should similarly be avoided in time-critical loops. For the Cortex-A8 processor, this is particularly true of register transfers from VFP registers to integer registers.
- Load/store multiple operations are preferred to the use of multiple, individual floating-point loads and stores, to make efficient use of available transfer bandwidth.

Chapter 8

Introducing NEON

This chapter provides an introduction to NEON technology, which can be used to accelerate the performance of multimedia applications running on ARM Cortex-A series processors.

NEON technology provides Single Instruction Multiple Data (SIMD) operations in ARM processors implementing the Advanced SIMD architecture extensions. These operations can significantly accelerate repetitive operations on large data sets. This can be useful in applications such as media codecs.

NEON is implemented as a separate hardware unit that is available as an option on Cortex-A series processors. Allowing the NEON hardware to be optional enables ARM SoCs to be optimized for specific markets. In most general-purpose applications processor SoCs, NEON will probably be included. However, for an embedded application such as a network router, NEON can be omitted, allowing a small saving in silicon area which translates to a small cost saving.

We provide only a brief introduction to NEON technology here. A complete description could fill a book of its own!

8.1 SIMD

SIMD is a computational technique for processing a number of data values (generally a power of two) using a single instruction, with the data for the operands packed into special wide registers. One instruction can therefore do the work of many separate instructions. For code which can be parallelized, large performance improvements can be achieved. SIMD extensions exist on many 32-bit architectures – PowerPC has AltiVec, while x86 has several variants of MMX/SSE. SIMD is described further in [Integer SIMD instructions on page 6-14](#).

Many software programs operate on large datasets. The data items can be less than 32 bits in size (8-bit pixel data is common in video, graphics and image processing, 16-bit samples in audio codecs). In such cases, the operations to be performed are simple, repeated many times and have little need for control code. SIMD can offer considerable performance improvements for this type of data processing. It is particularly beneficial for digital signal processing or multimedia algorithms, such as:

- Block-based data processing.
- Audio, video, and image processing codecs.
- 2D graphics based on rectangular blocks of pixels
- 3D graphics
- Color-space conversion.
- Physics simulations.

On a 32-bit microprocessors such as the Cortex-A series processors, it is relatively inefficient to perform large numbers of 8-bit or 16-bit single operations. The processor ALU, registers and datapath are designed for the 32-bit calculations. SIMD enables a single instruction to treat a register value as multiple data elements (for example, as four 8-bit values in a 32-bit register) and to perform multiple identical operations on those elements.




Figure 8-1 Comparing SIMD parallel add with 32-bit scalar add

To achieve four separate additions without using SIMD requires us to use four ADD instructions, as shown in [Figure 8-1](#), and additional instructions to prevent one result from overflowing into the adjacent byte. SIMD needs only one instruction to do this.

8.1.1 ARMv6 SIMD instructions

A set of basic 32-bit SIMD instructions was introduced in ARMv6, allowing the processor to operate on packed 8-bit or 16-bit values in 32-bit registers. These instructions have an 8 or 16 as a suffix, to indicate the size of data to operate on. These were described in [Integer SIMD instructions on page 6-14](#). These should not be confused with the separate NEON instructions described in this chapter.

8.2 NEON architecture overview

NEON was designed as an additional load/store architecture to provide good vectorizing compiler support from languages such as C/C++. A rich set of NEON instructions operate on wide 64-bit and 128-bit vector registers, enabling a high level of parallelism. The NEON instructions are straightforward and easy to understand, which also makes hand-coding easy for applications that need the very highest performance.

A key advantage of NEON technology is that instructions form part of the normal ARM or Thumb code, making programming simpler than with an external hardware accelerator. There are NEON instructions available to read and write external memory, move data between NEON registers and other ARM registers and to perform SIMD operations.

The NEON architecture uses a 32×64 -bit register file. These are actually the same registers used by the floating-point unit (VFPv3). It does not matter that the floating-point registers are re-used as NEON registers. All compiled code and subroutines will conform to the EABI, which specifies which registers can be corrupted and which registers must be preserved. The compiler is free to use any NEON or VFPv3 registers for floating-point values or NEON data at any point in the code.

The NEON architecture allows for 64-bit or 128-bit parallelism. This choice was made to keep the size of the NEON unit manageable (a vector ALU can easily become quite large), while still offering good performance benefits from vectorization. The NEON architecture also does not specify instruction timings and may require different numbers of cycles to execute the same instruction on different processors.

8.2.1 Commonality with VFP

NEON and Vector Floating-Point (VFP) instructions are both defined as coprocessor CP10 and CP15 operations (see [ARM VFP on page 7-4](#)). The ARM architecture can support a wide range of different NEON and VFP options, but in practice we see only the combinations:

- No NEON or VFP.
- VFP only.
- NEON and VFP.

These are vendor implementation options for the architecture, and so are fixed for a particular implementation of an ARM-based design.

The key differences between NEON and VFP are that NEON only works on vectors, does not support double-precision floating-point (double-precision is supported by the VFP), and does not support certain complex operations such as square root and divide. NEON has a register bank of thirty-two 64-bit registers. If both NEON and VFPv3 are implemented, this register bank is shared between them in hardware. This means that VFPv3 must be present in its VFPv3-D32 form, which has 32 double-precision floating-point registers. This makes support for context switching simpler. Code which saves and restores VFP context also saves and restores NEON context.

8.2.2 Data types

NEON instructions operate on elements of the following types:

- 32-bit single precision floating-point
- 8, 16, 32 and 64-bit unsigned and signed integers
- 8 and 16-bit polynomials.

Data type specifiers in NEON instructions comprise a letter indicating the type of data and a number indicating the width. They are separated from the instruction mnemonic by a point, for example, `VMLAL.S8`. So we have the following possibilities:

- unsigned integer U8 U16 U32 U64.
- signed integer S8 S16 S32 S64.
- integer of unspecified type I8 I16 I32 I64.
- floating-point number F16 F32 .
- polynomial over $\{0,1\}$ P8.

Note

F16 is not supported for data processing operations. It is only supported as a format to be converted to, or from.

The polynomial type is to help with anything that needs to use power-of-two finite fields or simple polynomials over $\{0,1\}$. Normal ARM integer code would typically use a lookup table for finite field arithmetic, but large lookup tables cannot be vectorized. Polynomial operations are hard to synthesize out of other operations, so it is useful having a basic multiply operation (add is EOR) out of which larger multiplies or other operations can be synthesized.

Applications of polynomial arithmetic can include error correction, such as Reed Solomon codes, and CRCs (the field operations typically use a polynomial multiply + mod, but this can be reduced to just multiplies using the Montgomery reduction method). Also elliptic curve cryptography often uses power-of-type fields.

NEON technology is IEEE 754-1985 compliant, but only supports round-to-nearest rounding mode. This is the rounding mode used by most high-level languages, such as C and Java. Additionally, NEON instructions always treats denormals as zero.

8.2.3 NEON registers

The register bank may be viewed as either sixteen 128-bit registers (Q0-Q15) or as thirty-two 64-bit registers (D0-D31). Each of the Q0-Q15 registers maps to a pair of D registers, as shown in [Figure 8-2](#).




Figure 8-2 NEON register bank

The view of registers in [Figure 8-3 on page 8-6](#) is determined by form of the instruction used, so that the software does not have to explicitly change state.




Figure 8-3 NEON registers

Individual elements can also be accessed as scalars. The advantage of the dual view is that it accommodates mathematical operations which widen or narrow the result. For example multiplying two D registers gives a Q register result. The dual-view enables the register bank to be used more efficiently.

NEON data processing instructions are typically available in Normal, Long, Wide, Narrow and Saturating variants.

- *Normal* instructions can operate on any vector types, and produce result vectors the same size, and usually the same type, as the operand vectors.
- *Long* instructions operate on doubleword vector operands and produce a quadword vector result. The result elements are usually twice the width of the operands, and of the same type. Long instructions are specified using an L appended to the instruction. [Figure 8-4](#) shows this, with input operands being promoted before the operation.

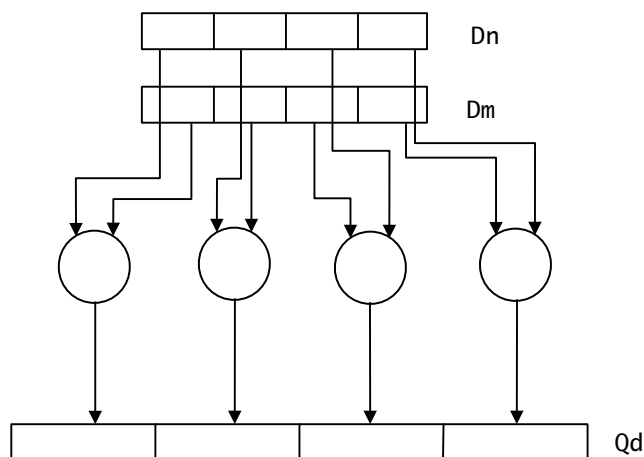


Figure 8-4 NEON long instructions

- *Wide* instructions operate on a doubleword vector operand and a quadword vector operand, producing a quadword vector result. The result elements and the first operand are twice the width of the second operand elements. Wide instructions have a W appended to the instruction. [Figure 8-5 on page 8-7](#) shows this, with the input doubleword operands being promoted before the operation.




Figure 8-5 NEON wide instructions

- *Narrow* instructions operate on quadword vector operands, and produce a doubleword vector result. The result elements are usually half the width of the operand elements. Narrow instructions are specified using an N appended to the instruction. Figure 8-6 shows this, with input operands being demoted before the operation.




Figure 8-6 NEON narrow instructions

Some NEON instructions act on scalars together with vectors. The scalars can be 8-, 16-, 32-, or 64-bit. Instructions which use scalars can access any element in the register bank, although there are differences for multiply instructions. The instruction uses an index into a doubleword vector to specify the scalar value. Multiply instructions only support 16-bit or 32-bit scalars, and can only access the first 32 scalars in the register bank (that is, D0-D7 for 16-bit scalars or D0-D15 for 32-bit scalars).

8.2.4 NEON instruction set

All mnemonics for NEON instructions (as with VFP) begin with the letter “V”. Instructions are generally able to operate on different data types, with this being specified in the instruction encoding. The size is indicated with a suffix to the instruction. The number of elements is indicated by the specified register size.

For example, looking at the instruction

VADD.I8 D0, D1, D2

VADD indicates a NEON ADD operation.

the I8 suffix indicates that 8-bit integers are to be added

D0, D1 and D2 specify the 64-bit registers used (D0 for the result destination, D1 and D2 for the operands).

So this instruction performs eight additions in parallel.

There are operations which have different size registers for input and output.

VMULL.S16 Q2, D8, D9

This instruction performs four 16-bit multiplies of data packed in D8 and D9 and produces four 32-bit results packed into Q2.

The VCVT instruction converts elements between single-precision floating-point and 32-bit integer, fixed-point, and (if implemented) half-precision floating-point.

NEON includes load and store instructions which can load or store individual or multiple values to a register. In addition, there are instructions which can transfer blocks of data between multiple registers and memory. It is also possible to interleave or de-interleave data during such multiple transfers.

The following modifiers can be used with certain Advanced SIMD instructions (some modifiers can be used only with a small subset of the available instructions):

- Q** The instruction uses saturating arithmetic, so that the result is saturated within the range of the specified data type. The sticky QC bit in the FPSCR is set if saturation occurs in any lane. VQADD is an example of such an instruction.
- H** The instruction will halve the result. It does this by shifting right by one place (effectively a divide by two with truncation). VHADD is an example of such an instruction – it could be used to calculate the mean of two inputs.
- D** The instruction doubles the result and saturates. This is commonly required when multiplying numbers in Q15 format, where an additional doubling is needed to get the result into the correct form.
- R** The instruction will perform rounding on the result, equivalent to adding 0.5 to the result before truncating. VRHADD is an example of this.

Instructions have the following general format:

V{<mod>}<op>{<shape>}{<cond>}{<.dt>}{<dest>}, src1, src2

where:

<mod> is one of the previously described Modifiers (Q, H, D, R)

<op> - operation (for example, ADD, SUB, MUL)

<shape> - Shape (L, W or N, as described in [NEON registers on page 8-5](#))

<cond> - Condition, used with IT instruction

<.dt> - Data type

<dest> - Destination

<src1> - Source operand 1

<src2> - Source operand 2.

The NEON instruction set includes a range of vector addition and subtraction operations, including pairwise adding, which adds adjacent vector elements together.

There are a number of multiply operations, including multiply-accumulate and multiply-subtract and doubling and saturating options. There is no SIMD division operation, but such an operation can be performed using the VRECPE (Vector Reciprocal Estimate) and VCREPS (Vector Reciprocal Step) instructions to perform Newton-Raphson iteration.

Similarly, there is no vector square root instruction, but VRSQRTE, VRSQRTS, and multiplies to compute square roots. Shift left, right and insert operations are also available, along with instructions which select minimum or maximum values. Common logic operations (AND, OR, EOR, AND NOT and OR NOT) can be performed. The instruction set also includes the ability to count numbers of bits set in an element or to count leading zeros or sign bits.

As SIMD instructions perform multiple operations, they cannot use the standard ALU flags for comparison instructions. Instead two elements can be compared, with the result of the comparison in the destination register. The destination element is set to all 0's if the tested condition is false or all 1's if the tested condition is true. This bit mask can then be used to control subsequent instructions. A number of different comparison operations are supported. There are bitwise select instructions which can be used in conjunction with these bit masks.

There are a number of different instructions to move data between registers, or between elements. It is also possible for instructions to swap or duplicate registers, to perform reversal, matrix transposition and extract individual vector elements.

[Appendix B](#) lists all NEON and VFP instructions and gives a brief description of their function. Full details for each instruction can also be found in the *ARM Architecture Reference Manual*.

Chapter 9

Caches

The word cache derives from the French verb *cacher*, “to hide”. A cache is a hidden storage place. The application of this word to a processor is obvious – a cache is a place where the processor can store instructions and data, hidden from the programmer and system. In many cases, it would be true to say that the cache is transparent to, or hidden from the programmer. But very often, as we shall see, it is important to understand the operation of the cache in detail.

When the ARM architecture was first developed, the clock speed of the processor and the access speeds of memory were broadly similar. Today’s processors are much more complicated and can be clocked orders of magnitude faster. However, the frequency of the external buses and of memory devices has not scaled to the same extent. It is possible to implement small blocks of on-chip SRAM which can operate at the same speeds of the processor, but such RAM is very expensive in comparison to standard DRAM blocks, which can have thousands of times more capacity. In many ARM processor based systems, access to external memory will take tens or even hundreds of processor cycles.

Essentially, a cache is a small, fast block of memory which (conceptually at least) sits between the processor core and main memory. It holds copies of items in main memory. Accesses to the cache memory happen significantly faster than those to main memory. As the cache holds only a subset of the contents of main memory, it must store both the address of the item in main memory and the associated data. Whenever the processor wants to read or write a particular address, it will first look for it in the cache. Should it find the address in the cache, it will access the data in the cache, rather than having to perform an access to main memory. This significantly increases the potential performance of the system, by reducing the effect of slow external memory access times. It also reduces the power consumption of the system, by avoiding the need to drive external signals.

Cache sizes are small relative to the overall memory used in the system. Larger caches make for more expensive chips. In addition, making an internal processor cache larger can potentially limit the maximum speed of the processor. Significant research has gone into identifying how hardware can determine what it should keep in the cache. Efficient use of this limited resource is a key part of writing efficient applications to run on a processor.

On-chip SRAM can be used to implement caches, which hold temporary copies of instructions and data from main memory. Code and data have the properties of temporal and spatial locality. This means that programs tend to re-use the same addresses over time (temporal locality) and tend to use addresses which are near to each other (spatial locality). Code, for instance, can contain loops, meaning that the same code gets executed repeatedly or a function can be called multiple times. Data accesses (for example, to the stack) can be limited to small regions of memory. It is this fact that access to RAM by the processor exhibits such locality, and is not truly random, that enables caches to be successful.

Access ordering rules obey the *weakly ordered* model. See [ARM memory ordering model on page 11-4](#).

We will also look at the write buffer. This is a block which decouples writes being done by the processor when executing store instructions from the external memory bus. The processor places the address, control and data values associated with the store into a set of hardware buffers. This is the write buffer. Like the cache, it sits between the processor core and main memory. This enables the processor to move on and execute the next instructions without having to stop and wait for the slow main memory to actually complete the write operation.

9.1 Why do caches help?

Caches speed things up, as we have seen, because program execution is not random. Programs access the same sets of data repeatedly and execute the same sets of instructions repeatedly. By moving code or data into faster memory when it is first accessed, following accesses to that code or data become much faster. The initial access which provided the data to the cache is no faster than normal. It is any subsequent accesses to the cached values which are faster, and it is from this the performance increase derives. The processor hardware will check all instruction fetches and data reads or writes in the cache, although obviously we need to mark some parts of memory (those containing peripheral devices, for example) as non-cacheable. As the cache holds only a subset of main memory, we need a way to determine (quickly) whether the address we are looking for is in the cache.

9.2 Cache drawbacks

It may seem that caches and write buffers are automatically a benefit, as they speed up program execution. However, they also add some problems which are not present in an uncached processor. One such drawback is that program execution time can become non-deterministic.

What this means is that, because the cache is small and holds only a subset of main memory, it fills rapidly as a program executes. When the cache is full, existing code or data is replaced, to make room for new items. So, at any given time, it is not normally possible for an application to be certain whether or not a particular instruction or data item is to be found in the cache.

This means that the execution time of a particular piece of code can vary significantly. This can be something of a problem in hard real-time systems where strongly deterministic behavior is needed.

Furthermore, as we shall see, we need a way to control how different parts of memory are accessed by the cache and write buffer. In some cases, we want the processor to read up-to-date data from an external device, such as a peripheral. It would not be sensible to use a cached value of a timer peripheral, for example. Sometimes we want the processor to stop and wait for a store to complete. So caches and write buffers give the programmer some extra work to do.

Occasionally data in the cache and data in external memory may not be the same. This is a problem with *coherency*. This can be a particular problem when we have multiple processors or memory agents like an external DMA controller. We will consider such coherency issues in greater detail later in the book.

9.3 Memory hierarchy

In computer science, a memory hierarchy refers to a hierarchy of memory types, with faster and smaller memories closer to the processor and slower and larger memory further away. In most systems, we can have secondary storage, such as disk drives and primary storage such as flash, SRAM and DRAM. In embedded systems, we typically further sub-divide this into on-chip and off-chip memory. Memory which is on the same chip (or at least in the same package) as the processor will typically be much faster.

A cache can be included at any level in the hierarchy and should improve system performance where there is an access time difference between different parts of the memory system.

In ARM processor based systems, we typically have level 1 (L1) caches, which are connected directly to the processor logic that fetches instructions and handles load and store instructions. These are Harvard caches (that is, there are separate caches for instructions and for data) in all but the lowest performing members of the ARM processor family and effectively appear as part of the processor. Over the years, the size of L1 caches has increased, due to SRAM size and speed improvements. At the time of writing, 16KB or 32KB cache sizes are most common, as these are the largest RAM sizes capable of providing single cycle access at a processor speed of 1GHz or more.

Many ARM systems have, in addition, a level 2 (L2) cache. This is larger than the L1 cache (typically 256KB, 512KB or 1MB), but slower and unified (holding both instructions and data). It can be inside the processor itself, as in the Cortex-A7, Cortex-A8 and Cortex-A15 processors, or be implemented as an external block, placed between the processor and the rest of the memory system. The ARM L2C-310 is an example of such an external L2 cache controller block.

In addition, we have processors that can be implemented in multi-processor clusters in which each processor has its own cache. Such systems require mechanisms to maintain coherency between caches, so that when one processor changes a memory location, that change is made visible to other processors which share that memory. We describe this in more detail when we look at multi-processing.

9.4 Cache architecture

In a von Neumann architecture, there is a single cache used for instruction and data (a unified cache). A modified Harvard architecture has separate instruction and data buses and therefore there are two caches, an instruction cache (I-cache) and a data cache (D-cache). In many ARM systems, we can have distinct instruction and data level 1 caches backed by a unified level 2 cache.

Let's consider how a cache memory is constructed. The cache needs to hold an address, some data and some status information. The top bits of the 32-bit address tells the cache where the information came from in main memory and is known as the *tag*. The total cache size is a measure of the amount of data it can hold; the RAMs used to hold tag values are not included in the calculation. The tag does, however, take up physical space in the cache.

It would be inefficient to hold one word of data for each tag address, so we typically group several locations together under the same tag. This logical block is commonly known as a cache line. The middle bits of the address, or *index*, identify the line. The index is used as address for the cache RAMs and does not need to be stored as a part of the tag. This will be covered in more detail later in this chapter. A cache line is said to be valid when it contains cached data or instructions, and invalid when it does not.

This means that the bottom few bits of the address (the offset) are not required to be stored in the tag – we need to record the address of a line, not of each byte within the line, so the five or six least significant bits will always be 0.

Associated with each line of data are one or more status bits. Typically, we will have a valid bit, which marks the line as containing data that can be used. (This means that the address tag represents some real value.) In a data cache we will also have one or more dirty bits which mark whether the cache line (or part of it) holds data which is not the same as (newer than) the contents of main memory. We will treat this in more detail later in the chapter.

9.4.1 Cache controller

This is a hardware block which has the task of managing the cache memory, in a way which is (largely) invisible to the program. It automatically writes code or data from main memory into the cache. It takes read and write memory requests from the processor and performs the necessary actions to the cache memory or the external memory.

When it receives a request from the processor it must check to see whether the requested address is to be found in the cache. This is known as a cache look-up. It does this by comparing a subset of the address bits of the request with tag values associated with lines in the cache. If there is a match (a hit) and the line is marked valid then the read or write will happen using the cache memory.

When the processor requests instructions or data from a particular address, but there is no match with the cache tags, or the tag is not valid, a cache *miss* results and the request must be passed to the next level of the memory hierarchy – an L2 cache, or external memory. It can also cause a cache linefill. A cache linefill causes the contents of a piece of main memory to be copied into the cache. At the same time, the requested data or instructions are streamed to the processor. This process happens transparently and is not directly visible to the programmer.

The processor need not wait for the linefill to complete before using the data. The cache controller will typically access the critical word within the cache line first. For example, if we perform a load instruction which misses in the cache and triggers a cache linefill, the first read to external memory will be that of the actual address supplied by the load instruction. This critical data is supplied to the processor pipeline, while the cache hardware and external bus interface then read the rest of the cache line, in the background.

9.4.2 Direct mapped caches

We now look at various different ways of implementing caches starting with the simplest, a *direct mapped* cache.

In a direct mapped cache, each location in main memory maps to a single location in the cache. As main memory is many times larger than the cache, many addresses map to the same cache location. [Figure 9-1](#) shows a small cache, with four words per line and four lines. This means that the cache controller will use two bits of the address (bits [3:2]) as the offset to select a word within the line and two bits of the address (bits [5:4]) as the index to select one of the four available lines. The remaining bits of the address (bits [31:6]) will be stored as a tag value.




Figure 9-1 Direct mapped cache operation

To look up a particular address in the cache, the hardware extracts the index bits from the address and reads the tag value associated with that line in the cache. If the two are the same and the valid bit indicates that the line contains valid data, it has a hit. It can then extract the data value from the relevant word of the cache line, using the offset and byte portion of the address. If the line contains valid data, but does not generate a hit (that is, the tag shows that the cache holds a different address in main memory) then the cache line is removed and is replaced by data from the requested address.

It should be clear that all main memory addresses with the same value of bits [5:4] will map to the same line in the cache. Only one of those lines can be in the cache at any given time. This means that we can easily get a problem called *thrashing*. Consider a loop which repeatedly accesses address 0x00, 0x40 and 0x80, as in the code below:

```
void add_array(int *data1, int *data2, int *result, int size)
{
    int i;

    for (i=0 ; i<size ; i++) {
        result[i] = data1[i] + data2[i];
    }
}
```

```
}

```

In this code example, if `result`, `data1`, and `data2` are pointers to `0x00`, `0x40` and `0x80` respectively then this loop will cause repeated accesses to memory locations that all map to the same line in the basic cache, as shown in [Figure 9-1 on page 9-7](#).

- When we first read address `0x40`, it will not be in the cache and so a linefill takes place putting the data from `0x40` to `0x4F` into the cache.
- When we then read address `0x80`, it will not be in the cache and so a linefill takes place putting the data from `0x80` to `0x8F` into the cache – and in the process we lose the data from address `0x40` to `0x4F` from the cache.
- The result is written to `0x00`. Depending on the allocation policy this may cause another line fill. The data from `0x80` to `0x8F` may be lost.
- The same thing will happen on each iteration of the loop and our software will perform poorly. Direct mapped caches are therefore not typically used in the main caches of ARM processors, but we do see them in some places – for example in the branch target address cache of the ARM1136 processor.

Processors can have hardware optimizations for situations where the whole cache line is being written to. This is a condition which can take a significant proportion of total cycle time in some systems. For example, this can happen when `memcpy()`- or `memset()`-like functions which perform block copies or zero initialization of large blocks are executed. In such cases, there is no benefit in first reading the data values which will be over-written. This can lead to situations where the performance characteristics of the cache are different to what might normally be expected.

Cache allocate policies act as a hint to the processor, they do not guarantee that a piece of memory will be read into the cache, and as a result, programmers should not rely on that.

9.4.3 Set associative caches

The main cache(s) of ARM processors are always implemented using a set associative cache. This significantly reduces the likelihood of the cache thrashing seen with direct mapped caches, thus improving program execution speed and giving more deterministic execution. It comes at the cost of increased hardware complexity and a slight increase in power (because multiple tags are compared on each cycle).

With this kind of cache organization, we divide the cache into a number of equally-sized pieces, called ways. The index field of the address continues to be used to select a particular line, but now it points to an individual line in each way. Commonly, there are 2- or 4-ways, but some ARM implementations have used higher numbers. Level 2 cache implementations (such as the ARM L2C-310) can have larger numbers of ways (higher associativity) due to their much larger size. The cache lines with the same index value are said to belong to a set. To check for a hit, we must look at each of the tags in the set.

In [Figure 9-2 on page 9-9](#), a cache with 2-ways is shown. Data from address `0x00` (or `0x40`, or `0x80`) may be found in line 0 of either (but not both) of the two cache ways.




Figure 9-2 A 2-way set-associative cache

Increasing the associativity of the cache reduces the probability of thrashing. The ideal case is a fully associative cache, where any main memory location can map anywhere within the cache. However, building such a cache is impractical for anything other than very small caches (for example, those associated with MMU TLBs – see [Chapter 10](#)). In practice, performance improvements are minimal for Level 1 caches above 4-way associativity, with 8-way or 16-way associativity being more useful for larger level 2 caches.

9.4.4 Cache terminology

A brief summary of some of the terms used may be helpful:




Figure 9-3 Cache terminology

- A *line* refers to the smallest loadable unit of a cache, a block of contiguous words from main memory.
- The *index* is the part of a memory address which determines in which line(s) of the cache the address can be found.
- A *way* is a subdivision of a cache, each way being of equal size and indexed in the same fashion. The line associated with a particular index value from each cache way grouped together forms a *set*.

- The *tag* is the part of a memory address stored within the cache which identifies the main memory address associated with a line of data.

9.4.5 A real-life example

Before we move on to look at write buffers, let's consider an example which is more realistic than those shown in the previous two diagrams. [Figure 9-4](#) is a 4-way set associative 32KB data cache, with an 8-word cache line length. This kind of cache structure might be found on the Cortex-A9 or Cortex-A5 processors.

The cache line length is eight words (32 bytes) and we have 4-ways. 32KB divided by 4, divided by 32 gives us a figure of 256 lines in each way. This means that we need the eight bits to index a line within a way (bits [12:5]). We need to use bits [4:2] of the address to select from the eight words within the line. The remaining bits [31:13] will be used as a tag.




Figure 9-4 A 32KB 4-way set associative cache

9.4.6 Virtual and physical tags and indexes

This section assumes some knowledge of the address translation process. Readers unfamiliar with virtual addressing may wish to revisit this section after reading [Chapter 10](#).

In *A real-life example*, we were a little imprecise about specification of exactly which address is used to perform cache lookups. Early ARM processors (for example, the ARM720T or ARM926EJ-S processors) used virtual addresses to provide both the index and tag values. This has the advantage that the processor can do a cache look-up without the need for a virtual to physical address translation. The drawback is that changing the virtual to physical mappings in the system means that the cache must first be cleaned and invalidated, and this can have a significant performance impact. We will go into more detail about these terms in [Invalidating and cleaning cache memory on page 9-16](#).

ARM11 family processors use a different cache tag scheme. Here, the cache index is still derived from a virtual address, but the tag is taken from the physical address. The advantage of a physical tagging scheme is that changes in virtual to physical mappings do not now require the cache to be invalidated. This can have significant benefits for complex multi-tasking operating systems which can frequently modify page table mappings. Using a virtual index has some hardware advantages. It means that the cache hardware can read the tag value from the appropriate line in each way in parallel without actually performing the virtual to physical address translation, giving a fast cache response. Such a cache is often described as *Virtually Indexed, Physically Tagged* (VIPT). Cache properties of Cortex-A series processors, including the use of these tagged caches are given in [Table 9-1](#). Other properties of the Cortex-A series processors are listed in [Table 2-3 on page 2-10](#).

Table 9-1 Cache features of Cortex-A series processors

	Processor				
	Cortex-A5	Cortex-A7	Cortex-A8	Cortex-A9	Cortex-A15
L2 Cache	External	Integrated	Integrated	External	Integrated
L2 Cache size	-	128KB to 1MB ^a	0KB to 1MB ^a	-	512KB to 4MB ^a
Cache Implementation (Data)	PIPT	PIPT	PIPT	PIPT	PIPT
Cache Implementation (Instruction)	VIPT	VIPT	VIPT	VIPT	PIPT
L1 Cache size (data) ^a	4K to 64K ^a	8KB to 64KB ^a	16/32KB ^a	16KB/32KB/64KB ^a	32KB
Cache size (Inst) ^a	4K to 64K ^a	8KB to 64KB ^a	16/32KB ^a	16KB/32KB/64KB ^a	32KB
L1 Cache Structure	2-way set associative (Inst) 4-way set associative (Data)	2-way set associative (Inst) 4-way set associative (Data)	4-way set associative	4-way set associative (Inst) 4-way set associative (Data)	2-way set associative (Inst) 2-way set associative (Data)
L2 Cache Structure	-	8-way set associative	8-way set associative	-	16-way associative
Cache line (words)	8	8	16	8	16

a. Configurable

However, there is a drawback to a VIPT implementation. For a 4-way set associative 32KB or 64KB cache, bits [12] and [13] of the address are needed to select the index. If 4KB pages are used in the MMU, bits [13:12] of the virtual address may not be equal to bits [13:12] of the physical address. There is therefore scope for potential cache coherency problems if multiple virtual address mappings point to the same physical address. This is resolved by placing certain restrictions on such multiple mappings which kernel page table software must obey. This is described as a *page coloring* issue and exists on other processor architectures for the same reasons.

This problem is avoided by using a *Physically Indexed, Physically Tagged* (PIPT) cache implementation. The Cortex-A series of processors described in this book use such a scheme for their data caches. It means that page coloring issues are avoided, but at the cost of hardware complexity.

9.5 Cache policies

There are a number of different choices which can be made in cache operation. We need to consider what causes a line from external memory to be placed into the cache (*allocation policy*). We need to look at how the controller decides which line within a set associative cache to use for the incoming data (*replacement policy*). And we need to control what happens when the processor performs a write which hits in the cache (*write policy*).

9.5.1 Allocation policy

When the processor does a cache look-up and the address it wants is not in the cache, it must determine whether or not to perform a cache linefill and copy that address from memory.

- A *read allocate* policy allocates a cache line only on a read. If a write is performed by the processor which misses in the cache, the cache is not affected and the write goes to the next level of the hierarchy.
- A *write allocate* policy allocates a cache line for *either* a read or write which misses in the cache (and so might more accurately be called a *read-write cache allocate* policy). For both memory reads which miss in the cache and memory writes which miss in the cache, a cache linefill is performed. This is typically used in combination with a *write-back* write policy on current ARM processors, as we shall see in [Write policy](#).

9.5.2 Replacement policy

When there is a cache miss, the cache controller must select one of the cache lines in the set for the incoming data. The cache line selected is called the *victim*. If the victim contains valid, dirty data, the contents of that line must be written to main memory before new data can be written to the victim cache line. This is called *eviction*.

The *replacement policy* is what controls the victim selection process. The index bits of the address are used to select the set of cache lines, and the replacement policy selects the specific cache line from that set which is to be replaced.

- *Round-robin* or *cyclic* replacement means that we have a counter (the *victim counter*) which cycles through the available ways and cycles back to 0 when it reaches the maximum number of ways.
- *Pseudo-random* replacement randomly selects the next cache line in a set to replace. The victim counter is incremented in a pseudo-random fashion and can point to any line in the set.

Most ARM processors support both policies.

A round-robin replacement policy is generally more predictable, but can suffer from poor performance in certain use cases and for this reason, the pseudo-random policy is often preferred.

9.5.3 Write policy

When the processor executes a store instruction, a cache lookup on the address(es) to be written is performed. For a cache hit on a write, there are two choices.

- *Write-through*. With this policy writes are performed to both the cache and main memory. This means that the cache and main memory are kept coherent. As there are more writes to main memory, a write-through policy is slower than a write-back policy if the write buffer fills and therefore is less commonly used in today's systems (although it can be useful for debug).

- *Write-back.* In this case, writes are performed only to the cache, and not to main memory. This means that cache lines and main memory can contain different data. The cache line holds newer data, and main memory contains older data (said to be *stale*). To mark these lines, each line of the cache has an associated *dirty* bit (or bits). When a write happens which updates the cache, but not main memory, the dirty bit is set. If the cache later evicts a cache line whose dirty bit is set (a *dirty line*), it writes the line out to main memory. Using a write-back cache policy can significantly reduce traffic to slow external memory and therefore improve performance and save power. However, if there are other agents in the system which can access memory at the same time as the processor, we may need to worry about coherency issues. This is described in more detail later.

9.6 Write and Fetch buffers

A write buffer is a hardware block inside the processor (but sometimes in other parts of the system as well), implemented using a number of buffers. It accepts address, data and control values associated with processor writes to memory. When the processor executes a store instruction, it may place the relevant details (the location to write to, the data to be written, the transaction size and so forth) into the buffer. The processor does not have to wait for the write to be completed to main memory. It can proceed with executing the next instructions. The write buffer itself will drain the writes accepted from the processor, to the memory system.

A write buffer can increase the performance of the system. It does this by freeing the processor from having to wait for stores to complete. In effect, provided there is space in the write buffer, the write buffer is a way to hide latency. If the number of writes is low or well spaced, the write buffer will not become full. If the processor generates writes faster than they can be drained to memory, the write buffer will eventually fill and there will be little performance benefit.

Some write buffers support write merging (also called write combining). They can take multiple writes (for example, a stream of writes to adjacent bytes) and merge them into one single burst. This can reduce the write traffic to external memory and therefore boost performance.

It will be obvious to the experienced programmer that sometimes the behavior of the write buffer is not what we want when accessing a peripheral, we might want the processor to stop and wait for the write to complete before proceeding to the next step. Sometimes we really want a stream of bytes to be written and we don't want the stores to be combined. In [ARM memory ordering model on page 11-4](#), we'll look at memory types supported by the ARM architecture and how to use these to control how the caches and write buffers are used for particular devices or parts of the memory map.

Similar components, called fetch buffers, can be used for reads in some systems. In particular, processors typically contain prefetch buffers which read instructions from memory ahead of them actually being inserted into the pipeline. In general, such buffers are transparent to the programmer. We will consider some possible hazards associated with this when we look at memory ordering rules.

9.7 Cache performance and hit rate

The *hit rate* is defined as the number of cache hits divided by the number of memory requests made to the cache during a specified time (normally calculated as a percentage). Similarly, the *miss rate* is the number of total cache misses divided by the total number of memory requests made to the cache. One may also calculate the number of hits or misses on reads or writes only.

Clearly, a higher hit rate will generally result in higher performance. It is not really possible to quote example figures for typical software, the hit rate is very dependent upon the size of the critical parts of the code or data operated on and of course, the size of the cache.

There are some simple rules which can be followed to give better performance. The most obvious of these is to enable caches and write buffers and to use them wherever possible (typically for all parts of the memory system which contain code and more generally for RAM and ROM, but not peripherals). Performance will be considerably increased in Cortex-A series processors if instruction memory is cached. Placing frequently accessed data together in memory can also be helpful. For example, a frequently accessed array will benefit from having a base address at the start of a cache line.

Fetching a data value in memory involves fetching a whole cache line; if none of the other words in the cache line will be used, there will be little or no performance gain. Smaller code may cache better than larger code and this can sometimes give (seemingly) paradoxical results. For example, a piece of C code may fit entirely within cache when compiled for Thumb (or for the smallest size) but not when compiled for ARM (or for maximum performance) and as a consequence can actually run faster than the more optimized version. We describe cache considerations in much more detail in [Chapter 19 Optimizing Code to Run on ARM Processors](#).

9.8 Invalidating and cleaning cache memory

Cleaning and invalidation can be required when the contents of external memory have been changed and the programmer wishes to remove stale data from the cache. It can also be required after MMU related activity such as changing access permissions, cache policies, or virtual to physical address mappings.

The word *flush* is often used in descriptions of clean and invalidate operations. ARM generally uses only the terms *clean* and *invalidate*.

- Invalidation of a cache (or cache line) means to clear it of data. This is done by clearing the valid bit of one or more cache lines. The cache always needs to be invalidated after reset as its contents will be undefined. If the cache might contain dirty data, it is generally incorrect to invalidate it. Any updated data in the cache from writes to write-back cacheable regions would be lost by simple invalidation.
- Cleaning a cache (or cache line) means to write the contents of dirty cache lines out to main memory and clear the dirty bit(s) in the cache line. This makes the contents of the cache line and main memory coherent with each other. Clearly, this is only applicable for data caches in which a write-back policy is used. Cache invalidate, and clean and invalidate operations can be performed by cache set, or way, or by virtual address.

Copying code from one location to another (or other forms of self-modifying code) may require the programmer either to clean and/or to invalidate the cache. The memory copy code will use load and store instructions and these will operate on the data side of the processor. If the data cache is using a write-back policy for the area to which code is written, it will be necessary to clean that data from the cache before the code can be executed. This ensures that the instructions stored as data go out into main memory and are then available for the instruction fetch logic. In addition, if the area to which code is written was previously used for some other program, the instruction cache could contain stale code (from before main memory was re-written). Therefore, it may also be necessary to invalidate the instruction cache before branching to the newly copied code.

The commands to either clean or invalidate the cache are CP15 operations. They are available only to privileged code and cannot be executed in User mode. In systems where the TrustZone Security Extensions are in use, there can be hardware limitations applied to non-secure use of some of these operations.

CP15 instructions exist which will clean, invalidate, or clean *and* invalidate level 1 data or instruction caches. Invalidation without cleaning is safe only when it is known that the cache cannot contain dirty data – for example a Harvard instruction cache. The programmer can perform the operation on the entire cache, or just on individual lines. These individual lines can be specified either by giving a virtual address to be cleaned or to be invalidated, or by specifying a line number in a particular set, in cases where the hardware structure is known. The same operations can be performed on the L2 or outer caches and we will look at this in [Level 2 cache controller on page 9-18](#).

Of course, these operations will be accessed through kernel code – in Linux, you will use the `__clear_cache()` function implemented in `arch/arm/mm/cache-v7.S`. Equivalent functions exist in other operating systems – Google Android has `cacheflush()`, for example.

A common situation where cleaning or invalidation can be required is DMA (Direct Memory Access). When it is required to make changes made by the processor visible to external memory, so that it can be read by a DMA controller, it might be necessary to clean the cache. When external memory is written by a DMA controller and it is necessary to make those changes visible to the processor, the affected addresses will need to be invalidated in the cache.

9.9 Point of coherency and unification

For set/way based clean and invalidate, the operation is performed on a specific level of cache. For operations which use a virtual address, the architecture defines two conceptual points.

- *Point of Coherency (PoC)*. For a particular address, the PoC is the point at which all blocks (for example, processors, DSPs, or DMA engines) which can access memory are guaranteed to see the same copy of a memory location. Typically, this will be the main external system memory.
- *Point of Unification (PoU)*. The PoU for a processor is the point at which the instruction and data caches and the page table walks of the processor are guaranteed to see the same copy of a memory location. For example, a unified level 2 cache would be the point of unification in a system with Harvard level 1 caches and a TLB for cacheing translation table entries.

Readers unfamiliar with the terms hardware page table walk or Translation Lookaside Buffer (TLB) will find these described in [Chapter 10](#). If no external cache is present, main memory would be the PoU.

In the case of an MPCore processor, or a big.LITTLE combination (an example of a Inner Shareable shareability domain, using the terminology from [Chapter 11 Memory Ordering](#)), the PoU is where instruction and data caches and page table walks of all the processors within the MPCore cluster are guaranteed to see the same copy of a memory location.

Knowledge of the PoU enables self-modifying code to ensure future instruction fetches are correctly made from the modified version of the code. They can do this by using a two-stage process:

- Clean the relevant data cache entries (by address).
- Invalidate instruction cache entries (by address).

In addition, the use of memory barriers will be required here, the *ARM Architecture Reference Manual* provides detailed examples of the necessary code sequences.

Similarly, we can use the clean data cache entry and invalidate TLB operations to ensure that all writes to the page tables are visible to the MMU.

9.10 Level 2 cache controller

At the start of this chapter, we briefly described the partitioning of the memory system and explained how many systems have a multi-layer cache hierarchy. The Cortex-A5 and Cortex-A9 processors do not have an integrated level 2 cache. Instead, the system designer can opt to connect the ARM L2 cache controller (L2C-310) outside of the processor instance.

The L2C-310 cache controller can support a cache of up 8MB in size, with a set associativity of between four and sixteen ways. The size and associativity are fixed by the SoC designer. The level 2 cache can be shared between multiple processors (or indeed between the processor and other agents, for example a graphics processor). It is possible to lockdown cache data on a per-master per-way basis, enabling management of cache sharing between multiple components.

9.10.1 Level 2 cache maintenance

We saw in [Virtual and physical tags and indexes on page 9-10](#) how the programmer may need the ability either to clean or invalidate some or all of a cache. This can be done by writing to memory-mapped registers within the L2 cache controller in the case where the cache is external to the processor, or through CP15, where the level 2 cache is implemented inside the processor. The registers themselves are not cached, which makes this feasible. Where such operations are performed by having the processor perform memory-mapped writes, the processor needs a way of determining when the operation is complete. It does this by polling a further memory-mapped register within the L2 cache controller.

The ARM L2C-310 Level 2 cache controller operates only on physical addresses. Therefore, to perform cache maintenance operations, it may be necessary for the program to perform a virtual to physical address translation. The L2C-310 provides a *cache sync* operation which forces the system to wait for pending operations to complete.

9.11 Parity and ECC in caches

So-called *soft errors* are increasingly a concern in today's systems. Smaller transistor geometries and lower voltages give circuits an increased sensitivity to perturbation by cosmic rays and other background radiation, alpha particles from silicon packages, or from electrical noise. This is particularly true for memory devices which rely on storing small amounts of charge and which also occupy large proportions of total silicon area. In some systems, mean-time-between-failure could be measured in seconds if appropriate protection against soft errors was not employed.

The ARM architecture provides support for parity and *Error Correcting Code* (ECC) in the caches (and tightly coupled memories). Parity means that we have an additional bit which marks whether the number of bits with the value one is even (or odd, depending upon the scheme chosen). This provides a simple check against single bit errors. An ECC scheme enables detection of multiple bit failures and possible recovery from soft errors, but recovery calculations can take several cycles. Implementing a processor which is tolerant of level 1 cache RAM accesses taking multiple clock cycles significantly complicates the processor design. ECC is therefore more commonly used only on blocks of memory (for example, the Level 2 cache), outside the processor. The Cortex-A15, however, supports ECC and parity inside the processor.

Parity is checked on reads and writes, and can be implemented on both tag and data RAMs. Parity mismatch generates a prefetch or data abort exception, and the fault status address registers are updated appropriately.

Chapter 10

Memory Management Unit

This chapter describes the main concepts of virtual memory systems and memory management. An important function of a *Memory Management Unit* (MMU) is to allow us to manage tasks as independent programs running in their own private virtual memory space. A key feature of such a virtual memory system is address relocation, which is the translation of the (virtual) address issued by the processor to a different (physical) address in main memory. The ARM MMU translates addresses of code and data from the virtual view of memory to the physical addresses in the real system. The translation is carried out by the MMU hardware and is transparent to the application (ignoring for the moment any performance issues). In addition, it controls such things as memory access permissions, memory ordering and cache policies for each region of memory.

In multi-tasking embedded systems, we typically need a way to partition the memory map and assign permissions and memory attributes to these regions of memory. In situations where we are running more complex operating systems, like Linux, we need even greater control over the memory system.

The MMU enables tasks or applications to be written in a way which requires them to have no knowledge of the physical memory map of the system, or about other programs which may be running simultaneously. This enables us to use the same virtual memory address space for each program. It also lets us work with a contiguous virtual memory map, even if the physical memory is fragmented. This virtual address space is separate from the actual physical map of memory in the system. Applications are written, compiled and linked to run in the virtual memory space. Virtual addresses are those used by the programmer, compiler and linker when placing code in memory. Physical addresses are those used by the actual hardware system.

It is the responsibility of the operating system to program the MMU to translate between these two views of memory. [Figure 10-1 on page 10-2](#) shows an example system, illustrating the virtual and physical views of memory. Different processors and/or devices in a single system may have different virtual and physical address maps, for example, some multi-core boards and PCI devices.

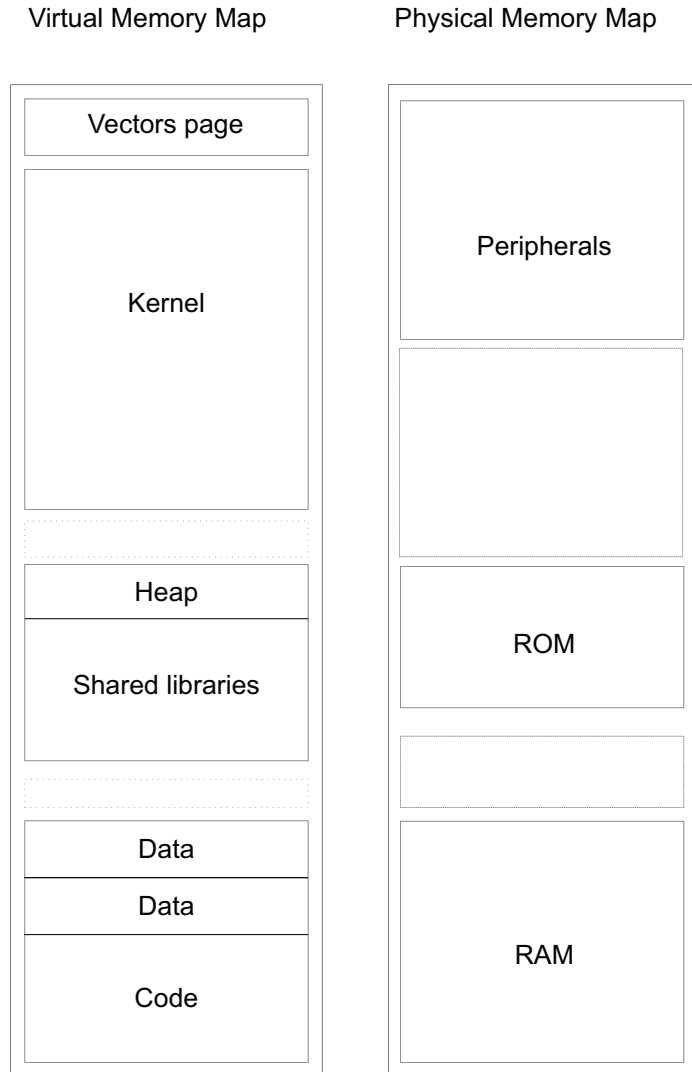


Figure 10-1 Virtual and physical memory

When the MMU is disabled, all virtual addresses map directly to the corresponding physical address (*a flat mapping*). If the MMU cannot translate an address, it generates an abort exception on the processor and provides information to the processor about what the problem was. This feature can be used to map memory or devices on-demand, one page at a time.

10.1 Virtual memory

The MMU allows us to build systems with multiple virtual address maps. Each task can have its own virtual memory map. The OS kernel places code and data for each application in physical memory, but the application itself does not need to know the location.

The key feature of the MMU hardware is address translation. It does this using *page tables*. These contain a series of entries, each of which describes the physical address translation for part of the memory map. Page table entries are organized by virtual address. In addition to describing the translation of that virtual page to a physical page, they also provide access permissions and memory attributes necessary for that page.

Note

In the ARM architecture, the concept referred to in generic computer terminology as *page tables* has a more specific meaning. The ARM architecture uses multi-level page tables, and defines *translation tables* as a generic term for all of these. It then reserves the use of the term page tables to describe second level translation tables when using the *short-descriptor format* (not making use of *Large Physical Address Extension* (LPAE)).

Addresses generated by the processor are virtual addresses. When the MMU is not enabled, there is a direct mapping between the physical and virtual addresses. The MMU essentially replaces the most significant bits of this virtual address with some other value, to generate the physical address (effectively defining a base address of a piece of memory). The lower bits are the same in both addresses (defining an offset in physical memory from the base address). This smallest amount of memory that can be mapped independently is known as a page. The ARM MMU supports a multi-level page table architecture with two levels of page table: level 1 (L1) and level 2 (L2). We will describe the meaning of level 1 and level 2 in a moment. The same translation tables are used to define the translations and memory attributes that apply to both instruction fetches and to data accesses. The process in which the MMU accesses page tables to translate addresses is known as *page table walking*.

10.2 Level 1 page tables

Let's take a look at the process by which a virtual address is translated to a physical address using level 1 page table entries on an ARM processor. The first step is to locate the page table entry associated with the virtual address.

The first stage of translation uses a single level 1 page table (sometimes called a master page table). It can contain two basic types of page table entries. The L1 page table divides the full 4GB address space of a 32-bit processor into 4096 equally sized sections, each of which describes 1MB of virtual memory space. The L1 page table therefore contains 4096 word-sized entries.

Each entry can either hold a pointer to the base address of a level 2 page table or a page table entry for translating a 1MB section. If the page table entry is translating a 1MB section (determined by the encoding, see [Figure 10-3 on page 10-5](#)), it gives the base address of the 1MB page in physical memory.

The base address of the L1 page table is known as the Translation Table Base Address and is held within a register in CP15 c2. It must be aligned to a 16KB boundary.

When the MMU performs a translation, the top 12 bits of the requested virtual address acts as the index into the translation table. This is illustrated in [Figure 10-2](#).

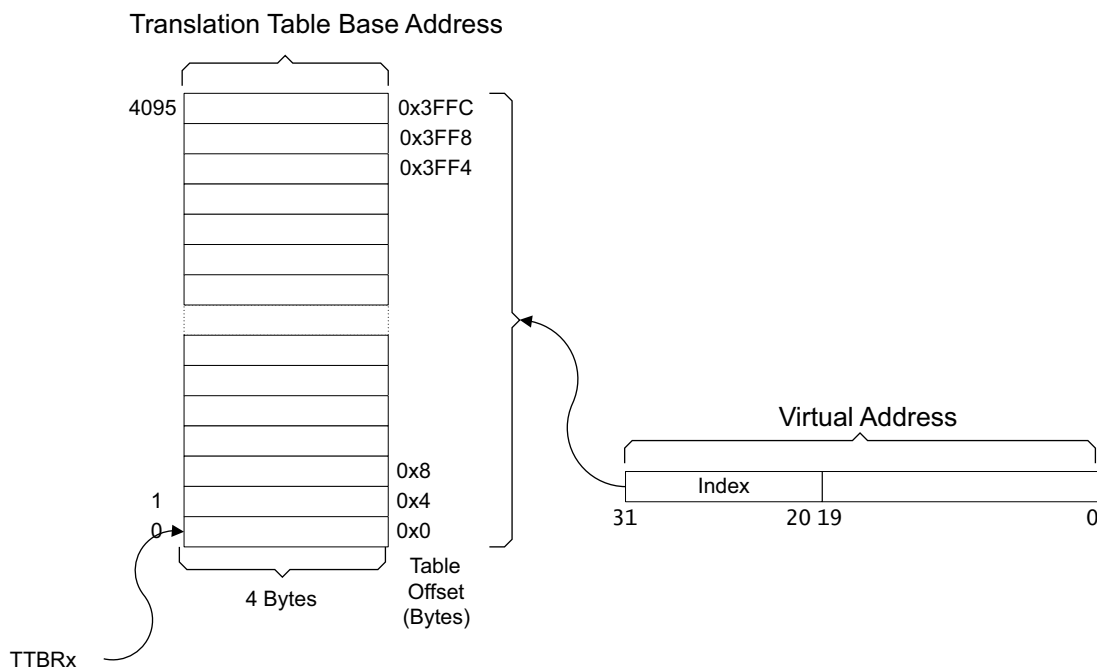


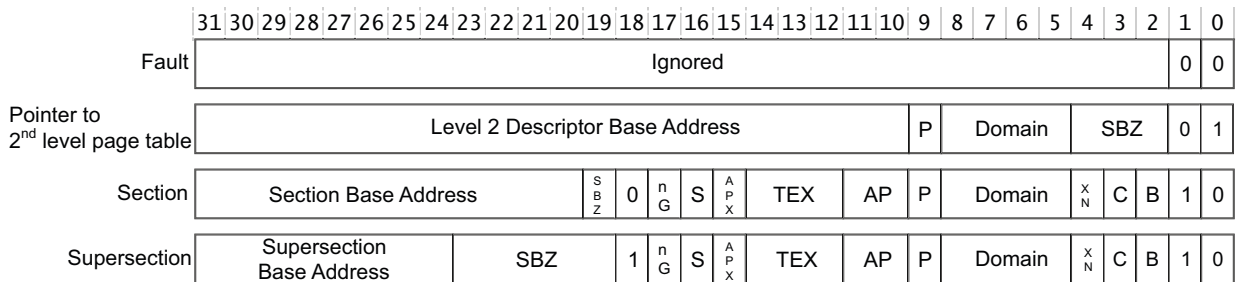
Figure 10-2 Finding the address of the level 1 page table entry

To take a simple example, suppose we place our L1 page table at address 0x12300000. The processor issues virtual address 0x00100000. The top 12 bits [31:20] define which 1MB of virtual address space is being accessed. In this case 0x001, so we need to read table entry 1. To get the offset into the table we must multiply the entry number by entry size:

$$0x001 * 4 \text{ bytes} = \text{address offset of } 0x004$$

The address of the entry the MMU reads the word from is $0x12300000 + 0x004 = 0x12300004$.

[Figure 10-3](#) shows the format of a L1 page table entry.

**Figure 10-3 Level 1 page table entry format**

To summarize, an L1 page table entry can be one of four possible types:

- A fault entry that generates an abort exception. This can be either a prefetch or data abort, depending on the type of access. This effectively indicates virtual addresses that are unmapped.
- A 1MB section translation entry, mapping a 1MB region to a physical address.
- An entry that points to an L2 page table. This enables a 1MB piece of memory to be further sub-divided into pages.
- A 16MB supersection. This is a special kind of 1MB section entry, which requires 16 entries in the page table, but can reduce the number of entries allocated in the Translation Lookaside Buffer for this region.

The least significant two bits [1:0] in the entry define whether the entry is a fault entry, a page table entry, or a section entry. Bit [18] is used to distinguish between a normal section and supersection.

The page table entry for a section (or supersection) contains the physical base address used to translate the virtual address. Note that many other bits are given in the page table entry, including the Access Permissions (AP) and Cacheable (C) or Bufferable (B) types, which we will consider in [Memory attributes on page 10-12](#). This is all of the information required to access the corresponding physical address and in these cases, the MMU does not need to look beyond the L1 table.

[Figure 10-4 on page 10-6](#) summarizes the translation process for an address translated by a section entry in the L1 page table.

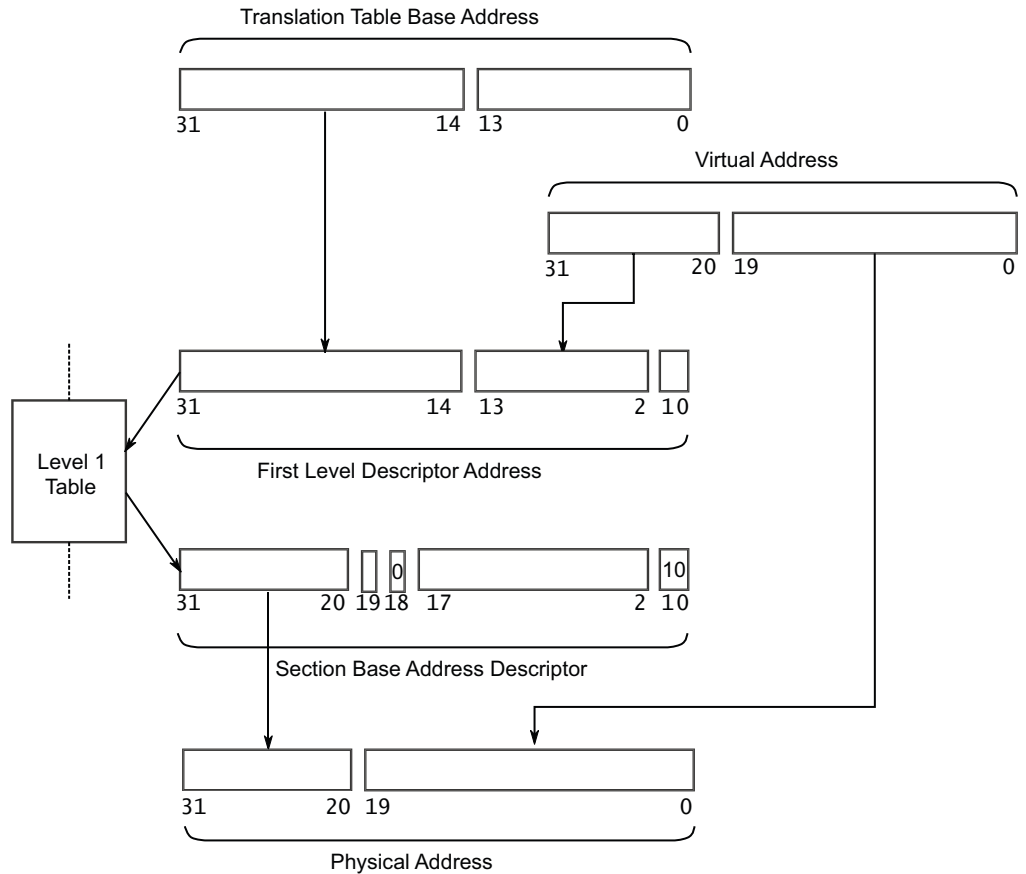


Figure 10-4 Generating a physical address from a level 1 page table entry

In a page table entry for a 1MB section of memory, the upper 12 bits of the page table entry replace the upper 12 bits of the virtual address when generating the physical address, as [Figure 10-3 on page 10-5](#) shows.

A supersection is a 16MB piece of memory, which must have both its virtual and physical base address aligned to a 16MB boundary. As L1 page table entries each describe 1MB, we need 16 consecutive, identical entries within the table to mark a supersection. In [Choice of page sizes on page 10-11](#), we describe why supersections can be useful.

10.3 Level 2 page tables

An L2 page table has 256 word-sized (4 byte) entries, requires 1KB of memory space and must be aligned to a 1KB boundary. Each entry translates a 4KB block of virtual memory to a 4KB block in physical memory. A page table entry can give the base address of either a 4KB or 64KB page.

There are three types of entry used in L2 page tables, identified by the value in the two least significant bits of the entry:

- A large page entry points to a 64KB page.
- A small page entry points a 4KB page.
- A fault page entry generates an abort exception if accessed.

Figure 10-5 shows the format of L2 page table entries.




Figure 10-5 Format of a level 2 page table entry

As with the L1 page table entry, a physical address is given, along with other information about the page. Type extension (TEX), Shareable (S), and Access Permission (AP, APX) bits are used to specify the attributes necessary for the ARMv7 memory model. Along with TEX, the C and B bits control the cache policies for the memory governed by the page table entry. The nG bit defines the page as being global (applies to all processes) or non-global (used by a specific process). We will describe all of these bits in more detail in [Memory attributes on page 10-12](#).




Figure 10-6 Generating the address of the level 2 page table entry

In [Figure 10-6](#) we see how the address of the L2 page table entry that we need is calculated by taking the (1KB aligned) base address of the level 2 page table (given by the level 1 page table entry) and using 8 bits of the virtual address (bits [19:12]) to index within the 256 entries in the L2 page table.

[Figure 10-7](#) summarizes the address translation process when using two layers of page tables. Bits [31:20] of the virtual address are used to index into the 4096-entry L1 page table, whose base address is given by the CP15 TTB register. The L1 page table entry points to an L2 page table, which contains 256 entries. Bits [19:12] of the virtual address are used to select one of those entries which then gives the base address of the page. The final physical address is generated by combining that base address with the remaining bits of the virtual address.

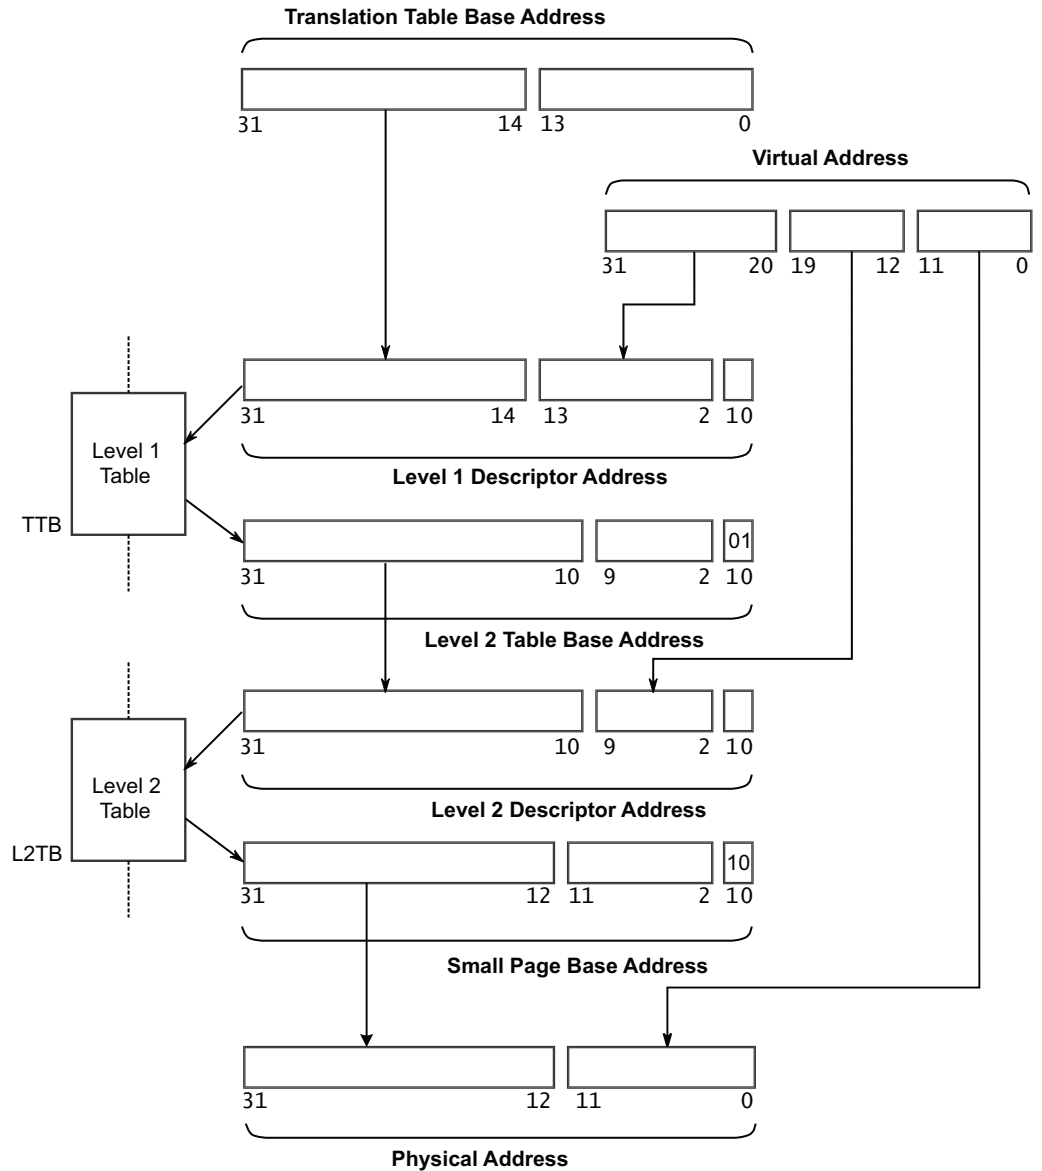


Figure 10-7 Summary of generation of physical address using the L2 page table entry

10.4 The Translation Lookaside Buffer

We have seen that a single memory request from the processor can result in a total of three memory requests to external memory – one for the level one page table walk, a second access for the L2 page table walk and then finally the original request from the processor. This seems like it would be ruinous for performance of the system. Fortunately, page table walks are relatively uncommon events in the majority of systems, due to another part of the MMU.

The *Translation Lookaside Buffer* (TLB) is a cache of page translations within the MMU. On a memory access, the MMU first checks whether the translation is cached in the TLB. If the requested translation is available, we have a TLB hit, and the TLB provides the translation of the physical address immediately. If the TLB does not have a valid translation for that address, we have a TLB miss and an external page table walk is required. This newly loaded translation can then be cached in the TLB for possible reuse.

The exact structure of the TLB differs between implementations of the ARM processors. What follows is a description of a typical system, but individual implementations may vary from this. There are one or more micro-TLBs, which are situated close to the instruction and data caches. Addresses with entries which hit in the micro-TLB require no additional memory look-up and no cycle penalty. However, the micro-TLB has only a small number of mappings (typically eight on the instruction side and eight on the data side). This is backed by a larger main TLB (typically 64 entries), but there may be some penalty associated with accesses which miss in the micro-TLB but which hit in the main TLB. [Figure 10-8](#) shows how each TLB entry contains physical and virtual addresses, but also attributes (such as memory type, cache policies and access permissions) and potentially an ASID value (described in [Address Space ID on page 10-15](#)).

The TLB is like other caches and so has a TLB line replacement policy, but this is effectively transparent to the programmer (although see the next section on maintaining TLB coherency). If the page table entry is a valid one, the virtual address, physical address and other attributes for the whole page or section are stored as a TLB entry. If the page table entry is not valid, the TLB will not be updated. The ARM architecture requires that only valid page table descriptors are cached within the TLB.




Figure 10-8 Illustration of TLB structure

10.5 TLB coherency

When the operating system changes page table entries, it is possible that the TLB could contain stale translation information. The OS should therefore take steps to invalidate TLB entries. There are several CP15 operations available which allow a global invalidate of the TLB or removal of specific entries. As speculative instruction fetches and data reads may cause page table walks, it is essential to invalidate the TLB when a valid page table entry is changed. Invalid page table entries cannot be cached in the TLB so they can be changed without invalidation.

The Linux kernel has a number of functions which use these CP15 operations, including `flush_tlb_all()` and `flush_tlb_range()`. Such functions are not typically required by device drivers.

10.6 Choice of page sizes

This is essentially controlled by the operating system, but it is worth being aware of the considerations involved when selecting page sizes. Smaller page sizes allow finer control of a block of memory and potentially can reduce the amount of unused memory in a page. If a task needs 7KB of data space, there is less unused space if it is allocated two 4KB pages as opposed to a 64KB page or a 1MB section. Smaller page sizes also allow finer control over permissions, cache properties and so forth.

However, with increased page sizes, each entry in the TLB holds a reference to a larger piece of memory. It is therefore more likely that a TLB hit will occur on any access and so there will be fewer page table walks to slow external memory. For this reason, 16MB supersections can be used with large pieces of memory which do not require detailed mapping. In addition, each L2 page table requires 1KB of memory.

10.7 Memory attributes

We have seen how page table entries allow the MMU hardware to translate virtual to physical addresses. However, they also specify a number of attributes associated with each page, including access permissions, memory type and cache policies.

10.7.1 Memory Access Permissions

The Access Permission (AP and APX) bits in the page table entry give the access permission for a page. See [Table 10-1](#).

An access which does not have the necessary permission (or which faults) will be aborted. On a data access, this will result in a precise data abort exception. On an instruction fetch, the access will be marked as aborted and if the instruction is not subsequently flushed before execution, a prefetch abort exception will be taken. Faults generated by an external access will not, in general, be precise.

Information about the address of the faulting location and the reason for the fault is stored in CP15 (the fault address and fault status registers). The abort handler can then take appropriate action – for example, modifying page tables to remedy the problem and then returning to the application to retry the access. Alternatively, the application which generated the abort may have a problem and need to be terminated.

Table 10-1 Summary of Access Permission encodings

APX	AP	Privileged	Unprivileged	Description
0	00	No access	No access	Permission fault
0	01	Read/Write	No access	Privileged Access only
0	10	Read/Write	Read	No user-mode write
0	11	Read/Write	Read/Write	Full access
1	00	-	-	Reserved
1	01	Read	No access	Privileged Read only
1	10	Read	Read	Read only
1	11	-	-	Reserved

10.7.2 Memory types

Earlier ARM architecture versions enabled the programmer to specify the memory access behavior of pages by configuring whether the cache and write buffer could be used for that location. This simple scheme is inadequate for today's more complex systems and processors, where we can have multiple levels of caches, hardware managed coherency between multiple processors sharing memory and processors which can speculatively fetch both instructions and data. The new memory types added to the ARM architecture in ARMv6 and extended in the ARMv7 architecture are designed to meet these needs.

Three mutually exclusive memory types are defined in the ARM architecture. All regions of memory are configured as one of these three types:

- Strongly-ordered
- Device

- Normal.

These are used to describe the memory regions. A summary of the memory types is shown in [Table 10-2](#).

Table 10-2 Memory attributes

Memory type	Shareable/ Non-shareable	Cacheable	Description
Normal	Shareable	Yes	Designed to handle normal memory that is shared between several processors.
	Non-shareable	Yes	Designed to handle normal memory that is used only by a single processor.
Device	-	No	Designed to handle memory-mapped peripherals. ^a
Strongly ordered	-	No	All memory accesses to Strongly-ordered memory occur in program order. All Strongly-ordered accesses are assumed to be shared.

a. Shared memory was originally used to distinguish between accesses directed to the “peripheral private port” found on several ARM11 processors. This use is now deprecated and processors implementing LPAE treat all device accesses as Shareable.

Further descriptions of these memory types can be found in [ARM memory ordering model on page 11-4](#).

[Table 10-3](#) shows how the TEX, C and B bits within the page table entry are used to set the memory types of a page and also the cache policies to be used. The meaning of each of the memory types is described in [Chapter 11](#), while the cache policies were described in [Chapter 9](#).

Table 10-3 Memory type and cacheable properties encoding in page table entry

TEX	C	B	Description	Memory type
000	0	0	Strongly-ordered	Strongly-ordered
000	0	1	Sharable device	Device ^a
000	1	0	Outer and Inner write-through, no allocate on write	Normal
000	1	1	Outer and Inner write-back, no allocate on write	Normal
001	0	0	Outer and Inner non-cacheable	Normal
001	-	-	Reserved	-
010	0	0	Non-shareable device	Device ^a
010	-	-	Reserved	-
011	-	-	Reserved	-
1XX	Y	Y	Cached memory XX = Outer policy YY = Inner policy	Normal

a. LPAE treats all device accesses as Shareable

The final entry within the table needs further explanation. For normal cacheable memory, the two least significant bits of the TEX field are used to provide the outer cache policy (perhaps for level 2 or level 3 caches) while the C and B bits give the inner cache policy (for level 1 and any other cache which is to be treated as inner cache). This enables us to specify different cache policies for both the inner and outer cache. For the Cortex-A15 and Cortex-A8 processors, inner cache properties set by the page table entry apply to both L1 and L2 caches. On some older processors, outer cache may support write allocate, while the L1 cache may not. Such processors should still behave correctly when running code which requests this cache policy, of course.

10.7.3 Domains

The ARM architecture has an unusual feature which enables regions of memory to be tagged with a domain ID. There are 16 domain IDs provided by the hardware and CP15 c3 contains the *Domain Access Control Register* (DACR) which holds a set of 2-bit permissions for each domain number. This enables each domain to be marked as no-access, *manager* mode or *client* mode. No-access causes an abort on any access to a page in this domain, irrespective of page permissions. Manager mode ignores all page permissions and enables full access. Client mode uses the permissions of the pages tagged with the domain.

———— **Note** ————

The use of domains is deprecated in the ARMv7 architecture, and will eventually be removed, but in order for access permissions to be enforced, it is still necessary to assign a domain number to a section and to ensure that the permission bits for that domain are set to client.

10.8 Multi-tasking and OS usage of page tables

In most systems using Cortex-A series processors, we will have a number of applications or tasks running concurrently. Each task will have its own unique page tables residing in physical memory. Typically, much of the memory system is organized so that the virtual-to-physical address mapping is fixed, with page table entries that never change. This typically is used to contain operating system code and data, and also the page tables used by individual tasks.

Whenever an application is started, the operating system will allocate it a set of page table entries which map both the code and data used by the application to physical memory. If the application needs to map in code or extra data space (for example through a `malloc()` call), the kernel can subsequently modify these tables. When a task completes and the application is no longer running, the kernel can remove any associated page table entries and re-use the space for a new application. In this way, multiple tasks can be resident in physical memory. Upon a task switch, the kernel switches page table entries to the page in the next thread to be run. In addition, the dormant tasks are completely protected from the running task. This means that the MMU does not allow the running task to access the code or data of the kernel or of other user privilege tasks.

When the page table entries are changed, an access by code to a particular virtual address can now translate to different location in physical memory. This can give rise to several possible problems. The ARM architecture has some features to try to mitigate the performance impact of these.

Older ARM processors (from the ARM7 and ARM9 family) have cache tags which store virtual addresses. When page table mappings are changed, the caches can contain invalid data from the old page table mapping. To ensure memory coherency, the caches would need to be cleaned and invalidated. This can have a significant performance impact, as often instructions and data from a location which has just been invalidated would then need to be re-fetched from external memory. However, all Cortex-A series processors use physically tagged caches. As a result no coherency problems are created by changing page table entries.

In addition, the TLB may also have cached old page table entries and these will need to be invalidated.

10.8.1 Address Space ID

When we described the page table bits in [Level 2 page tables on page 10-7](#) we noted a bit called nG (non-global). If the nG bit is set for a particular page, the page is associated with a specific application. When the MMU performs a translation, it uses both the virtual address and an ASID value.

The ASID is a number assigned by the OS to each individual task. This value is in the range 0-255 and the value for the current task is written in the ASID register (accessed via CP15 c13). When the TLB is updated and the entry is marked as non-global, the ASID value will be stored in the TLB entry in addition to the normal translation information. Subsequent TLB look-ups will only match on that entry if the current ASID matches with the ASID that is stored in the entry. We can therefore have multiple valid TLB entries for a particular page (marked as non-global), but with different ASID values. This significantly reduces the software overhead of *context switches*, as it avoids the need to flush the on-chip TLBs. The ASID forms part of a larger (32-bit) process ID register that can be used in task-aware debugging.

————— Note —————

A *context switch* denotes the scheduler transferring execution from one process to another. This typically requires saving the current process state and restoring the state of the next process waiting to be run.

Figure 10-9 illustrates this. Here, we have multiple applications (A, B and C), each of which is linked to run from virtual address 0. Each application is located in a separate address space in physical memory. There is an ASID value associated with each application so we can have multiple entries within the TLB at any particular time, which will be valid for virtual address 0.

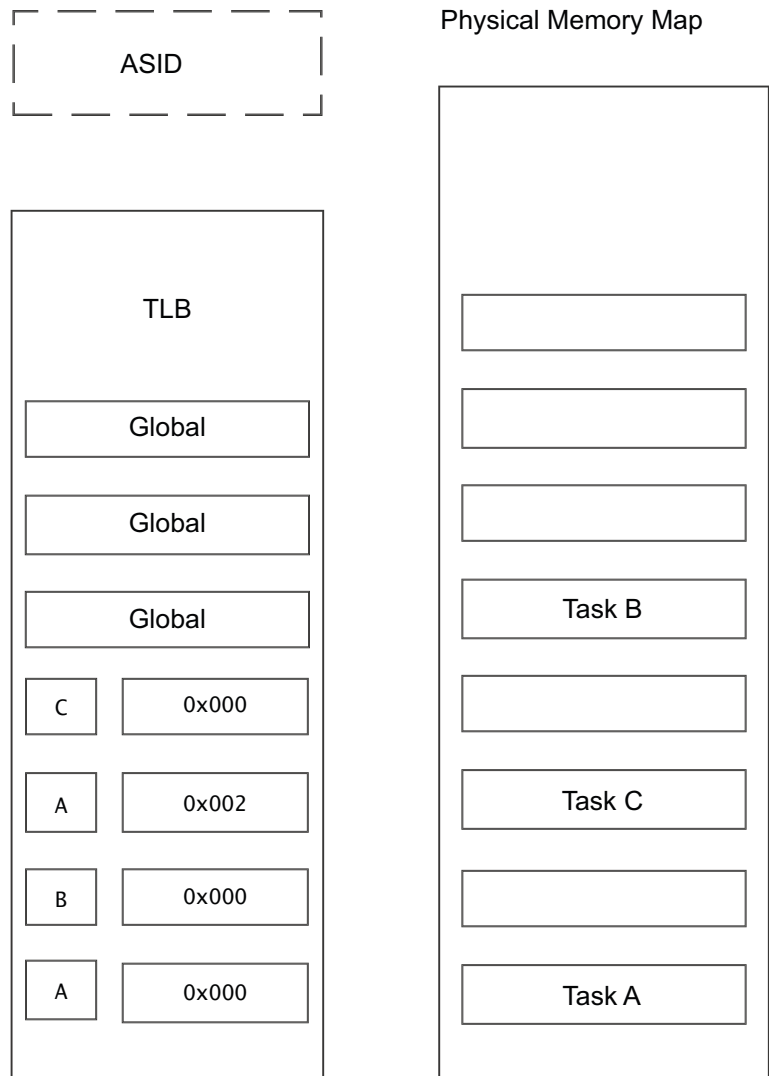


Figure 10-9 ASIDs in TLB mapping the same virtual address

10.8.2 Page Table Base Register 0 and 1

A further potential difficulty associated with managing multiple applications with their individual page tables is that there may need to be multiple copies of the L1 page table, one for each application. Each of these will be 16KB in size. Most of the entries will be identical in each of the tables, as typically only one region of memory will be task-specific, with the kernel space being unchanged in each case. Furthermore, if there is a need to modify a global page table entry, the change will be needed in each of the tables.

To help reduce the effect of these problems, a second page table base register can be used. CP15 contains two *Translation Table Base Registers*, TTBR0 and TTBR1. A control register (the TTB Control Register) is used to program a value in the range 0 to 7. This value (denoted by N) tells the MMU how many of the upper bits of the virtual address it should check to determine which of the two TTB registers to use.

When N is 0 (the default), all virtual addresses are mapped using TTBR0. With N in the range 1-7, the hardware looks at the most significant bits of the virtual address. If the N most significant bits are all zero, TTBR0 is used, otherwise TTBR1 is used.

For example, if N is set to 7, any address in the bottom 32MB of memory will use TTBR0 and the rest of memory will use TTBR1. As a result, the application-specific page table pointed to by TTBR0 will contain only 32 entries (128 bytes). The global mappings are in the table pointed to by TTBR1 and only one table needs to be maintained.

When these features are used, a context switch will typically require the operating system to change the TTBR0 and ASID values, using CP15 instructions. However, as these are two separate, non-atomic operations, some care is needed to avoid problems associated with speculative accesses occurring using the new value of one register together with the older value of the other. OS programmers making use of these features should become familiar with the sequences recommended for this purpose in the *ARM Architecture Reference Manual*.

10.8.3 The Fast Context Switch Extension

The *Fast Context Switch Extension* (FCSE) was added to the ARMv4 architecture but has been deprecated since. It enabled multiple independent tasks to run in a fixed, overlapping area at the bottom of the virtual memory space without the need to clean the cache or TLB on a context switch. It did this by modifying virtual addresses by substituting a process ID value into the top seven bits of the virtual address (but only if that address lay within the bottom 32MB of memory). Some ARM documentation distinguishes *Modified Virtual Addresses* (MVA) from *Virtual Addresses* (VA). This distinction is useful only when the FCSE is used. Since it is deprecated, we will not discuss the FCSE any further.

10.9 Large Physical Address Extensions

The Cortex-A15 and Cortex-A7 processors implement the ARMv7-A architecture *Large Physical Address Extension* (LPAE), which introduced a number of new features:

- Address mapping for a 40-bit address space, giving a range of accessible physical addresses from 4GB to 1024GB (a terabyte), with a granularity of 4KB. An additional page table entry format, called the *long-descriptor format*, was added. The existing *Virtual Memory Systems Architecture* (VMSA-v7) short-descriptor format page tables are still supported.
- Support is available for hierarchical permissions.
- A contiguous page hint bit for page table entries. This is set to show that the entry is one of 16 contiguous entries which point to a contiguous output address range. If set, the TLB need cache only one translation for the group of 16 pages.
- An additional access permission setting – *Privileged eXecute Never* (PXN). This marks a page as containing code that can be executed only in a non-privileged (user) mode. This setting is also added to the legacy descriptor format. There is also a privileged execute setting (PX) which means that code can be executed only in privileged mode.
- The ASID is stored in the TTBR. This allows atomic ASID changes, which reduce the overhead of context switches.
- Simplified fault status encoding.

The Cortex-A15 and Cortex-A7 processor MMUs also support virtualization. This provides a second stage of address translation when running virtual machines. The first stage of this translation produces an *Intermediate Physical Address* (IPA) and the second stage then produces the *Physical Address* (PA). TLB entries may also have an associated *Virtual Machine ID* (VMID), in addition to an ASID. This will be covered in more detail in [Chapter 27 Virtualization](#). It is possible to disable the stage 2 MMU and have a flat mapping from IPA to PA.

There are certain similarities between the short-descriptor format and the long-descriptor format:

- Both have an input 32-bit virtual address from the processor.
- Both use TTBR0/TTBR1 (See [Table 6-5 on page 6-20](#)) to point to level 1 translation tables (with TTBCR selecting which part of the address space is covered by each).
- Memory attributes are the same between the two formats, but the encoding is different.
- Long descriptors also provides a way to indicate “Outer-shareable” or “Inner-shareable” as opposed to just “Shareable”.

However, there are also some significant differences between the two:

- Long-descriptors have 64-bit sized translation table entries, while short-descriptor format uses 32-bit entries.
- Long-descriptors can have up to three levels of translation tables (only two were supported in short-descriptor format).
- Long-descriptors support up to 40-bit physical addresses and blocks (of 1GB or 2MB) and 4KB pages (short-descriptors support 16MB or 1MB blocks and 64KB or 4KB pages).
- Cache policies are given by the memory attributes index field of a long-descriptor (short-descriptor format uses TEX [2:0], C & B bits).

- Device memory is always considered shareable with long descriptor format.

Figure 10-10 shows the format of Level 1 and Level 2 descriptors.

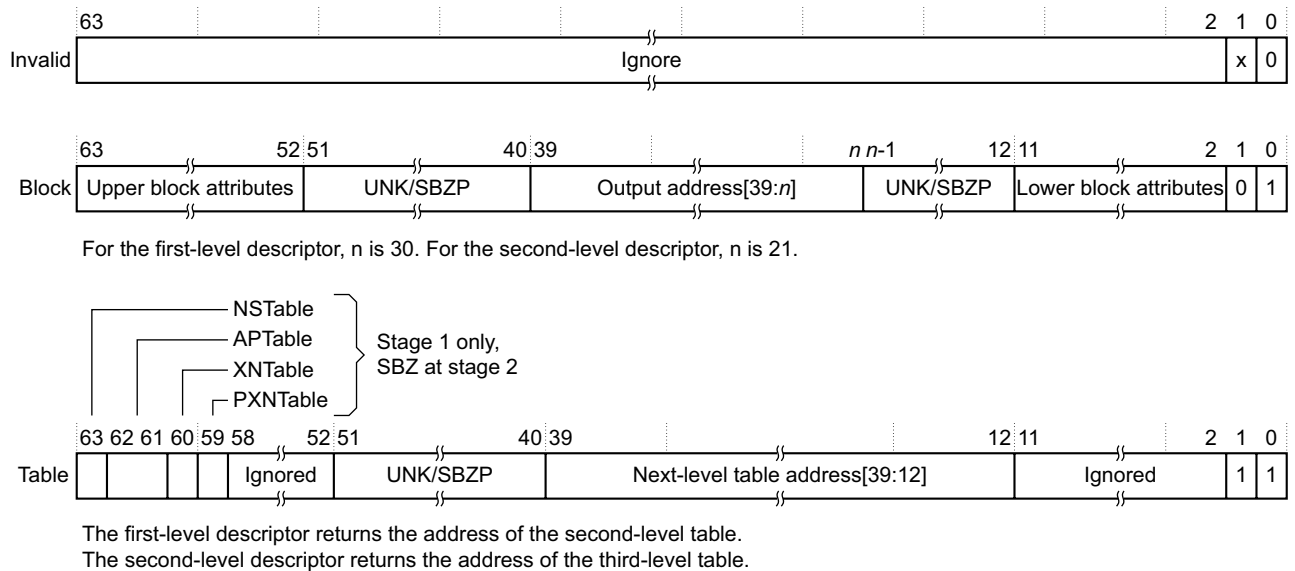


Figure 10-10 Format of Level 1 and Level 2 descriptors

where UNK is ‘Unknown’ and SBZP is ‘Should Be Preserved’


The TTBR points to a four entry L1 translation table, with bits [31:30] of the virtual address being used to index this table. This first level can be used for 1GB blocks. If only 30-bit virtual addresses are used in the system, then this L1 translation table may be omitted. This is configured by a TTBCR bit and in this case, the TTBR then points directly to the first entry in the L2 translation table.

The L2 translation table contains 512 entries, each of which might be for a 2MB block or point to a third level (L3) translation table. These L3 translation tables contain 2048 entries, with bits [20:12] of the virtual address used as an index. Each entry points to a 4KB page.

The page table entries have two possible formats. The first, a table entry, is used to provide the 28-bit address of a lower level translation table (given in bits [39:12] of the descriptor). Bits [1:0] are set to “11” to mark this as a table entry. The attribute bits (bits [63:59]) provide information such as access permissions, memory type, cache policies, and security status.

The other format, marked with bits [1:0]= “01” is used to give the physical address for the memory access. First level block descriptors point to a 1 GB block of memory, whose base address is given in descriptor bits [39:30]. Bits [39:21] are used to point to a 2MB block of memory, in a second level block descriptor, with bits [39:12] used to specify a 4KB page in a third level page descriptor. The lower attribute bits are used to define access permissions (for example shareable, global or non-global) similar to the VMSAv7 short-descriptor format.

Figure 10-11 on page 10-20 shows the format of Level 3 descriptors.

**Figure 10-11 Format of Level 3 descriptors**


Remember that long-descriptor page tables may output a 40-bit physical address. The physical base address given by this translation will be of the form[39:0], with bit 39 of the descriptor producing bit [39] of the physical address. It is not necessary to map all of the physical address space, but all of the 4GB virtual address space has to be taken care of. There must be a descriptor for each virtual address – although it might simply be a *fault* descriptor.

Chapter 11

Memory Ordering

Older implementations of the ARM architecture execute all instructions in program order and each instruction is completely executed before the next instruction is started.

Newer micro-processors employ a number of optimizations which relate to the way memory accesses are performed. As we have seen, the speed of execution of instructions by the processor is significantly higher than the speed of external memory. Caches and write buffers are used to partially hide the latency associated with this difference in speed. One potential effect of this is to re-order memory accesses. The order in which load and store instructions are executed by the processor will not necessarily be the same as the order in which the accesses are seen by external devices.

**Figure 11-1 Memory ordering example**

In [Figure 11-1](#), we have three instructions listed in program order. The first instruction performs a write to external memory which in this example, misses in the cache (Access 1). It is followed in program order by two reads, one which misses in the cache (Access 2) and one which hits in the cache (Access 3). Both of the read accesses could complete before the write buffer completes the write associated with Access 1. *Hit-under-miss* behaviors in the cache mean that a load which hits in the cache (like Access 3) can complete before a load earlier in the program which missed in the cache (like Access 2).

It is still possible to preserve the illusion that the hardware executes instructions in the order written by the programmer. There are generally only a few cases where the programmer has to worry about such effects. For example, when modifying CP15 registers or when copying or otherwise changing code in memory, it may be necessary for the programmer to explicitly make the processor wait for such operations to complete.

For very high performance processors, which support speculative data accesses, multi-issuing of instructions, cache coherency protocols and out-of-order execution in order to make further performance gains, there are even greater possibilities for re-ordering. In general, the effects of this re-ordering are invisible to the programmer, in a single processor system. The processor hardware takes care of many possible hazards for us. It will ensure that data dependencies are respected and ensure the correct value is returned by a read, allowing for potential modifications caused by earlier writes.

However, in cases where we have multiple processors which communicate through Shareable memory (or share data in other ways), memory ordering considerations become more important. In general, we are most likely to care about exact memory ordering at points where multiple execution threads must be synchronized.

Processors which conform to the ARM v7-A architecture employ a weakly-ordered model of memory, this means that the order of memory accesses is not expected to be the same as the program order for load and store operations. The model can reorder memory read operations (from LDR, LDM and LDD instructions) with respect to each other, to store operations, and certain other instructions. Reads and writes to Normal memory can be re-ordered by hardware, with such re-ordering being subject only to data dependencies and explicit memory barrier instructions. In cases where we need stronger ordering rules to be observed, we must communicate this to the processor through the memory type attribute of the page table entry

which describes that memory. Enforcing ordering rules on the processor limits the possible hardware optimizations and therefore reduces performance and increases power consumption. The programmer therefore needs to understand when to apply such ordering constraints.

In this chapter, we will consider the memory ordering model of the ARM architecture. We will look at the different types of memory (and other memory attributes) which are assigned to pages using page table entries. We will look at the barrier instructions in ARM assembly language (and accessible through C compiler intrinsics). Finally we will look at the support for coherency between SMP clusters and the concept of coherency domains, as well as considering some common cases where memory ordering problems can be encountered.

11.1 ARM memory ordering model

As we have said, Cortex-A series processors employ a weakly ordered memory model. However, within this model specific regions of memory can be marked as Strongly-ordered. In this case memory transactions are guaranteed to occur in the order they are issued,

Three mutually exclusive memory types are defined. All regions of memory are configured as one of these three types:

- Strongly-ordered.
- Device.
- Normal.

In addition, for Normal and Device memory, it is possible to specify whether the memory is Shareable (accessed by other agents) or not. For Normal memory, inner and outer cacheable properties can be specified.

11.1.1 Strongly-ordered and Device memory

Accesses to Strongly-ordered and Device memory have the same memory-ordering model. Access rules for this memory are as follows:

- The number and size of accesses will be preserved. Accesses will be atomic, and will not be interrupted part way through.
- Both read and write accesses can have side-effects on the system. Accesses are never cached. Speculative accesses will never be performed.
- Accesses cannot be unaligned.
- The order of accesses arriving at Device memory is guaranteed to correspond to the program order of instructions which access Strongly-ordered or Device memory. This guarantee applies only to accesses within the same peripheral or block of memory. The size of such a block is implementation defined, but has a minimum size of 1KB.
- In the ARMv7 architecture, the processor can re-order Normal memory accesses around Strongly-ordered or Device memory accesses.

The only difference between Device and Strongly-ordered memory is that:

- A write to Strongly-ordered memory can complete only when it reaches the peripheral or memory component accessed by the write.
- A write to Device memory is permitted to complete before it reaches the peripheral or memory component accessed by the write.

System peripherals will almost always be mapped as Device memory.

Regions of Device memory type can further be described using the Shareable attribute.

On some ARMv6 processors, the Shareable attribute of Device accesses is used to determine which memory interface will be used for the access, with memory accesses to areas marked as Device, Non-Shareable performed using a dedicated interface, the private peripheral port. This mechanism is not used on ARMv7 processors.

Note

These memory ordering rules provide guarantees only about explicit memory accesses (those caused by load and store instructions). The architecture does not provide similar guarantees about the ordering of instruction fetches or page table walks with respect to such explicit memory accesses.

11.1.2 Normal memory

Normal memory is used to describe most parts of the memory system. All ROM and RAM devices are considered to be Normal memory. All code to be executed by the processor must be in Normal memory. The architecture does not permit code to be in a region of memory which is marked as Device or Strongly-ordered.

The properties of Normal memory are as follows:

- The processor can repeat read and some write accesses.
- The processor can pre-fetch or speculatively access additional memory locations, with no side-effects (if permitted by MMU access permission settings). The processor will not perform speculative writes, however.
- Unaligned accesses can be performed.
- Multiple accesses can be merged by processor hardware into a smaller number of accesses of a larger size. Multiple byte writes could be merged into a single double-word write, for example.

Regions of Normal memory must also have cacheability attributes described (see [Chapter 9](#) for details of the supported cache policies). The ARM architecture supports cacheability attributes for Normal memory for two levels of cache, the inner and outer cache. The mapping between these levels of cache and the implemented physical levels of cache is implementation defined.

Inner refers to the innermost caches, and always includes the processor level 1 cache. An implementation might not have any outer cache, or it can apply the outer cacheability attribute to an L2 or L3 cache. For example, in a system containing a Cortex-A9 processor and the L2C-310 level2 cache controller, the L2C-310 is considered to be the outer cache. The Cortex-A8 L2 cache can be configured to use either inner or outer cache policy.

Normal memory must also be identified either as Shareable or Non-Shareable. A region of Normal memory with the Non-Shareable attribute is one which is used only by this processor. There is no requirement for the processor to make accesses to this location coherent with other processors. If other processors do share this memory, any coherency issues must be handled in software. For example, this can be done by having individual processors perform cache maintenance and barrier operations.

The Outer Shareable attribute enables the definition of systems containing multiple levels of coherency control. For example, an Inner Shareable domain could consist of a multiple core Cortex-A15 processor and a multiple core Cortex-A7 processor. Within a cluster, the data caches of the processors are coherent for all data accesses which have the inner Shareable attribute. The Outer Shareable domain, meanwhile, might consist of this cluster of processors and a graphics processor with multiple cores. An Outer Shareable domain can consist of multiple Inner Shareable domains, but an Inner Shareable domain can only be part of one Outer Shareable domain.

Shareability

A region with the Shareable attribute set is one which can be accessed by other processors in the system. Accesses to memory in this region by other processors within the same shareability domain are coherent. This means that the programmer does not need to take care of the effects of data or caches. Without the Shareable attribute, in situations where cache coherency is not maintained between processors for a region of shared memory, the programmer would have to explicitly manage coherency themselves.

The ARMv7 architecture enables the programmer to specify Shareable memory as *inner Shareable* or *outer Shareable* (this latter case means that the location is both inner and outer Shareable).

11.2 Memory barriers

A memory barrier is an instruction which requires the processor to apply an ordering constraint between memory operations which occur before and after the memory barrier instruction in the program. Such instructions may also be known as *memory fences* in other architectures.

The term memory barrier can also be used to refer to a compiler mechanism which prevents the compiler from scheduling data access instructions across the barrier when performing optimizations. For example in GCC, we can use the inline assembler memory *clobber*, to indicate that the instruction changes memory and therefore the optimizer cannot re-order memory accesses across the barrier. The syntax is as follows:

```
asm volatile("" ::: "memory");
```

ARM RVCT includes a similar intrinsic, called `__schedule_barrier()`.

Here, however, we are looking at hardware memory barriers, provided through dedicated ARM assembly language instructions. As we have seen, processor optimizations such as caches, write buffers and out-of-order execution can result in memory operations occurring in an order different from that specified in the executing code. Normally, this re-ordering is invisible to the programmer and application developers do not normally need to worry about memory barriers. However, there are cases where we may need to take care of such ordering issues, for example in device drivers or when we have multiple observers of the data which need to be synchronized.

The ARM architecture specifies memory barrier instructions, which allow the programmer to force the processor to wait for memory accesses to complete. These instructions are available in both ARM and Thumb code, in both user and privileged modes. In older versions of the architecture, these were performed using CP15 operations in ARM code only. Use of these is now deprecated, although preserved for compatibility.

Let's start by looking at the practical effect of these instructions in a uni-processor system. Note that this description is a simplified version of that given in the *ARM Architecture Reference Manual*, what is written here is intended to introduce the usage of these instructions. The term explicit access is used to describe a data access resulting from a load or store instruction in the program. It does not include instruction fetches.

Data Synchronization Barrier (DSB)

This instruction forces the processor to wait for all pending explicit data accesses to complete before any further instructions stages can be executed. There is no effect on pre-fetching of instructions.

Data Memory Barrier (DMB)

This instruction ensures that all memory accesses in program order before the barrier are observed in the system before any explicit memory accesses that appear in program order after the barrier. It does not affect the ordering of any other instructions executing on the processor, or of instruction fetches.

Instruction Synchronization Barrier (ISB)

This flushes the pipeline and prefetch buffer(s) in the processor, so that all instructions following the ISB are fetched from cache or memory, after the instruction has completed. This ensures that the effects of context altering operations (for example, CP15 or ASID changes or TLB or branch predictor operations), executed before the ISB instruction are visible to any instructions fetched after the ISB. This does not in itself cause synchronization between data and instruction caches, but is required as a part of such an operation.

Several options can be specified with the DMB or DSB instructions, to provide the type of access and the shareability domain it should apply to, as follows:

SY	This is the default and means that the barrier applies to the full system, including all processors and peripherals.
ST	A barrier which waits only for stores to complete.
ISH	A barrier which applies only to the Inner Shareable domain.
ISHST	A barrier which combines the above two (that is, it only stores to the Inner Shareable).
NSH	A barrier only to the Point of Unification (PoU). (See Point of coherency and unification on page 9-17).
NSHST	A barrier that waits only for stores to complete and only out to the point of unification.
OSH	Barrier operation only to the Outer Shareable domain.
OSHST	Barrier operation that waits only for stores to complete, and only to the Outer Shareable domain.

To make sense of this, we need to use a more general definition of the DMB and DSB operations in a multi-processor system. The use of the word processor (or agent) in the following text does not necessarily mean a processor and also could refer to a DSP, DMA controller, hardware accelerator or any other block that accesses shared memory.

The DMB instruction has the effect of enforcing memory access ordering within a shareability domain. All processors within the shareability domain are guaranteed to observe all explicit memory accesses before the DMB instruction, before they observe any of the explicit memory accesses after it.

The DSB instruction has the same effect as the DMB, but in addition to this, it also synchronizes the memory accesses with the full instruction stream, not just other memory accesses. This means that when a DSB is issued, execution will stall until all outstanding explicit memory accesses have completed. When all outstanding reads have completed and the write buffer is drained, execution resumes as normal.

It may be easier to appreciate the effect of the barriers by considering an example. Consider the case of a Cortex-A9 MPCore containing four processors. These processors operate as an SMP cluster and form a single Inner Shareable domain. When a single processor within the cluster executes a DMB instruction, that processor will ensure that all data memory accesses in program order before the barrier complete, before any explicit memory accesses that appear in program-order after the barrier. This way, it can be guaranteed that all processors within the cluster will see the accesses on either side of that barrier in the same order as the processor that performs them. If the DMB ISH variant is used, the same is not guaranteed for external observers such as DMA controllers or DSPs.

11.2.1 Memory barrier use example

Consider the case where we have two processors A and B and two addresses in Normal memory (Addr1 and Addr2) held in processor registers. Each processor executes two instructions as shown in [Example 11-1](#):

Example 11-1 Code example showing memory ordering issues

Processor A:

```
STR R0, [Addr1]
LDR R1, [Addr2]
```

Processor B:

```
STR R2, [Addr2]
LDR R3, [Addr1]
```

Here, there is no ordering requirement and we can make no statement about the order in which any of the transactions occur. The addresses Addr1 and Addr2 are independent and there is no requirement on either processor to execute the load and store in the order written in the program, or to care about the activity of the other processor.

There are therefore four possible legal outcomes of this piece of code, with four different sets of values from memory ending up in processor A, register R1 and processor B, register R3:

- A gets the old value, B gets the old value.
- A gets the old value, B gets the new value.
- A gets the new value, B gets the old value.
- A gets the new value, B gets the new value.

If we were to involve a third processor, C, we should also note that there is no requirement that it would observe either of the stores in the same order as either of the other processors. It is perfectly permissible for both A and B to see an old value in Addr1 and Addr2, but for C to see the new values.

So, let's consider the case where the code on B looks for a flag being set by A and then reads memory – for example if we are passing a message from A to B. We might now have code similar to that shown in [Example 11-2](#):

Example 11-2 Possible ordering hazard with postbox

Processor A:

```
STR R0, [Msg] @ write some new data into postbox
STR R1, [Flag] @ new data is ready to read
```

Processor B:

```
Poll_loop:
    LDR R1, [Flag]
    CMP R1, #0 @ is the flag set yet?
    BEQ Poll_loop
    LDR R0, [Msg] @ read new data.
```

Again, this might not behave in the way that is expected. There is no reason why processor B is not allowed to speculatively perform the read from [Msg] before the read from [Flag]. This is normal, weakly ordered memory and the processor has no knowledge about a possible dependency between the two. The programmer must explicitly enforce the dependency by inserting a memory barrier. In this example, we actually need *two* memory barriers. Processor A needs a DMB between the two store operations, to make sure they happen in the programmer specified order. Processor B needs a DMB before the LDR R0, [Msg] to be sure that the message is not read until the flag is set.

11.2.2 Avoiding deadlocks with a barrier

Here is another case which can cause a deadlock if barrier instructions are not used. Consider a situation where a processor writes to an address and then polls for an acknowledge value to be applied by a peripheral.

[Example 11-3](#) shows the type of code which can cause a problem.

Example 11-3 Deadlock

```
STR R0, [Addr] @ write a command to a peripheral register
DSB
Poll_loop:
    LDR R1, [Flag]
    CMP R1, #0 @ wait for an acknowledge/state flag to be set
    BEQ Poll_loop
```

The ARMv7 architecture without multiprocessing extensions does not strictly require the processor's store to [Addr] to ever complete (it could be sitting in a write buffer while the memory system is kept busy reading the flag), so both processors could potentially deadlock, each waiting for the other. Inserting a DSB after the STR of the processor forces its store to be observed before it will read from Flag.

Processors which implement the multiprocessing extensions are required to complete accesses in a finite time (that is, their write buffers must drain) and so the barrier instruction is not required.

11.2.3 WFE and WFI Interaction with barriers

The WFE (Wait For Event) and WFI (Wait For Interrupt) instructions, described further in [Chapter 25 Power Management](#), allow us to stop execution and enter a low-power state. If we need to ensure that all memory accesses prior to executing WFI or WFE have been completed (and made visible to other processors), we must insert a DSB instruction.

A further consideration relates to usage of WFE and SEV (Send Event) in an MP system. These instructions allow us to reduce the power consumption associated with a lock acquire loop (a spinlock). A processor which is attempting to acquire a mutex may find that some other processor already has the lock. Instead of having the processor repeatedly poll the lock, we can suspend execution and enter a low-power state, using the WFE instruction. We wake either when an interrupt or other asynchronous exception is recognized, or another processor sends an event (with the SEV instruction). The processor that had the lock will use the SEV instruction to wake-up other processors in the WFE state after the lock has been released. For the purposes of memory barrier instructions, the event signal is not treated as an explicit memory access. We therefore need to take care that the update to memory which releases the lock is actually visible to other processors before the SEV instruction is executed. This requires the use of a DSB. DMB is not

sufficient as it only affects the ordering of memory accesses without synchronizing them to a particular instruction, whereas DSB will prevent the SEV from executing until all preceding memory accesses have been seen by other processors.

11.2.4 Linux use of barriers

In this section, we will look at the implications of barriers in multi-core systems and take a much more detailed look at SMP operation.

Barriers are needed to enforce ordering of memory operations. Most programmers will not need to understand, or explicitly use memory barriers. This is because they are already included within kernel locking and scheduling primitives. Nevertheless, writers of device drivers or those seeking an understanding of kernel operation may find a detailed description useful.

Both the compiler and processor micro-architecture optimizations permit the order of instructions and associated memory operations to be changed. Sometimes, however, we wish to enforce a specified order of execution of memory operations. For example, we can write to a memory mapped peripheral register. This write can have side effects elsewhere in the system. Memory operations which are in before or after this operation in our program can appear as if they can be re-ordered, as they operate on different locations. In some cases, however, we wish to ensure that all operations complete before this peripheral write completes. Or, we may want to make sure that the peripheral write completes before any further memory operations are started. Linux provides some functions to do this, as follows:

- We instruct the compiler that re-ordering is not permitted for a particular memory operation. This is done with the `barrier()` function call. This controls only the compiler code generation and optimization and has no effect on hardware re-ordering.
- We call a memory barrier function which maps to ARM processor instructions that perform the memory barrier operations. These enforce a particular hardware ordering. The available barriers are as follows (in a Linux kernel compiled with Cortex-A SMP support):
 - the read memory barrier `rmb()` function ensures that any read that appears before the barrier is completed before the execution of any read that appears after the barrier
 - the write memory barrier `wmb()` function ensures that any write that appears before the barrier is completed before the execution of any write that appears after the barrier
 - the memory barrier `mb()` function ensures that any memory access that appears before the barrier is completed before the execution of any memory access that appears after the barrier.
- There are corresponding SMP versions of these barriers, called `smp_mb()`, `smp_rmb()` and `smp_wmb()`. These are used to enforce ordering on Normal cacheable memory, between cores inside the same SMP processor. For example, each processor inside a Cortex-A15 MPCore. They can be used with devices and they work even for normal non-cacheable memory. When the kernel is compiled without `CONFIG_SMP`, each invocation of these are expanded to `barrier()` statements.

All of the locking primitives provided by Linux include any required barrier.

For these memory barriers, it is almost always the case that a pair of barriers is required. For further information, see <http://www.kernel.org/doc/Documentation/memory-barriers.txt>.

11.3 Cache coherency implications

The caches are largely invisible to the application programmer. However they can become visible when memory locations are changed elsewhere in the system or when memory updates made from the application code must be made visible to other parts of the system.

A system containing an external DMA device and a processor provides a simple example of possible problems. There are two situations in which a breakdown of coherency can occur. If the DMA reads data from main memory while newer data is held in the processor's cache, the DMA will read the old data. Similarly, if a DMA writes data to main memory and stale data is present in the processor's cache, the processor can continue to use the old data.

Therefore dirty data which is in the ARM processor data cache must be explicitly cleaned before the DMA starts. Similarly, if the DMA is copying data to be read by the ARM processor, it must be certain that the ARM processor data cache does not contain stale data (the cache will not be updated by the DMA writing memory and this may need the ARM processor to clean or invalidate the affected memory areas from the cache(s) before starting the DMA). As all ARMv7-A processors can do speculative memory accesses, it will also be necessary to invalidate after using the DMA.

11.3.1 Issues with copying code

Boot code, kernel code or JIT compilers can copy programs from one location to another, or modify code in memory. There is no hardware mechanism to maintain coherency between instruction and data caches. The programmer must invalidate stale code from the instruction cache by invalidating the affected areas, and ensure that the code written has actually reached the main memory. Specific code sequences including instruction barriers are needed if the processor is then intended to branch to the modified code.

11.3.2 Compiler re-ordering optimizations

It is important to understand that memory barrier instructions apply only to hardware re-ordering of memory accesses. Inserting a hardware memory barrier instruction may not have any direct effect on compiler re-ordering of operations. The `volatile` type qualifier in C tells the compiler that the variable can be changed by something other than the currently executing code that is accessing it. This is often used for C language access to memory mapped I/O, allowing such devices to be safely accessed through a pointer to a `volatile` variable. The C standard does not provide rules relating to the use of `volatile` in systems with multiple processors. So, although we can be sure that `volatile` loads and stores will happen in program specified order with respect to each other, there are no such guarantees about re-ordering of accesses relative to non-volatile loads or stores. This means that `volatile` does not provide a shortcut to implement mutexes.

Chapter 12

Exception Handling

In this chapter, we look at how ARM processors respond to exceptions. An exception is any condition that needs to halt normal execution and instead run software (known as an exception handler) associated with each exception type. Other architectures may refer to what ARM calls exceptions as traps or interrupts, however, in the ARM architecture, these terms are reserved for specific types of exceptions, as described in [Types of exception on page 12-3](#). We begin by introducing exceptions and see how the ARM processor handles each of the different types and what they are used for. We then look in more detail at interrupts and describe mechanisms of interrupt handling on ARM processors and standard interrupt handling schemes.

All microprocessors need to respond to external asynchronous events, such as a button being pressed, or a clock reaching a certain value. Normally, there is specialized hardware which activates input lines to the processor. This causes the microprocessor to temporarily stop the current program sequence and execute a special handler routine. The speed with which a processor can respond to such events may be a critical issue in system design. Indeed in many embedded systems, there is no main program as such – all of the functions of the system are handled by code which runs from interrupts, and assigning priorities to these is a key area of design. Rather than have the processor constantly polling the flags from different parts of the system to see if there is something to be done, we instead allow the system to tell the processor that something needs to happen, by generating an interrupt. Complex systems have very many interrupt sources with different levels of priority and requirements for nested interrupt handling (where a higher priority interrupt can interrupt a lower priority one).

In normal program execution, the program counter increments through the address space, with branches in the program modifying the flow of execution (for example, for function calls, loops, and conditional code). When an exception occurs, this sequence is interrupted.

In addition to responding to external interrupts, there are a number of other things which can cause the processor to take an exception, both external (reset, external aborts from the memory system) and internal (MMU generated aborts or OS calls using the SVC instruction). Dealing with exceptions causes the ARM processor to switch between modes and copy some registers into others. Readers new to the ARM architecture may wish to refresh their understanding of the modes and registers described in [Chapter 4](#), before continuing with this chapter.

12.1 Types of exception

As we have seen from [Table 4-1 on page 4-3](#), the A classes and R classes of architecture support seven processor modes, six privileged modes called *FIQ*, *IRQ*, *Supervisor*, *Abort*, *Undefined* and *System*, and the non-privileged *User* mode. The *Hyp* mode and *Monitor* mode registers can be added to the list if the Virtualization Extensions and Security Extensions are implemented. The current mode can change under software control or when processing an exception.

However, the unprivileged *User* mode can switch to another mode only by generating an exception.

When an exception occurs, the processor saves the current status and the return address, enters a specific mode and possibly disables hardware interrupts. Execution resumes from a fixed memory address called an exception vector. This happens automatically and is not under direct control of the programmer.

The following types of exception exist:

Interrupts There are two types of interrupts provided on ARMv7-A processors, called *IRQ* and *FIQ*.

FIQ is higher priority than *IRQ*. *FIQ* also has some potential speed advantages owing to its position in the vector table and the higher number of banked registers available in *FIQ* mode. This potentially saves processor clock cycles on pushing registers to the stack within the handler. Both of these kinds of exception are typically associated with input pins on the processor – external hardware asserts an interrupt request line and the corresponding exception type is raised when the current instruction finishes executing, assuming that the interrupt is not disabled.

Aborts Aborts can be generated either on failed instruction fetches (prefetch aborts) or failed data accesses (data aborts). They can come from the external memory system giving an error response on a memory access (indicating perhaps that the specified address does not correspond to real memory in the system). Alternatively, the abort can be generated by the Memory Management Unit (MMU) of the processor. An operating system can use MMU aborts to dynamically allocate memory to applications.

An instruction can be marked within the pipeline as aborted, when it is fetched. The prefetch abort exception is taken only if the processor then actually tries to execute it. The exception takes place before the instruction actually executes. If the pipeline is flushed before the aborted instruction reaches the execute stage of the pipeline, the abort exception will not occur. A data abort exception happens when a load or store instruction executes and is considered to happen after the data read or write has been attempted.

An abort is described as *synchronous* if it is generated as a result of execution or attempted execution of the instruction stream, and where the return address will provide details of the instruction that caused it. An *asynchronous* abort is not generated by executing instructions, while the return address may not always provide details of what caused the abort.

The ARMv7 architecture distinguishes between precise and imprecise asynchronous aborts. Aborts generated by the MMU are always synchronous. The architecture does not require particular classes of externally aborted accesses to be synchronous.

For example, on a particular processor implementation, it may be the case that an external abort reported on a page table walk is treated as precise, but this is not required to be the case for all processors. For precise asynchronous aborts, the abort handler can be certain which instruction caused the abort and that no further

instructions were executed after that instruction. This is in contrast to an imprecise asynchronous abort, which results when the external memory system reports an error on an unidentifiable access.

In this case, the abort handler cannot determine which instruction caused the problem (or further instructions may have executed after the one which generated the abort). For example, if a buffered write receives an error response from the external memory system, further instructions will have been executed after the store. This means that it will be impossible for the abort handler to fix the problem and return to the application. All it can do is to kill the application which caused the problem. Device probing therefore needs special handling, as externally reported aborts on reads to non-existent areas will generate imprecise synchronous aborts even when such memory is marked as Strongly-ordered, or Device.

Detection of asynchronous aborts is controlled by the CPSR A bit. If the A bit is set, asynchronous aborts from the external memory system will be recognized by the processor, but no abort exception will be generated immediately. Instead, the processor keeps the abort pending until the A bit is cleared and takes an exception at that time. Kernel code will typically ensure (through the use of a barrier instruction) that pending asynchronous aborts are recognized against the correct application. If a thread has to be killed due to an imprecise abort, it needs to be the correct one!

Reset All processors have a reset input and will take the reset exception immediately after they have been reset. It is the highest priority exception and cannot be masked.

Exceptional instructions

There are two classes of instruction which can cause exceptions on an ARM processor. The first is the Supervisor Call (SVC), previously known as Software Interrupt (SWI). This is typically used to provide a mechanism by which User mode programs can pass control to privileged, kernel code in the OS to perform OS-level tasks. The second is an undefined instruction. The architecture defines certain bit-patterns as corresponding to undefined opcodes. Trying to execute one of these causes an Undefined Instruction exception to be taken. In addition, executing coprocessor instructions for which there is no corresponding coprocessor hardware will also cause this trap to happen. Some instructions can be executed only in a privileged mode and executing these from User mode will cause an undefined instruction exception.

When an exception occurs, code execution passes to an area of memory called the vector table. Within the table just one word is allocated to each of the various exception types and this will usually contain a branch instruction to the actual exception handler. This behavior is different from the majority of other architectures, which usually store an execution address in the exception table, rather than an instruction.

You can write the exception handlers in either ARM or Thumb code. The CP15 SCTLR.TE bit is used to specify whether exception handlers will use ARM or Thumb. When handling exceptions, the prior mode, state, and registers of the processor must be preserved so that the program can be resumed after the exception has been handled.

12.2 Exception mode summary

[Table 12-1](#) lists the state of the interrupt disabling I and F bits of the CPSR upon entering an exception handler.

Table 12-1 Summary of exception behavior

Vector offset	Exception	Mode	Event	CPSR interrupt mask	Return address offset
0x0	Reset	Supervisor	Reset input asserted	F = 1 I = 1	Not applicable
0x4	UNDEFINED instruction	Undefined	Executing undefined instruction	I = 1	0 (if emulating the instruction) 4 (if re-executing after, for example, enabling VFP)
0x8	Supervisor Call	Supervisor	SVC instruction	I = 1	0
0xC	Prefetch Abort	Abort	Instruction fetch from invalid address	I = 1	4
0x10	Data Abort	Abort	Data Read/Write to invalid address	I = 1	8 (if retry of the aborting instruction is wanted)
0x14	Hypervisor Call ^a	HYP	Hypervisor entry	-	0
0x18	Interrupt	IRQ	IRQ input asserted	I = 1	4
0x1C	Fast Interrupt	FIQ	FIQ input asserted	F = 1 I = 1	4

a. Hypervisor entry exception (described in [Chapter 27 Virtualization](#)) is available only in processors which support Virtualization Extensions and is unused in other processors.

12.2.1 Exception priorities

As some of the exception types can occur simultaneously, the processor assigns a fixed priority for each exception, as shown in [Table 12-1 on page 12-5](#). The Undefined instruction, prefetch abort and Supervisor Call exceptions are due to execution of an instruction (there are specific bit patterns for undefined and SVC opcodes) and so can never happen together. They therefore have the same priority.

Table 12-2 Exception priorities

Priority	Exception
Highest	1 Reset
	2 Data Abort
	3 FIQ
	4 IRQ
Lowest	6 Prefetch Abort
	7 Undefined instruction SVC, SMC, HVC

Note

The ARM architecture does not define when asynchronous exceptions are taken. Therefore the prioritization of asynchronous exceptions relative to other exceptions, both synchronous and asynchronous, is implementation defined.

It is important to distinguish between prioritization of exceptions, which happens when multiple exceptions are required at the same time, and the actual exception handler code. You will notice that [Table 12-1 on page 12-5](#) contains columns explaining how FIQ and IRQ are automatically disabled by some exceptions. (All exceptions disable IRQ, only FIQ and reset disable FIQ.) This is done by the processor automatically setting the CPSR I (IRQ) and F (FIQ) bits.

So, an FIQ exception can interrupt an abort handler or IRQ exception. In the case of a data abort and FIQ occurring simultaneously, the data abort (which has higher priority) is taken first. This lets the processor record the return address for the data abort. But as FIQ is not disabled by data abort, we then take the FIQ exception immediately. At the end of the FIQ we return back to the data abort handler.

More than one exception can potentially be generated at the same time, but some combinations are mutually exclusive. A prefetch abort marks an instruction as invalid and so cannot occur at the same time as an undefined instruction or SVC (and of course, an SVC instruction cannot also be an undefined instruction). These instructions cannot cause any memory access and therefore cannot cause a data abort. The architecture does not define when asynchronous exceptions, FIQ, IRQ or asynchronous aborts must be taken, but the fact that taking an IRQ or data abort exception does not disable FIQ exceptions means that FIQ execution will be prioritized over IRQ or asynchronous abort handling.

12.3 Entering an exception handler

When an exception occurs, the ARM processor automatically does the following things:

- Preserves the address of the next instruction, in the Link Register (LR) of the new mode.
- Copies CPSR to the SPSR, one of the banked registers specific to each (non-user) mode of operation.
- Modifies the CPSR mode bits to a mode associated with the exception type. The other CPSR mode bits are set to values determined by bits in the CP15 System Control Register. The T bit is set to the value given by the CP15 TE bit. The J bit is cleared and the E bit (Endianness) is set to the value of the EE (Exception Endianness) bit. This enables exceptions to always run in ARM or Thumb state and in little or big-endian, irrespective of the state the processor was in before the exception.
- Forces the PC to point to the relevant instruction from the exception vector table.

Once in the new mode the processor will access the view associated with that mode, as shown in [Figure 4-1 on page 4-4](#).

It will almost always be necessary for the exception handler software to save registers onto the stack immediately upon exception entry. (FIQ mode has more dedicated registers and so a simple handler may be able to be written in a way which needs no stack usage.)

A special assembly language instruction is provided to assist with saving the necessary registers, called SRS (Store Return State). This instruction pushes the LR and SPSR onto the stack of any mode; which stack should be used is specified by the instruction operand.

12.4 Exit from an exception handler

To return from an exception handler, two separate operations must take place atomically:

- Restore the CPSR from the saved SPSR.
- Set the PC to the return address offset, see [Table 12-1 on page 12-5](#).

In the ARM architecture this can be achieved either by using the RFE instruction or any flag-setting data processing operation (with the S suffix) with the PC as the destination register, such as SUBS PC, LR, #offset (note the S). The Return From Exception (RFE) instruction pops the link register and SPSR off the current mode stack.

12.5 Vector table

The first column in [Table 12-1 on page 12-5](#) gives the vector offset within the vector table associated with the particular type of exception. This is a table of instructions that the ARM processor jumps to when an exception is raised. These instructions are located in a specific place in memory. The normal vector base address is `0x00000000`, but most ARM processors allow the vector base address to be moved to `0xFFFF0000` (or `HIVECS`). All Cortex-A series processors permit this, and it is the default address selected by the Linux kernel. Processors that implement the Security Extensions can additionally set the vector base address, separately for Secure and Non-secure states, using the CP15 Vector base Address registers.

You will notice that there is a single word address associated with each exception type. Therefore, only a single instruction can be placed in the vector table for each exception (although, in theory, two 16-bit Thumb instructions could be used). FIQ is different. Therefore, the vector table entry almost always contains one of the various forms of branches.

B<label>

This performs a PC-relative branch. It is suitable for calling exception handler code which is close enough in memory that the 24-bit field provided in the branch instruction is large enough to encode the offset.

LDR PC, [PC, #offset]

This loads the PC from a memory location whose address is defined relative to the address of the exception instruction. This lets the exception handler be placed at any arbitrary address within the full 32-bit memory space (but takes some extra cycles relative to the simple branch above).

12.6 Return instruction

The Link Register (LR) is used to store the appropriate return address for the PC after the exception has been handled. Its value needs to be modified as shown in [Table 12-1 on page 12-5](#), depending upon the type of exception occurred. The *ARM Architecture Reference Manual* defines the LR values that are appropriate (the definition derives from the values that were convenient for early hardware implementations).

Chapter 13

Interrupt Handling

In this chapter, we will look at a range of methods by which interrupts are handled in ARM processors, and briefly look at the Generic Interrupt Controller (GIC) architecture.

Older versions of the ARM architecture allowed implementers considerable freedom in their design of an external interrupt controller, with no agreement over the number or types of interrupts or the software model to be used to interface to the interrupt controller block. The GIC architecture provides a much more tightly controlled specification, with a greater degree of consistency between interrupt controllers from different manufacturers. This enables interrupt handler code to be more portable.

13.1 External interrupt requests

As we discussed in [Types of exception on page 12-3](#), all ARM processors have two external interrupt requests, FIQ and IRQ. Both of these are level-sensitive active-LOW inputs. Individual implementations have interrupt controllers which accept interrupt requests from a wide variety of external sources and map them onto FIQ or IRQ, causing the processor to take an exception.

In general, an interrupt exception can be taken only when the appropriate CPSR disable bit (the F and I bits respectively) is clear.

The CPS assembly language instruction provides a simple mechanism to enable or disable the exceptions controlled by CPSR A, I and F bits (asynchronous abort, IRQ and FIQ respectively). CPS can be used additionally to change mode, as shown below.

```
CPS #<mode>
CPSIE <if>
CPSID <if>
```

where <mode> is the number of the mode to change to. If this option is omitted, no mode change occurs. The values of these modes are listed in [Table 4-1 on page 4-3](#).

IE or ID will enable or disable exceptions respectively. The exceptions to be enabled or disabled are specified using one or more of the letters A, I and F. Exceptions whose corresponding letters are omitted will not be modified.

In Cortex-A series processors, it is possible to configure the processor so that FIQs cannot be masked by software. This is known as Non-Maskable FIQ and is controlled by a hardware configuration input signal which is sampled when the processor is reset. They will still be masked automatically upon taking an FIQ exception.

13.1.1 Assigning interrupts

A system will always have an interrupt controller which accepts interrupt requests from multiple pieces of external hardware. This typically contains a number of registers enabling software running on the ARM processor to mask individual interrupt sources, to acknowledge interrupts from external devices and to determine which interrupt sources are currently active.

This interrupt controller can be a design specific to the system, or it can be an implementation of the ARM *Generic Interrupt Controller* (GIC) architecture, which is described in [Generic Interrupt Controller on page 13-4](#).

13.1.2 Simplistic interrupt handling

This represents the simplest kind of interrupt handler. An interrupt occurs and while processing that interrupt, further interrupts are disabled. We can only handle further interrupts at the completion of the first interrupt request and there is no scope for a higher priority or more urgent interrupt to be handled during this time. This is not generally suitable for complex embedded systems, but it is useful to examine before proceeding to a more realistic example.

The steps taken to handle an interrupt are as follows:

1. An IRQ exception is raised by external hardware. The processor performs several steps automatically. The contents of the PC in the current execution mode are stored in LR_IRQ. The CPSR register is copied to SPSR_IRQ. The bottom byte of the CPSR is updated to change to IRQ mode, and to disable IRQ which prevent further exceptions from occurring. The PC is set to IRQ entry in the vector table.

2. The instruction at the IRQ entry in the vector table (a branch to the interrupt handler) is executed.
3. The interrupt handler saves the context of the interrupted program (that is, it pushes onto the stack any registers which will be corrupted by the handler).
4. The interrupt handler determines which interrupt source needs to be processed and calls the appropriate device driver.
5. Finally, the SPSR_IRQ is copied back into the CPSR, which switches the system back to the previous execution mode. At the same time, the PC is restored from the LR_IRQ.

13.1.3 Nested interrupt handling

In a nested handler, we re-enable interrupts before the handler has fully served the current interrupt. This allows us to prioritize interrupts and make significant improvements to the latency of high priority events at the cost of additional complexity.

A reentrant interrupt handler must save the IRQ state and then switch processor modes, and save the state for the new processor mode, before it branches to a nested subroutine or C function with interrupts enabled. This is because a fresh interrupt could occur at any time, which would cause the processor to store the return address of the new interrupt in LR_IRQ, overwriting the original interrupt. When the original interrupt attempts to return to the main program, it will cause the system to fail. The nested handler must change into an alternative kernel mode before re-enabling interrupts in order to prevent this.

———— Note ————

A computer program is *reentrant* if it can be interrupted in the middle of its execution and then be called again before the previous version has completed.

A reentrant interrupt handler must therefore take the following steps after an IRQ exception is raised and control is transferred to the interrupt handler in the way previously described.

1. The interrupt handler saves the context of the interrupted program (that is, it pushes onto the alternative kernel mode stack any registers which will be corrupted by the handler, including the return address and SPSR_IRQ).
2. It determines which interrupt source needs to be processed and clears the source in the external hardware (preventing it from immediately triggering another interrupt).
3. The interrupt handler changes the processor SVC mode, leaving the CPSR I bit set (interrupts are still disabled).
4. The interrupt handler saves the exception return address on the stack (a stack for the new mode, located in kernel memory) and re-enables interrupts.
5. It calls the appropriate device driver.
6. Upon completion, the interrupt handler disables IRQ and pops the exception return address from the stack.
7. It restores the context of the interrupted program directly from the alternative kernel mode stack. This includes restoring the PC, and the CPSR which switches back to the previous execution mode. If the SPSR does not have the I bit set then the operation also re-enables interrupts.

13.2 Generic Interrupt Controller

The GIC architecture defines a Generic Interrupt Controller (GIC) which comprises a set of hardware resources for managing interrupts in a single or multi-core system. The GIC provides memory-mapped registers which can be used to manage interrupt sources and behavior and (in multi-core systems) for routing interrupts to individual processors. It enables software to mask, enable and disable interrupts from individual sources, to prioritize (in hardware) individual sources and to generate software interrupts. It also provides support for the TrustZone Security Extensions described in [Chapter 26 Security](#). The GIC accepts interrupts asserted at the system level and can signal them to each processor it is connected to, potentially resulting in an IRQ or FIQ exception being taken.

From a software perspective, a GIC has two major functional blocks:

Distributor

which is shared between all processors in a multiprocessor system. This is used for configuring things such as prioritization and routing, as well as providing global enabling or disabling of individual interrupts.

CPU Interface

which is each processor's private channel for handling interrupts. This is where you find out which interrupt has been triggered and notify when you have completed processing an interrupt.

Each interrupt can be considered to be in one of four states:

- *Inactive*
- *Pending* – this means that the interrupt source has been asserted, but is waiting to be handled by a processor
- *Active* – this describes an interrupt that has been acknowledged by a processor and is currently being serviced
- *Active and pending* – this describes the situation where a processor is servicing the interrupt and the GIC also has a pending interrupt from the same source.

Interrupts can be of a number of different types:

Software Generated Interrupt (SGI)

This is generated by writing to a dedicated register, the Software Generated Interrupt Register (ICDSGIR). It is most commonly used for inter-processor communication.

Private Peripheral Interrupt (PPI)

This is generated by a peripheral that is private to an individual processor.

Shared Peripheral Interrupt (SPI)

This is generated by a peripheral that the Interrupt Controller can route to more than one processor.

Interrupts can either be edge-triggered (considered to be asserted when the Interrupt Controller detects a rising edge on the relevant input – and to remain asserted until cleared) or level-sensitive (considered to be asserted only when the relevant input to the Interrupt Controller is HIGH).

The source of an individual interrupt can be determined using an ID number. The GIC assigns a specified range of ID values to different types of interrupt.

The processor hardware then compares the interrupt priority with the current interrupt priority for the processor. If the interrupt has sufficient priority, an interrupt exception request is signaled.

We will return to the subject of the GIC in [Chapter 22](#), where we describe its implementation within an ARM MPCore processor. More detailed information on GIC behavior can be found in the TRMs for the individual processors and in the *ARM Generic Interrupt Controller Architecture* specification.

13.2.1 Configuration

Access to the Generic Interrupt Controller registers is memory-mapped. In an MPCore processor which uses the GIC, this is done using an interface private to each processor (see [Handling interrupts in an SMP system on page 22-6](#) for further details).

The distributor contains a set of configuration registers, the number depending on how many external interrupts are implemented:

- The Interrupt Configuration Registers configure individual interrupt sources as edge or level sensitive.
- The Interrupt Priority Registers set priority values for individual interrupts. A lower value indicates a higher priority.

The CPU Interfaces provide per-processor instances of registers for interrupt handling, as opposed to configuration:

- The Priority Mask Register that prevents interrupts below a certain priority level from being delivered to a processor. Interrupts of the lowest priority are never delivered.
- The Binary Point Register that enables a configurable number of the least significant bits of the priority value to be ignored for pre-emption purposes. This enables interrupts to be configured such that groups of interrupts with similar levels of priority do not pre-empt each other, but are still handled in priority order if triggered simultaneously.

Initialization sequence

This section describes the initialization sequence.

1. Enable the distributor using the Distributor Control Register.
2. Enable and set priority of SPIs using the Enable Set and Priority Level registers. The reset priority level of interrupts might be too low for them to be delivered.
3. Enable the CPU Interface using the CPU Interface Control Register.
4. Set Priority Mask Register (the reset value prevents all interrupts from being delivered).
5. Enable and set the priority of the Private Peripheral Interrupts and Software Generated Interrupts using the Enable Set and Priority Level registers. These operations are performed in the Distributor Priority Level registers. These registers are banked providing a separate copy for each processor.
6. Ensure that a valid entry exists in the vector table for the interrupt.
7. Clear the I bit in the CPSR.

13.2.2 Interrupt handling

When the processor takes the interrupt exception, it reads the *Interrupt Acknowledge Register* (ICCIAR), to acknowledge the interrupt. This read returns an interrupt ID, which is used to select the correct interrupt handler. When the GIC sees this read, it changes the state of the interrupt from pending to active or to active and pending, as appropriate. If no interrupt is currently pending, a pre-defined ID for the spurious interrupt is returned.

If an interrupt was made active, the interrupt controller then de-asserts the IRQ input to the processor. This means that the interrupt service routine can now re-enable interrupts. This enables the arrival of a higher-priority interrupt to pre-empt processing of the current one.

When the interrupt service routine has completed handling the interrupt, it signals this by writing to the *End of Interrupt Register* (ICCEOIR) in the GIC. Until this is done, new signaling of that interrupt (and any interrupts of lower priority) will not be detected.

- Interrupt numbers ID1020-ID1023 are reserved for special purposes, including signaling spurious interrupts.
- Interrupt ID values in the range ID32-ID1019 are used for SPIs.
- Interrupt numbers ID0-ID31 are used for interrupts that are private to a processor interface. In MPCore systems (or on a multi-core GIC implementation) these numbers are banked on a per core basis. ID0-ID15 are used for SGIs and ID16-ID31 are used for PPIs.

Chapter 14

Other Exception Handlers

In this chapter, we will briefly look at handlers for aborts, undefined instructions and SVC instructions and look at how interrupts are handled by the Linux kernel. Reset handlers are covered in depth in [Chapter 15 Boot Code](#).

14.1 Abort handler

Abort handler code can vary significantly between systems. In many embedded systems, an abort indicates an unexpected error and the handler will record any diagnostic information, report the error and have the application (or system) quit gracefully.

In systems which support virtual memory using an MMU, the abort handler can load the required virtual page into physical memory. In effect, it tries to fix the cause of the original abort and then return to the instruction which aborted and re-execute it. [Chapter 10 Memory Management Unit](#) gives further information about how Linux does this.

CP15 registers provide the address of the memory access which caused an abort (the Fault Address Register) and the reason for the abort (Fault Status Register). The reason might be lack of access permission, an external abort or a page table translation fault. In addition, the link register (with a -8 or -4 adjustment, depending on whether the abort was caused by an instruction fetch or a data access), gives the address of the instruction executing before the abort exception. By examining these registers, the last instruction executed and possibly other things in the system (for example, page table entries), the abort handler can determine what action to take.

14.2 Undefined instruction handling

An undefined instruction exception is taken if the processor tries to execute an instruction with an opcode which is described in the ARM architecture specification as UNDEFINED, or when a coprocessor instruction is executed but no coprocessor recognizes it as an instruction that it can execute.

In some systems, it is possible that code includes instructions for a coprocessor (such as a VFP coprocessor), but that no corresponding VFP hardware is present in the system. In addition, it may be that the VFP hardware cannot handle the particular instruction and wishes to call software to emulate it. Alternatively, the VFP hardware is disabled, and we take the exception so that we can enable it and re-execute the instruction.

Such emulators are called through the undefined instruction vector. They examine the instruction opcode which caused the exception and determine what action to take (for example, perform the appropriate floating-point operation in software). In some cases, such handlers may need to be daisy-chained together (for example, there might be multiple coprocessors to emulate).

If there is no software which makes use of undefined or coprocessor instructions, the handler for the exception should record suitable debug information and kill the application which failed due to this unexpected event.

An additional use for the undefined instruction exception in some cases is to implement user breakpoints, see [Chapter 29 Debug](#) for more information on breakpoints. (See also the description of the Linux context switch for VFP in [Chapter 7](#).)

14.3 SVC exception handling

A supervisor call (SVC) is typically used to allow User mode code to access OS functions. For example, if user code wishes to access privileged parts of the system (for example to perform file I/O) it will typically do this using an SVC instruction.

Parameters can be passed to the SVC handler either in registers or (less frequently) by using the comment field within the opcode.

Code to illustrate SVC usage with the Linux kernel is shown in [Example 14-1](#).

Example 14-1 Linux kernel SVC usage

```

_start:
    MOV     R0, #1           @ STDOUT
    ADR     R1, msgtxt       @ Address
    MOV     R2, #13          @ Length
    MOV     R7, #4           @ sys_write
    SVC     #0
    ....
    .align 2
msgtxt:
    .asciz "Hello World\n"

```

The SVC #0 instruction makes the ARM processor take the SVC exception, which is the mechanism to access a kernel function. Register R7 defines which system call we want (in this case, sys_write). The other parameters are passed in registers; for sys_write we have R0 telling where to write to, R1 pointing to the characters to be written and R2 giving the length of the string.

Another example of the use of the SVC instruction can be seen by application developers. Tools developed by ARM use SVC 0x123456 (ARM state) or SVC 0xAB (Thumb) to represent semi-hosting debug functions (for example, outputting a character on a debugger window).

14.4 Linux exception program flow

Linux utilizes a cross-platform framework for exception handling which does not distinguish between different privileged processor modes when handling exceptions. Therefore, the ARM implementation uses an exception handler stub to allow the kernel to handle all exceptions in SVC mode. All exceptions other than SVC and FIQ use the stub to switch to SVC mode and invoke the correct exception handler.

14.4.1 Boot process

During the boot process, the kernel will allocate a 4KB page as the vector page. It maps this to the location of the exception vectors, virtual address `0xFFFF0000` or `0x00000000`. This is done by `devicemaps_init()` in the file `arch/arm/mm/mmu.c`. This is invoked very early in the ARM system boot. After this, `trap_init` (in `arch/arm/kernel/traps.c`), copies the exception vector table, exception stubs and kuser helpers into the vector page. The exception vector table obviously has to be copied to the start of the vector page, the exception stubs being copied to address `0x200` (and kuser helpers copied to the top of the page, at `0x1000 - kuser_sz`), using a series of `memcpy()` operations, as shown in [Example 14-2](#).

Example 14-2 Copying exception vectors during Linux boot

```
unsigned long vectors = CONFIG_VECTORS_BASE;

memcpy((void *)vectors, __vectors_start, __vectors_end - __vectors_start);
memcpy((void *)vectors + 0x200, __stubs_start, __stubs_end - __stubs_start);
memcpy((void *)vectors + 0x1000 - kuser_sz, __kuser_helper_start, kuser_sz);
```

When the copying is complete, the kernel exception handler is in its runtime dynamic status, ready to handle exceptions

14.4.2 Interrupt dispatch

There are two different handlers, `__irq_usr` and `__irq_svc`. These save all of the processor registers and use a macro `get_irqnr_and_base` which indicates if there is an interrupt pending. The handlers loop around this code until no interrupts remain. If there is an interrupt, the code will branch to `do_IRQ` which exists in `arch/arm/kernel/irq.c`.

At this point, the code is the same in all architectures and we call an appropriate handler written in C.

There is, however, a further point to consider. When the interrupt is completed, we would normally need to check whether or not the handler has done something which needs the kernel scheduler to be called. If the scheduler decides to go to a different thread, the one that was originally interrupted stays dormant until it is selected to run again.

Chapter 15

Boot Code

In this chapter, we will look at the work which needs to be undertaken within the boot code running in an ARM processor based system. We will focus on two distinct areas:

- Code to be run immediately after the processor comes out of reset, on a so-called bare-metal system, that is, one in which code is run without the use of an operating system. This is a situation which is often encountered when first starting up a chip or system.
- How a bootloader loads and runs the Linux kernel.

15.1 Booting a bare-metal system

When the processor has been reset, it will commence execution at the location of the reset vector within the exception vector table (at either address 0x00000000 or 0xFFFF0000). The reset handler code will need to do some, or all of the following:

- In a multi-processor system, put non-primary processors to sleep (see [Booting SMP systems on page 22-10](#)).
- Initialize exception vectors.
- Initialize the memory system, including the MMU.
- Initialize processor mode stacks and registers.
- Initialize any critical I/O devices.
- Perform any necessary initialization of NEON or VFP.
- Enable interrupts.
- Change processor mode or state.
- Handle any set-up required for the Secure world (see [Chapter 26](#)).
- Call the main() application.

The first consideration is placement of the exception vector table. We need to make sure that it contains a valid set of instructions which branch to the appropriate handlers. The `_start` directive in the GNU Assembler tells the linker to locate code at a particular address and can be used to place code in the vector table. The initial vector table will be in non-volatile memory and can contain branch to self instructions (other than the reset vector) as no exceptions are expected at this point. Typically, the reset vector contains a branch to the boot code in ROM. The ROM can be aliased to the address of the exception vector. The ROM then writes to some memory-remap peripheral, which maps RAM into address 0 and the real exception vector table is copied into RAM. This means the part of the boot code which handles remapping should be position-independent, as only PC-relative addressing can be used. [Example 15-1](#) shows an example of typical code which will be placed in the exception vector table.

Example 15-1 Typical exception vector table code

```

start
    B    Reset_Handler
    B    Undefined_Handler
    B    SWI_Handler
    B    Prefetch_Handler
    B    Data_Handler
    NOP  @ Reserved vector
    B    IRQ_Handler

@ FIQ_Handler will follow directly after this table

```

We may then need to initialize stack pointers for the various modes that our application can make use of. [Example 15-2 on page 15-3](#) gives a simple example, showing code which initializes the stack pointers for FIQ and IRQ modes.

Example 15-2 Code to initialize the stack pointers

```

LDR    R0, stack_base
@ Enter each mode in turn and set up the stack pointer
MSR    CPSR_c, #Mode_FIQ:OR:I_Bit:OR:F_Bit ;
MOV    SP, R0
SUB    R0, R0, #FIQ_Stack_Size
MSR    CPSR_c, #Mode_IRQ:OR:I_Bit:OR:F_Bit ;
MOV    SP, R0

```

The next step is to set up the caches, MMU and branch predictors. An example of such code is shown in [Example 15-3](#). We begin by disabling the MMU and caches and invalidating the caches and TLB. The example code is for the Cortex-A9 processor. For the Cortex-A8 and Cortex-A5 processors, the cache invalidation can be done automatically by a hardware state machine at reset, but for the Cortex-A9 processor, boot code must explicitly cycle through the lines of the cache and invalidate them. The MMU TLBs must be invalidated. The branch target predictor hardware may not need to be explicitly invalidated, but it should be enabled by boot code. Branch prediction can safely be enabled at this point; this will improve performance.

Example 15-3 Setting up caches, MMU and branch predictors

```

@ Disable MMU
MRC    p15, 0, r1, c1, c0, 0 @ Read Control Register configuration data
BIC    r1, r1, #0x1
MCR    p15, 0, r1, c1, c0, 0 @ Write Control Register configuration data

@ Disable L1 Caches
MRC    p15, 0, r1, c1, c0, 0 @ Read Control Register configuration data
BIC    r1, r1, #(0x1 << 12) @ Disable I Cache
BIC    r1, r1, #(0x1 << 2) @ Disable D Cache
MCR    p15, 0, r1, c1, c0, 0 @ Write Control Register configuration data

@ Invalidate L1 Caches
@ Invalidate Instruction cache
MOV    r1, #0
MCR    p15, 0, r1, c7, c5, 0

@ Invalidate Data cache
@ to make the code general purpose, we calculate the
@ cache size first and loop through each set + way

MRC    p15, 1, r0, c0, c0, 0 @ Read Cache Size ID
MOV    r3, #0x1ff
AND    r0, r3, r0, LSR #13 @ r0 = no. of sets - 1

MOV    r1, #0 @ r1 = way counter way_loop
way_loop:
MOV    r3, #0 @ r3 = set counter set_loop
set_loop:
MOV    r2, r1, LSL #30 @
ORR    r2, r3, LSL #5 @ r2 = set/way cache operation format
MCR    p15, 0, r2, c7, c6, 2 @ Invalidate line described by r2
ADD    r3, r3, #1 @ Increment set counter
CMP    r0, r3 @ Last set reached yet?
BNE    set_loop @ if not, iterate set_loop
ADD    r1, r1, #1 @ else, next
CMP    r1, #4 @ Last way reached yet?
BNE    way_loop @ if not, iterate way_loop

```

```

@ Invalidate TLB
MCR    p15, 0, r1, c8, c7, 0

@ Branch Prediction Enable
MOV    r1, #0
MRC    p15, 0, r1, c1, c0, 0 @ Read Control Register configuration data
ORR    r1, r1, #(0x1 << 11) @ Global BP Enable bit
MCR    p15, 0, r1, c1, c0, 0 @ Write Control Register configuration data

```

After this, we can create some page tables, as shown in the example code of [Example 15-4](#). The variable `ttb_address` is used to denote the address to be used for the initial page table. This should be a 16KB area of memory (whose start address is aligned to a 16KB boundary), to which an L1 page table can be written by this code.

Example 15-4 Create page tables

```

@ Enable D-side Prefetch
MRC    p15, 0, r1, c1, c0, 1 @ Read Auxiliary Control Register
ORR    r1, r0, #(0x1 << 2) Enable D-side prefetch
MCR    p15, 0, r1, c1, c0, 1 ;@ Write Auxiliary Control Register
DSB
ISB
@ DSB causes completion of all cache maintenance operations appearing in program
@ order before the DSB instruction
@ An ISB instruction causes the effect of all branch predictor maintenance
@ operations before the ISB instruction to be visible to all instructions
@ after the ISB instruction.
@ Initialize PageTable

@ We will create a basic L1 page table in RAM, with 1MB sections containing a flat
(VA=PA) mapping, all pages Full Access, Strongly Ordered

@ It would be faster to create this in a read-only section in an assembly file

LDR    r0, =2_00000000000000000000000000000000110111100010 @ r0 is the non-address part of
descriptor
LDR    r1, ttb_address
LDR    r3, = 4095 @ loop counter
write_pte
ORR    r2, r0, r3, LSL #20 @ OR together address & default PTE bits
STR    r2, [r1, r3, LSL #2] @ write PTE to TTb
SUBS   r3, r3, #1 @ decrement loop counter
BNE    write_pte

@ for the very first entry in the table, we will make it cacheable, normal,
write-back, write allocate
BIC    r0, r0, #2_1100 @ clear CB bits
ORR    r0, r0, #2_0100 @ inner write-back, write allocate
BIC    r0, r0, #2_1110000000000000 @ clear TEX bits
ORR    r0, r0, #2_1010000000000000 @ set TEX as write-back, write allocate
ORR    r0, r0, #2_1000000000000000 @ shareable
STR    r0, [r1]

@ Initialize MMU
MOV    r1, #0x0
MCR    p15, 0, r1, c2, c0, 2 @ Write Translation Table Base Control Register
LDR    r1, ttb_address
MCR    p15, 0, r1, c2, c0, 0 @ Write Translation Table Base Register 0

```



```

@ In this simple example, we don't use TRE or Normal Memory Remap Register.
@ Set all Domains to Client
LDR r1, =0x55555555
MCR p15, 0, r1, c3, c0, 0 @ Write Domain Access Control Register

@ Enable MMU
MRC p15, 0, r1, c1, c0, 0 @ Read Control Register configuration data
ORR r1, r1, #0x1 @ Bit 0 is the MMU enable
MCR p15, 0, r1, c1, c0, 0 @ Write Control Register configuration data

```

The L2 cache, if present, and if running without an operating system, may also need to be invalidated and enabled at this point. NEON or VFP access must also be enabled. If the system makes use of the TrustZone Security Extensions, it may need to switch to the Normal world when the Secure world is initialized. See [Chapter 26 Security](#) for details of this,

The next steps will depend upon the exact nature of the system. It may be necessary, for example, to zero-initialize memory which will hold uninitialized C variables, copy the initial values of other variables from a ROM image to RAM, and set up application stack and heap spaces. It may also be necessary to initialize C library functions, call top-level constructors (for C++ code) and other standard embedded C initialization.

We will examine MPCore programming in more detail in [Chapter 21](#) - [Chapter 24](#), but we can touch briefly on how these are booted here. A common approach is to allow a single processor within the MPCore to perform system initialization, while the same code, if run on a different processor, will cause it to sleep (that is, enter WFI state, as described in [Chapter 25](#)). The other processors might be woken after processor 0 has created a simple set of L1 page table entries, as these could be used by all processors in the system. [Example 15-5](#) shows example code which determines which processor it is running on and either branches to initialization code (if running on processor 0), or goes to sleep otherwise. The secondary processors are typically woken up later by an SMP OS.

Example 15-5 Determining which processor is running

```

@ Only CPU 0 performs initialization. Other CPUs go into WFI
@ to do this, first work out which CPU this is
@ this code typically is run before any other initialization step

MRC p15, 0, r1, c0, c0, 5 @ Read Multiprocessor Affinity Register
AND r1, r1, #0x3 @ Extract CPU ID bits
CMP r1, #0
BEQ initialize @ if we're on CPU0 goto the start

wait_loop:
@ Other CPUs are left powered-down
....
....
....
initialize:
@ next section of boot code goes here

```

15.2 Configuration

There are a number of control register bits within the ARM processor which will typically be set by boot code. In all cases, for best performance, code should run with the MMU, instruction and data caches and branch prediction enabled. Page table entries for all regions of memory which are not peripheral I/O devices should be marked as L1 Cacheable and (by default) set to read-allocate, write-back cache policy. For MPCore systems, pages should be marked as Shareable and the broadcasting feature for CP15 maintenance operations should be enabled.

In addition to the CP15 registers mandated by the ARM architecture, processors typically also have registers which control implementation-specific features. Programmers of boot code should refer to the relevant technical reference manual for the correct usage of these.

15.3 Booting Linux

It is useful to understand what happens from the ARM processor coming out of reset and executing its first instruction at the exception base address `0x00000000` or `0xFFFF0000` if HIVECS (known as high vectors) is selected, until the Linux command prompt appears. (See [Vector table on page 12-9](#).)

When the kernel is present in memory, the sequence on an ARM processor based system is similar to what might happen on a desktop computer. However, the bootloading process can be very different, as ARM processor based phones or more deeply embedded devices can lack a hard drive or PC-like BIOS.

Typically, what happens when we power the system on is that hardware specific boot code runs from flash or ROM. This code initializes the system, including any necessary hardware peripheral code and then launches the bootloader (for example U-Boot). This initializes main memory and copies the compressed Linux kernel image into main memory (from a flash device, memory on a board, MMC, host PC or elsewhere). The bootloader passes certain initialization parameters to the kernel. The Linux kernel then decompresses itself and initializes its data structures and running user processes, before starting the command shell environment. Let's take a more detailed look at each of those processes.

15.3.1 Reset handler

There is typically a small amount of system-specific boot monitor code, which configures memory controllers and performs other system peripheral initialization. It sets up stacks in memory and typically copies itself from ROM into RAM, before changing the hardware memory mapping so that RAM is mapped to the exception vector address, rather than ROM. In essence this code is independent of which operating system is to be run on the board and performs a function similar to a PC BIOS. When it has completed execution, it will call a Linux bootloader, such as U-Boot.

15.3.2 Bootloader

Linux needs a certain amount of code to be run out of reset, to initialize the system. This performs the basic tasks needed to allow the kernel to boot:

- Initializing the memory system and peripherals.
- Loading the kernel image to an appropriate location in memory (and possibly also an initial RAM disk).
- Generate the boot parameters to be passed to the kernel (including machine type).
- Set up a console (video or serial) for the kernel.
- Enter the kernel.

The exact steps taken vary between different bootloaders, so for detailed information, please refer to documentation for the one that you wish to use. U-Boot is a widely used example, but other bootloader possibilities include Apex, Blob, Bootldr and Redboot.

When the bootloader starts, it is typically not present in main memory. It must start by allocating a stack and initializing the processor (for example invalidating its caches) and installing itself to main memory. It must also allocate space for global data and for use by `malloc()` and copy exception vector entries into the appropriate location.

15.3.3 Initialize memory system

This is very much a board or system specific piece of code. The Linux kernel has no responsibility for the configuration of the RAM in the system. It is presented with the physical memory layout, but has no other knowledge of the memory system. In many systems, the available RAM and its location are fixed and the bootloader task is straightforward. In other systems, code must be written which discovers the amount of RAM available in the system.

15.3.4 Kernel images

The kernel image from the build process is typically compressed in zImage format (the conventional name given to the bootable kernel image). Its head code contains a magic number, used to verify the integrity of the decompression, plus start and end addresses. The kernel code is position independent and can be located anywhere in memory. Conventionally, it is placed at a 0x8000 offset from the base of physical RAM. This gives space for the parameter block placed at a 0x100 offset (used for page tables etc).

Many systems need an initial RAM disk (initrd), as this lets us have a root filesystem available without other drivers being setup. The bootloader can place an initial ramdisk image into memory and pass the location of this to the kernel using ATAG_INITRD2 (a tag which describes the physical location of the compressed RAM disk image) and ATAG_RAMDISK.

The bootloader will typically setup a serial port in the target, allowing the kernel serial driver to detect the port and use it for a console. In some systems, another output device such as a video driver can be used as a console. The kernel command line parameter `console=` can be used to pass the information.

15.3.5 Kernel parameters

Historically, the parameters passed to the kernel are in the form of a tagged list, placed in physical RAM with register R2 holding the address of the list. Tag headers hold two 32-bit unsigned ints, with the first giving the size of the tag in words and the second providing the tag value (indicating the type of tag). For a full list of parameters which can be passed, consult the appropriate documentation. Examples include ATAG_MEM to describe the physical memory map and ATAG_INITRD2 to describe where the compressed ramdisk image is located. The bootloader must also provide an ARM Linux machine type number (MACH_TYPE). This can be a hard-coded value, or the boot code can inspect the available hardware and assign a value accordingly.

There is currently work underway to implement a more flexible, or generic method for passing this information using *Flattened Device Trees* (FDTs).

15.3.6 Kernel entry

Kernel execution must commence with the ARM processor in a fixed state. The bootloader calls the kernel image by branching directly to its first instruction – the start label in `arch/arm/boot/compressed/head.S`. The MMU and data cache must be disabled. The processor must be in Supervisor mode, with CPSR I and F Bits set (IRQ and FIQ disabled). R0 must contain 0, R1 the MACH_TYPE value and R2 the address of the tagged list of parameters.

The first step in getting the kernel working is to decompress it. This is mostly architecture independent. The parameters passed from the bootloader are saved and the caches and MMU are enabled. Checks are made to see if the decompressed image will overwrite the compressed image, before calling `decompress_kernel()` in `arch/arm/boot/compressed/misc.c`. The cache is then cleaned and invalidated before being disabled again. We then branch to the kernel startup entry point in `arch/arm/kernel/head.S`.

15.3.7 Platform-specific actions

A number of architecture specific tasks are now undertaken. The first checks processor type using `__lookup_processor_type()` which returns a code specifying which processor it is running on. The function `__lookup_machine_type()` is then used (unsurprisingly) to look up machine type. A basic set of page tables is then defined which map the kernel code. The cache and MMU are initialized and other control registers set. The data segment is copied to RAM and `start_kernel()` is called.

15.3.8 Kernel start-up code

In principle, the rest of the startup sequence is the same on any architecture, but in fact some functions are still hardware dependent.

IRQ interrupts are disabled with `local_irq_disable()`, while `lock_kernel()` is used to stop FIQ interrupts from interrupting the kernel. It initializes the tick control, memory system and architecture-specific subsystems and deals with the command line options passed by the bootloader. Stacks are setup and the Linux scheduler is initialized. The various memory areas are set-up and pages allocated. The interrupt and exception table and handlers are setup, along with the GIC. The system timer is setup and at this point IRQs are enabled. Further memory system initialization occurs and then a value called *BogoMips* is used to calibrate the processor clock speed. Internal components of the kernel are set up, including the filesystem and the initialization process, followed by the thread daemon which creates kernel threads. The kernel is unlocked (FIQ enabled) and the scheduler started. The function `do_basic_setup()` is called to initialize drivers, `sysctl`, work queues and network sockets. At this point, the switch to User mode is performed.




Figure 15-1 Linux virtual memory view

The memory map used by Linux is shown in [Figure 15-1](#). ZI refers to zero-initialized data. There is a broad split between kernel memory, above address 0xBF000000 and user memory, below that address. Kernel memory uses global mappings, while user memory uses non-global mappings, although both code and data can be shared between processes. As already mentioned, application code starts at 0x1000, leaving the first 4KB page unused, to allow for trapping of NULL pointer references.

Chapter 16

Porting

New projects will normally use an existing operating system and re-use code from existing applications. New code might be targeted at the ARMv7-A architecture, but might eventually need porting to a different board. In this chapter, we will mainly cover issues associated with porting code from a different architecture to run on an ARM processor, or from older versions of the ARM architecture to ARMv7-A.

For many applications (particularly those coded with portability issues in mind), this will mean recompiling the source code. For example, a large amount of Linux application software is designed to run in many different environments and tends to make fewer assumptions about the underlying hardware. However, there are a number of areas where C code is not fully portable. We shall look at these further and see what problems this causes and how to resolve them. In particular, low level, hardware-specific code such as device drivers may require more effort than porting applications.

There is a further consideration when porting code between processors: that of efficiency. It may be the case that optimizations applied to code running on another processor, or to older versions of the ARM architecture, do not apply to that code when running on ARMv7-A. Equally, there may be scope for making code smaller or faster on ARMv7-A class processors. Optimizations may differ between processors or systems. Code which is optimal for one processor may not be optimal for others. We will consider the area of ARM-specific optimization in [Chapter 19](#).

16.1 Endianness

The use of the terms *little-endian* and *big-endian* was introduced by Danny Cohen in his 1980 paper “On Holy Wars and a Plea for Peace”. Cohen has also been responsible for many advances in the fields of networks and computer graphics. It is a reference to *Gulliver’s Travels*, a famous satire from the early 18th century, by Irish writer Jonathan Swift, in which a war is fought between the fictional countries of Lilliput and Blefuscu over the correct end to open a boiled egg.

There are two basic ways of viewing bytes in memory – little-endian and big-endian. On big-endian machines, the most significant byte of an object in memory is stored at the least significant (closest to zero) address. On little-endian machines, the least significant byte is stored at the address closest to zero (lowest address).

The term byte-ordering can also be used rather than endian. Other kinds of endianness do exist, notably middle-endian and bit-endian, but we will not discuss these further.

Consider the following simple piece of code ([Example 16-1](#)):

Example 16-1 Endian access

```
int i = 0x44332211;
unsigned char c = *(unsigned char *)&i;
```

On a 32-bit big-endian machine, *c* is given the value of the most significant byte of *i*: 0x44. On little-endian machines, *c* is the least significant byte of *i*: 0x11.

[Figure 16-1 on page 16-3](#) illustrates the two differing views of memory. It should be stated at this point that many people find endianness confusing and that even the act of drawing a diagram to illustrate it can reveal a personal bias! The diagram shows a 32-bit value in a register being written to address 0x1000, using a STR instruction. The processor then performs a read of a byte, using a LDRB instruction. A different value will be returned by this instruction sequence depending upon whether we have a little- or big-endian memory system.




Figure 16-1 Different endian behaviors

ARM processors support both modes, but are most commonly used in, and typically default to little-endian mode. Most Linux distributions for ARM tend to be little-endian only. The x86 architecture is little-endian. The PowerPC or the venerable 68K, on the other hand, are generally big-endian (although the Power architecture can also handle little-endian). Several common file formats and networking protocols specify different endianness. For example, .BMP and .GIF files are little-endian, while .JPG is big-endian, and TCP/IP is big-endian, but USB and PCI are little-endian.

So, there are two issues to consider – code portability and data sharing. Systems are built from multiple blocks and can include one or more processors, DSPs, peripherals, memory, network connections and so forth. Whenever data is shared between these elements, there is a potential endianness conflict. If code is being ported from a system with one endianness to a system with different endianness, it may be necessary to modify that code, either to make it endian-neutral or to work with the opposite byte-ordering.

Cortex-A series processors provide support for systems of either endian configuration, controlled by the CPSR E bit, which enables software to switch dynamically between viewing data as little or big-endian. Instructions in memory are always treated as little-endian. The REV instruction (see [Byte reversal on page 6-23](#)) can be used to reverse bytes within an ARM register, providing simple conversion between big and little-endian formats.

In principle, it is straightforward to support mixed endian systems (typically this means the system is natively of one endian configuration, but there are peripherals which are of the opposite endianness). The CPSR E bit can be modified dynamically by software, and there is a SETEND assembly instruction provided to do this. The CP15:SCTLR (System Control Register, c1), contains the EE bit (see [Coprocessor 15 on page 6-19](#)), which defines the endian mode to switch to upon an exception and also the endianness of translation table lookups. It would be

difficult if exception code had to worry about which endian state the processor was in upon arrival at the handler. In practice, however, it can be difficult to tell the compiler that part of the system is of a different endian configuration to the rest of memory.

Modern ARM processors support a big-endian format known architecturally as BE8. Older ARM processors used a different format known as BE-32, which applied to both instructions and data. BE8 corresponds to what most other computer architectures simply call big-endian.

Let's look at a simple piece of code ([Example 16-2](#)) which will behave differently when run on architectures with different endianness.

Example 16-2 Non-portable code

```
int i= 0x12345678;
char *buf = (char*)&i;
char i0, i1, i2, i3;

i0 = buf[0];
i1 = buf[1];
i2 = buf[2];
i3 = buf[3];
```

The values of `i0...i3` are not guaranteed to be the same if the system endianness changes. This kind of code is therefore inherently non-portable.

When inspecting code in which you suspect endianness problems, you should look for the following potential causes of problems:

Unions A union can hold objects of different types and sizes. The programmer must keep track of what the data member represents at any particular time. Code which uses unions should be carefully checked. If the union is used to access the same data, but with different data types, there exists a possible endianness, alignment, and packing problem. Any time that halfword, word (or longer) data types are combined or viewed as an array of bytes is a potential issue.

Casting of data types

Anywhere that data is accessed in a way outside of its native data type is a potential problem. Similarly, if there are arrays of bytes, they should not be accessed other than as a byte data type. Casting of pointers changes how data is addressed and can be endian sensitive.

Bitfields To avoid endianness problems code that defines bitfields or performs bit operations should not be used in code that is intended to be portable.

Data sharing

Any code which reads shared data from another block, or exports data to another block, should be checked to see whether the two blocks agree endian definitions. If the two are different, it may be necessary to implement byte swapping at one location.

Network code

Code which accesses networking or other I/O devices needs to be reviewed to see if there is any endian dependency. Again, it may be necessary to re-write code for greater efficiency, or swap bytes at the interface.

16.2 Alignment

The alignment of accesses is significant on ARM processors. On older ARM processors, accesses to addresses which are not aligned are possible, but with a different behavior to those using the ARMv7 architecture. On ARM7 and ARM9 processors, an unaligned LDR is performed in the memory system in the same way as an aligned access, but with the data returned being rotated so that the data at the requested address is placed in the least significant byte of the loaded register. Some older compilers and operating systems were able to use this behavior for clever optimizations. This can represent a portability problem when moving code from an ARMv4 or ARMv5 to ARMv7 architecture.

ARM MMUs can be configured to automatically detect such unaligned accesses and abort them (using the CP15:SCTL A bit), see [Coprocesor 15 on page 6-19](#).

For the Cortex-A series of processors, unaligned accesses are supported (although one must enable this by setting the U bit in the CP15:SCTL register, indicating that unaligned accesses are permitted). This means that instructions to read or write words or halfwords can access addresses which are not aligned to word or halfword boundaries. However, load and store multiple instructions (LDM and STM) and load and store double-word (LDRD or STRD) must be aligned to at least a word boundary. Furthermore, loads and stores of floating-point values must always be aligned. Additional alignment constraints may be imposed by the ABI that are stronger than those imposed by the ARM architecture.

These unaligned accesses can take additional cycles in comparison with aligned accesses and therefore alignment is also a performance issue. In addition, such accesses are not guaranteed to be atomic. This means that a external agent (another processor in the system) might perform a memory access which appears to occur part way through the unaligned access. For example, it might read the accessed location and see the *new* value of some bytes and the *old* value of others.




Figure 16-2 Aligned words at address 0 or 8

A word aligned address is one which is a multiple of four, for example 0x100, 0x104, 0x108, 0x10C, 0x110. [Figure 16-2](#) shows examples of aligned accesses.




Figure 16-3 An unaligned word

An unaligned word at address 1 is shown in [Figure 16-3](#). It takes three bytes from the word at address 0 and one byte from the word at address 4.

A simple example where alignment effects can have significant performance effects is the use of `memcpy()`. Copying small numbers of bytes between word aligned addresses will be compiled into LDM or STM instructions. Copying larger blocks of memory aligned to word boundaries will typically be done with an optimized library function which will also use LDM or STM. Copying blocks of memory whose start or end points do not fall on a word boundary can result in a call to a generic `memcpy()` function which can be significantly slower. Although, if the source and destination are similarly unaligned then only the start and end fragments are non-optimal. Whenever explicit typecasting is performed, that cast always carries alignment implications.

16.3 Miscellaneous C porting issues

In this section we consider some other possible causes of problems when porting C code.

16.3.1 unsigned char and signed char

The piece of code in [Example 16-3](#) illustrates a very simple example of a possible issue when porting code to ARM.

Example 16-3 Use of unsigned char

```
char c = -1;
if (c > 0) printf("c is positive \n");
else     printf("c is negative \n");
```

When this is compiled for some architectures (for example, for x86) the result is the one you might intuitively expect, which is that it reports the variable `c` as negative, but compiling the code on an ARM Compiler will produce code which reports `c` as positive (and typically a warning will be emitted by the compiler, too).

The ANSI C standard specifies a range for both signed (at least -127 to +127) and unsigned (at least 0 to 255) chars. Simple chars are not specifically defined and it is compiler dependent whether they are signed or unsigned. Although the ARM architecture now has the instruction `LDRSB` (which loads a signed byte into a 32-bit register with sign extension), the earliest versions of the processor did not. And so it made sense for the compiler to treat simple chars as unsigned, whereas on the x86 simple chars are, by default, treated as signed.

One workaround for users of GCC is to use the `-fsigned-char` command line switch (or `--signed-chars` for RVCT), which forces all chars to become signed, but a better practice is to write portable code by declaring char variables appropriately. Unsigned char should be used for accessing memory as a block of bytes or for small unsigned integers. Signed char should be used for small signed integers and simple char should be used only for ASCII characters and strings. In fact, on an ARM processor, it is usually better to use ints rather than chars, even for small values, for performance reasons. You can read more on this in [Chapter 19 Optimizing Code to Run on ARM Processors](#).

A second piece of code, in [Example 16-4](#), illustrates another possible problem with chars. Here we compare EOF with an unsigned char. On the ARM, the while loop will never complete. The value of EOF is defined as -1 and when it is converted to be compared with a char (which is unsigned and therefore in the range 0 to 255), it can never be equal and so the loop does not terminate.

Example 16-4 Use of EOF

```
char c;
while ((c = getchar()) != EOF) putchar(c);
```

Here, we should declare the variable as `int` instead of `char` to avoid the problem – in fact, this is how the functions in `stdio.h` are defined.

Similar cases to look out for include the use of `getopt()` and `getc()` – both are defined as returning an `int`.

16.3.2 Compiler packing of structures

Compilers are not allowed to re-order members of a structure and have to follow the alignment restrictions of the processor architecture. This means that compilers may have to add unused bytes into user defined structures, for best performance and code size. Such padding is architecture specific and can therefore lead to portability problems if assumptions have been made about the location and size of this padding.

Marking a structure as `__packed` in the ARM Compiler or using the attribute `__packed__` in GCC, will remove any padding. This reduces the size of the structure and can be useful when porting code or for structures being passed from external hardware, but can reduce performance and increase code size, although generally it will be relatively efficient on Cortex-A series processors.

If we have some simple struct code, as shown in [Example 16-5](#):

Example 16-5 A typical C struct

```
struct test
{
    unsigned char c;
    unsigned int i;
    unsigned short s;
}
```

Then the arrangement of data within the struct will be as in [Figure 16-4](#).




Figure 16-4 Packed versus unpacked

If we now mark the structure as packed, as in [Example 16-6](#):

Example 16-6 Packed structure

```
struct test
{
    unsigned char c;
    unsigned int i;
    unsigned short s;
} __attribute__((__packed__));
```

The layout is now that byte 0 holds `c`, bytes 1-4 hold `i` and bytes 5-6 will hold `s`. To access `s` or `i` will require the processor to perform an unaligned access, which is also shown in [Figure 16-4 on page 16-8](#).

16.3.3 Use of the stack

Code which makes assumptions about the stack structure can require porting effort. For example, functions with a variable number and type of arguments can receive their variables through the stack. The `<stdarg.h>` macros dealing with accessing these arguments will walk through the stack frame and provide compatibility between systems, but code which does not use standard libraries or macros can have a problem.

16.3.4 Other issues

A function prototype is a declaration that omits the function body but gives the function name, argument and return types. It effectively gives a way to specify the interface of a function separately from its definition. Incorrectly prototyped functions can behave differently between different compilers.

Compilers can allocate different numbers of bytes to `enum`. Care is therefore required when enumerations are used; cross-linking of code and libraries between different compilers may not be possible if `enums` are used.

Code written for 8-bit or 16-bit microprocessors must assume that integer variables are 16-bit. On ARM processors, they will always be 32 bits. The program may rely on 16-bit behaviors. In general, this is easily fixed by the use of the C `short` type. Use of `short ints` can be less efficient, as we shall see in the chapter on optimization and in cases where the code does not rely on 16-bit behavior, it is usually better to promote these variables to a 32-bit `int`.

16.4 Porting ARM assembly code to ARMv7

So far in this chapter, we have looked at porting C code from other architectures to ARM. It is sometimes necessary to port assembly code from older ARM processors to the Cortex-A series. Sometimes, it can be difficult to determine which ARM architecture variant your code was originally targeting. GCC has a series of macros, with names like `__ARM_ARCH_6__` which are mutually exclusive. The ARM Compiler has a similar set of macros such as `__TARGET_ARCH_7_A`. In general, ARM assembler code is backward compatible and will work unmodified.

There are a few special cases to look for:

CP15 Operations

The architecture specifies a consistent mapping of CP15 to designed system control operations. In general, one should attempt to understand the purpose of code which performs CP15 instructions and ensure that this code is appropriate for the ARMv7-A Architecture. In addition, a number of CP15 registers (for example CP15:ACTL, the Auxiliary Control Register) are implementation specific. Code which references such registers will always need attention when being ported.

SWP

The SWP (or SWPB) instruction was used to implement atomic operations in older versions of the ARM architecture, but is deprecated and its use strongly discouraged in the ARMv7 Architecture. There is no encoding for SWP in Thumb at all, so SWP is not allowed when building for Thumb. In the ARMv7 Architecture, SWP is disabled by default, but can be re-enabled by setting CP15:SYSCTL bit [10]. Code which uses SWP should be rewritten to make use of LDREX or STREX (and possibly also barrier instructions – see [Chapter 11](#)). Alternatively, the GCC `__sync_...` intrinsics could be used. It is usually preferable to use library functions for such things as spinlocks, semaphores, and mutexes, rather than writing such primitives yourself. The mechanisms used are different from those used by SWP, so it is necessary to port all code accessing an atomic object, not just some of it. Usually an atomic object will be managed by a piece of library code shared between the threads which access it, so this is not typically a problem. Cortex-A series processors treat SWP as Undefined out of reset.

16.4.1 Memory access ordering and memory barriers

The ARMv7 architecture has a weakly-ordered memory model. Code for older processors which makes implicit assumptions about ordering of memory accesses may not operate correctly on ARMv7 devices. This may be particularly true for code which interacts with other devices, such as a DMA controller or other bus master. Such code should be inspected and modified, possibly by the insertion of appropriate barrier instructions or by the use of suitable atomic primitives. See [Chapter 11](#) for a more detailed description of memory ordering and memory barriers.

16.5 Porting ARM code to Thumb

We also consider problems associated with porting code written for the ARM instruction set to the Thumb instruction set. As we have seen, use of Thumb is often preferred due to its combination of small code size with higher performance relative to the older 16-bit only Thumb instruction set.

16.5.1 Use of PC as an operand

Instructions which use PC (R15) as an explicit operand can cause problems when assembling in Thumb. As we have seen, it is not possible to encode any arbitrary 32-bit address into any instruction as an operand. For this reason, it is common to address data stored in a literal pool using offsets relative to the current instruction location. In ARM code, the PC value can be used (with some adjustment) to determine the address of the currently executing instruction, for this purpose and this enables position-independent coding. However, the PC value obtained in this way can show some inconsistencies between ARM and Thumb states and can also depend upon the type of instruction executed. For this reason, ARM assembly code which directly references the PC register may need to be modified in order to work correctly in Thumb.

It is better to avoid explicit PC arithmetic and instead to use an instruction like:

```
LDR    r0, =<value>
```

This automatically puts <value> somewhere in the text section and assembles an appropriate PC-relative LDR instruction. You can still do a load from a local text section label which you declare explicitly, as shown below, but again, the assembler or linker should be allowed to perform the PC offset calculation:

```
LDR    r0, data
...
data:  .long  <value>
```

Sometimes the required PC-relative address offset is too large to encode in a single LDR instruction, causing the assembler to complain that a literal pool is out of range. This can be resolved by explicitly placing a literal pool, using the `.ltorg` directive, which tells the assembler where to insert literal data. The programmer must ensure that the literal data is not located where it might be executed as code. This typically means the `.ltorg` directive is placed after an unconditional branch or function return instruction.

16.5.2 Branches and interworking

When using Thumb, the system will typically have both ARM and Thumb functions (even if we compile our application for Thumb, we may still need to think about such things as libraries and prebuilt binaries). The processor needs to know which instruction set is to be used for the code being executed after a branch, procedure call or return. This interworking between instruction sets was described in [Interworking on page 5-11](#). When writing C code, the linker takes care of this for us, but a little more care is needed when porting assembly code.

The target instruction set state is determined in different ways depending on the type of branch. We can consider a number of different instructions:

Function return

Code written for ARMv4 can use the MOV PC, LR instruction. This is unsafe for systems which contain a mix of ARM and Thumb code and should be replaced by BX LR for code running on all later architectures.

Function return from the stack

This is done using the LDMFD SP!, {registers, pc} instruction, which will continue to work correctly in the ARMv7-A architecture (although a newer, equivalent form, POP {<registers>, pc} is also available). This is used when the LR and other registers which need to be preserved by the function are PUSHed at the start of the function.

Branch

A simple B instruction will work in the same fashion on all ARM architectures. If ARM and Thumb instructions are mixed in a single source file (this is unusual), there is no automatic instruction set switch for local symbols. The assembler may or may not introduce a *veneer* depending on whether it knows that the destination is in a different instruction set and is definitely a code symbol (such as a .type <symbol>, %function or .thumb_func). Note that just because a symbol appears in a code section it is not assumed to be a code symbol unless specifically tagged in this way. If the label is in a different file, the linker will take care of any necessary instruction set change. Similar considerations apply for a function call (BL).

———— Note ————

Veneers are small pieces of code which are automatically inserted by the linker when it detects that a branch target is out of range or is a conditional branch to code in the other state, for example, from Thumb to ARM or vice-versa. The veneer becomes an intermediate target of the original branch with the veneer itself then being a branch to the target address. Often these veneers can be inlined. The linker can reuse a veneer generated for a previous call, for other calls to the same function if it is in range from both calls. Occasionally, such veneers can be a performance factor. If you have a loop which calls multiple functions through veneers, you will get many pipeline flushes and therefore sub-optimal performance. Placing related code together in memory can avoid this. The ARM linker can be made to export information on this by specifying the --info veneers option.

PC modification

Care may be needed with other instructions which modify the PC and produce a branch effect. For example, MOV PC, register must be replaced with BX register in systems which contain both ARM and Thumb code.

Function call to register address

If code contains a sequence like MOV LR, PC followed by MOV PC, register, this will not work in a system which has both ARM and Thumb code. It should be replaced with the single instruction BLX <register>.

When a destination or return address is variable or calculated at run-time, care is needed to appropriately set the Thumb bit (bit [0]) in the address correctly and to do the correct type of branch, to make sure that the call (and return, if applicable) switches instruction set appropriately.

If an external label or function defined in another object is referenced, the linker will produce an address with the Thumb bit (bit [0]) set appropriately. However, if you reference a symbol internal to the object, things are more complicated. For C functions, or code tagged as Thumb, bit [0] will be set appropriately, but it will not be set appropriately for other symbols. In particular, GNU Assembler local labels will not have the Thumb bit set appropriately, nor will the GNU current assembly location symbol “.”.

Therefore, when coding in assembler, if an address will be passed to any other function or object, for example, as a return address, method address or callback, you must handle the Thumb bit setting yourself, setting bit [0] of the address where required.

16.5.3 Operand combinations

Thumb and ARM Assembly code have different restrictions on instruction operands. You may therefore find that existing ARM code can produce assembler errors when targeted for Thumb.

Branch out of range errors occur when the distance between the current instruction and the branch target is too large to be encoded in a Thumb instruction. To resolve this, it may be necessary to use a different type of branch, move code sections or to use two separate branches, a so-called *trampoline*.

Similarly, *index out of range* errors may be produced on load and store operations, and to resolve these it may be necessary to manually add part (or all) of the required index offset to the base register in a separate explicit instruction.

Generally, use of SP should be limited to stack operations – other usage may not be permitted in Thumb code. This means that PUSH, POP, LDMFD SP!, STMFD SP!, ADD, SUB or MOV instructions which use the SP are allowed, but other operations should be treated as possible problems. Similarly, operations which directly operate on the PC should be checked (other than the usual function or exception return operations, or literal pool loads).

16.5.4 Other ARM/Thumb differences

There are a number of other differences which can need attention by the assembly language programmer.

- The RSC instruction is not available in Thumb. Therefore, code which uses RSC needs to be re-coded using RSB or SBC, or be built in ARM state.
- Most ARM instructions can optionally be made conditional. This is not the case in Thumb, other than for branches. Instead, small instruction sequences can be executed conditionally by preceding them with the IT instruction. For compatibility with both ARM and Thumb, the IT block construct is always understood when using unified assembler syntax. Manually modifying code to use IT instructions can be tedious. Fortunately, the assembler command-line option `-mimplicit-it=<when>`, where `<when>` can be one of `never`, `arm`, `thumb` or `always`. In this case we don't need to add IT instructions, the assembler will work out the right thing to do for the given target. When assembling for Thumb, however, it is sensible to use `-mimplicit-it=thumb`.

Chapter 17

Application Binary Interfaces

The C compiler is able to generate code from many separately compiled modules. For programs to execute successfully these modules must be able to work together with each other, with the operating system code and any code which is written in assembler or any other compiled language. For that reason, we must define a set of conventions to govern inter-operability between separate pieces of code.

The *Application Binary Interface* (ABI) for the ARM architecture specification describes a set of rules that an ARM executable must adhere to in order to execute in a specific environment. It specifies conventions for executables, including file formats and ensures that objects from different compilers or assemblers can be linked together successfully. There are variants of the ABI for specific purposes, for example, we may consider the Linux ABI for the ARM architecture or the Embedded ABI (EABI).

The *ARM Architecture Procedure Call Standard* (AAPCS) is part of the ABI and specifies conventions for register and stack usage by the compiler and during subroutine calls. Knowledge of this is vital for inter-working C and assembly code and can be useful for writing optimal code. The AAPCS supersedes the previous *ARM-Thumb Procedure Call Standard* (ATPCS).

The AAPCS mandates specific things that must be done by a caller to allow a callee function to run and what the called routine must do to preserve the program state of the caller through a function. It describes the way that data is laid out in memory and how the stack is laid out, plus permitted variations for processor extensions. It defines how code which has been separately compiled or assembled works together.

17.1 Procedure Call Standard

As we have seen, there are 16 integer registers available in the processor, each of size 32-bits. These are labeled R0-R15. [Table 17-1](#) shows the role assigned to registers within the procedure call standard.

Table 17-1 APCS registers

Register	PCS name	PCS role
R0	a1	argument 1/scratch register/result
R1	a2	argument 2/scratch register/result
R2	a3	argument 3/scratch register/result
R3	a4	argument 4/scratch register/result
R4	v1	register variable
R5	v2	register variable
R6	v3	register variable
R7	v4	register variable
R8	v5	register variable
R9	tr/sb/v6	static base/ register variable
R10	s1/v7	stack limit/stack chunk handle/register variable
R11	FP/v8	frame pointer/register variable
R12	IP	scratch register/new -sb in inter-link-unit calls
R13	SP	Lower end of the current stack frame
R14	LR	link register/scratch register
R15	PC	program counter

For the purposes of function calls, the registers are divided into three groups:

- *Argument* registers R0-R3 (a1-a4). These can be used as scratch registers or as caller-saved register variables which can hold intermediate values within a routine, between calls to other functions.
- Callee-saved registers, normally used as register variables. Typically, the registers R4-R8, R10 and R11 (v1-v5, v7 and v8) are used for this purpose.
- Registers which have a dedicated role.

The function of the program counter, link register and stack pointer should by now be clear. If not, you should refer to [Registers on page 4-4](#).

The IP (R12) register can be used by the linker, as a scratch register between a routine and any subroutine it calls, or as an additional local variable within a function. As the BL instructions cannot address the full 32-bit address space, the linker may need to insert a veneer between the caller and callee. Veneers can also be used for ARM-Thumb inter-working or dynamic linking. Veneers are permitted to modify the contents of IP (R12).

Register R9 has a role which is specific to a particular environment. It can be used as the static base register (SB) to point to position-independent data, or as the thread register (TR) where thread-local storage is used. In code that has no need for such a special register, it can be used as an extra callee-saved variable register, v6.

The first four word-sized parameters passed to a function will be transferred in registers R0-R3. Sub-word sized arguments (for example, char) will still use a whole register. Arguments larger than a word will be passed in multiple registers. If more arguments are passed, the fifth and subsequent words will be passed on the stack. Passing arguments on the stack always requires additional instructions and memory accesses and therefore reduces performance. For optimal code, therefore, the programmer should always try to limit arguments to four words or fewer. If this is not possible, the most commonly used parameters should be defined in the first four positions of the function definition. If the arguments are part of a structure then it is more optimal to pass a pointer to the structure instead. C++ uses the first argument to pass the this pointer to member functions, so only three arguments can be passed in registers.

There are additional rules about 64-bit types. 64-bit types must always be 8-byte aligned in memory (recall that in [Alignment on page 16-5](#) we described how there are limitations on use of LDRD and STRD double-word instructions to unaligned addresses). In addition, 64-bit arguments to functions must be passed in an even + consecutive odd register pair (for example, R0 + R1 or R2 + R3).

If 64-bit arguments are passed on the stack, they must be at an 8-byte aligned location. Again, this is because of restrictions on LDRD and STRD instructions. If such 64-bit arguments are listed in a sub-optimal fashion, there can be wasted space in registers or on the stack. When considering such issues, it is important to take into account the this pointer in R0 present in all non-static C++ member functions.




Figure 17-1 Efficient parameter passing

[Figure 17-1](#) shows some examples of how sub-optimal listing of arguments can cause unnecessary spilling of variables to the stack. The figure shows how two different function calls, which pass identical parameters, make use of the registers and stack.

The first function passes an int, a double and a further int. The first parameter is passed in R0. The second argument is 64-bits and must be passed in an even and consecutive odd register (or in an 8-byte aligned location on the stack). It is therefore passed in R2 and R3. This means that the final argument is passed on the stack.

As the stack pointer needs to be 8-byte aligned (a constraint imposed by the ABI to simplify the interface at function boundaries), there will be an additional unused word pushed and popped. In the second function call, we can pass the two `int` values in registers R0 and R1 and the double value in R2 and R3. This means that no values are spilled to the stack, which gives more efficient code, requiring both fewer instructions and fewer memory accesses.

The registers R4-R11 (v1-v8) are used to hold the values of the local variables of a subroutine. A subroutine is required to preserve (on the stack), the contents of the registers R4-R8, R10, R11 and SP (and R9 in PCS variants that designate R9 as v6), if they are used by that subroutine.

A caller function will have code like this:

```
@ may need to preserve r0-r3
@ does not need to preserve r4-r11
BL Func
```

while the callee function will have code like this:

```
Func:

@ Must preserve r4-r11, lr (if used)
@ May corrupt r0-r3, r12
PUSH {r4-r11, lr}
...
...
...
POP {r4-r11, pc}
@ Returns value in r0 - char, short or int
@ Returns value in r0 & r1 - double
```

The PUSH/POP instructions must maintain 8-byte stack alignment, and hence use an even number of registers. Leaf functions do not have to do so. The example callee code as shown pushes or pops R4-R11 and LR and PC, which would not preserve an 8-byte aligned stack. It is shown like this to indicate which registers need to be saved. In practice, the compiler will normally push an extra register, depending on whether the function is leaf and which registers are modified by the function. Actual instructions will usually be one of PUSH/POP {r4, lr/pc}, PUSH/POP {r4-r8, lr/pc}, PUSH/POP {r4-r10, lr/pc} or PUSH/POP {r4-r12, lr/pc}. In each case, we would PUSH lr and POP pc.

17.1.1 VFP and NEON register usage

Readers unfamiliar with ARM floating-point may wish to refer to [Chapter 7 Floating-Point](#) before reading this section.

VFPv3 has 32 single-precision registers s0-s31, which can also be accessed in pairs as double-precision registers d0-d15. There are a further 16 double-precision registers, d16-d31. NEON can also view these as quadword registers q0-q15. Registers s16-s31 (d8-d15, q4-q7) must be preserved across subroutine calls; registers s0-s15 (d0-d7, q0-q3) do not need to be preserved (and can be used for passing arguments or returning results in standard procedure-call variants). Registers d16-d31 (q8-q15), do not need to be preserved.

The Procedure Call Standard specifies two ways in which floating-point parameters can be passed. For software floating-point, they will be passed using ARM registers R0-R3 and on the stack, if required. An alternative, where floating-point hardware exists in the processor, is to pass parameters in the VFP or NEON registers.

This hardware floating-point variant behaves in the following way:

- Integer arguments are treated in exactly the same way as in softfp. So, if we consider the function *f* below, we see that the 32-bit value *a* will be passed to the function in R0, and because the value *b* must be passed in an even or odd register pair, it will go into R2 or R3, leaving R1 unused.

```
void f(uint32_t a, uint64_t b)
    r0: a
    r1: unused
    r2: b[31:0]
    r3: b[63:32]
```

- FP arguments fill d0-d7 (or s0-s15), independently of any integer arguments. This means that integer arguments can flow onto the stack and FP arguments will still be slotted into FP registers (if there are enough available).
- FP arguments are able to back-fill, so it's less common to get the unused slots that we see in integer arguments. Consider the following examples:

```
void f(float a, double b)
    d0:
        s0: a
        s1: unused
    d1: b
```

Here, *b* is aligned automatically by being assigned to d1 (which occupies the same physical registers as VFP s2 or VFP s3).

```
void f( float a, double b, float c)
    d0:
        s0: a
        s1: c
    d1: b
```

In this example, the compiler is able to place *c* into s1; it does not need to be placed into s4.

In practice, this is implemented (and described) by using separate counters for s, d and q arguments, and the counters always point at the next available slot for that size. In the second FP example above, *a* is allocated first because it's first in the list, and it goes into first available s register, which is s0. Next, *b* is allocated into the first available d register, which is d1 because *a* is using part of d0. When *c* is allocated, the first available s register is s1. A subsequent double or single argument would go in d2 or s4, respectively.

- There is a further case when filling FP registers for arguments. When an argument must be spilled to the stack, no back-filling can occur, and stack slots are allocated in exactly the same way for further parameters as they are for integer arguments.

```
void f( double a, double b, double c, double d,
        double e, double f, float g, double h,
        double i, float j)
    d0: a
    d1: b
    d2: c
    d3: d
    d4: e
    d5: f
    d6:
        s12: g
        s13: unused
    d7: h
    *sp: i
    *sp+8: j
    *sp+12: unused (4 bytes of padding for 8-byte sp alignment)
```

Arguments a-f are allocated to d0-d5 as expected.

The single-precision `g` is allocated to `s12`, and `h` goes to `d7`.

The next argument, `i`, can't fit in registers, so it is stored on the stack. It would be interleaved with stacked integer arguments if there were any. However, while `s13` is still unused, `j` must go on the stack because we cannot back-fill to registers when FP arguments have hit the stack.

- Double-precision and quad-precision registers can also be used to hold vector data. This would not occur in typical C code.
- No VFP registers are used for variadic procedures, that is, a procedure which does not have a fixed number of arguments. They are instead treated as in `softfp`, in that they are passed in integer registers (or the stack). Note that single-precision variadic arguments are converted to doubles, as in `softfp`.

17.1.2 Linkage

If the platform has hardware support for NEON technology or an FPU, the highest performance is achieved by passing NEON and FPU parameters and return values in NEON and FPU registers. This is called *hardware floating-point linkage*. In some situations, using the general-purpose registers for parameter passing might be preferred, to simplify software compatibility between platforms with and without hardware floating-point support. This is called *software floating-point linkage*.

You cannot mix objects with different floating-point linkage in a single image. Any dynamic libraries loaded while the application is running must also use the same linkage.

Any system that supports both NEON and VFP instructions uses a common register bank for these instructions, therefore configuration options that affect the floating-point calling convention also affect how NEON parameters are passed and returned.

17.1.3 Stack and heap

The stack implementation is full-descending, with the current top of the stack pointed to by `R13` (`SP`). The stack must be contiguous in the virtual memory space. Detection of a stack overflow is usually handled with memory management. The stack must always be aligned to a word boundary, except at an external interface, when it must be double-word aligned.

The heap is an area (or areas) of memory that are managed by the process itself (for example, with the `C malloc()` function). It is typically used for the creation of dynamic data objects.

17.1.4 Returning results

A function which returns a `char`, `short`, `int` or single-precision float value will do so using `R0`. A function which returns a 64-bit value (a double-precision float or `long long`) does so in `R0` and `R1`. As mentioned in [VFP and NEON register usage on page 17-4](#), floating-point and NEON return values will be in `s0`, `d0`, or `q0` when using hardware linkage. If the software uses hardware linkage it will return floating-point values in `s0`, `d0`, or `q0`. If the software uses software linkage, it will return single-precision float in `r0`.

17.2 Mixing C and assembly code

One example of why it can be useful to understand the AAPCS is to write assembly code which is compatible with C code. One way to do this is write separate modules and assemble them with GNU Assembler. They can be defined as extern functions in C and called; provided the AAPCS rules are followed, there should be no problem.

We can also insert assembly code into our C code. Let's look at some inline or embedded assembly, through the GCC `asm` statement. This is very simple to use. For example, we can implement a NOP as shown in [Example 17-1](#).

Example 17-1 NOP

```
asm("nop");
```

The time taken to carry out a NOP is undefined.

In fact, it is likely that this NOP will have no effect, because the C compiler will optimize it away, or the processor will discard the instruction. When you include assembly language code with inline assembler, the resulting code is still subject to optimization by the C compiler. This is a very important point which must be taken into account whenever inline assembly is used. Even if the compiler does not optimize-away the NOP instruction, the processor itself may filter out the NOP from the instruction stream so that it never reaches the execute stage.

Inline assembly code has a different syntax to regular assembly code. Registers and constants have to be specified differently, if they refer to C expressions.

Let's look at a slightly more complicated example. We take the value of an `int` and use the `USAD8` assembly instruction to calculate the sum of difference of bytes and then store the result in a different `int`. [Example 17-2](#) shows the relevant code.

Example 17-2 Using the `USAD8` instruction

```
asm volatile ("usad8 %0, %1, %2" : "=r" (result): "r"(value1), "r"(value2));
```

The general format of such inline assembly code is:

```
asm volatile (assembler instructions : output operands (optional) : input operands
(optional) : clobbered registers (optional) );
```

The colons divide the statement up into parts. The first part `"usad8 %0, %1, %2"` is the actual assembly instruction. The second part is an (optional) list of output values from the sequence. If more than one output is needed, commas are used to separate the entries. We may then optionally have a list of input values for the sequence, with the same format as the output. If you don't specify an output operand for an assembly sequence, it is quite likely that the C compiler optimizer will decide that it is not serving any useful purpose and optimize it away! A way to avoid this is to use the `volatile` attribute, which tells GCC not to optimize the sequence.

In the actual assembly language statement, operands are referenced by a percent sign followed by the symbolic name in square brackets. The symbolic name references the item with the same name in either the input or output operand list. This name is completely distinct from any other symbol within your C code (although clearly it is less confusing to use symbols which do have a meaning within your code). Alternatively, the name can be omitted and the operand can be specified using a percent sign followed by a digit indicating the position of the operand in the list (that is, `%0, %1 ... %9`), as shown in the example.

There is an optional fourth part to an asm statement, called the *clobber* list. This enables us to specify to the compiler what will be changed by the assembly code. We can specify registers (for example, R0), the condition code flags (cc) or memory.

This makes the compiler store affected values before and reload them after executing the instruction.

The constraints mentioned when we talked about input and output operand lists relate to the fact that assembly language instructions have specific operand type requirements. When passing parameters to inline assembly statements, the compiler must know how they should be represented. For example, the constraint “r” specifies one of the registers R0-R15 in ARM state, while “m” is a memory address and “w” is a single precision floating-point register. These characters have an = placed before them to indicate a write-only output operand, a “+” for a read/write output operand (that is, one that is both input and output to the instruction). The “&” modifier instructs the compiler not to select any register for the output value, which is used for any of the input operands.

You can force the inline assembler to use a particular register to hold a local variable by using something like the code shown in [Example 17-3](#).

Example 17-3 Inline assembler local variable usage

```
void func (void) {
    register unsigned int regzero asm("r0");
```

and later

```
    asm volatile("rev r0, r0");
```

This usage can interfere with the compiler optimization and does not guarantee that the register will not be re-used, for example when the local variable is no longer referenced. Hard coding register usage is always bad practice. It is almost always better to use local variables instead.

[Example 17-4](#) gives a longer example of inline assembler, taken from the Linux kernel. It shows how a series of inline assembly language instructions can be used. The code manipulates the CPSR, to change modes. This would not be possible using C code.

Example 17-4 Inline assembler

```
void __naked get_fiq_regs(struct pt_regs *regs)
{
    register unsigned long tmp;
    asm volatile (
        "mov    ip, sp\n\
        stmfd  sp!, {fp, ip, lr, pc}\n\
        sub    fp, ip, #4\n\
        mrs    %0, cpsr\n\
        msr    cpsr_c, %2 @ select FIQ mode\n\
        mov    r0, r0\n\
        stmia  %1, {r8 - r14}\n\
        msr    cpsr_c, %0 @ return to SVC mode\n\
        mov    r0, r0\n\
        ldmfd  sp, {fp, sp, pc}"
        : "=r" (tmp)
        : "r" (&regs->ARM_r8), "I" (PSR_I_BIT | PSR_F_BIT | FIQ_MODE));
}
```

The ARM compiler tools have a similar concept, albeit with different syntax. In addition to inline assembly, they also support *embedded* assembly which is assembled separately from the C code and produces a compiled object which is combined with the object from the C compilation.

A global index gives unique data for each thread within a process. A thread allocates the index when the process starts. The other threads use this to retrieve the unique data associated with the index. When a thread starts, it allocates a block of dynamic memory and stores a *Thread Local Storage* (TLS) pointer to this memory. This can either be located in memory, or in ARM MPCore systems, a dedicated CP15 register.

Chapter 18

Profiling

Donald Knuth once famously observed that “Premature optimization is the root of all evil”. However, Knuth's comment does not argue against *appropriate* optimization.

Code optimization is an important part of the work of software engineers— modifying software so that it runs more quickly, uses less power and makes less use of memory or other resources. To do this, we must first identify which part (or parts) of the code should be optimized. It is worth noting, however, that programmer’s best guesses about where the time is spent are notoriously inaccurate - hence the importance of profiling.

Profiling is a technique that lets us identify functions (or other pieces of code) which consume large proportions of the total execution time. It is usually more productive to focus optimization efforts on code segments which are executed very frequently, or which take a long time to execute than to optimize rarely used functions or code which takes little time to execute. A profiler will tell us which parts of the code are frequently executed and which occupy the most processor cycles. A profiler can help us identify *bottlenecks*, situations where the performance of the system is constrained by a small number of functions (or just one). This data is collected using instrumentation, an execution trace or sampling, and can often be a good method for finding hidden bugs too.

When you have identified some slow part of your code it's important to consider first whether you can change the algorithm, before trying to improve the existing code. For example, if the time is being spent searching a linked list, it's probably much more beneficial to change to using a tree or hash table instead of spending effort to speed up the linked list search.

Profiling can be considered as a form of dynamic code analysis. Profiling tools can gather information in a number of different ways. We can distinguish two basic approaches to gathering information.

Time based sampling

Here, the state of the system is sampled at a periodic, time-based interval. The size of this interval can affect the results – a smaller sampling interval can increase execution time but produce more detailed data.

Event based sampling

Here, sampling is driven by occurrences of an event, which means that the time between sampling intervals is usually variable. Events can often be hardware related – for example, cache misses.

It is also important to understand that profilers typically operate on a statistical basis – they may not necessarily produce absolute counts of events. In complex systems, it may be necessary to control profiling information by use of annotation options to specify which events are to be recorded, which events shown, or thresholds to avoid displaying large numbers of functions with low count numbers.

18.1 Profiler output

Profiler tools normally provide two kinds of information:

Call graph The call graph tells us the number of times each function was called. This can help point out which function calls can be eliminated or replaced and shows inter-relations between different functions. Viewing a call graph can suggest code to optimize and reveal hidden bugs (for example, if code is unexpectedly calling an error function many times). Collecting call graph information can require building the code with special options.

Flat profile A flat profile, as in [Example 18-1](#) shows how much processor time each function uses and the number of times it was called. This enables a simple identification of which functions consume large fractions of run-time and should therefore be considered first for possible optimizations.

Example 18-1 Example flat profile

Flat profile:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
33.34	0.02	0.02	6275	0.00	0.00	start
16.67	0.03	0.01	192	0.07	0.21	func1
16.67	0.04	0.01	15	1.20	1.20	memcpy
16.67	0.05	0.01	7	1.41	1.41	write

It may be useful at this point to consider some example profiling tools which can be used in Cortex-A series processors.

18.1.1 Gprof

GProf is a GNU tool which provides an easy way to profile your C/C++ application and find the locations that need work.

Using GCC, you can generate profile information by compiling with a special flags. The source code has to be compiled with the `-pg` option. For line-by-line profiling, the `-g` option would also be needed. You then execute the compiled program and the profiling data is collected. You then run `gprof` on the resulting statistics file to view the data, in a variety of convenient formats.

When you compile with `-pg` the compiler adds profiling instrumentation that collects data at function entry and exits at runtime. It therefore profiles only the user application code and not anything that happens in code that has not been built for profiling (for example, `libc`) or the kernel. Gprof can give misleading results if the performance limitations of the code come from kernel or I/O issues (memory fragmentation or file accesses for example).

It may be necessary to remove GCC optimization flags, as some compiler optimizations can cause problems while profiling. It is also the case that the use of the profiling flags will actually slow the program down. This can be an important consideration in some types of real-time system, where it may be that the interaction of real-time events has a significant effect on the performance of the profiled code. A binary file called `gmon.out` containing profiling information is generated. This file can then be operated on by the `gprof` tool.

18.1.2 OProfile

OProfile is a whole system profiling tool which runs on Linux and includes the kernel in its metrics. Unlike gprof, it works using a statistical sampling method. OProfile can examine the system at regular intervals, determine what code is running, and update appropriate counters. If a long enough profile is taken, with a sufficient sample rate, an accurate picture of the execution is obtained. Like other profilers which make use of interrupts, code which disables interrupts can cause inaccuracies. For this reason, the Linux function `spinlock_irq_restore()`, which re-enables interrupts after a spinlock has been relinquished can erroneously appear to be a major system bottleneck, as the time for which interrupts were disabled can be counted against it. OProfile can also be made to trigger on hardware events and will record all system activity including kernel and library code execution.

OProfile does not need code to be recompiled with any special flags (provided symbol information is available). It provides useful hardware information about such things as clock cycles and cache misses. Call graphs are statistically generated, so may not be completely accurate.

18.1.3 DS-5 Streamline

DS-5 Streamline is a graphical performance analysis tool, which can be used to analyze the performance of an Linux or Android system. It is a component of ARM DS-5 and combines a kernel driver, target daemon and Eclipse-based user interface. (For more information about DS-5, see [ARM DS-5 on page 3-11](#).)

DS-5 Streamline takes sampling data and system trace information and produces reports that present the data visually and in a statistical form. It uses hardware performance counters and kernel metrics to provide an accurate representation of system resources. DS-5 Streamline allows the application source code to be annotated to augment its graphical display with additional textual or visual information, such as labelled timing markers or screen shots from the target.

For example, DS-5 Streamline has a display mode called X-ray mode. In this mode the process trace view changes from an intensity map of time to a view that highlights processor affinity. [Figure 18-1 on page 18-5](#) shows the DS-5 Streamline view of a multi-threaded Linux application running on a multi-core processor. Threads are being assigned to the processors by the Linux kernel:

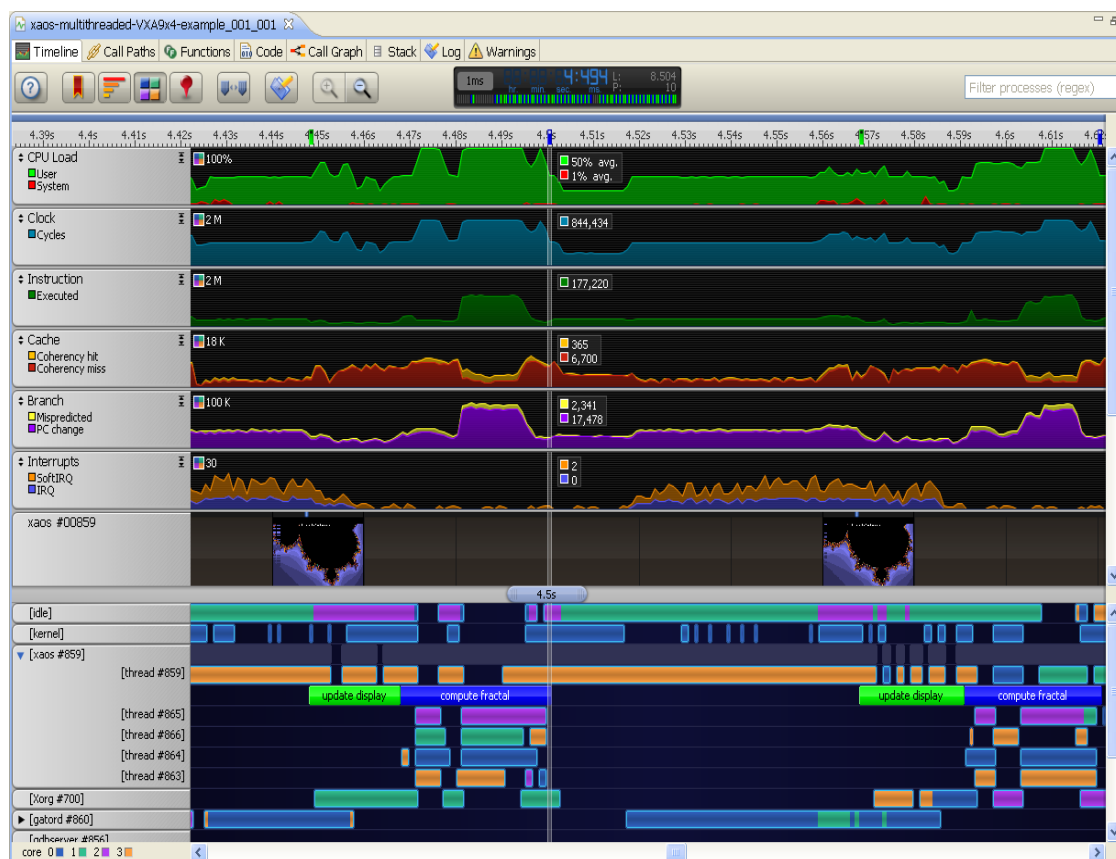


Figure 18-1 X-ray mode in ARM DS-5 Streamline

18.1.4 ARM performance monitor

The performance monitor hardware is able to count several events, using multiple counters. Normally, we combine together multiple values to generate useful parameters to optimize. For example, we can choose to count the total number of clock cycles and the number of instructions executed and use this to derive a cycles per instruction figure which is a useful proxy for the efficiency with which the processor is operating. We can generate information about cache hit or miss rates (separately for both L1 data and instruction caches) and examine how code changes can affect these.

Cortex-A series processors contain event counting hardware which can be used to profile and benchmark code, including generation of cycle and instruction count figures and to derive figures for cache misses and so forth. The performance counter block contains a cycle counter which can count processor cycles, or be configured to count every 64 cycles. There are also a number of configurable 32-bit wide event counters which can be set to count instances of events from a wide-ranging list (for example, instructions executed, or MMU TLB misses). These counters can be accessed through debug tools, or by software running on the processor, through the CP15 *Performance Monitoring Unit* (PMU) registers. They provide a non-invasive debug feature and do not change the behavior of the processor. CP15 also provides a number of controls for enabling and resetting the counters and to indicate overflows (there is an option to generate an interrupt on a counter overflow). The cycle counter can be enabled independently of the event counters.

It is important to understand that information generated by such counters may not be exact. In a superscalar, out-of-order processor, for example, it can be difficult to guarantee that the number of instructions executed is precise at the time any other counter is updated.

The standard countable events, common to all ARMv7-A processors are listed in [Table 18-1](#). The TRM for the specific processor being used provides further information on the lists of events which can be monitored, which can include a large number of additional possibilities in addition to those listed here.

Table 18-1 Performance monitor events

Number	Event counted
0x00	Software increment of the Software Increment Register
0x01	Instruction fetch that causes a Level 1 instruction cache refill
0x02	Instruction fetch that causes a Level 1 instruction TLB refill
0x03	Memory Read or Write operation that causes a Level 1 instruction TLB refill
0x04	Memory Read or Write operation that causes a Level 1 data cache access
0x05	Memory Read or Write operation that causes a Level 1 data TLB refill
0x06	Memory-reading instruction executed
0x07	Memory-writing instruction executed
0x09	Exception taken
0x0A	Exception return executed
0x0B	Instruction that writes to the Context ID register
0x0C	Software change of program counter
0x0D	Immediate branch instruction executed
0x0F	Unaligned load or store
0x10	Branch mispredicted or not predicted
0x11	Cycle count; the register is incremented on every cycle
0x12	Predictable branch speculatively executed
0x13	Data memory access
0x14	Level 1 instruction cache access
0x15	Level 1 data cache write-back
0x16	Level 1 data cache write-back
0x17	Level 2 data cache refill
0x18	Level 2 data cache write-back
0x19	Bus access
0x1A	Local memory error
0x1B	Instruction speculatively executed
0x1C	Instruction write to TTBR
0x1D	Bus cycle
0x1E-0x3F	Reserved

MPCore processors include a significant number of additional events relating to SMP capabilities, described in [Chapter 22 SMP Architectural Considerations](#). These include numbers of barrier instructions executed, measures of coherent cache activity and counts of exclusive access failures.

In Linux, these counters are normally accessed through the kernel via the Linux OProfile tool or the Linux perf events framework. However, User mode access to these counters can be enabled for direct access if required.

18.1.5 Linux perf events

Linux contains patches (as of Linux-2.6.34) to support the Ingo Molnar Perf Events framework. It handles multiple events, including processor cycles and cache-misses but also measures kernel events like context switches. It provides simple, platform independent access to the ARM performance counter registers and also to software events. Through this framework, tools can produce graphs and statistics about such things as function calls. You can also record execution traces, profile on a per-processor, per-application and per-thread basis and generate summaries of events.

18.1.6 Ftrace

Ftrace is an increasingly widely used trace tool on Linux. It provides visualization of the flow within the kernel by tracing each function call. It enables interactions between parts of the kernel to be viewed and enables developers to investigate potential problem paths where interrupts or pre-emption are disabled for long periods.

18.1.7 Valgrind and Cachegrind

Valgrind is a widely used tool, commonly used for detection of memory leaks and other memory handling problems; however, it can also be used for profiling memory usage and is potentially able to give much more detailed results than OProfile. Valgrind translates a program into a processor-neutral intermediate representation. Other tools associated with Valgrind are able to operate on this and then Valgrind translates the code back into native machine code. This process makes the code run an order of magnitude slower, but enables checks for accesses to undefined memory, off-by-one errors and so forth to be performed by tools like Memcheck.

For memory access optimization, the Cachegrind tool can be used. As Valgrind simulates the execution of the program, we can use Cachegrind to record all uses of program memory. By simulating the operation of the processor caches, we can generate statistics about cache usage. Some care is needed – this is a simulation of the cache and it is possible that it does not represent the real cache hardware completely accurately. Nevertheless, it can be a very useful tool to examine cache and memory usage by an application. When writing high level applications, it can be difficult for the programmer to have much appreciation for (or control over) addresses used by a program. The linker will typically generate many of the virtual addresses used within an image, while the runtime loader will control positioning of libraries and so forth. Finally, the kernel is responsible for placement of code and data in physical memory. Memory profiling tools can therefore be a useful aid.

Chapter 19

Optimizing Code to Run on ARM Processors

In [Chapter 18](#), we looked at how to identify which parts of the program are suitable for optimization. In this chapter we will look at some possible software optimizations. Much of what is presented will apply equally to any system, but some of the suggested optimizations are specific to ARM systems, or indeed to particular ARM processors.

Optimization does not necessarily mean optimizing to make programs faster. In embedded systems, we may prefer to optimize for battery life, code density or memory footprint, for example. Writing code which is more efficient delivers not only higher levels of performance, but can also be crucial in conserving battery life. If we can get a job done faster, in fewer cycles, we can turn off the power for longer periods.

Many compilers provide options to help with this. For example, both the ARM Compiler and the GNU GCC compiler have `-O` flags to specify the level of optimization. We will spend some time looking at the effects of these and how to trade off optimizations of code density, speed and debug visibility.

We will continue by looking at ways in which you can modify your program in order to generate more optimized code. In a book of this size, we cannot cover the areas of algorithm design which can lead to significant performance gains. We will instead concentrate our efforts on areas where an understanding of the underlying ARM architecture can help to optimize code.

Although cycle timing information can be found in the *Technical Reference Manual* (TRM) for the processor that you are using, it is very difficult to work out how many cycles even a trivial piece of code will take to execute. The movement of instructions through the pipeline is dependent on the progress of the surrounding instructions and can be significantly affected by memory system activity. Pending loads or instruction fetches which miss in the cache can stall code for tens of cycles. Standard data processing instructions (logical and arithmetic) will take only one or two

cycles to execute, but this does not give the full picture. Instead, we must use profiling tools, or the system performance monitor built-in to the processor, to extract useful information about performance.

19.1 Compiler optimizations

The ARM Compiler and GNU GCC give us a wide range of options which aim to increase the speed, or reduce the size, of the executable files it generates. For each line in the source code there are generally many possible choices of assembly instructions that could be used. The compiler must trade-off a number of resources, such as registers, stack and heap space, code size (number of instructions), compilation time, ease of debug, and number of cycles per instruction in order to produce the best image file.

We'll begin by looking at some of the source-level optimization options available to the compiler.

19.1.1 Function inlining

When a function is called, there is a certain overhead. A caller function must store its own return address on the stack. Instructions may be required to place arguments in the appropriate registers and push registers on the stack, in accordance with the Procedure Call Standard. There is a corresponding overhead when returning to the original point of execution when the function ends, again requiring a branch (and corresponding instruction pipeline flush) and possibly popping registers from the stack. This function-call overhead can become important when there are functions which contain only a few instructions, and where these functions represent a significant amount of the total run-time. Also, executing branches uses branch predictor resources, which can affect overall program performance. Function inlining eliminates this overhead by replacing calls to a function by a copy of the actual code of the function itself (known as placing the code *inline*).

Inlining for critical code paths is always a worthwhile optimization if there is only one place where the function is called. It is always worthwhile if calling the function requires more instructions (memory) than inlining the function body. A further consideration is that inlining can help permit other optimizations. Clearly, increasing the number of times that a function is called will increase the number of inlined copies of the function that are made and this will increase the cost in code size.

GCC performs inlining only within each compilation unit. The `inline` keyword can be used to request that a specific function should be inlined wherever possible, even in other files. The GCC documentation gives more details of this (and how its use can be combined with `static` and `extern`).

We will look at inlining in a little more detail when we consider cache optimizations.

19.1.2 Eliminating common sub-expressions

Another simple source-level optimization is re-using already computed results in a later expression. This *common sub-expression elimination* is performed automatically when optimization command line switches are used and can make code both smaller and faster. However, the compiler may not necessarily catch all cases, and it can sometimes be more useful to do this by hand.

[Example 19-1](#) illustrates how this works:

Example 19-1 Common sub-expression

```
i = a * b + c;
j = a * b * d;
```

The compiler can treat this code as if it had been written as in [Example 19-2](#).

Example 19-2 Common sub-expression elimination

```
tmp = a * b;
i = tmp + c;
j = tmp * d;
```

This reduces both the instruction count and cycle count.

19.1.3 Loop unrolling

Every iteration of a loop has a certain penalty associated with it. Every conditional loop must include a test for the end of loop on each iteration. Furthermore, there is a branch instruction to iterate over the loop, which can take a number of cycles to execute. We can avoid this penalty by unrolling loops, partially or fully.

Consider the simple code shown in [Example 19-4](#), to initialize an array.

Example 19-3 Initializing an array

```
for (i = 0; i < 10; i++)
{
    x[i] = i;
}
```

Each iteration of the loop contains an assembler sequence of the form in [Example 19-4](#).

Example 19-4 Loop termination assembly code

```
CMP i, #10
BLT for_loop
```

A large proportion of the total run time will have been spent checking if the loop has terminated and in executing a branch to re-execute the loop.

The same code can be written by *unrolling the loop*, as shown in [Example 19-5](#).

Example 19-5 Unrolled loop

```
x[0] = 0;
x[1] = 1;
x[2] = 2;
```

```

x[3] = 3;
x[4] = 4;
x[5] = 5;
x[6] = 6;
x[7] = 7;
x[8] = 8;
x[9] = 9;

```

When the code is written in this way, we remove the compare and branch instruction and have a sequence of stores and adds. This is clearly larger than the original code but can execute considerably faster.

Conventionally, loop unrolling is often considered to increase the speed of the program but at the expense of an increase in code size (except for very short loops). However, in practice this may not always be the case on many hardware platforms. In many systems, an access to external memory takes significant numbers of cycles and an instruction cache is provided. Code which loops will typically fit into the cache very well. The code is fetched into the cache during the first loop iteration and is executed directly from cache after that. Unrolling the loop can mean that the code is executed only once and so does not cache so well. This is more likely to be the case for functions which are executed only once – loops which are frequently executed may well be cached whether they are unrolled or not. A further consideration is that modern ARM processors typically include branch prediction logic which can hide the effect of pipeline flushes from the programmer by speculatively predicting whether a branch will or will not be taken ahead of the actual evaluation of a condition. In some cases, the branch instruction can be folded, so that it does not require an actual processor cycle to execute.

Cortex-A series processors can have long, complex instruction pipelines, with interdependencies between instructions, particularly loads and instructions which set condition code flags. The compiler understands the rules associated with a particular processor and can often re-arrange instructions so that pipeline interlocks are avoided. This is called *scheduling* and typically involves re-arranging the order of instructions in ways which do not alter the logical correctness of the program or its size, but which reduce its execution time. This can significantly increase the compiler effort, increasing both the time and memory required for the compilation. It can also restrict the ability to perform source level debug. There may no longer be a strict one-to-one link between a line of C source and a sequence of assembly instructions. We can instead have a couple of instructions from a C statement followed by instructions for the next statement and then some more instructions for the first statement.

19.1.4 GCC optimization options

GCC has a range of optimization levels, plus individual options to enable or disable particular optimizations.

The overall compiler optimization level is controlled by the command line option `-O n` , where n is the required optimization level, as follows:

- `-O0`. (default). No optimization is performed. Each source code command relates directly to the corresponding instructions in the executable file. This gives the clearest view for source level debugging.
- `-O1`. This enables most common forms of optimization that requires no size versus speed decisions, including function inlining. It can often actually produce a faster compile than `-O0`, due to the resulting files being smaller.
- `-O2`. This enables further optimizations, such as instruction scheduling. Again, optimizations which can have speed versus size implications will not be used.

- `-O3`. This enables further optimizations, such as aggressive function inlining and can therefore increase the speed at the expense of image size. Furthermore, this option enables `-ftree-vectorize` - causing the compiler to attempt to automatically generate NEON code from standard C or C++. See [Chapter 20 Writing NEON Code](#).
- `-funroll-loops`. This option is independent of the above and enables loop unrolling. As previously described, loop unrolling can increase code size and may not have a beneficial effect in all cases.
- `-Os`. This selects optimizations which attempt to minimize the size of the image, even at the expense of speed.

Higher levels of optimization can restrict debug visibility and increase compile times. It is usual to use `-O0` for debugging, and `-O2` for finished code. When using the above optimization options with the `-g` (debug) switch, it can be difficult to see what is happening. The optimizations can change the order of statements or remove (or add) temporary variables among other things. But an understanding of the kinds of things the compiler will do means that satisfactory debug is normally still possible with `-O2 -g`.

For optimal code, it is important to specify to the compiler as much detailed information about the target platform as practically possible.

For the compiler to optimize for a specific processor or instruction set version, you also need to provide some information about your target platform. Many useful options are documented on <http://gcc.gnu.org/onlinedocs/gcc/ARM-Options.html>.

The main platform-specifying parameters are:

`-march=<arch>`

where `<arch>` is the architecture version to compile for. This defines the instruction set supported. It can make a significant difference to performance to specify `-march=armv7-a` if this is supported by your platform but is not used by default by your compiler.

`-mcpu=<cpu>`

More specific than `-march`, `-mcpu` specifies which processor to optimize for, including scheduling instructions in the way most efficient for that processor's pipeline.

`-mtune=<cpu>`

This option provides processor specific tuning options for code, even when only an architecture version is specified on the command line. For instance, the command line might contain `-march=armv5te -mtune=cortex-a8`. This selects instructions for the architecture ARMv5TE but tunes the selected instructions for execution on a Cortex-A8 processor.

`-mfpu=<fpu>`

If your target platform supports hardware floating-point or NEON, specify this to ensure that the compiler can make use of these instructions. For a Cortex-A5 target, you would specify `-mfpu=neon-vfpv4`.

`-mfloat-abi=<name>`

This option specifies the floating-point ABI to use. Values for `<name>` are:

- | | |
|---------------------|---|
| <code>soft</code> | causes GCC to generate code containing calls to the software floating-point library for floating-point operations. |
| <code>softfp</code> | enables GCC to generate code containing hardware floating-point instructions, but still uses the software floating-point linkage. |

hard enables GCC to generate code containing hardware floating-point instructions and uses FPU-specific hardware floating-point linkage.

The default depends on the target configuration. You must compile your entire program with the same ABI, and link with a compatible set of libraries.

Table 19-1 shows a few examples of code generation for floating-point operations.

Table 19-1 Floating-point code generation

-mfpu	-mfloat-abi	Resultant code
Any value	soft	Floating-point emulation using software floating-point library
vfpv3	softfp	VFPv3 floating-point code
vfpv3-d16	softfp	VFPv3 floating-point code
neon	hard	VFPv3 and Advanced SIMD code, where the floating-point and SIMD types use the hardware FP registers

19.1.5 armcc optimization options

The armcc compiler enables you to compile your C and C++ code. It is an optimizing compiler with a range of command-line options to enable you to control the level of optimization.

The command line option gives a choice of optimization levels, as follows:

- **-Ospace.** This option instructs the compiler to perform optimizations to reduce image size at the expense of a possible increase in execution time.
- **-Otime.** This option instructs the compiler to perform optimizations to reduce execution time at the expense of a possible increase in image size.
- **-O0.** Turns off most optimizations. It gives the best possible debug view and the lowest level of optimization.
- **-O1.** Removes unused inline functions and unused static functions. Turns off optimizations that seriously degrade the debug view. If used with `--debug`, this option gives a satisfactory debug view with good code density.
- **-O2 (default).** High optimization. If used with `--debug`, the debug view might be less satisfactory because the mapping of object code to source code is not always clear.
- **-O3.** performs the same optimizations as -O2 however the balance between space and time optimizations in the generated code is more heavily weighted towards space or time compared with -O2. That is:
 - **-O3 -Otime** aims to produce faster code than -O2 -Otime, at the risk of increasing your image size
 - **-O3 -Ospace** aims to produce smaller code than -O2 -Ospace, but performance might be degraded.

19.2 ARM memory system optimization

Writing code which is optimal for the system it will run on is a key part of the art of programming. It requires the programmer to understand how the compiler and underlying hardware will carry out the tasks described in the lines of code. If we can do the job with less access to external memory, we can save power by keeping everything on-chip. Furthermore, by accessing the external memory less frequently, we improve the performance of the system, allowing software to run faster (or the processor to be clocked more slowly or for shorter periods, to save power).

19.2.1 Data cache optimization

In most Cortex-A series processors, there is a significant gap in performance between memory accesses which hit in the cache and those that do not. Cache misses can take tens of cycles to resolve. Cache hits return data within a few cycles and the compiler can often schedule instructions in a way that hides latency. For most algorithms, therefore, ensuring that cache misses are minimized is the most important possible optimization. The most important improvements are those which affect the level 1 cache.

Let us consider the problem of data cache misses first. Optimization is particularly significant for pieces of code which use a dataset larger than the available cache size. It is important for the programmer of such code to understand the arrangement of data in memory and how that corresponds to data cache accesses. Code should be structured in a way which ensures maximum re-use of data already loaded into the cache. It is this principle of *data locality*, the degree to which accesses to the same cache line are concentrated during program execution, in both space and time, which gives best performance.

Several techniques to improve this locality can be considered.

19.2.2 Loop tiling

Loop tiling divides loop iterations into smaller pieces, in a way which promotes data cache re-use. Large arrays are divided into smaller blocks (tiles), which match the accessed array elements to the cache size. The classic example to illustrate this approach is a large matrix vector product.

Consider two square matrices *a* and *b*, each of size 1024×1024 . [Example 19-6](#) shows code to compute a matrix vector product. This requires us to multiply each element in each array with each element in the other array.

Example 19-6 Matrix vector product code

```
for (i = 0; i < 1024; i++)
  for (j = 0; j < 1024; j++)
    for (k = 0; k < 1024; k++)
      result[i][j] = result[i][j] + a[i][k] * b[k][j];
```

In this case, the contents of matrix *a* are accessed sequentially, but matrix *b* advances in the inner loop, by row. It is therefore, highly probable that we will encounter a cache miss for each multiply operation.

It is obvious that the order in which the additions for each element of the result matrix are calculated does not change the result, ignoring the effect of such things as overflows. It should therefore be possible to rewrite the code in a way which improves the cache hit rate. The elements of matrix *b* are accessed in the following way (0,0), (1,0), (2,0)... (1023, 0), (0,1), (1,1)... (1023,1). The elements are stored in memory in the order (0,0), (0,1) etc. For word sized

elements, it means that the elements (0,0), (0,1)...(0,7) will be stored in the same cache line. (For simplicity, we will assume that the start address of the matrix is aligned to a cache line. Alignment will be mentioned again in [Structure alignment on page 19-10](#).) Therefore, elements (0,0), (0,1), (0,2) etc. will be in the same cache line; when we load (0,0) into the cache, we get (0,1...7) too. By the time the inner loop completes, it is likely that this cache line will be evicted.

If we modify the code so that two (or indeed four, or eight) iterations of the middle loop are performed at once while executing the inner loop, as in [Example 19-7](#) we can make a big improvement. Similarly, we should unroll the outer loop two (or four, or eight) times as well.

Example 19-7 Code using tiles

```
for (io = 0; io < 1024; io += 8)
  for (jo = 0; jo < 1024; jo += 8)
    for (ko = 0; ko < 1024; ko += 8)
      for (ii = 0, rresult = &result[io][jo],
           ra = &a[io][ko]; ii < 8;
           ii++, rresult += 1024, ra += 1024)
        for (ki = 0, rb = &b[ko][jo];
             ki < 8; ki++, rb += 1024)
          for (ji = 0; ji < 8; ji++)
            rresult[ji] += ra[ki] * rb[ji];
```

There are now six nested loops. The outer loops iterate with steps of 8, representing the fact that eight int sized elements are stored in each line of the level 1 cache. Some further optimizations have also been introduced. The order of ji and ki has been reversed as only one expression uses ki, but two use ji. In addition, we optimize by removing common expressions from the inner loops. All pointer accesses are potential sources of aliasing in C, so by using `result`, `ra` and `rb` to access array elements, the array indexing is speeded up. We cover this in more detail in [Source code modifications on page 19-13](#). [Figure 19-1](#) illustrates the changing cache access pattern which results from changes to the C code.




Figure 19-1 Effect of tiling on cache usage

19.2.3 Loop interchange

In many programs, there will be nested loops – a very simple example would be code that stepped through the items in a 2-dimensional array. For reasonably complex code, we can sometimes get better performance by re-arrangement of the loops. It is better to have the loop with the smaller number of iterations as the outer loop and the one with the highest iteration count as the innermost loop.

This gives two potential advantages. One is that the compiler can potentially unroll the inner loop. Perhaps more importantly for complex loops where the size of the nested loop is sufficiently large that it may not all be held in level 1 cache at the same time, the overall cache hit rate will be improved by this change. Some compilers can make this change automatically at higher levels of optimization. For example, GCC 4.4 adds the switch `-floop-interchange` to do this.

19.2.4 Structure alignment

Efficient placement of structure elements and alignment are not the only aspects of data structures which influence cache efficiency. Where code has a large working set, it is important to make efficient use of the available cache space. To achieve this, it might be necessary to rearrange data structures.

It is common to have data structures which span multiple cache lines, but where the program uses only a few parts of the structure at any particular time. If there are many objects of this type, it can make sense to try to split the structure so that it fits within a cache line. For example, you can split an array of structures into two or more arrays of smaller structures. This only makes sense if the object itself is aligned to a cache boundary. For example, consider the case where we have a very large array of instances of a 64-byte structure (much larger than the cache size). Within that structure, we have a byte-sized quantity and we have a commonly used function which iterates through the array looking only at that byte-sized quantity. This function would probably make inefficient use of the cache, as we would need to load an entire cache line to read the 8-bit value. If instead those 8-bit values were stored in their own array (rather than as part of a larger structure), we would get 32 or 64 values per cache linefill.

Furthermore, as we saw in [Chapter 16 Porting](#), unaligned accesses are supported, but can take extra cycles in comparison to aligned accesses. For performance reasons, therefore, it can be sensible to remove or reduce unaligned accesses.

19.2.5 Associativity effects

As we have seen, ARM L1 caches are normally 4-way set-associative, but L2 caches typically have 8- or 16-way associativity. There can be performance problems if more than four of the locations in the data fall into the same cache set, as there can be repeated cache misses, even though other parts of the cache can be unused. The ARM L1 Cache uses physical rather than virtual addresses, so it can be difficult for programmers operating in User mode to take care of this.

A particularly common cause of this problem is arranging data so that it is on boundaries of powers of two. If the cache size is 16KB, each way is 4KB in size. If you have multiple blocks of data arranged on boundaries which are multiples of 4KB, the first access to each block will go into line 0 of a way. If code accesses the first line in several such blocks then we can get cache misses even if only five cache lines in total are being used. Unaligned accesses can increase the likelihood of this, as each access might require two cache lines rather than one.

19.2.6 Optimizing instruction cache usage

The C programmer does not directly have control over how the instruction cache is used by code. Code is linear between branch instructions and this pattern of sequential accesses uses the cache efficiently. The branch prediction logic of the processor will try to minimize the stalls due to branches, so there is little the programmer can do to assist. The main goal for the programmer is to reduce the code footprint. Many of the compiler optimizations enabled at `-O2` and `-O3` for the ARM Compiler and GCC deal with loop optimizations and function inlining. These optimizations will improve performance if the code accounts for a significant part of the total program execution. In particular, function inlining has multiple potential benefits. Obviously, it

can reduce branch penalties by removing branches on both function call and exit, and potentially also stack usage. Equally importantly, it enables the compiler to optimize over a larger block of code which can lead to better optimizations for value range propagation and elimination of unused code.

However, modifications intended for speed optimizations which increase code size can actually reduce performance due to cache issues. Larger code is less likely to fit in the L1 cache (or indeed the L2 cache) and the performance lost by the additional cache linefills can well outweigh any benefits of the optimization. It is often better to use the `armcc -Ospace` or `gcc -Os` option to optimize for code density rather than speed. Clearly, using Thumb code will also improve code density and cache efficiency.

There are some interesting decisions to be made around function inlining and in some cases human judgment can improve on that of the compiler. A function which is only ever called from one place will always give a benefit if inlined. One might think that inlining very small functions always gives a benefit, but this is not the case. An instance of a tiny function which is called from many places is likely to be re-used many times within the instruction cache. If the same function is repeatedly inlined, it is much more likely that it will cause a cache miss (and also evict other potentially useful code from the cache). The branch prediction logic within Cortex-A series processors is efficient and an unconditional function call and return consumes few cycles – much less than would be used for a cache linefill. The programmer may wish to use the GCC function attributes `noinline` and `always_inline` to control such cases.

This is a general problem and not specific to inlining functions. Whenever conditional execution is used and it is lopsided (that is, the expression far more often leads to one result than the other) there is the potential for false static branch prediction and thus bubbles in the pipeline. It is usually better to order conditional blocks so that the often-executed code is linear, while the less commonly executed code has to be branched to and does not get pre-fetched unless it is actually used. The GCC attribute `__builtin_expect` used with the `-freorder-blocks` optimization option can help with this.

The performance monitor block of the processor (and OProfile) can be used to measure branch prediction rates in code. Note that there are two effects at play here. Correct branch prediction saves clock cycles by avoiding pipeline flushes, but taking fewer conditional branches which skip forward over code can help performance by making more of the program fit within the L1 cache.

19.2.7 Optimizing L2 and outer cache usage

Everything said about optimizations for using the L1 cache also applies to the L2 cache accesses. Best performance results from having a working dataset which is smaller than the L2 cache size, and where the data is used more than once; there is little benefit caching data which is used only once, other than possibly producing more optimal bus accesses. If the dataset is larger than the cache size, the programmer should consider similar techniques to those already described for the L1 cache. There is, however, a further point to consider with outer caches, which is that they may well be shared with other processors and therefore the effective size for an individual processor can be less than the actual size. In addition, when writing generic code to run on a number of ARM families, it can be difficult to make optimal use of the L2 cache. The presence of such a cache is not guaranteed and its size can vary significantly between systems.

19.2.8 Optimizing TLB usage

In general, the scope for optimizing usage of the Translation Lookaside Buffer (see [Chapter 10](#)) is much less than for optimizing cache accesses. The key points are to minimize the number of pages in use (this obviously gives fewer TLB misses) and to use large MMU mappings (supersections or sections in preference to 4KB pages) as this reduces the cost of individual page

table walks (one external memory access rather than two) and also means that a larger amount of memory is represented within an individual TLB entry (thus also giving fewer TLB misses). In practice, however, an operating system like Linux uses 4KB pages everywhere, so the main optimization technique available is to separate the frequently accessed code and data from the infrequently accessed code and data (for example exception handling code can be moved to a different page) and trying to limit the number of frequently accessed pages to below the maximum number supported by the processor hardware. The main optimization would be to try to process multiple cache lines' worth of data per page, so that the L1 cache is the limiting factor rather than TLB entries.

19.2.9 Data abort optimization

As we saw in [Chapter 10 Memory Management Unit](#), in the context of Linux, data aborts will be generated by page faults on the first time that a memory page is accessed and again when the page is first written to. This means that the kernel abort handler is called to take appropriate action and there is a certain performance overhead to this. Simplistically, we can reduce this overhead by using fewer pages. Again, code optimizations which make code smaller will help, as will reducing the size of the data space.

19.2.10 Prefetching a memory block access

ARM Cortex-A series processors contain sophisticated cache systems, and support for speculation and out of order execution which can hide latencies associated with memory accesses. However, accesses to the external memory system are usually sufficiently slow that there will still be some penalty. If we can prefetch instructions or data into the cache before we need them, we can hide this latency.

Many ARM processors provide support for preloading of data, using the PLD instruction. The PLD instruction is a hint that enables the programmer to request that data is loaded to the data cache in advance of it actually being read or written by the application. The PLD operation may generate a cache linefill on a data cache miss, independent of load and store instruction execution, while the processor continues to execute other instructions. If supported and used correctly, PLD can significantly improve performance by hiding memory access latencies. There is also a PLI instruction, which enables the programmer to hint to the processor that an instruction load from a particular address is likely in the near future. This can cause the processor to preload the instructions to its cache.

In addition to this programmer-initiated prefetch, the processor may also support automatic data prefetching. Essentially, the processor looks out for a series of sequential accesses to memory. When it sees this, it automatically requests the following cache lines speculatively, in advance of the program actually using them.

In many systems, significant numbers of cycles are consumed initializing or moving blocks of memory, using the `memset()` or `memcpy()` functions. Optimized ARM libraries will typically implement such functions by using Store Multiple instructions, with each store aligned to a cache line boundary.

19.3 Source code modifications

We have seen how profiling tools allow us to identify code segments or functions which can benefit from optimization and how different compiler options can enable compiler optimizations to our code. We will now look at a variety of source code modifications which can yield faster or smaller code on the ARM.

19.3.1 Loop termination

For inner loops that have been identified by the profiler, it may be appropriate to have integer loop counters which end at 0 (zero), rather than start from 0 (zero). This is because a compare with zero comes for free with the ADD or SUB instruction used to update the loop counter, whereas a compare with a non-zero value will typically require an explicit CMP instruction.

Replace a loop which counts up to a terminating value:

```
for (i = 1; i <= total; i++)
```

with one which counts down to zero:

```
for (i = total; i != 0; i--)
```

This will remove a CMP instruction from each iteration of the loop.

It is also good practice to use `int` (32-bit) variables for loop counters. This is because the ARM is natively a 32-bit machine. Its ADD assembly language instruction operates on two 32-bit registers. If it does an ADD (or other data processing operation) with a smaller quantity, the compiler may need to insert additional instructions to handle overflow (see also [Variable selection on page 19-14](#)).

19.3.2 Loop fusion

This is one of a variety of other possible loop techniques which can be employed either by the programmer, or by an optimizing compiler. It essentially means merging loops which have the same iteration count and no interdependencies ([Example 19-8](#) and [Example 19-9](#)).

Example 19-8 Loop fusion

```
for (i = 0; i < 10; i++)
{
    x[i] = 1;
}
for (j = 0; j < 10; j++)
{
    y[j] = j;
}
```

It should be immediately apparent that this can be optimized to:

Example 19-9 Fused loops

```
for (i = 0; i < 10; i++)
{
    x[i] = 1;
    y[i] = i;
}
```

 }

It is worth mentioning that this approach can sometimes lead to a reduction in performance because of cache effects such as thrashing, depending on the cache associativity and the addresses of the data being accessed.

19.3.3 Reducing stack and heap usage

In general, it is a good idea to try to minimize memory usage by code. The ARM processor has a register set which provides a relatively limited set of resources for the compiler to keep variables in. When all registers are allocated with currently live variables, further variables will be spilled to the stack – causing memory operations and extra cycles for the code to execute. There are a number of ways available to the programmer to try to help. A key rule is to try to limit the number of live variables at any one time.

When we looked at the procedure call standard in [Chapter 17](#), it was stated that up to four parameters can be passed to a function. Further parameters are passed on the stack. It is therefore significantly more efficient to pass four or fewer parameters than to pass five or more. Of course, the ARM registers in question are 32-bits in size and therefore if we pass a 64-bit variable, it will take two of our four register slots. For similar reasons, recursive functions do not typically yield efficient processor register usage. Remember also that non-static C++ functions also consume one argument slot with the this pointer.

19.3.4 Variable selection

ARM integer registers are 32-bit sized and optimal code is therefore produced most readily when using 32-bit sized variables, as this avoids the need to provide extra code to deal with the case where a 32-bit result overflows an 8- or 16-bit sized variable.

Consider the following code:

```
unsigned int i, j, k;
i = j+k;
```

The compiler would typically emit assembly code similar to:

```
ADD R0, R1, R2
```

If these variables were instead short (16-bit) or char (8-bit), the compiler must ensure the result does not overflow the halfword or byte.

The same code might need to be as shown in [Example 19-10](#), for signed halfwords (shorts).

Example 19-10 Addition of 2 signed shorts (assembly code)

```
ADD    R0, R1, R2
SXTH   R0, R0
```

Or for unsigned halfwords as in [Example 19-11](#).

Example 19-11 Addition of 2 unsigned shorts (assembly code)

```
ADD    R0, R1, R2
BIC    R0, R0, #0x10000
```

This has the effect of clipping the result to the defined size.

Although the compiler can cope with such things as an incorrect type specification for a loop counter variable, it is generally best to use the correct type in the first place.

19.3.5 Pointer aliasing

If a function has two pointers *pa* and *pb*, with the same value, we say the pointers *alias* each other. This introduces constraints on the order of instruction execution. If two write accesses which alias occur in program order, they must happen in the same order on the processor and cannot be re-ordered. This is also the case for a write followed by a read, or a read followed by a write. Two read accesses to aliases are safe to re-order. As any function pointer argument could alias any other function pointer argument in C, the compiler must assume that memory regions accessed through these pointers can overlap, which prevents many possible optimizations. C++ enables more optimizations, as pointer arguments will not be treated as possible aliases if they point to different types.

C99 introduces the `restrict` keyword which specifies that a particular pointer argument does not alias any other. If you know that pointers do not overlap, using this keyword to give the compiler this information can yield significant improvements. However, misusing it can lead to incorrect program function. The `restrict` keyword qualifies the pointer and not the object being pointed to. This consideration is not specific to the ARM architecture and applies to all processors. When using GCC, you can enable the C99 standard by adding `-std=c99` to your compilation flags. In code that cannot be compiled with C99, use either `__restrict` or `__restrict__` to enable the keyword as a GCC extension.

Consider the following simple code sequence:

```
void foo(unsigned int *ptr1, unsigned int *ptr2, unsigned int *i)
{
    *ptr1 += *i;
    *ptr2 += *i;
}
```

The pointers could possibly refer to the same memory location and this causes the compiler to generate code which is less efficient. In this example, it must read the value `*i` from memory twice, once for each add, as it cannot be certain that changing the value of `*ptr1` does not also change the value of `*i`.

If the function is instead declared as:

```
void foo(unsigned int *restrict ptr1, unsigned int *restrict ptr2, unsigned int
*restrict i)
```

This means that the compiler can assume that the three pointers may not refer to the same location and optimize accordingly. The programmer must ensure that the pointers do not ever overlap.

19.3.6 Division and modulo

Not all ARM processors have hardware support for division. For these processors, C division typically calls a library routine which takes tens of cycles to run for divides of 32-bit integers,

Note

Division is slower than multiplication even in hardware. In performance-critical code it is almost always worth replacing it if possible. This should be done as a trade-off against code maintainability.

Where possible, divides should be avoided, or removed from loops. Division with a fixed divisor is faster than dividing two variable quantities. The compiler can replace a divide by a shift-multiply pair in this case. (A 32×32 multiply by fixed constant, then shift right to adjust the most significant word).

Modulo arithmetic is another case to be aware of, as this will also use division library routines.

The code

```
minutes = (minutes + 1) % 60;
```

will run significantly faster on machines with no hardware divide, if coded as

```
if (++minutes == 60) minutes=0;
```

which substitutes a two cycle add and compare in place of a call to a library function.

19.3.7 Extern data

Accessing external variables requires the processor to execute a series of load instructions to acquire the address of the variable through a base pointer and then read the actual variable value. If multiple variables are defined as members of a structure, they can share a base pointer, saving cycles and instructions. It is therefore good practice to define the variables inside the same struct.

19.3.8 Inline or embedded assembler

In some cases, it can be a worthwhile optimization to use assembly code, in addition to C. The general principle here is for the programmer to code in a high level language, use a profiler to determine which sections will produce the most benefit if optimized and then inspect the compiler-produced assembly code to look for possible improvements.

If a code section is identified as being a performance bottleneck, one should not immediately reach for the assembly language manual. Improvements to the algorithm should first be sought and then compiler optimizations tried before considering use of assembly code. Even then, it is often the case that poor performance is due to cache misses and memory access delays rather than the actual assembly code.

The ARM Compiler and GCC (and most other C compilers) uses the `-s` flag to tell the compiler to produce assembly code output. The `-fverbose-asm` command line option can also be useful in gcc. Interleaved source and assembler can be produced by the ARM Compiler with the `--interleave` option.

[Chapter 16 Porting](#), of this book gives further information about the use of an inline assembler.

19.3.9 Complex addressing modes

It is often better to avoid complex addressing modes. In cases where the address to be used for a load or store requires a complex calculation, dual-issue of instructions is not possible. Only the addressing mode which uses a base register plus an offset (specified either by a register or an immediate value) with an optional shift left by an immediate value of two is fast. Other, less commonly used, addressing modes can be executed more quickly by splitting into two instructions which might be dual-issued. For example, `MOV R2, R1 LSL#3; LDR R2, [R0, R2]` can be faster than `LDR R2, [R0, R1 LSL #3]`. `LDRH` and `LDRB` have no extra penalty, but `LDRSH` and `LDRSB` have a single cycle load-use penalty, but no early forwarding path and can therefore incur additional latency if a subsequent instruction uses the loaded value.

19.3.10 Unaligned access

Unaligned LDRs have an extra cycle penalty compared with aligned loads, but unaligned LDRs which cross cache-lines have many cycles of additional penalty. In general, stores are less likely to stall the system compared to loads. STRB and STRH have similar performance to STR, due to the merging write buffer. As there are four slots in the load/store unit, more than four consecutive pending loads will always cause a pipeline stall.

19.3.11 Linker optimizations

Some code optimizations can be performed at the link, rather than the compile stage of the build, for example, unused section elimination and linker feedback. Multi-file optimization can be carried out across multiple C files, and unused sections can be removed. Similarly, multi-file compilation enables the compiler to perform optimization across multiple files instead of on individual files.

Chapter 20

Writing NEON Code

Support for NEON is provided in the ARM Compiler toolchain and in GNU tools. When writing NEON assembly code with the GNU tools, you must use the `-mfpu=neon` option to specify that the hardware supports NEON. In the ARM Compiler, you specify a target which supports NEON, for example, by using `--cpu=Cortex-A7`.

20.1 NEON C Compiler and assembler

Code targeted at NEON hardware can be written in C or in assembly language and a range of tools and libraries is available to support this.

In many cases, it may be preferable to use NEON code within a larger C/C++ function, rather than in a separate file to be processed by the assembler. This can be done using the intrinsics, described in Section 20.1.3.

20.1.1 Vectorization

A vectorizing compiler can take your C or C++ source code and parallelize it in a way which enables efficient usage of NEON hardware. This means you can write portable C code, while still obtaining the levels of performance made possible by NEON. The C language does not specify parallelizing behavior, so it can be necessary to provide hints to the compiler about this. For example, it may be necessary to use the `__restrict` keyword when defining pointers. This has the effect of guaranteeing that pointers will not overlap regions of memory. It can also be necessary to ensure that the number of loop iterations is a multiple of four or eight. Automatic vectorization is specified with the GCC option `-ftree-vectorize` (along with `-mfpu=neon`). Using the ARM Compiler, you must specify optimization level `-O2` (or `-O3`), `-Otime` and `--vectorize`.

20.1.2 NEON libraries

There is much free open source software which makes use of NEON, for example:

- OpenMAX, a set of APIs for processing audio, video, and still images. It is part of a standard created by the Khronos group. There is a free ARM implementation of the OpenMAX DL layer for NEON at <http://www.khronos.org/openmax/>
- ffmpeg, a collection of codecs for many different audio and video standards under LGPL license at <http://ffmpeg.org/>
- Eigen3, a linear algebra, matrix math C++ template library at eigen.tuxfamily.org/
- Pixman, a 2D graphics library (part of Cairo graphics) at <http://pixman.org/>
- x264, a rights-free GPL H.264 video encoder at <http://www.videolan.org/developers/x264.html>
- Math-neon at <http://code.google.com/p/math-neon/>

20.1.3 Intrinsics

NEON C/C++ intrinsics are available in `armcc`, `GCC/g++`, and `llvm`. They use the same syntax, so source code that uses intrinsics can be compiled by any of these compilers.

NEON intrinsics are a way to write NEON code that is more easily maintained than NEON assembler instructions, while still keeping control of the NEON instructions which are generated. There are new data types that correspond to NEON registers (both D-registers and Q-registers) containing different sized elements, allowing C variables to be created that map directly onto NEON registers. These variables are passed to NEON intrinsic functions. The compiler will generate NEON instructions directly instead of performing an actual subroutine call.

NEON intrinsics provide low-level access to NEON instructions but with the compiler doing some of the hard work normally associated with writing assembly language, such as:

- Register allocation.

- Code scheduling, or re-ordering instructions to get highest performance. The C compilers can be told which processor is being targeted, and they can reorder code to ensure the minimum number of stalls.

The main disadvantage with intrinsics is that it is not possible to get the compiler to output exactly the code you want, so there is still some possibility of improvement when moving to NEON assembler code.

20.1.4 NEON types in C

The *ARM C Language Extensions* document contains a full list of NEON types. The format is:

<basic type>x<number of elements>_t

Table 20-1 NEON type definitions

64-bit type (D-register)	128-bit type (Q-register)
int8x8_t	int8x16_t
int16x4_t	int16x8_t
int32x2_t	int32x4_t
int64x1_t	int64x2_t
uint8x8_t	uint8x16_t
uint16x4_t	uint16x8_t
uint32x2_t	uint32x4_t
uint64x1_t	uint64x2_t
float16x4_t	float16x8_t
float32x2_t	float32x4_t
poly8x8_t	poly8x16_t
poly16x4_t	poly16x8_t

There are also further types, which combine two, three or four of each of the above, into a larger struct type. These types are used to map the registers accessed by NEON load/store operations, which can load/store up to four registers with a single instruction. For example:

```
struct int16x4x2_t
{
    int16x4_t val[2];
}<var_name>;
```

These types are only used by loads, stores, transpose, interleave and de-interleave instructions; to perform operations on the actual data, the individual registers would be selected using, for example, <var_name>.val[0] and <var_name>.val[1] in the example above.

20.1.5 Variables and constants

In this section we show some example code to access variable or constant data using NEON.

Declaring a variable

Declaring a new variable is as simple as declaring any variable in C:

```
uint32x2_t vec64a, vec64b; // create two D-register variables
```

Using constants

Using constants is straightforward. The following code will replicate a constant into each element of a vector:

```
uint8x8 start_value = vdup_n_u8(0);
```

To load a general 64-bit constant into a vector, use:

```
uint8x8 start_value = vreinterpret_u8_u64(vcreate_u64(0x123456789ABCDEFULL));
```

Moving results back to normal C variables

To access a result from a NEON register, we can either store it to memory using VST, or move it back to ARM using a *get lane* type operation:

```
result = vget_lane_u32(vec64a, 0); // extract lane 0
```

This MRC operation can have some performance impact, particularly on the Cortex-A8 processor.

Accessing two D-registers of a Q-register

This can be done using *vget_low* and *vget_high*, as below:

```
vec64a = vget_low_u32(vec128); // split 128 bit vector
vec64b = vget_high_u32(vec128); // into 2x 64 bit vectors
```

Casting NEON variables between different types

NEON intrinsics are strongly typed, so they must be cast between vectors of different types. This is done using *vreinterpret*, which doesn't actually generate any code, but just enables you to cast the NEON types:

```
uint8x8_t byteval;
uint32x2_t wordval;
byteval = vreinterpret_u8_u32(wordval);
```

Note that the destination type, *u8* is listed first after *vreinterpret*.

20.1.6 Generating NEON instructions from C/C++ code

To generate NEON vector instructions from C/C++ code, the following ARM compiler options must be specified:

- A target `--cpu` that has NEON capability, such as the Cortex-A9 or Cortex-A15 processors.
- `--vectorize` to enable NEON vectorization.
- `--02` (default) or `--03` optimization level.
- `--0time` to optimize for performance rather than code size.

In addition, use `--diag_warning=optimizations` to obtain useful diagnostics from the compiler on what it could or could not optimize/vectorize.

For example:

```
armcc --cpu Cortex-A9 --vectorize -03 -0time --diag_warning_optimizations source.c
```


20.1.7 NEON assembler and ABI restrictions

For the very highest performance, hand-coded NEON assembler is the best approach for experienced programmers. This requires familiarity with the pipeline and cycle timings of the specific ARM processor that is being used, to ensure maximum throughput. Both GNU Assembler (gas) and ARM `armasm` support assembly of NEON instructions.

When writing assembler functions, you must be aware of the ARM EABI which defines how registers can be used.

The ARM EABI (Embedded Application Binary Interface) specifies which registers are used to pass parameters, return results, or must be preserved. This specifies the usage of 32 D-registers in addition to the ARM integer registers and is summarized in [Table 20-2](#).

Table 20-2 ARM EABI NEON register use

D0-D7	Argument registers / return value (can be corrupted if not required)
D8-D15	preserved by called functions
D15-D31	not preserved by called functions

Typically, data values are not actually passed to a NEON subroutine, so the best order to use NEON registers is D0-D7, then D15-D31 (then finally D8-D15 only if saved). See also [VFP and NEON register usage on page 17-4](#).

20.1.8 Detecting NEON

As NEON hardware can be omitted from a processor implementation, it may be necessary to test for its presence.

Build-time NEON selection

This is the easiest way to select NEON. In `armcc` (RVCT 4.0 and later), or GCC, the predefined macro `__ARM_NEON__` is defined when a suitable set of processor and FPU options is provided to the compiler. The `armasm` equivalent predefined macro is `TARGET_FEATURE_NEON`.

This could be used to have a C source file which has both NEON and non-NEON optimized versions.

Run-time NEON detection

To detect NEON at run-time requires help from the operating system, since the ARM architecture intentionally does not expose processor capabilities to user-mode applications.

Under Linux, `/proc/cpuinfo` contains this information in human-readable form.

On Tegra2 (a dual-core Cortex-A9 processor with FPU), `cat /proc/cpuinfo` reports:

```
...
Features      : swp half thumb fastmult vfp edsp thumbee vfpv3 vfpv3d16
...
```

The ARM quad-core Cortex-A9 processor with NEON gives a slightly different result:

```
...
Features      : swp half thumb fastmult vfp edsp thumbee neon vfpv3
...
```

As the `/proc/cpuinfo` output is text based, it is often preferred to look at the auxiliary vector `/proc/self/auxv`. This contains the kernel `hwcap` in a binary format. The `/proc/self/auxv` file can be easily searched for the `AT_HWCAP` record, to check for the `HWCAP_NEON` bit (4096).

Some Linux distributions (for example, Ubuntu 09.10, or later) take advantage of NEON transparently. The `ld.so` linker script is modified to read the `hwcap` via `glibc`, and add an additional search path for NEON-enabled shared libraries. In the case of Ubuntu, a new search path `/lib/neon/vfp` contains NEON-optimized versions of libraries from `/lib`.

20.2 Optimizing NEON assembler code

To obtain best performance from hand-written NEON code, it is necessary to be aware of some underlying hardware features. In particular, the programmer should be aware of pipelining and scheduling issues, memory access behavior and scheduling hazards. We will briefly describe each of these here.

20.2.1 Memory access optimizations

It is likely NEON will be used to process large amounts of data, such as digital images. One important optimization is to make sure the algorithm is accessing the data in the most cache-friendly way possible. It is also important to consider the number of active memory locations. Under Linux, each 4KB page will require a separate TLB entry. A typical optimization is to arrange the algorithm to process image data in suitably sized *tiles* in order to maximize the cache and TLB hit rate. In general, the memory optimization considerations described in [Chapter 19](#) will also apply to NEON code.

The instructions which support interleaving and de-interleaving can provide significant scope for performance improvements. VLD1 loads registers from memory, with no de-interleaving. However, the other VLD operations allow us to load, store and de-interleave structures containing two, three or four equally sized 8, 16 or 32-bit elements. VLD2 loads two or four registers, de-interleaving even and odd elements. This could be used, for example, to split left and right channel stereo audio data as in [Figure 20-1](#). Similarly, VLD3 could be used to split RGB pixels into separate channels and correspondingly, VLD4 might be used with ARGB or CMYK images.




Figure 20-1 Labeled example of load de-interleaving

[Figure 20-1](#) shows loading two NEON registers with VLD2.16, from memory pointed to by R1. This produces four 16-bit elements in the first register, and four 16-bit elements in the second, with adjacent paired left and right values separated out to each register.

20.2.2 Alignment

The NEON architecture provides full unaligned support for NEON data access. However, the instruction opcode contains an *alignment hint* which permits implementations to be faster when the address is aligned and a hint is specified.

The base address specified as [`<Rn>:align`]

Note

It is a programming error to specify a hint, but to then use an incorrectly aligned address. This will result in a data abort.

An alignment hint can be one of :64, :128 or :256 bits depending on the number of D-regs. The *ARM Architecture Reference Manual* uses an “@” symbol to describe this, but this acts as a comment marker in GNU assembler and so should not be used.

20.2.3 Scheduling

To get the very best performance with NEON, you should be aware of how to schedule code for the specific ARM processor you are using. Careful hand-scheduling is recommended to get the best out of any NEON assembler code you write, especially for performance-critical applications such as video codecs.

If writing C or NEON intrinsics, the compiler (GCC or armcc) will automatically schedule code from NEON intrinsics or vectorizable C source code, but it can still help to make the source code as friendly as possible to scheduling optimizations.

Result-use scheduling

This is the main performance optimization when writing NEON code. NEON instructions typically issue in one cycle, but the result is not always ready in the next cycle except for the simplest NEON instructions, such as VADD and VMOV.

In some cases there can be a considerable latency, particularly VMLA multiply-accumulate (five cycles for an integer; seven cycles for a floating-point). Code using these instructions should be optimized to avoid trying to use the result value before it is ready, otherwise a stall will occur. Despite having a few cycles result latency, these instructions do fully pipeline so several operations can be *in flight* at once.

The resulting latency is the same between the Cortex-A8 and Cortex-A9 processors for the majority of instructions. The Cortex-A5 and Cortex-A7 processors use a simplified NEON design more tailored to reduced power and area implementation, and most NEON instructions have a three-cycle result latency. The Cortex-A15 processor has significantly higher performance, with 128-bit data paths for NEON and the ability to issue multiple operations per cycle.

20.3 NEON power saving

Not all code can benefit from NEON. It depends to a large extent how data-intensive an application is and whether the algorithm can operate on multiple values in parallel. In a typical mobile device, NEON will be in a different power domain to the main processor, which allows it to be powered off when not in use. This can save significant power because NEON is often used intensively by applications (for example, real-time multi-media processing) or not at all. Mobile devices can spend significant amounts of time in standby with the clocked stopped, and it makes sense to ensure NEON is powered down at the same time.

The power saving mechanism is implemented by the OS, and is expected to be transparent to the user application. Typically the OS will turn off NEON after a period of non-activity, and turn it on again when there is an exception caused by attempting to execute a NEON instruction.

NEON power-down sequence (performed by the operating system):

1. Disable NEON (for example, by setting the Cortex-A9 processor ASEDIS to 1).
2. Remove power to the NEON region (SoC specific).

If a NEON instruction is executed when NEON is disabled, the processor will take an *undefined instruction* exception (as would any coprocessor if the coprocessor is disabled). The OS undefined instruction handler would detect that a NEON instruction (CP10 or CP11) has caused the exception and would execute the power-on sequence:

1. Restore power supply (SoC specific) and wait for the power to stabilize.
2. Enable NEON (for example, by setting the Cortex-A9 processor ASEDIS to 0).
3. Re-execute the instruction.

This on-demand control over NEON by the operating system means that use of NEON should be restricted to sufficiently large workloads and/or bursts for greater power efficiency.

Chapter 21

Introduction to Multi-processing

In the chapters that follow we look at systems with multiple processors and examine programming techniques for these. We distinguish between systems that contain multiple separate processor elements (a majority of today's embedded systems) and true multi-processors, with multiple processors closely coupled together in hardware. We will introduce terminology to describe and categorize such systems, in terms of both hardware and software and we look at some example ARM implementations.

Multi-processing can be defined as running two or more sequences of instructions simultaneously within a device containing two or more processors. The concept of multi-processing has been a subject of research for a number of decades, and has seen widespread commercial use over the past 15 years. Multi-processing is now a widely adopted technique in both systems intended for general-purpose application processors and in areas more traditionally defined as embedded systems.

Multi-processing systems can deliver higher performance, due to the simple fact that more processing units are available. This allows multiple tasks to be executed in parallel, potentially reducing the amount of time required to perform the allocated task.

The overall energy consumption of a multi-core system can be significantly lower than that of a system based on a single processor core. Multiple cores allow execution to be completed faster and so some elements of the system might be completely powered down for longer periods. Alternatively, a system with multiple cores may be able to operate at a lower frequency than that required by a single processor to achieve the same throughput. A lower power silicon process or a lower supply voltage can result in lower power consumption and reduced energy usage. Note that most current systems do not allow the frequency of cores to be changed independently. However, each core can be dynamically clock gated, giving further power and energy savings.

Multi-core systems also add flexibility and scalability to system designs. A system that contains one or two cores could be scaled up for more performance by adding further cores, without requiring redesign of the whole system or significant changes to software.

Having multiple cores at our disposal also allows more options for system configuration. For example, we might have a system which uses separate cores – one to handle a hard real-time requirement and another for an application requiring high, uninterrupted performance. These could be consolidated into a single multi-processor system.

A multi-core device is also likely to be more responsive than one with a single core. When interrupts are distributed between cores there will be more than one core available to respond to an interrupt and fewer interrupts per core to be serviced. Multiple cores will also allow an important background process to progress simultaneously with an important but unrelated foreground process.

Multi-core systems can also extract more performance from high latency memory systems (for example, DDR memory) by allowing a memory controller to queue up and optimize requests to the memory. Processors working on coherent data can benefit from reductions in linefills and evictions. When the data is not shared, it is likely that performance will be adversely affected. An L2 cache can mean improved utilization for shared memory regions (including file caches), shared libraries and kernel code. Additionally, if the number of cores is increased, then without a corresponding increase in memory bandwidth performance will also deteriorate.

In the past, much software was written to operate within the context of a single processor. Some operating systems provide support for time-slicing. This gives the illusion of multiple processes or tasks running simultaneously. It is important to clearly understand the difference between multi-threading (for example POSIX threads, or Java) and multi-processing. A multi-threaded application can be run on a single core, but only with multi-processing can the threads truly execute in parallel.

Migrating multi-threaded software from a single processor system to a multi-processing one can trigger problems with incorrect programs that could not be exposed by running the same program time-sliced on a single processor. It can also cause very infrequent bugs to become very frequently triggered. What it cannot do is to cause correctly written multi-threaded programs to misbehave – only expose previously unnoticed errors.

21.1 Multi-processing ARM systems

From early in the history of the architecture, ARM processors were likely to be implemented in systems which contained other processors. This commonly meant a *heterogeneous* system, perhaps containing an ARM processor plus a separate DSP processor. Such systems have different software executing on different cores and the individual processors can have differing privileges and views of memory. Many widely used ARM systems (for example, TI's OMAP series, or Freescale's i.MX) are examples of this.




Figure 21-1 Example of a system with multiple processors

We can distinguish between systems which contain:

- a single processor
- multiple processors (such as that shown in [Figure 21-1](#)), which contain multiple independent processing elements
- a multi-core processor, which contain multiple cores capable of independent instruction execution, which can be externally viewed like a single unit, either by the system designer or by an operating system that can abstract the underlying resources from the application layer. In this chapter (and those which follow), we will focus on such processors.

ARM was among the first companies to introduce such multi-core processors to the System-on-Chip market, when it introduced the ARM11 MPCore processor in 2004. All the processors described in this book, with the exception of the Cortex-A8 processor are examples of such systems.

An ARM MPCore processor can contain between one and four processing cores. Each core can be individually configured to take part (or not) in a data cache coherency management scheme. A *Snoop Control Unit* (SCU) device inside the processor has the task of automatically maintaining cache coherency, between cores within the processor without software intervention.

ARM MPCore processors include an integrated interrupt controller. Multiple external interrupt sources can be independently configured to target one or more of the individual processor cores. Furthermore, each core is able to signal (or *broadcast*) any interrupt to any other core or set of cores in the system, from software (software triggered Interrupts). These mechanisms allow the OS to share and distribute the interrupts across all cores and to coordinate activities using the low-overhead signaling mechanisms provided.

Cortex-A MPCore processors also provide hardware mechanisms to accelerate OS kernel operations such as system wide cache and TLB maintenance operations. (This feature is not found in the ARM11 MPCore.)

In summary, each of the Cortex-A series MPCore processors have the following features:

- configurable between one and four cores (at design time)
- data cache coherency
- integrated interrupt controller
- local timers and watchdogs
- (optional) Accelerator Coherency Port (ACP).

Figure 21-2 shows the structure of the Cortex-A9 MPCore processor, though a generalized form of this description also applies to other ARM MPCore processors. We will describe each of these features in more detail in the course of this chapter and those which follow.




Figure 21-2 Cortex-A9 MPCore processor

21.2 Symmetric multi-processing

Symmetric multi-processing (SMP) is a software architecture which dynamically determines the roles of individual processors. Each core within the processor has the same view of memory and of shared hardware. Any application, process or task can run on any core and the operating system scheduler can dynamically migrate tasks between cores to achieve optimal system load.

We expect that readers will be familiar with the fundamental operating principles of an OS, but OS terminology will be briefly reviewed here. An application which executes under an operating system is known as a *process*. It performs many operations through calls to the system library, which provides certain functions from library code, but also acts a wrapper for system calls to kernel operations.

For the purposes of describing SMP operation, we will use the term *kernel* to represent that portion of the operating system which contains exception handlers, device drivers and other resource and process management code. We will assume also the presence of a *scheduler*, which typically is called using a timer interrupt and which has the task of time-slicing the available cycles on processors between multiple tasks, dynamically determining the priority level of individual processes and deciding which process to run next. Individual processes have associated resources (including stack, heap and constant data areas) and properties (such as scheduling priority settings). The kernel view of a process is called a *task*.

Threads are separate tasks executing within the same process space, which allow separate parts of the application to execute in parallel on different cores. They also permit one part of an application to keep executing while another part is waiting for a resource.

In general, all threads within a process share a number of global resources (including the same memory map and access to any open file and resource handles). Threads also have their own local resources, including their own stacks and register usage, which will be saved and restored by the kernel on a context switch. The fact that these resources are local does not, however, mean that the local resources of any thread are guaranteed to be protected from incorrect accesses by other threads. Threads are scheduled individually and can have different priority levels even within a single process.

An SMP capable operating system provides an abstracted view of the available processor resources to the application. Multiple applications can run concurrently in an SMP system without re-compilation or source code changes. A conventional multitasking OS enables the system to perform several task or activities at the same time, in both single-core and multi-core processors. There is, however, a key distinction. In the single (uniprocessor) case, multitasking is performed through time-slicing of a single processor, giving the illusion of many tasks being performed at the same time. In the multi-core system, we can have true concurrency; multiple task are actually carried out at the same time, in parallel, on separate cores. The role of managing the distribution of such tasks across the available cores is performed by the operating system.

To provide the best possible energy usage, it may be desirable for the OS to distribute the concurrent tasks evenly across the available processors. This enables us to maximize the benefits of Dynamic Voltage and Frequency Scaling (DVFS) power management techniques, since, typically, DVFS applies to the whole multi-core cluster. Such activity is known as *load balancing*. It attempts to ensure that there is an even spread of workloads across all available cores. To do this, it must be able to dynamically re-prioritize tasks and have the ability to monitor the workload of each core.

Interrupt handling can also be load balanced. By default all interrupts on ARM processors are sent to CPU0. This allows for using user space software to rebalance interrupts. Under Linux this can be performed by the *irqbalance* daemon, <http://irqbalance.org/>. Irqbalance monitors the interrupt handling load on each core and tries to shift interrupts around, using the `irq_set_affinity` system call, to achieve an even load across the system. This is not as trivial a task as it might first appear. For example, a single device might have more than one interrupt

output, or separate devices might be feeding data into the same kernel subsystems. Both these situations, if ignored, could lead to a lot of data being shuffled back and forth between cores by the cache coherency hardware. To combat this, `irqbalance` tries to make clever decisions about which interrupts should be grouped together on a specific core. However, it provides additional facilities for controlling this behavior if the default parameters are not optimal for your system. It is possible to manually configure interrupt targets for load balancing, for example, you may want to do this for interrupts associated with an individual peripheral. On ARM and Intel processors, the default in-kernel behavior is NOT to do interrupt load balancing.

The scheduler in an SMP system can dynamically reprioritize tasks. This *dynamic task prioritization* allows other tasks to run while the current task sleeps. In Linux, for example, tasks whose performance is bound by I/O activity can have their priority decreased in favor of tasks whose performance is limited by processor activity. The I/O-bound task will typically schedule I/O activity and then sleep pending such activity.

21.3 Asymmetric multi-processing

We need to distinguish between SMP (Symmetric) and AMP (Asymmetric) multi-processing. In an AMP system, the programmer statically assigns individual roles to a core within a processor, so that in effect, we have separate cores, each performing separate jobs, within each processor. This can be referred to as a function-distribution software architecture and typically means that we have separate operating systems running on the individual cores. The system can appear to the programmer as a single-core system with dedicated accelerators for certain critical system services. In general AMP does not refer to systems in which tasks or interrupts are associated with a particular core.

In an AMP system, each task can have a different view of memory and there is no scope for a core which is highly loaded to pass work to one which is lightly loaded. There is no need for hardware cache coherency in such systems, although there will typically be mechanisms for communication between the cores through shared resources, possibly requiring dedicated hardware. The Cortex-A9 processor cache coherence management hardware described in [Cache coherency on page 22-2](#), can help reduce the overheads associated with sharing data between the systems.

Reasons for implementing an AMP system using a multi-core processor might include security, need for guaranteeing meeting of real-time deadlines, or because individual cores are dedicated to perform specific tasks.

There are classes of systems that have both SMP and AMP features. This means that we have two or more cores running an SMP operating system, but the system has additional elements which do not operate as part of the SMP system. The SMP sub-system can be regarded as one element within the AMP system. Cache coherency is implemented between the SMP cores, but not necessarily between SMP cores and AMP elements within the system. In this way, independent subsystems can be implemented within the same multi-core processor.

It is entirely possible (and normal) to build AMP systems in which individual processors are running different operating systems (these are called Multi-OS systems).

The selection of software MP model is determined by the characteristics of the applications running in the system. In networking systems, for example, it can be convenient to provide a separation between control-plane (AMP) and data-plane (SMP) sections of the system. In other systems, it may be desirable to isolate those parts of the system which require hard real-time response from applications which require raw processor performance and to implement these on separate processors.

Note

Where synchronization is required between these separate processors, it can be provided through message passing communication protocols, for example, the *Multicore Association API* (MCAP). These can be implemented by using shared memory to pass data packets and by use of software triggered interrupts to implement a so-called door-bell mechanism.

Chapter 22

SMP Architectural Considerations

An SMP system will by definition have shared memory between the processors. To maintain the desired level of abstraction to application software, the hardware must take care of providing a consistent and coherent view of memory for the programmer.

Changes to shared regions of memory must be visible to all processors without any explicit software coherency management. Likewise, any updates to the memory map (for example due to demand paging, allocation of new memory or mapping a device into the current virtual address space) of either the kernel or applications must be consistently presented to all processors.

The ARM MPCore technology provides a number of features that reduce some of the overheads traditionally associated with writing software for SMP systems. This chapter describes these features, as well as going through other important aspects of system software running on an ARM MPCore processor.

22.1 Cache coherency

Cache coherency is vitally important in an SMP system. It means that changes to data held in the cache of one processor are visible to the other processors, making it impossible for processors to see stale copies of data (the old data from before it was changed by the first processor).

This activity is handled in ARM MPCore systems by a device known as the *Snoop Control Unit* (SCU). This device maintains coherency between the cores' L1 data caches. Since executable code changes much less frequently, this functionality is not extended to the L1 instruction caches. The coherency management is implemented using a MESI-based protocol, optimized to decrease the number of external memory accesses. In order for the coherency management to be active for a memory access, all of the following must be true:

- The SCU is enabled, through its control register located in the private memory region. See [Private memory region on page 22-12](#). (The SCU has configurable access control, allowing restriction of which processors can configure it.)
- The core performing the access is configured to participate in the inner Shareable domain (configured by the operating system at boot time, by setting the somewhat misleadingly named SMP bit in the CP15:ACTLR, Auxiliary Control Register).
- The MMU is enabled.
- The page being accessed is marked as Normal Shareable, with a cache policy of *write-back*, *write-allocate*. Device and Strongly-ordered memory, however, are not cacheable, and write-through caches behave just like uncached memory from the point of view of the core.

Note that the SCU can only maintain coherency within the MP block. If there are additional processors or other bus masters in the system, explicit software synchronization is required when these share memory with the MP block. Alternatively, the Accelerated Coherency Port can be used (described in [Accelerator Coherency Port \(ACP\) on page 22-3](#))

22.1.1 MESI protocol

There are a number of standard ways by which cache coherency schemes can operate. As mentioned in [Cache coherency](#), ARM MPCore processors use a modified version of the MESI protocol. Let's take a look at how that works and then describe some of the modifications made to it for performance reasons:

The SCU marks each line in the cache with one of the following attributes: M (Modified), E (Exclusive), S (Shared) or I (Invalid). These are described below:

Modified	The most up-to-date version of the cache line is within this cache. No other copies of the memory location exist within other caches. The contents of the cache line are no longer coherent with main memory.
Exclusive	The cache line is present in this cache and coherent with main memory. No other copies of the memory location exist within other caches.
Shared	The cache line is present in this cache and coherent with main memory. Copies of it can also exist in other caches in the coherency scheme.
Invalid	The cache line is invalid.

Now let's look at the standard implementation of the protocol. The rules are as follows:

- A write can only be done if the cache line is in the Modified or Exclusive state. If it is in the Shared state, all other cached copies must be invalidated first. A write moves the line into the Modified State.
- A cache can discard a Shared line at any time, changing to the Invalid state. A Modified line is written back first.
- If a cache holds a line in the Modified state, reads from other caches in the system will get the updated data from the cache. This is conventionally done by first writing the data to main memory and then changing the cache line to the Shared state, before performing a read.
- A cache that has a line in the Exclusive state must move the line to the Shared state when another cache reads that line.
- The Shared state may not be precise. If one cache discards a Shared line, another cache may not be aware that it could now move the line to Exclusive status.

ARM MPCores also implement optimizations that can copy clean data and move dirty data directly between participating L1 caches, without having to access (and wait for) external memory.

22.1.2 MOESI protocol

The Cortex-A5 MPCore processor implements a slightly different coherency protocol. It uses the MOESI protocol, which introduces a fifth state for a shareable line in an L1 data cache. This is *Owned*, which describes a line that is dirty and in possibly more than one cache. A cache line in the owned state holds the most recent, correct copy of the data. Only one core can hold the data in the owned state. The other cores can hold the data in the shared state.

22.1.3 Accelerator Coherency Port (ACP)

The *Accelerator Coherency Port* (ACP) is an optional feature of the Cortex-A9 and Cortex-A5 processors. It provides an AXI slave interface into the snoop control unit of the processor. (The AXI bus interface is defined in ARM's AMBA specification.)

This slave interface can be connected to an external uncached AXI master, such as a DMA engine, for example, and is able to initiate both read and write memory transfers to the ACP. It enables such a device to snoop the L1 caches of all cores, avoiding the need for synchronization through external memory. A cached device can also be connected, but this requires manual coherency management through software.

The behavior of accesses performed on the ACP is as follows:

- addresses used by the ACP are physical addresses that can be snooped by the SCU to be fully coherent
- ACP reads can hit in any processor's L1 D-cache
- writes on the ACP invalidate any stale data in L1 and write-through to L2.

The ACP allows an external device to see processor-coherent data without knowledge of where the data is in the memory hierarchy. Memory transfers are automatically coherent in the same way as happens between the processor's L1 D-caches.

Use of the ACP can both increase performance and save power, as there will be reduced traffic to external memory and faster execution.

Programmers writing device drivers which use ACP do not need to be concerned with coherency issues, as no cache cleaning or invalidation is required to ensure coherency. However, the use of memory barriers (DMB) or external cache synchronization operations can still be necessary, if a particular ordering must be enforced.

22.2 TLB and cache maintenance broadcast

An individual processor can broadcast TLB and cache maintenance operations to other processors in the inner Shareable coherency domain. This can be required whenever shared page tables are modified, for example. This behavior may need to be enabled by the programmer. For example, in the Cortex-A9 processor, this is controlled by the FW bit in the *Auxiliary Control Register* (ACTLR). Maintenance operations can only be broadcast and received when the processor is configured to participate in the inner Shareable domain (using the SMP bit in ACTLR). Only inner Shareable operations are broadcast, for example:

- invalidate TLB entry by virtual address
- clean or invalidate data cache line by virtual address
- invalidate instruction cache line by virtual address.

Maintenance broadcasting was not available in the ARM11 MPCore, requiring broadcasting of the above operations to be implemented in software.

Some care is needed with cache maintenance activity in multi-core systems which include a L2C-310 L2 cache (or similar). Cleaning or invalidating the L1 cache and L2 cache will not be a single atomic operation. A processor may therefore perform cache maintenance on a particular address in both L1 and L2 caches only as two discrete steps. If another processor were to access the affected address between those two actions, a coherency problem can occur. Such problems can be avoided by following two simple rules.

- when cleaning, always clean the innermost (L1) cache first and then clean the outer cache(s)
- when invalidating, always invalidate the outermost cache first and the L1 cache last.

22.3 Handling interrupts in an SMP system

ARM MPCore processors include an integrated interrupt controller, which is an implementation of the GIC architecture (see [Generic Interrupt Controller on page 13-4](#) for further details). This controller provides 32 private interrupts per core, of which the lower 16 are *Software Generated Interrupts* (SGI) that can be generated only through software operations, and the rest are *Private Peripheral Interrupts* (PPI). It also provides a configurable number of *Shared Peripheral Interrupts* (SPI) – up to 224 in current MPCore implementations). It supports interrupt prioritization, pre-emption and routing to different cores.

In an ARM MPCore processor, the GIC control registers are memory mapped and located within the Private memory region, see [Private memory region on page 22-12](#).

The Interrupt Processor Targets registers configure which cores individual interrupts are routed to. They are ignored for private interrupts.

The registers controlling the private interrupts (0-31) are banked, so that each core can have its own configuration for these. This includes priority configuration and the enabling or disabling of individual interrupts.

The Software Generated Interrupt (SGI) Register can assert a private software generated interrupt on any core, or a groups of cores. The priority of a software interrupt is determined by the priority configuration of the receiving core, not the one that sends the interrupt. The interrupt acknowledge register bits [12:10] will provide the ID of the core that made the request. The target list filter field within this register provides shorthand for an SGI to be sent to all processors, all but self or to a target list.

For a software generated interrupt in an MPCore processor, the Interrupt Acknowledge Register also contains a bitfield holding the ID of the core that generated it. When the interrupt service routine has completed, it must write-back the value previously read from the Interrupt Acknowledge Register into the End Of Interrupt Register. For an SGI, this must also match the ID of the signalling core.

In both AMP and SMP systems, it is likely that cores will need to trigger interrupts on other cores (or themselves) - a so called *softirq*.

These Inter-Processor Interrupts can be used for kernel synchronization operations, or for communicating between AMP processors. For operations requiring more information passed than a raised interrupt, you can use a shared buffer to store messages. Before the core can receive an interrupt, some initialization steps are required both in the distributor and in the processor interface.

22.4 Exclusive accesses

In an SMP system, data accesses frequently have to be restricted to one modifier at any particular time. This can be true of peripheral devices, but also for global variables and data structures accessed by more than one thread or process. Furthermore, library code which is used by multiple threads must be designed to ensure that concurrent access or execution is possible – it must be *reentrant*.

Note

A computer program is *reentrant* if it can be interrupted in the middle of its execution and then be called again before the previous version has completed.

Protection of such shared resources is often through a method known as *mutual exclusion*. The section of code which is being executed by a core while accessing such a shared resource is known as the *critical section*.

In a single processor system, mutual exclusion can be achieved by disabling interrupts when inside critical sections. This is not sufficient in a multi-core system, as disabling interrupts on one core will not prevent others from entering the critical section. It is also not ideal since interrupts cannot be disabled from within User mode. Of course, there are other problems with this technique, including reduced responsiveness to real-time events, particularly if the critical section is long.

In a multi-core system, we can use a *spinlock* – effectively a shared flag with an atomic (indivisible) mechanism to test and set its value. We perform this operation in a tight loop to wait until another thread (or core) clears the flag. We require hardware assistance in the form of special assembly language instructions to implement this. Most application developers do not need to understand the low-level implementation detail, but should instead become familiar with the lock and unlock calls available in their OS or threading library API. Nevertheless, in the next section, we will examine how this is implemented on the ARM architecture.

Three assembly language instructions relating to exclusive access are provided in the ARM processor architecture (in ARMv6 onwards). Variants of these instructions which operate on byte, halfword, word or doubleword sized data are also provided. The instructions rely on the ability of the core or memory system to tag particular addresses to be monitored for exclusive access by that core, using an *exclusive access monitor*.

- LDREX (Load Exclusive) performs a load of memory, but also tags the physical address to be monitored for exclusive access by that processor.
- STREX (Store Exclusive) performs a conditional store to memory, succeeding only if the target location is tagged as being monitored for exclusive access for that core. This instruction returns the value of 1 in a general purpose register if the store does not take place, and a value of 0 if the store is successful.
- CLREX (Clear Exclusive) clears any exclusive access tag for that processor.

Load Exclusive and Store Exclusive operations must be performed only on Normal memory (see [Normal memory on page 11-5](#)) and have slightly different effect depending on whether the memory is marked as Shareable or not. If the core reads from Shareable memory with an LDREX, the load happens and that physical address is tagged to be monitored for exclusive access by that processor. If any other processor writes to that address and the memory is marked as Shareable, the tag is cleared.

If the memory is not Shareable then any attempt to write to the tagged address by the one that tagged it results in the tag being cleared. If the processor does a further LDREX to a different address, the tag for the previous LDREX address is cleared. Each processor can only have one address tagged.

STREX can be considered as a conditional store. The store is performed only if the physical address is still marked as exclusive access (this means it was previously tagged by this core and no other core has since written to it). STREX returns a status value showing if the store succeeded. STREX always clears the exclusive access tag.

The use of these instructions is not limited to multi-core systems. In fact, they are frequently employed in single processor systems, to implement synchronization operations between threads running on the same processor.

Let's take a very simple example ([Example 22-1](#)) of how the load and store exclusive operations can be used to implement a spinlock.

Example 22-1 Example spin lock implementation with LDREX and STREX

```

MOV    R1, #0      @ the LOCK value
LDR     R2, <spinlock_addr>
spin_lock:
LDREX   R0, [R2]    @ load value
CMP     R0, R1      @ if not locked already
STREXNE R0, R1, [R2] @ try to claim the lock
CMPNE   R0, #1      @ check if we succeeded
BEQ     spin_lock   @ retry if failed

```

In this code, we perform an LDREX (exclusive load) from the memory address which holds the lock value (address held in R2). This does two things. It returns the value of the lock into register R0. It also tags the address as being exclusively accessed. We then check to see if the value at the address is equal to 0 (the value we have chosen to represent that the lock is in a locked state). If it is not, we proceed to try to claim the lock. We execute the STREX (exclusive store) instruction. This first checks that the address is still marked as exclusive. If it is, this means that we check that no other thread on this processor, or on another, has touched the address in the time since the read was done and the address was tagged as exclusive. If this check succeeds then the memory location is updated and the exclusive access tag is cleared. In this case the value 0 is written, indicating that the resource is now locked. The register R0 is used to hold the value returned by the STREX instruction. If R0 is 1 then the STREX failed. The BEQ at the end covers both cases (that is, LDREX returned a value of 0 indicating the lock is taken or that STREX returned a 1 indicating that someone touched the lock since we read a non-zero value in the LDREX).

It should be clear from this that LDREX and STREX provide the necessary primitives to allow us to build complex synchronization objects. Now let's take a look at how the exclusive tagging actually works.

In hardware, the processor includes a device named the *local monitor*. This monitor observes the processor. When the core performs an exclusive load access, it records that fact in the local monitor. When it performs an exclusive store, it checks that a previous exclusive load was performed and fails the exclusive store if this was not the case. The architecture enables individual implementations to determine the level of checking performed by the monitor. The processor can only tag one physical address at a time. An LDREX from a particular address should be followed shortly after by an STREX to the same location, before an LDREX from a different address is performed. This is because the local monitor does not need to store the address of the exclusive tag (although it can do, if the processor implementer decides to do this). The

architecture enables the local monitor to treat any exclusive store as matching a previous LDREX address. For this reason, use of the CLREX instruction to clear an existing tag is required upon context switches.

Where exclusive accesses are used to synchronize with external masters outside the processor, or to regions marked as Sharable even between cores in the same processor, it is necessary to implement a global monitor within the hardware system. This acts as a wrapper to one or more memory slave devices and is independent of the individual processors. This is specific to a particular SoC and may or may not exist in any particular system. An LDREX/STREX sequence performed to a memory location which has no suitable exclusive access monitor will fail, with the STREX instruction always returning 1.

Finally, let's look at the specific case of Linux – the arch/arm/include/asm/spinlock.h file, which includes two functions arch_spin_lock() and arch_spin_unlock(). The first of these contains the inline assembly sequence shown in [Example 22-2](#).

Example 22-2 Linux spinlock code

```
static inline void arch_spin_lock(arch_spinlock_t *lock)
{
    unsigned long tmp;

    __asm__ __volatile__(
        "1:    ldrex    %0, [%1]\n"
        "      teq     %0, #0\n"
        "#ifdef CONFIG_CPU_32v6K\n"
        "      wfene\n"
        "#endif\n"
        "      strexeq %0, %2, [%1]\n"
        "      teqeq   %0, #0\n"
        "      bne     1b\n"
        "      : "=&r" (tmp)\n"
        "      : "r" (&lock->lock), "r" (1)\n"
        "      : "cc");

    smp_mb();
}
```

As you can see, this is very similar to the code in [Example 22-1 on page 22-8](#), a key difference being that Linux running on an MPCore processor can put a core which is waiting for a lock to become available into standby state, to save power. Of course, this relies on the other processor telling us when it has finished with the lock to wake this processor up, using the SEV instruction. More information on WFE and SEV can be found in [Chapter 25 Power Management](#).

The smp_mb() macro at the end of the sequence is required to ensure that external observers see the lock acquisition before they see any modifications of the protected resource, and also to ensure that accesses to the region before the acquisition, have completed before the lock holder reads from it. See [Linux use of barriers on page 11-11](#) for more information on the barrier macros used in the Linux kernel.

22.5 Booting SMP systems

Initialization of the external system may need to be synchronized between cores. Typically, only one of the cores in the system needs to run code which initializes the memory system and peripherals. Similarly, the SMP operating system initialization typically runs on only one core – the *primary core*. When the system is fully booted, the remaining cores are brought online and this distinction between the primary core and the others (secondary cores) is lost.

If all of the cores come out of reset at the same time, they will normally all start executing from the same reset vector. The boot code will then read the processor ID to determine which core is the primary. The primary core will perform the initialization described above and then signal to the secondary ones that everything is ready. An alternative method is to hold the secondary cores in reset while the primary core does the initialization. This requires hardware support to co-ordinate the reset.

In an AMP system, the bootloader code will determine the suitable start address for the individual cores, based on their processor ID (as each processor will be running different code). Care may be needed to ensure correct boot order in the case where there are dependencies between the various applications running on different cores.

22.5.1 Processor ID

Booting provides a simple example of a situation where particular operations need to be performed only on a specific core. Other operations need to perform different actions dependent on the core on which they are executing.

The CP15:MPIDR Multiprocessor Affinity Register provides a processor identification mechanism in a multiprocessor system.

This register was introduced in version 7 of the ARM architecture, but was in fact already used in the same format in the ARM11 MPCore. In its basic form, it provides up to three levels of affinity identification, with 8 bits identifying individual blocks at each level. In less abstract terms, you could say that there is:

- one 8-bit field showing which core you are executing on within an MPCore processor
- one 8-bit field showing which MPCore processor you are executing on within a cluster of MPCore processors
- one 8-bit field showing which cluster of MPCore processors you are executing on within a cluster of clusters of MPCore processors.

This information can also be of value to an operating system scheduler, as an indication of the order of magnitude of the cost of migrating a process to a different core, processor or cluster.

The format of the register was slightly extended with the ARMv7-A multiprocessing extensions implemented in the Cortex-A9 and Cortex-A5 processors. This extends the previous format by adding an identification bit to reflect that this is the new register format, and also adds the *U* bit which indicates whether the current core is the only core a uniprocessor implementation or not.

22.5.2 SMP boot in Linux

The boot process for the primary core is as described in [Boot process on page 14-5](#). The method for booting the secondary cores can differ somewhat depending on the SoC being used. The method that the primary core invokes in order to get a secondary core booted into the operating system is called `boot_secondary()` and needs to be implemented for each *mach* type that supports SMP. Most of the other SMP boot functionality is extracted out into generic functions in `linux/arch/arm/kernel`.

The method below describes the process on an ARM Versatile Express development board (mach-vexpress).

While the primary core is booting, the secondary cores will be held in a standby state, using the WFI instruction. It will provide a startup address to the secondary cores and wake them using an *Inter-Processor Interrupt* (IPI), meaning an SGI signalled through the GIC (see [Handling interrupts in an SMP system on page 22-6](#)). Booting of the secondary cores is serialized, using the global variable `pen_release`. Conceptually, we can think of the secondary cores being in a *holding pen* and being released one at a time, under control of the primary core. The variable `pen_release` is set by the kernel code to the ID value of the processor to boot and then reset by that core when it has booted. When an inter-processor interrupt occurs, the secondary core will check the value of `pen_release` against their own ID value using the MPIDR register.

Booting of the secondary processor will proceed in a similar way to the primary. It enables the MMU (setting the TTB register to the new page tables already created by the primary). It enables the interrupt controller interface to itself and calibrates the local timers. It sets a bit in `cpu_online_map` and calls `cpu_idle()`. The primary processor will see the setting of the appropriate bit in `cpu_online_map` and set `pen_release` to the next secondary processor.

22.6 Private memory region

In the Cortex-A5, and Cortex-A9 MPCore processors, all of the internal peripherals are mapped to the private address space. This is an 8KB region location within the memory map at an address determined by the hardware implementation of the specific device used (this can be read using the CP15 Configuration Base Address Register).

The registers in this region are fixed in little-endian byte order, so some care is needed if the CPSR E bit is set when accessing it. Some locations within the region exist as banked versions, dependent on the processor ID. The Private memory region is not accessible through the Accelerator Coherency Port. [Table 22-1](#) shows the layout of this Private memory region.

Table 22-1 Private memory region layout

Base Address offset	Function
0x0000	Snoop Control Unit (SCU)
0x0100	Interrupt controller CPU Interface
0x0200	Global Timer
0x0600	Local Timer/Watchdog
0x1000	Interrupt Controller Distributor

22.6.1 Timers and watchdogs

We looked at the SCU and interrupt control functions in [Cache coherency on page 22-2](#). In addition, each core in an ARM MPCore implements a standard timer and a watchdog, both private to that core.

These can be configured to trigger after a number of processor cycles, using a 32-bit start value and an 8-bit pre-scale. They can be operated using interrupts, or by periodic polling (supported with the timer/watchdog Interrupt Status Registers). They stop counting while the core is in debug state. The timer can be configured in *single-shot* or *auto-reload* mode. The watchdog can be operated in classic watchdog fashion, where it asserts the core reset signal (for that specific core) on timeout. Alternatively, it can be used as a second timer.

Revision 1 and later of the Cortex-A9 processor, and all versions of the Cortex-A5 processor also include a global timer, shared between all cores, but with banked comparator and auto-increment registers for each core. It is a single, incrementing 64-bit counter, accessible only through 32-bit accesses. It can be configured to trigger an interrupt when the comparator value is reached. The auto-increment feature causes the processor comparator register to be incremented after each match. This is typically used by the OS scheduler, to trigger the scheduler on each core, at different times.

Chapter 23

Parallelizing Software

In [Chapter 21 *Introduction to Multi-processing*](#), we described how an SMP system can allow us to run multiple threads efficiently and concurrently across multiple processors. In this case, the parallelization is, in effect, handled on our behalf by the OS scheduler.

In many cases, however, this is insufficient and the programmer must take steps to rewrite code to take advantage of speed-ups available through parallelization. An obvious example is where a single application requires more performance than can be delivered by a single processor. More commonly, we can have the situation where an application requires much more performance than all of the others within a system, when it is said to be *dominant*. This prevents efficient energy usage, as we cannot perform optimal load-balancing. An unbalanced load distribution does not allow efficient dynamic voltage and frequency scaling.

The operating system cannot automatically parallelize an application. It is limited to treating that application as a single scheduling unit. In such cases, the application itself has to be split into multiple smaller tasks by the programmer. Of course, this means each of these tasks must be able to be independently scheduled by the OS, as separate threads. A thread is a part of a program that can be run independently and concurrently with other parts of a program. If the programmer *decomposes* an application into smaller execution entities which can be separately scheduled, the OS can spread the threads of the application across multiple processors.

23.1 Decomposition methods

The best approach to decomposition of an application into smaller tasks capable of parallel execution depends on the characteristics of the original application. Large data-processing algorithms can be broken down into smaller pieces by sub-division into a number of similar threads which execute in parallel on smaller portions of a dataset. This is known as *data decomposition*.

Consider the example of color-space conversion, from RGB to YUV. We start with an array of pixel data. The output is a similar array giving chrominance and luminance data for each pixel. Each output value is calculated by performing a small number of multiplies and adds. Crucially, the output Y, U and V values for each pixel depend only upon the input R, G and B values for that pixel. There is no dependency on the data values of other pixels. Therefore, the image can be divided into smaller blocks and we can perform the calculation using any number of instances of our code. This does not require any change to our original algorithm – simply changes to the amount of data supplied to each thread.

We split the image into stripes ($1/N$ arrays, where we have N threads) and each thread works on a stripe. The level of detail of the stripes can be an important consideration (it is clearly better for cacheability if each thread works on a contiguous block of pixels in array order). The code does not have to be modified to take care of scheduling – it is the operating system which takes care of it. Color space conversion would be a task where the NEON unit could significantly improve performance. Splitting the task across several processors can provide further parallelization gains than using NEON instructions alone.

A different approach is that of *task decomposition*. Here, we identify areas of code which are independent of each other and capable of being executed concurrently. This is a little more difficult, as we need now to think about the discrete operations being carried out and the interactions between them. A simple example might be the start-up sequence of a program. One task might be to check that the user has a valid license for the software. Another task might be to display a start-up banner with a copyright message. These are independent tasks with no dependency on each other and can be performed in separate threads. Again, no change is required to the source code which carries out these isolated tasks. We have to supply them to the OS kernel Scheduler as separate execution threads.

Of course, not all algorithms are able to be handled through data or task decomposition. Instead, we must analyze the program with the aim of identifying functional blocks. These are independent pieces of code with defined inputs and outputs that have some scope to be parallelized. Such functional blocks often depend upon input from other blocks (they have a *serial dependency*), but do not have a corresponding dependency upon time (a *temporal dependency*). This is (in some respects) analogous to the hardware pipelining employed in the processor itself.

The software for an MPEG video encoder provides a good example of this. Input data, in the form of an analog video signal is sampled and processed through a pipeline of discrete functional blocks. First, both inter-frame and intra-frame redundancies are removed. Then, quantization takes place to reduce the number of bits required to represent the video. After this, motion vector compensation takes place, run length compression and finally the encoded sub-stream is stored.

At the same time that data from one frame is being run-length compressed and stored, we can also start to process the next frame. Within a frame, the motion vector compensation process can be parallelized. We can use multiple parallel threads to operate on a frame (an example of data decomposition).

When decomposing an application using these techniques, we must consider the overheads associated with task creation and management. An appropriate level of granularity is required for best performance. If we make our datasets too small, too big, or have too many datasets, it can reduce performance. In our example of color-space conversion, it would not be sensible to have a separate thread for each pixel, even though this is logically possible.

23.2 Threading models

When an algorithm has been analyzed to determine potential changes which can be made for parallelization, the programmer must modify code to map the algorithm to smaller, threaded execution units. There are two widely-used threading models, the *workers' pool* model and the *fork-join* model, not to be confused with the UNIX *fork* system call. The latter creates (spawns) a new thread whenever one is needed (that is, threads are created on-demand.) The operating system then schedules the various threads across the available processors.

Each of the newly spawned threads is typically considered to be either a *detached* thread, or a *joinable* thread. A detached thread executes in the background and terminates when it has completed, without any message to the parent process. Of course, communication to or from such processes can be implemented manually by the programmer, through the available signaling mechanisms, or using global variables. A joinable thread, in contrast, will communicate back to the main thread, at a point set by the programmer. The parent process might have to wait for all joinable threads to return before proceeding with the next execution step.

In the fork-join model, individual threads have explicit start and end conditions. There is an overhead associated with managing their creation and destruction and latencies associated with the synchronization point. This means that threads must be sufficiently long-lived to justify these costs.

If we know that some execution threads will be repeatedly required to consume input data, we can instead use the workers' pool threading model. Here, we create a pool of worker threads at the start of the application. The pool can consist of multiple instances of the same algorithm, where the distributor (also called producer or boss) will dispatch the task to the first available worker (consumer) thread. Alternatively, the workers' pool can contain several different data processing operators and data-items will be tagged to show which worker should consume the data.

The number of worker threads can be changed dynamically to handle peaks in the workload. Each worker thread performs a task until it is finished, then interrupts the boss to be assigned another task. Alternatively, the boss can periodically poll workers to see whether one is ready to receive another task. The work queue model is similar. The boss places tasks in a queue, and workers check the queue and take tasks to perform. A further variant is to have multiple bosses, sharing the pool of workers. The boss threads place tasks onto a queue, from where they are taken by the worker threads.

In each of these models, it should be understood that the amount of work to be performed by a thread can be variable and unpredictable. Even for threads which operate on a fixed quantity of data, it can be the case that data dependencies cause different execution times for similar threads. There is always likely to be some synchronization overhead associated with the need for a parent thread to wait for all spawned threads to return (in the fork-join model) or for a pool of workers to complete data consumption before execution can be resumed.

23.3 Threading libraries

We have looked at how to make our target application capable of concurrent execution. We must now consider actual source code modifications. This is normally done using a threading library, normally utilizing multi-threading support available in the OS. When modifying existing code, we must take care to ensure that all shared resources are protected by proper synchronization. This includes any libraries used by the code, as all libraries are not *reentrant*. In some cases, there can be separate reentrant libraries for use in multi-threaded applications. A library which is designed to be used in multi-threaded applications is called *thread-safe*. If a library is not known to be thread-safe, only one thread should be allowed to make calls to the library functions.

The most commonly used standard in this area is POSIX threads (Pthreads), a subset of the wider POSIX standard. POSIX (IEEE std. 1003) is the Portable Operating System Interface, a collection of OS interface standards. Its goal is to assure interoperability and portability of code between systems. Pthreads defines a set of API calls for creating and managing threads. Pthreads libraries are available for Linux, Solaris, and Windows.

There are several other multi-threading frameworks, such as OpenMP, which can simplify multi-threaded development by providing high-level primitives, or even automatic multi-threading. OpenMP is a multi-platform, multi-language API that supports shared memory multi-processing through a set of libraries and compiler directives plus environment variables which affect run-time behavior.

Pthreads provides a set of C primitives which allow us to create, manage, and terminate threads and to control thread synchronization and scheduling attributes. Let us examine, in general terms, how we can use Pthreads to build multi-threaded software to run on our SMP system. We'll deal with the following types:

- `pthread_t` – thread identifier
- `pthread_mutex_t` – mutex
- `sem_t` – semaphore.

We need to modify our code to include the appropriate header files.

```
#include <pthread.h>
#include <semaphore.h>
```

We must link our code using the pthread library with the switch `-lpthread`.

To create a thread, we must call `pthread_create()`, a library function which requires four arguments. The first of these is a pointer to a `pthread_t`, which is where we will store the thread identifier. The second argument is the attribute, which can point to a structure which modifies the thread's attributes (for example scheduling priority), or be set to NULL if no special attributes are required. The third argument is the function the new thread will start by executing. The thread will be terminated should this function return. The fourth argument is a void * pointer supplied to the thread. This can receive a pointer to a variable or data structure containing relevant information to the thread function.

A thread can complete either by returning, or calling `pthread_exit()`. Both will terminate the thread. A thread can be detached, using `pthread_detach()`. A detached thread will automatically have its associated data structures (but not explicitly allocated data) released on exit.

For a thread that has not been detached, this resource cleanup will happen as part of a `pthread_join()` call from another thread. The library function `pthread_join()` enables us to make a thread stall and wait for completion of another thread. Take care, as so-called *zombie* threads can be created by joining a thread which has already completed. It is not possible to join a detached thread (one which has called `pthread_detach()`).

Mutexes are created with the `pthread_mutex_init()` function. The functions `pthread_mutex_lock()` and `pthread_mutex_unlock()` are used to lock or unlock a mutex. `pthread_mutex_lock()` blocks the thread until the mutex can be locked. `pthread_mutex_trylock()` checks whether the mutex can be claimed and returns an error if it cannot, rather than just blocking. A mutex can be deleted when no longer required with the `pthread_mutex_destroy()` function.

Semaphores are created in a similar way, using `sem_init()` – one key difference being that we must specify the initial value of the semaphore. `sem_post()` and `sem_wait()` are used to increment and decrement the semaphore.

The GNU tools for ARM processors support full thread-local storage using the Native POSIX Thread library (NPTL), which enables efficient use of POSIX threads with the Linux kernel. There is a one-to-one correspondence between threads created with `pthread_create()` and kernel tasks

[Example 23-1](#) provides a simple example of using the Pthreads library.

Example 23-1 Pthreads code

```
void *thread(void *vargp);
int main(void)
{
    pthread_t tid;
    pthread_create(&tid, NULL, thread, NULL);
    /* Parallel execution area */
    pthread_join(tid, NULL);
    return 0;
}
/* thread routine */
void *thread(void *vargp)
{
    /* Parallel execution area */
    printf("Hello World from a POSIX thread!\n");
    return NULL;
}
```

23.3.1 Inter-thread communications

Semaphores can be used to signal to another thread. A simple example would be where one thread produces a buffer containing shared data. It could use a semaphore to indicate to another thread that the data can now be processed.

For more complex signaling, a message passing protocol can be needed. Threads within a process use the same memory space, so an easy way to implement message passing is by posting in a previously agreed-upon mailbox and then incrementing a semaphore.

23.3.2 Threaded performance

There are a few general points to consider when writing a multi-threaded application:

- Each thread has its own stack space and care may be needed with the size of this if large numbers of threads are in use.
- Multiple threads contending for the same mutex or semaphore creates contention and wasted processor cycles. There is a large body of research on programming techniques to reduce this performance loss.

- There is an overhead associated with thread creation. Some applications avoid this by creating a thread pool at startup. These threads are used on demand and then returned to the thread pool for later re-use, rather than being closed completely.

23.3.3 Thread affinity

Thread affinity refers to the practice of assigning a thread to a particular processor or processors. When the scheduler wants to run a particular thread, it will use only the selected processor(s) even if others are idle (this can be quite a problem if too many threads have an affinity set to a specific processor). By default, threads are able to run on any core in an SMP system.

ARM DS-5 Streamline is able to reveal a thread's affinity by using a display mode called X-ray mode. This mode can be used to visualize how tasks are divided up by the kernel and shared amongst several processors. See [DS-5 Streamline on page 18-4](#).

23.4 Synchronization mechanisms in the Linux kernel

When porting software from a uniprocessor environment to run on multiple processors, there can be situations where we need to modify code to enforce a particular order of execution or to control parallel access to shared peripherals or global data. The Linux kernel (like other operating systems) provides a number of different synchronization primitives for this purpose. Most such primitives are implemented using the same architectural features as application-level threading libraries like Pthreads. Understanding which of these is best suited for a particular case will give software performance benefits. Serialization and multiple threads contending for a resource can cause suboptimal use of the increased processing throughput provided by the multiple processors. In all cases, minimizing the size of the critical section provides best performance.

23.4.1 Completions

Completions are a feature provided by the Linux kernel, which can be used to serialize task execution. They provide a lightweight mechanism with limited overhead that essentially provides a flag to signal completion of an event between two tasks. The task which is waiting can sleep until it receives the signal, using `wait_for_completion (struct completion *comp)` and the task that is sending the signal typically uses either `complete (struct completion *comp)`, which will wake up one waiting process, or `complete_all (struct completion *comp)` which wakes all processes which are waiting for the event. Kernel version 2.6.11 added support for completions which can time out and for interruptible completions.

23.4.2 Spinlocks

A spinlock provides a simple binary locking mechanism, designed for protection of critical sections. It implements a busy-wait loop. A spinlock is a generic synchronization primitive that can be accessed by any number of threads. More than one thread might be spinning for obtaining the lock. However, only one thread can obtain the lock. The waiting task executes `spin_lock (spinlock_t *lock)` and the signaling task uses `spin_unlock (spinlock_t *lock)`. Spinlocks do not sleep and disable pre-emption.

23.4.3 Semaphores

Semaphores are a widely used method to control access to shared resources, and can also be used to achieve serialization of execution. They provide a counting locking mechanism, which can cope with multiple threads attempting to lock. They are designed for protection of critical sections and are useful when there is no fixed latency requirement. However, where there is a significant amount of contention for a semaphore, performance will be reduced. The Linux kernel provides a straightforward API with functions `down (struct semaphore *sem)` and `up (struct semaphore *sem)`; to lower and raise the semaphore.

Unlike spinlocks, which spin in a busy wait loop, semaphores have a queue of pending tasks. When a semaphore is locked, the task yields, so that some other task can run. Semaphores can be binary (in which case they are also mutexes) or counting.

23.4.4 Lock-free synchronization

The use of lock-free data structures, such as circular buffers, is widespread and can avoid the overheads associated with spinlocks or semaphores. The Linux kernel also provides two synchronization mechanisms which are lock-free, the *Read-Copy-Update* (RCU) and seqlocks. Neither of these mechanisms is normally used in device drivers.

If you have multiple readers and writers to a shared resource, using a mutex may not be very efficient. A mutex would prevent concurrent read access to the shared resource because only a single thread is allowed inside the critical section. Large numbers of readers might delay a writer from being able to update the shared resource. RCU can help in the case where the shared resource is mainly accessed by readers. Reader threads execute with little synchronization overhead. A thread which writes the shared resource has a much higher overhead, but is executed relatively infrequently. The writer thread must make a copy of the shared resource (access to shared resources must be done through pointers). When the update is complete, it publishes the new data structure, so that it is visible to all readers. The original copy is preserved until the next context switch on all processors. This guarantees that all current read operations can complete. RCUs are more complex to use than standard mutexes and are typically used only when traditional solutions are not suitable. Examples include shared file buffers or networking routing tables and garbage collection.

Seqlocks are also intended to provide quick access to shared resources, without use of a lock. They are optimized for short critical sections. Readers are able to access the shared resource with no overhead, but must explicitly check and re-try if there is a conflict with a write. Writes, of course, still require exclusive access to the shared resource. They were originally developed to handle things like system time – a global variable which can be read by many processes and is written only by a timer-based interrupt (on a frequent basis, of course!) The timer write has a high priority and a hard deadline, in order to be accurate. Using a seqlock instead of a mutex enables many readers to share access, without locking out the writer from accessing the critical section.

Chapter 24

Issues with Parallelizing Software

In this chapter, we will consider some of the problems and potential difficulties associated with making software concurrent. You might also at this point wish to revisit the explanation of barrier use in the Linux kernel described in [Linux use of barriers on page 11-11](#).

Amdahl's Law defines the theoretical maximum speedup achievable by parallelizing an application. The maximum speedup is given by the formula:

$$\text{Max speedup} = 1 / ((1-P) + (P/N))$$

where:

P = parallelizable proportion of program,

N = Number of processors.

This is, of course, an abstract, academic view. In practice, this provides a theoretical maximum speedup, as there are a number of overheads associated with concurrency. Synchronization overheads occur when a thread must wait for another task or tasks before it can continue execution. If a single task is slow, the whole program must wait. In addition, we will have critical sections of code, where only a single task is able to run at a time. We may also have occasions when all tasks are contending for the same resource or where no other tasks can be scheduled to run by the OS.

24.1 Thread safety and reentrancy

Functions which can be used concurrently by more than one thread concurrently must be both thread-safe and reentrant. This is particularly important for device drivers and for library functions.

For a function to be reentrant, it must fulfill the following conditions:

- all data must be supplied by the caller
- the function must not hold static or global data over successive calls
- the function cannot return a pointer to static data
- the function cannot itself call functions which are not reentrant.

For a function to be thread-safe, it must protect shared data with locks. (This means that the implementation needs to be changed by adding synchronization blocks to protect concurrent accesses to shared resources, from different threads.) Reentrancy is a stronger property, this means that not every thread-safe function is reentrant.

There are number of common library functions which are not reentrant. For example, the function `ctime()` returns a pointer to static data which is over-written on each call.

24.2 Performance issues

There are several multi-core specific issues relating to performance of threads:

Bandwidth The connection to external memory is shared between all processors within the MPCore. The individual cores run at speeds far higher than the external memory and so are potentially limited (in I/O intensive code) by the available bandwidth.

Thread dependencies and priority inversion

The execution of a higher priority thread can be stalled by a lower priority thread holding a lock to some shared data. Alternatively, an incorrect split in thread functionality can lead to a situation where no benefit is seen because the threads have fully serialized dependencies.

Cache contention and false sharing

If multiple threads are using data which reside within the same coherent cache lines, there can be cache line migration overhead even if the actual variables are not shared.

24.2.1 Bandwidth concerns

Bandwidth issues can be optimized in a number of ways. Clearly, the code itself must be optimized using the techniques described in [Chapter 19](#), to minimize cache misses and therefore reduce the bandwidth utilization.

Another option is to pay attention to thread allocation. The kernel scheduler does not pay any attention to data usage by threads; instead it makes use of priority to decide which threads to run. The programmer may be able to provide hints which allow more efficient scheduling through the use of thread affinity.

24.2.2 Thread dependencies

In real systems we can have threads with higher or lower priority which both access a shared resource. This gives scope for some potential difficulties. The term *starvation* is used to describe the situation where a thread is unable to get access to a resource after repeated attempts to claim it.

Priority inversion is said to occur when a lower priority task has a lock on a resource that a higher priority requires in order to be able to execute. In other words, a lower priority task prevents the higher priority task from executing. Priority inheritance resolves this by temporarily raising the priority of the task which has the lock to the highest level. This causes that task to execute as quickly as possible and relinquish the shared resource as soon as it can.

Operating systems (particularly real time operating systems) have ways to avoid such problems automatically. One method is not to allow lower-priority threads from directly accessing resources needed by higher-priority threads, they may need to use a higher-priority proxy thread to perform the operation. A similar approach is to temporarily increase the priority of the low-priority thread while it is holding the critical resource, ensuring that the scheduler will not pre-empt execution of that thread while in the critical selection.

A program that relies on threads executing in a particular sequence to work correctly may have a *race condition*. Single-core real-time systems often implicitly rely on tasks being executed in a priority based order. Tasks will then execute to completion, without pre-emption. Later tasks can rely on earlier tasks having completed. This can cause problems if such software is moved to a multi-core system without careful checking for such assumptions. A lower-priority task can run at the same time as a higher-priority task and the expected execution order of the original single-core system is no longer guaranteed. There are number of ways to resolve this. A simple

approach is to set task affinity to make those tasks run on the same processor. This requires little change to the legacy code, but does break the symmetry of the system and remove scope for load balancing. A better approach is to enforce serial execution through the use of the kernel synchronization mechanisms, which gives the programmer explicit control over the execution flow and better SMP performance, but does require the legacy code to be modified.

24.2.3 Cache thrashing

Processors implementing ARM architecture version 6 and later, including all ARM MPCore processors, use physically tagged caches which remove the need for flushing caches on context switch. In an SMP system, it is possible for tasks to migrate between the different processors in the system. The scheduler starts a task on a processor and it runs for a certain period and is then replaced by a different task. When that task is restarted at a later time by the scheduler, this could be on a different processor. This means that the task does not get the potential benefit of cache data already being present in the processor cache. Memory intensive tasks which quickly fill data cache might thrash each others' cached data. This has an impact on both performance (slower execution due to higher number of cache misses) and system energy usage (due to additional interaction with external memory). The ARM MPCore processor optimizations for cache line migration mitigate the effects of this. In addition, the OS scheduler can try to reduce the problem by aiming to keep tasks on the same processor. As we have seen, the programmer can also do this by setting processor affinity to threads and processes.

24.2.4 False sharing

This is a problem of systems with shared coherent caches and is effectively a form of involuntary memory contention. It can happen when a processor (or other block) regularly accesses data that is never changed by another processor and this data shares a cache line with data that will be altered by another processor. The MESI protocol can end up migrating data that is not truly shared between different parts of the memory system, costing clock cycles and power. Even though there is no actual coherency to be maintained, the MESI protocol invalidates the cache line, forcing it to be re-loaded on each write. However, the cache-to-cache migration capability of ARM MPCore processors reduces the overhead. Therefore, programmers should avoid having processors operating on independent data that is stored within the same cache line and increasing the level of detail for inner loop parallelization.

24.2.5 Deadlock and livelock

When writing code that includes critical sections, it is important to be aware of common problems that can break correct execution of the program:

- *Deadlock* is the situation where two (or more) threads are each waiting for another thread to release a resource. Such threads are effectively blocked, waiting for a lock that can never be released.
- *Livelock* occurs when multiple threads are able to execute, without blocking indefinitely (the deadlock case), but the system as a whole is unable to proceed, due to a repeated pattern of resource contention.

Both deadlocks and livelocks can be avoided either by correct software design, or by use of lock-free software techniques.

24.3 Profiling in SMP systems

ARM MPCores contain additional performance counter functions, which allow counting of the following SMP cache events:

- coherent linefill missed in all processors
- coherent linefill hit in other processor caches.

ARM DS-5 Streamline configures a default set of hardware performance counters that are a best-fit for optimizing applications. See [DS-5 Streamline on page 18-4](#) for more information.

Chapter 25

Power Management

Many ARM systems are mobile devices, powered by batteries. In such systems, optimization of power usage (in fact, it would be more accurate to look at total energy usage) is a key design constraint. Programmers often spend significant amounts of time trying to save battery life in such systems. Power-saving can also be of concern even in systems which do not use batteries. For example, we may wish to minimize energy usage for reduction of electricity costs to the consumer or for environmental reasons.

Built into ARM processors are many hardware design methods aimed at reducing power usage. In this chapter, we will focus on the features available to the programmer to reduce processor power consumption. Before we do so, let's review the components of processor energy usage.

Energy usage can be divided into two components – dynamic and static. Both are important. Static power consumption occurs whenever the processor logic or RAM blocks have power applied to them. In general terms, the leakage currents (any current that flows when the ideal current is zero) are proportional to the total silicon area – the bigger the chip, the more the leakage. The proportion of power consumption due to leakage gets significantly higher as we move to more advanced manufacturing process – they are much worse on fabrication geometries of 130nm and below. Dynamic power consumption occurs due to transistors switching and is a function of the processor clock speed and the numbers of transistors which change state per cycle. Clearly, higher clock speeds and more complex processors will consume more power.

25.1 Power and clocking

One way in which we can reduce energy usage is to remove power, which removes both dynamic and static currents (sometimes called *power gating*) or to stop the clock of the processor which removes dynamic power consumption only and can be referred to as *clock gating*.

ARM processors typically support a number of levels of power management, as follows:

- Run (Normal operation)
- Standby
- Shutdown
- Dormant.

In this chapter, we will examine each of these in turn and understand how power can be saved at each level. We will also look at the latency associated with switching between each of these modes. For certain operations, there is a requirement to save and restore state before and after removing power and both the time taken to do this and power consumed by this extra work can be an important factor in software selection of the appropriate power management activity.

The SoC device which includes the ARM processor can have further low power states, with names such as “STOP” and “Deep sleep.” These refer to the ability for the hardware *Phase Locked Loop* (PLL) and voltage regulators to be controlled by power management software.

25.1.1 Standby mode

In the standby mode of operation, the core is left powered-up, but most of its clocks are stopped, or *clock-gated*. This means that almost all parts of the processor are in a static state and the only power drawn is due to leakage currents and the clocking of the small amount of logic which looks out for the wake-up condition.

This mode is entered using either the WFI (Wait For Interrupt) or WFE (Wait For Event) instructions, which we’ll take a more detailed look at in a moment. We recommend use of a Data Synchronization Barrier (DSB) instruction before WFI or WFE, to ensure that pending memory transactions complete before changing state.

If a debug channel is active, it will remain active. The processor stops execution until a wakeup event is detected. The wakeup condition is dependent on the entry instruction. For WFI an interrupt or external debug request will wake the processor. For WFE, a number of specified events exist, including another processor in an MP system executing the SEV instruction. A request from the *Snoop Control Unit* (SCU) can also wake up the clock for a cache coherency operation in an MP system. This means that the cache of a processor which is in standby state will continue to be coherent with caches of other processors. A processor reset will always force the processor to exit from the standby condition.

Various forms of dynamic clock gating can also be implemented in hardware. For example the SCU, GIC, timers, CP15, instruction pipeline or NEON blocks can be automatically clock gated when an idle condition is detected, to save power.

Standby mode can be entered and exited quickly (typically in two-clock-cycles). It therefore has an almost negligible affect on the latency and responsiveness of the processor.

25.1.2 Dormant mode

In dormant mode, the actual processor logic is powered down, but the cache RAMs are left powered up. Often the RAMs will be held in a low-power retention state where they hold their contents but are not otherwise functional. This provides a far faster restart than complete shutdown, as live data and code persists in the caches. Again, in a multi-processor system, individual processors can be placed in dormant mode.

In an MPCore system which allows individual processors within the SMP cluster to go into dormant mode, there is no scope for maintaining coherency while the processor has its power removed. Such processors must therefore first isolate themselves from the coherence domain. They will clean all dirty data before doing this and will typically be woken up via another processor telling the external logic to re-apply power. The woken processor will then need to restore the original processor state before rejoining the coherency domain. As the memory state may have changed while the processor was in dormant mode, it might well need to invalidate the caches anyway. Dormant mode is therefore much more likely to be useful in a single processor environment rather than in a multi-processor set-up. This is due to the additional expense of leaving and rejoining the coherency domain. In an MPCore cluster, dormant mode is typically likely to be used only by the last processor when the other processors have already been shutdown.

25.1.3 Assembly language power instructions

ARM assembly language includes instructions which can be used to place the processor in a low power state. The architecture defines these instructions as *hints* – the processor is not required to take any specific action when it executes them. In the Cortex-A processor family, however, these instructions are implemented in a way which shuts down the clock to almost all parts of the processor. This means that the power consumption of the processor is significantly reduced – only static leakage currents will be drawn, and there will be no dynamic power consumption.

The WFI instruction has the effect of suspending execution until the processor is woken up by one of the following conditions:

- an IRQ interrupt (even if CPSR I-bit is set)
- an FIQ interrupt (even if CPSR F-bit is set)
- an asynchronous abort
- a Debug Entry request (even if JTAG Debug is disabled).

In the event of the processor being woken by an interrupt when the relevant CPSR interrupt flag is disabled, the processor will implement the next instruction after WFI. On older versions of the ARM architecture, the wait for interrupt function (also called standby mode) was accessed using a CP15 operation, rather than a dedicated instruction.

The WFI instruction is widely used in systems which are battery powered. For example, mobile telephones can place the processor in standby mode many times a second, while waiting for the user to press a button.

WFE is similar to WFI. It suspends execution until an event occurs. This can be one the events listed above, or an additional possibility – an event signaled by another processor in an MPCore. Other processors can signal events by executing the SEV instruction. SEV signals an event to all processors in a multi-processor system. We will describe WFE further when we look at multi-processor systems.

25.1.4 Dynamic Voltage and Frequency Scaling

Many systems operate under conditions where their workload is very variable. It would be useful to have the ability to reduce or increase the processor performance to match the expected processor workload. If we could clock the processor more slowly when it is less busy, we could save dynamic power consumption. The dynamic power consumption has a linear correlation with clock frequency, but a quadratic relationship with voltage (it is proportional to the square of the supply voltage).

If the processor is running more slowly, it is also the case that its supply voltage can be reduced somewhat. This hardware approach is called *Dynamic Voltage and Frequency Scaling* (DVFS). The advantage of reducing supply voltage is that it reduces both dynamic and static power. Compared to the alternative of running fast, then entering standby, then running fast and so forth, the approach of running slowly at a lower supply can save energy. To do this successfully requires two difficult engineering problems to be solved. The SoC needs a way in which software running on the ARM processor can reliably modify the clock speed and supply voltage of the processor, without causing problems in the system. (This needs things like voltage level-shifters and split power supplies on chip to cope with the variable supply, plus synchronizers between voltage domains to cope with timing changes.) Of equal importance is the ability of software running in the system to make accurate predictions about future workloads to set the voltage and clock speed accordingly.

Chapter 26

Security

The term *security* is used in the context of computer systems to cover a wide variety of features. For the purposes of this chapter, we will use a narrower definition. A secure system means one which protects assets (resources which need protecting, for example passwords, or credit card details) from a range of plausible attacks to prevent them from being copied or damaged or made unavailable (denial of service). Confidentiality is a key security concern for assets such as passwords and cryptographic keys. Defense against modification and proof of authenticity is vital for security software and on-chip secrets used for security. Examples of secure systems might include entry of Personal Identification Numbers (PINs) for such things as mobile payments, digital rights management, and e-Ticketing.

Security is harder to achieve in today's world of open systems where a wide range of software can be downloaded onto a platform. This gives the potential for malevolent or untrusted code to tamper with the system.

ARM processors include specific hardware extensions to allow construction of secure systems. Creating secure systems is outside the scope of this book. In the remainder of this chapter, we present the basic concepts behind the ARM Security Extensions (TrustZone). If your system is one which makes use of these extensions, you should be aware that this imposes some restrictions on the operating system and on unprivileged code (in other words, code which is not part of the secure system). TrustZone is of little or no use without memory system support.

It should, of course, be emphasized, that no security is absolute!

26.1 TrustZone hardware architecture

The TrustZone hardware architecture aims to provide resources that enable a system designer to build secure systems. It does this through a range of components and infrastructure additions. Low-level programmers need to have some awareness of the restrictions placed on the system by the TrustZone architecture, even if they are not intending to make use of the security features.

In essence, system security is achieved by dividing all of the device's hardware and software resources, so that they exist in either the Secure world for the security subsystem, or the Normal world for everything else. System hardware ensures that no Secure world resources can be accessed from the Normal world. A secure design places all sensitive resources in the Secure world, and has robust software running which can protect assets against a wide range of possible attacks.

Note that the use of the term *Non-Secure* is used in the *ARM Architecture Reference Manual* as a contrast to *Secure* state, but this does not imply that there is a security vulnerability associated with this state. We will refer to this as *Normal* operation here. The use of the word *world* is to emphasize the relationship between the Secure world and other states the device is capable of.

The additions to the processor enable a single physical processor to execute code from both the Normal world and the Secure world in a time-sliced fashion. The memory system is similarly divided. An additional bit, indicating whether the access is Secure or Non-Secure (the NS bit) is added to all memory system transactions, including cache tags and access to system memory and peripherals. This can be considered as an additional address bit, giving a 32-bit physical address space for the Secure world and a completely separate 32-bit physical address space for the Normal world.




Figure 26-1 Switching between Normal and Secure worlds

As the processor cores execute code from the two worlds in a time-sliced fashion, context switching between them is done using an additional core mode (like the existing modes for IRQ, FIQ etc.) called Monitor mode. A limited set of mechanisms by which the physical processor can enter Monitor mode from the Normal world is provided. Entry to monitor can be through a dedicated instruction, the Secure Monitor Call (SMC) instruction, or by hardware exception mechanisms. IRQ, FIQ and external aborts can all be configured to cause the processor to switch

into Monitor mode. In each case, this will appear as an exception to be dealt with by the Monitor mode exception handler. Figure 26-1 on page 26-2 provides a conceptual summary of this switching.

Figure 26-2 shows how, in many systems, FIQ is reserved for use by the secure world (it becomes, in effect, a non-maskable secure interrupt). An IRQ which occurs when in the Normal world is handled in the normal way, described in the chapters on exception handling. An FIQ which occurs while executing in the Normal world is vectored directly to Monitor mode. Monitor mode handles the transition to Secure world and transfers directly to the Secure world FIQ handler. If the FIQ occurs when in the Secure world, it is handled through the Secure vector table and routed directly to the Secure world handler. IRQs are typically disabled during execution in the Secure world.




Figure 26-2 Banked out registers

The software handler for Monitor mode is implementation specific, but will typically save the state of the current world and restore the state of the world being switched to, much like a normal context switch.

The NS bit in the Secure Configuration Register (SCR) in CP15 indicates which world the processor is currently in. In Monitor mode, the processor is always executing in the Secure world, regardless of the value of the SCR NS-bit, which is used to signal which world you were previously in. The NS-bit also enables code running in Monitor mode to snoop security banked registers, to see what is in either world.

TrustZone hardware also effectively provides two virtual MMUs, one for each virtual processor. This enables each world to have a local set of translation tables, with the Secure world mappings hidden and protected from the Normal world. The page table descriptions include a NS bit, which is used to determine whether accesses are made to the secure or non-secure physical address space. Although the page table entry bit is still present, the Normal virtual processor

hardware does not use this field, and memory accesses are always made with NS = 1. The Secure virtual processor can therefore access either Secure or Normal memory. Cache and TLB hardware permits Normal and Secure entries to co-exist.

It is good practice for code which modifies page table entries and which does not care about TrustZone based security, to always set the page table NS-bit to zero. This means that it will be equally applicable when the code is executing in the Secure or Normal worlds.

The ability to direct aborts, IRQ and FIQ directly to the monitor, enables trusted software to route the interrupt request accordingly, which permits a design to provide secure interrupt sources immune from manipulation by the Normal world software. Similarly, the Monitor mode routing means that from the point of view of Normal world code, an interrupt that occurs during Secure world execution appears to occur in the last Normal world instruction that occurred before the Secure world was entered.

A typical implementation is to use FIQ for the Secure world and IRQ for the Normal world. Exceptions are configured to be taken by the current world (whether Secure or Normal), or to cause an entry to the monitor. The monitor has its own vector table. As a result of this, the processor has three sets of exception vector tables. It has a table for the Normal world, one for the Secure world, and one for Monitor mode.

The hardware must also provide the illusion of two separate cores within CP15. Sensitive configuration CP15 registers can only be written by Secure world software. Other settings are normally banked in the hardware, or by the Monitor mode software, so that each world sees its own version.

Implementations which use TrustZone will typically have a light-weight kernel (Trusted Execution Environment) which hosts services (for example, encryption) in the Secure world. A full OS runs in the Normal world and is able to access the secure services via SMC. In this way, the Normal world gets access to functions of the service, without any ability to see keys or other protected data.

26.1.1 Multi-processor systems with security extensions

Each processor in a multi-core system has the programmer's model features described for uniprocessors earlier in this book. Any number of the processors in the cluster can be in the Secure world at any point in time, and processors are able to transition between the worlds independently of each other. The Snoop Control Unit is aware of security settings. Additional registers are provided to control whether Normal world code can modify SCU settings. Similarly, the generic interrupt controller which distributes prioritized interrupts across the Multi-processor cluster must also be modified to be aware of security concerns.

Theoretically, the Secure world OS on an SMP system could be as complicated as the Normal world OS. However, this is highly undesirable when aiming for security. In general, it is expected that a Secure world OS will actually only execute on one core of an SMP system (with security requests from the other cores being routed to this chosen core). This does provide some bottleneck issues. To some extent these will be balanced by the Normal world OS performing load balancing against the core that it will see as busy for unknown reasons. Beyond that this limitation has to be seen as one of the compromises that can be reached to hit a particular target level of security.

26.1.2 Interaction of Normal and Secure worlds

If you are writing code in a system which contains some secure services, it can be useful to understand how these are used. As we have seen, a typical system will have a light-weight kernel, *Trusted Execution Environment* (TEE) hosting services (for example, encryption) in the

Secure world. This interacts with a full OS in the Normal world, which can access the secure services using the SMC call. In this way, the Normal world is able to have access to functions of the service, without getting to see keys (for example).

Generally applications developers won't directly interact with TrustZone (or TEEs or Trusted Services). Instead, one makes use of a high level API (for example, it might be called `reqPayment()`) provided by a Normal world library. The library would be provided by the same vendor as the Trusted Service (for example, a credit card company), and would handle the low level interactions. Figure 26-3 shows this interaction and illustrates the flow from user application calling the API, which makes an appropriate OS call, which then passes to the TrustZone driver code, which passes execution into the TEE, through the Secure monitor.




Figure 26-3 Interaction with TrustZone

It is common to share data between the Secure and Normal worlds. For example, in the Secure world you might have a signature checker. The Normal world can request that the Secure world verifies the signature of a downloaded update, using the SMC call. The Secure world needs access to the memory used by the Normal world to store the package. The Secure world can use the NS-bit in its page table descriptors to ensure that it used non-secure accesses to read the data. This is important because data relating to the package might already be in the caches, due to the accesses done by the Normal world. These accesses with addresses marked as non-secure. As mentioned previously the security attribute can be thought of as an additional address bit. If the core used secure access to try to read the package, it would not hit on data already in the cache.

If you are a Normal world programmer, in general, you can ignore something happening in the Secure world, as its operation is hidden from you. One side effect is that interrupt latency can increase slightly, if an interrupt goes off in the Secure world, but this increase is small compared to the overall latency on a typical OS.

If you do need to access a secure application, you will need a driver-like function to talk to the Secure world OS and Secure applications, but the details of creating that Secure world OS and applications are beyond the scope of this book. Programmers writing code for the Normal world only need to know the particular protocol for the secure application being called.

Finally, the TrustZone System also controls availability of debug provision. Separate hardware over full JTAG debug and trace control is separately configurable for Normal and Secure software worlds, so that no information about the Secure system leaks.

Chapter 27

Virtualization

Modern compute subsystems are powerful but often under utilized. The domains that these systems address increasingly require multiple software environments working simultaneously on the same physical processor systems. Software applications and their environments may need to be completely separated for reasons of robustness, differing requirements for real-time behavior, or simply to have them isolated from each other. This is done by providing *virtual* processors for the software to execute on.

Implementing such virtual processors in an efficient fashion needs both dedicated hardware extensions (to accelerate switching between virtual machines) and software hypervisors. The volume of legacy software that needs to be ported to new hardware may make it more expedient to provide a virtual machine matching the legacy OS requirements than to port the OS and drivers to new hardware and OS versions. The Virtualization Extensions to the ARMv7 architecture provide a standard hardware accelerated implementation allowing the creation of high performance hypervisors. Writing a hypervisor is outside the scope of this book, as such software is typically produced by specialized teams. In this chapter, we will provide an overview of how the Virtualization Extensions work and the effects of this on OS and application code, in a fashion similar to our previous discussion of the TrustZone Security Extensions. The ARM Cortex-A15 processor is the first device implementing the Virtualization Extensions.

The term hypervisor actually dates back to the 1960s and was first used by IBM to describe software running on mainframes. Today's hypervisors are conceptually very similar and may be thought of as operating at a level higher than the supervisor or operating system. They typically allow multiple instances of one or more different operating systems, called Guest Operating Systems (*Guest OS*), to run on the system. The hypervisor provides a virtual system to each of the Guest operating systems and monitors their execution. The Guest OS does not typically need to be aware that it is running under a hypervisor or that other operating systems may be using the system. The term Virtual Machine Monitor (VMM) is sometimes used to describe a hypervisor. There are

two types of hypervisors being deployed on ARM. Type1 is as described with each Virtual Machine (VM) containing Guest OS. In Type 2, the hypervisor is an extension of the Host OS with each subsequent Guest OS contained in a separate VM.

One possible use of virtual machine hypervisors is in so-called cloud computing solutions, where software may be partitioned into client or server devices or where large amounts of data or content exist. This scenario is likely to increase the amount of addressable physical memory required in the system. For this reason, the Virtualization Extensions require that the processor also implement the *Large Physical Address Extension (LPAE)* which was described in [Chapter 10](#). This allows each of the multiple software environments, as well as different processes within, to access separate windows of physical addresses. LPAE provides an additional level of MMU translation table, so that each 32-bit virtual memory address can be mapped within a 40-bit physical memory range. In this scenario, this permits software to allocate enough physical memory to each virtual machine, even when total demands on memory exceed the range of 32-bit addressing. It would also be possible for a single operating system kernel to handle up to 40 bits of physical address space, with up to 4GB available at any given time. In theory, this can mean up to ~3GB per process.

27.1 ARMv7-A Virtualization Extensions

The ARMv7-A architecture Virtualization Extensions add a number of additional features to the programmer's model of the processor. These are:

- A hypervisor mode, in addition to the existing privileged modes.
- An additional memory translation stage is added. Previously, the MMU would translate the processor's virtual address to a physical address. Now, the virtual address is translated to an intermediate physical address (using the mappings provided by the OS itself). This intermediate physical address is then translated to a physical address, with access to real memory being supervised by the hypervisor.
- Hardware interrupts can be configured such that they are directed to the hypervisor. Registers are provided to hold virtual interrupts – called *guest IRQs*.

27.1.1 Privilege model in ARMv7-A Virtualization Extensions

The ARMv7-A Virtualization Extensions introduce a privilege model a little different to that of previous processors. The Cortex-A15 processor has three different levels of privilege:

- PL0 (referred to as Unprivileged elsewhere in this book). This is the privilege level of User mode.
- PL1 (referred to as Privileged elsewhere in this book). This is the privilege level of SVC, Undef, Abort, FIQ, IRQ, System and Monitor modes.
- PL2 (new level for processors which implement the Virtualization Extensions). This is the privilege level of Hyp mode.

These privilege levels are separate from the TrustZone Secure and Normal (Non-secure) settings.

27.1.2 Hypervisor mode

Hypervisor mode operates at the PL2 level of privilege. It has access to its own versions of R13 (SP) and SPSR. It uses the User mode link register for storing function return addresses, but has a dedicated register, called ELR_hyp, to store the exception return address. Hyp mode is available only in the Normal state and provides facilities for virtualization, which are only accessible in this mode. [Figure 27-1](#) shows a representation of systems with and without virtualization.




Figure 27-1 Hypervisor layer

When the processor is operating in Hyp mode, it uses Hyp mode vector entries, which are taken from a dedicated vector table belonging to the hypervisor. Hypervisor mode is entered through a special exception process called *Hyp trap entry* which makes use of the previously reserved 0x14 address within the vector table (see [Table 12-1 on page 12-5](#)). A dedicated register (the Hyp Syndrome Register) gives information to the hypervisor about the exception or other reason for entering the hypervisor (a trapped CP15 operation, for example).

27.1.3 Memory translation

The memory translation process in a processor which implements the Virtualization Extensions can be thought of as comprising four different MMUs (although the underlying hardware is likely to be a single actual block):

- A Normal world PL1&0 Stage 1 MMU. The operation of this MMU is typically under control of the Guest OS which may use two sets of translation tables, defined by the Normal world copies of TTBR1 and TTBR0, and controlled by the Normal world copy of TTBCR. This MMU supports both the Short-descriptor and Long-descriptor formats.
- A Normal world PL1&2 Stage 2 MMU. This performs translation between the Intermediate Physical Address (IPA) and the Physical Address (PA). The VTTBR defines the translation table for this MMU under control by the VTCR register. This MMU supports only the long-descriptor format.
- A Normal world PL2 Stage 1 MMU. This performs the virtual to physical address translations used for the hypervisor software itself. It is controlled by the HTCR and its translation table is located by the setting of the HTTBR. This MMU supports only the long-descriptor format.
- A Secure MMU used in secure state at both PL0 and PL1. Its operation is controlled by two sets of translation tables, whose use is governed by the secure versions of TTBCR, TTBR0 and TTBR1. This produces a physical address from a virtual address. This MMU supports both the short-descriptor and long-descriptor formats.

The MMU translation process now has two stages. The first stage translates to a physical address – either an Intermediate Physical Address when two stages are used or directly to a physical address when only one-stage translation is required. The second translation stage is used to translate an IPA to a PA.

Writers of device drivers that pass the physical address of buffers to other bus masters, such as DMA controllers or communications subsystems, will need to ensure that the other bus masters receive the correct physical address information when the driver is virtualized.

27.2 Hypervisor exception model

Asynchronous exceptions (IRQ, FIQ and asynchronous or asynchronous aborts) may be routed to Hypervisor mode. This behavior may be selected through dedicated control bits in the *Hyp Configuration Register* (HCR). A further bit in this register allows synchronous exceptions (undef, SVC or precise aborts) to also be routed to the Hypervisor Trap entry. The hypervisor may also be called through the HVC instruction. If this is performed while in privilege level PL1, it causes a hypervisor trap exception and entry into hypervisor mode. If performed when already in hypervisor mode, it causes an HVC exception.

The concept of virtual exceptions is introduced to the architecture – we have a virtual abort, a virtual IRQ and a virtual FIQ, each corresponding to their physical equivalent. These are taken in their conventional exception mode (that is, a virtual abort is taken in Abort mode, and so on). This concept is useful only when the corresponding physical exception is routed to Hyp mode. This means that the hypervisor has control of masking and generating a virtual exception. A real physical exception occurs, and is routed to the hypervisor software, which then signals the virtual exception to the current Guest OS. The Guest OS handles the exception as it would do for an equivalent physical exception and does not need to be aware that the hypervisor has been involved. The virtual exception is signalled through the virtual GIC, or by using dedicated HCR bits. When exceptions are routed to the hypervisor in this fashion, the CPSR A, I and F bits no longer mask physical exceptions, instead they mask the handling of virtual exceptions within the Guest OS.

27.3 Relationship between virtualization and ARM Security Extensions

We discussed the concept of privilege levels, PL0, PL1 and PL2, in [Privilege model in ARMv7-A Virtualization Extensions on page 27-3](#). Figure 27-2 illustrates how in the Normal world, we have PL0 (User mode), PL1 (for exception modes) and PL2 for the hypervisor, while in the Secure world, we have only PL0 and PL1, with secure monitor mode at PL1.




Figure 27-2 Privilege levels and security

In general, the hypervisor has no impact on the Secure world. It is unable to access secure parts of memory or otherwise interfere in the correct operation of the secure system.

Chapter 28

Introducing big.LITTLE

This chapter gives an introduction to ARM big.LITTLE technology, which provides new options for matching compute capacity to work load by appropriately distributing the load between different processors.

Traditionally, it has not been possible to design a processor capable of both high performance and high energy efficiency. Solutions typically involve integrating processors with different micro-architectures optimized for performance and energy efficiency respectively. For example a high performance application processor coupled with a low power ASIC.

This solution is often referred to as *Heterogeneous Multiprocessing* (HMP), and has been used for some time. Heterogeneous computing involves putting a number of specialized processors together – applications processors and graphics processors, for example – and sharing out tasks to each processor for the work that suits it best.

The ARM solution to this problem is to couple a high performance, low power processor design with a lower power, energy efficient processor while still retaining full *Instruction Set Architecture* (ISA) compatibility. This solution provides a wider dynamic range of power and performance control than would be possible using designs with only one type of processor.

28.1 big.LITTLE configuration

The ARM big.LITTLE design uses a Cortex-A7 processor as the LITTLE component and a Cortex-A15 processor as the big component, as in [Figure 28-1](#). These two processors are implementations of the same version of the ARM architecture, but with different performance levels. The two processors are architecturally identical, implementing the ARMv7-A architecture, including Virtualization and LPAE. An application compiled for one processor will execute in an architecturally consistent way on the other processor.

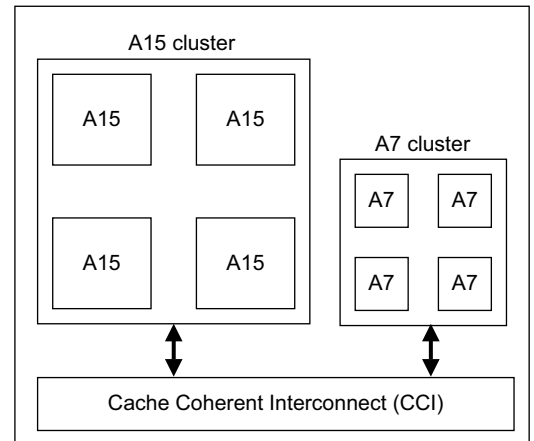


Figure 28-1 A big.LITTLE system

The Cortex-A7 processor is capable of handling most low intensity tasks such as audio playback, web-page scrolling, operating system events, and other always on, always connected tasks. As such, it is likely that the Cortex-A7 cluster is where the software stack will remain until intensive tasks such as gaming or video processing are run.

The Cortex-A15 processor can be utilized for heavy workloads such as high performance graphics. A coupling of these two processors provides opportunities to save energy as well as satisfy the increasing performance demands of applications stacks in mobile devices.

28.2 Structure of a big.LITTLE system

To support a big.LITTLE configuration requires a Cache Coherent Interconnect (CCI), such as the CoreLink CCI-400, that maintains cache memory coherency between the two processors, as well as IO coherency for components such as a graphics processor. In addition, such a system also requires a shared generic interrupt controller, such as the CoreLink GIC-400, allowing interrupts to be migrated between any cores in the clusters. Firmware takes care of task switching so the process is invisible to the operating system and applications. A typical system is shown in Figure 28-2.

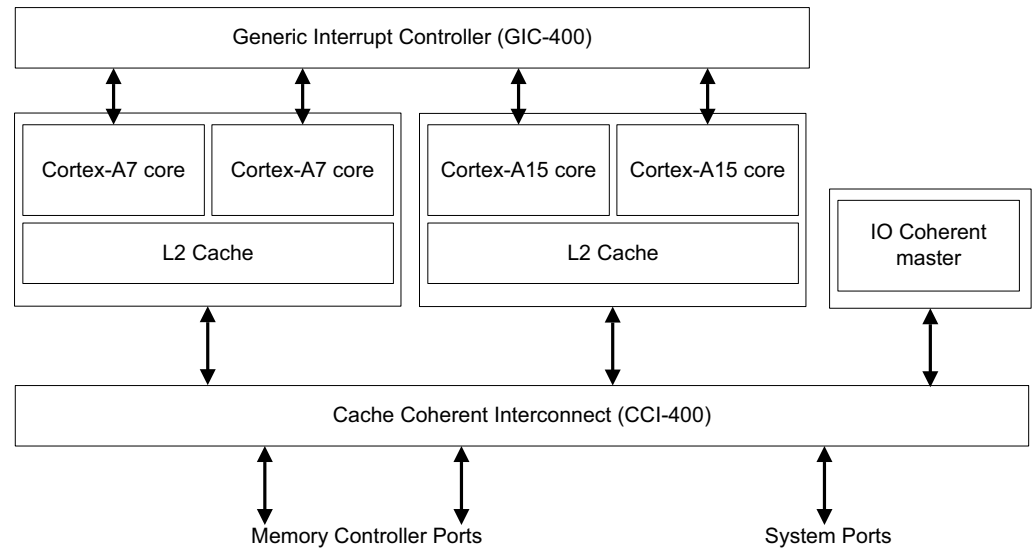


Figure 28-2 Basic big.LITTLE system

As a fully coherent system can create a significant volume of coherent transactions, the processor clusters and the CCI-400 have been designed to cope with worst case snooping scenarios. This includes the case where a graphics processor is connected to one of the I/O coherent CCI-400 ports and every transaction is snooping the processors caches at the same time as they are snooping each other.

From the perspective of trace and debug, both Cortex-A15 and Cortex-A7 processors offer trace solutions and are both compliant with the Debug v7.1 architecture. Full support for big.LITTLE debug and trace is provided through a CoreSight™ SoC.

28.3 Execution models in big.LITTLE

Since the same application can be made to run on a Cortex-A7 or Cortex-A15 processor without modification, this opens up the possibility to map applications to the right processor on an opportunistic basis. This fact is the basis for a number of execution models, namely:

- [big.LITTLE migration models](#)
- [big.LITTLE MP operation on page 28-9](#)

As the name suggests, migration models enable the capture and restoration of software context from one processor type to another

The MP operation model enables the software stack to be distributed across processors in both clusters. All processors may be in operation simultaneously (although in practice this is unlikely).

28.3.1 big.LITTLE migration models

Migration models are further divided into two types, namely:

- [Cluster migration on page 28-6](#)
- [CPU migration on page 28-7](#)

Cluster migration involves software context switching of the complete cluster while CPU migration involves software context switching of individual cores within the processor. The software stack only ever runs on a single cluster in the case of cluster migration. In CPU migration, each core in a cluster is paired with its counterpart on the other cluster and the software context is migrated opportunistically between clusters on a per-core basis. CPU migration is a finer grained form of cluster migration.

These execution models have different trade-offs in terms of ease of adoption and flexibility of operation.

Migration models are a natural extension to power-performance management techniques such as DVFS, (see [Dynamic Voltage and Frequency Scaling on page 25-4](#)). A migration action is similar to a DVFS operating point transition. Operating points on a processor's DVFS curve will be traversed in response to load variations. When the current processor (or cluster) has attained the highest operating point, if the software stack requires more performance, a processor (or cluster) migration action is effected (see [Figure 28-3 on page 28-5](#)). Execution then continues on the other processor (or cluster) with the operating points on this processor (or cluster) being traversed. When performance isn't needed, execution can switch back.




Figure 28-3 Processor DVFS curves

In a big.LITTLE system these operating points are applied both to the Cortex-A15 and Cortex-A7 processors. When the Cortex-A7 processor is executing the OS can tune the operating points as it would for an existing platform with a single applications processor. When the Cortex-A7 processor is at its highest operating point (Figure 28-3), if more performance is required a switch is invoked that transfers the OS and applications to the Cortex-A15 processor. Further DVFS tuning takes place on the Cortex-A15 processor if required, as the operating load increases.

Migration requires rapid context switching capability. Coherency is clearly a critical enabler in achieving a fast task migration time as it allows the state that has been saved on the *outbound* (migrated from) processor to be snooped and restored on the *inbound* (migrated to) processor rather than going via main memory. Additionally, for Cluster migration, (or for CPU migration when all processors have been switched) because the L2 cache of the outbound processor is coherent it can remain powered up after a task migration to improve the cache warming time of the inbound processor through snooping of data values. However, since the L2 cache of the outbound processor cannot be allocated, it will eventually need to be cleaned and powered off to save leakage power. The switching sequence is described in Figure 28-4 on page 28-6.





Figure 28-4 Migration process

28.3.2 Cluster migration

Only one cluster, either big or LITTLE, is active at any one time, except very briefly during a cluster context switch to the other cluster. To achieve the best power and performance efficiency, the software stack runs mostly on the energy-efficient Cortex-A7 cluster and only runs for short time periods on the Cortex-A15 processor.

**Figure 28-5 Cluster switching**

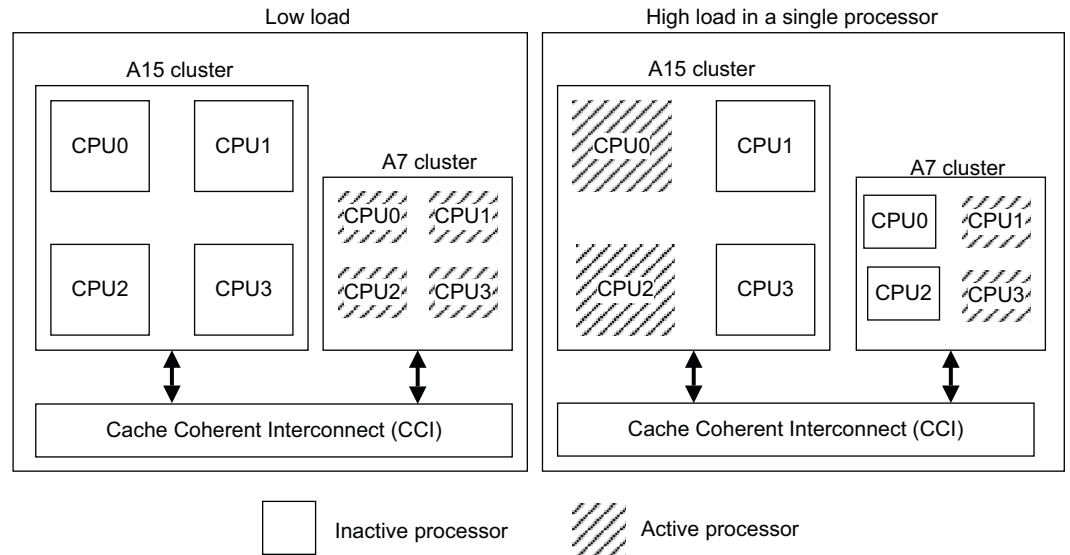
The system will consider the overall performance loading on the currently active cluster. Typically this would be the load of the processor that is busiest in the cluster. If the load warrants a change from big to LITTLE, or vice versa, the system synchronizes all the cores and then transfers all compute context to the other cluster. As part of this process, the system software on every processor has to save its state whilst still operating on the old (outbound) cluster. When the new (inbound) cluster boots up, the operating systems need to restore their state on each processor. Once the inbound cluster has restored context, the outbound cluster is switched off.

The mode works most efficiently with a symmetric big.LITTLE system (with same number of cores in the two clusters). An asymmetric system would require additional system software involvement to scale execution down to the least common number of processors before a switch could take place and while this is very much possible, it will increase the switching latency.

Cluster migration can be implemented alongside existing power management functionality (such as idle management) with about the same complexity.

28.3.3 CPU migration

CPU migration involves the use of both clusters at the same time with restrictions on which processors are switched on at any given moment. In CPU migration, each processor on the LITTLE cluster is paired with a processor on the big cluster. Processors are divided in CPU pairs (CPU0 on the Cortex-A15 and Cortex-A7 processors, CPU1 on the Cortex-A15 and Cortex-A7 processors and so on). When using CPU migration only one CPU per CPU pair can be used at any one time.

**Figure 28-6 CPU migration**

The system actively monitors the load on each processor. High load causes the execution context to be moved to the big processor, and equally, when the load is low, the execution is moved to the LITTLE processor. Only one processor in the pairing can be active at any time. When the load is moved from an outbound processor to an inbound processor, the former is switched off. This model allows a mix of big and LITTLE processors to be active at any one time, and support the use of asymmetric topologies (an unequal number of big and LITTLE processors).

28.4 big.LITTLE MP operation

Since a big.LITTLE system is fully coherent through the CCI-400, another approach is to allow both Cortex-A15 and Cortex-A7 processors to be powered on and simultaneously executing code. This is called big.LITTLE MP (which is essentially Heterogeneous Multi-Processing). This is the most sophisticated and flexible mode for a big.LITTLE system, involving scaling a single execution environment across both clusters. In this use model, a Cortex-A15 processor is powered on and executing simultaneously with a Cortex-A7 processor if there are threads that need such a level of processing performance. If not, only the Cortex-A7 processor needs to be powered on. Since processors are not explicitly matched, asymmetric topologies are simpler to support with MP operation.

With big.LITTLE MP, the operating system requires a higher degree of customization to extract maximum benefit from the design. For example, the scheduler subsystem will need to be aware of the power-performance capabilities of the different processors and will need to map tasks to suitable processors. Here the operating system runs on all processors in all clusters, which may be operating simultaneously. There is no explicit switching of complete software context, as in the two migration modes. Instead, the operating system migrates tasks between processors in the two clusters. Since the scheduler is involved in the directed placement of suitable tasks on suitable processors, this is a complex mode. The OS attempts to map tasks to processors that are best suited to running those tasks, and will power off unused, or under-used processors.

Multi-processor support in the Linux kernel assumes that all processors have identical performance capabilities and therefore a task can be allocated to any available processor. Supporting full big.LITTLE MP requires significant changes to the scheduling and power management components of Linux.

Chapter 29

Debug

Debugging is a key part of software development and is often considered to be the most time consuming (and therefore expensive) part of the process. Bugs can be difficult to detect, reproduce and fix and it can be difficult to predict how long it will take to resolve a defect. The cost of resolving problems grows significantly when the product is delivered to a customer. In many cases, when a product has a small time window for sales, if the product is late, it can miss the market opportunity. Therefore, the debug facilities provided by a system are a vital consideration for any developer.

Many embedded systems using ARM processors have limited input/output facilities. This means that traditional desktop debug methods (such as use of `printf()`) may not be appropriate. In such systems in the past, developers might have used expensive hardware tools like logic analyzers or oscilloscopes to observe the behavior of programs. The processors described in this book have caches and are part of a complex system-on-chip containing memory and many other blocks. There may be no processor signals which are visible off-chip and therefore no ability to monitor behavior by connecting up a logic analyzer (or similar). For this reason, ARM systems typically include dedicated hardware to provide wide-ranging control and observation facilities for debug.

29.1 ARM debug hardware

Cortex-A series processors provide hardware features which enable debug tools to provide significant levels of control over processor activity and to non-invasively collect large amounts of data about program execution. We can sub-divide the hardware features into two broad classes, *invasive* and *non-invasive*.

Invasive debug provides facilities which enable us to stop programs and step through them line by line (either at the C source level, or stepping through assembly language instructions). This can be by means of an external device which connects to the processor using the chip JTAG pins, or (less commonly) by means of debug monitor code in system ROM. JTAG stands for Joint Test Action Group and refers to the IEEE-1149.1 specification, which was originally designed to standardize testing of electronic devices on boards, but is now widely re-used for processor debug connection. A JTAG connection typically has five pins – two inputs, plus a clock, a reset and an output.

The debugger gives the ability to control execution of the program, allowing us to run code to a certain point, halt the processor, step through code and resume execution. We can set breakpoints on specific instructions (causing the debugger to take control when the processor reaches that instruction). These work using one of two different methods. Software breakpoints work by replacing the instruction with the opcode of the BKPT instruction. Obviously, these can only be used on code which is stored in RAM, but have the advantage that they can be used in large numbers. The debug software must keep track of where it has placed software breakpoints and what opcodes were originally located at those addresses, so that it can put the correct code back when we wish to execute the breakpointed instruction. Hardware breakpoints use comparators built into the processor and stop execution when execution reaches the specified address. These can be used anywhere in memory, as they do not require changes to code, but the hardware provides limited numbers of hardware breakpoint units (typically four in the Cortex-A family). Debug tools can support more complex breakpoints (for example stopping on any instruction in a range of addresses, or only when a specific sequence of events occurs or hardware is in a specific state). Data watchpoints give debugger control when a particular data address or address range is read or written. These can also be called data breakpoints.

Upon hitting a breakpoint, or when single-stepping, we can inspect and change the contents of ARM registers and memory. A special case of changing memory is code download. Debug tools typically enable the user to change our code, recompile and then download the new image to the system.

29.1.1 Debug events

A debug event is some part of the process being debugged that causes the system to notify the debugger. Debug events can be synchronous or asynchronous. Breakpoints, the BKPT instruction, and Watchpoints are all synchronous debug events. When any of these events occur, the processor can respond in one of a number of ways:

- It can ignore the debug event.
- It can take a debug exception.
- It will enter one of two debug modes, depending on the setup of the *Debug Status and Control Register* (DSCR):
 - Monitor debug mode.
 - Halt Debug mode.

Both of these are examples of invasive debug.

Halt debug mode

In Halt Debug mode, a debug event causes the processor to enter Debug state. The processor is halted and isolated from the rest of the system. This means that the debugger displays memory as seen by the processor, and the effects of memory management and cache operations will become visible.

In Debug state, the processor stops executing instructions from the location indicated by the program counter, and is instead controlled through the external debug interface, in particular using the *Debug Instruction Transfer Register* (DBGITR). This enables an external agent, such as a debugger, to interrogate processor context and control all subsequent instruction execution. Both the processor and system state can be modified. Because the processor is stopped, no interrupts will be handled until execution is restarted by the debugger.

Monitor debug-mode

In Monitor debug-mode, a debug event causes a debug exception to occur, either related to the instruction execution that generates a Prefetch Abort exception, or a data access that generates a Data Abort exception. Both of these must be handled by the software debug monitor. Since the processor is still operating, interrupts can still be serviced.

29.2 ARM trace hardware

Non-invasive debug (often called trace in ARM documentation) enables observation of the processor behavior while it is executing. It is possible to record memory accesses performed (including address and data values) and generate a real-time trace of the program, seeing peripheral accesses, stack and heap accesses and changes to variables. For many real-time systems, it is not possible to use invasive debug methods. Consider, for example, an engine management system – while we may be able to stop the processor at a particular point, the engine will keep moving and we will not be able to do useful debug. Even in systems with less onerous real-time requirements, trace can be very useful.

Trace is typically provided by an external hardware block connected to the processor. This is known as an *Embedded Trace Macrocell* (ETM) or *Program Trace Macrocell* (PTM) and is an optional part of an ARM processor based system. System-on-chip designers can omit this block from their silicon to reduce costs. These blocks observe (but do not affect) the processor behavior and are able to monitor instruction execution and data accesses.

There are two main problems with capturing trace. The first is that with today's very high processor clock speeds, even a few seconds of operation can mean trillions of cycles of execution. Clearly, to look at this volume of information would be extremely difficult. The second, related problem is that today's processors can potentially perform one or more 64-bit cache accesses per cycle, and to record both the data address and data values can require a large bandwidth. This presents a problem in that typically, only a few pins might be provided on the chip and these outputs can be switched at significantly lower rates than the processor can be clocked. If the processor generates 100 bits of information every cycle at a speed of 1GHz, but the chip can only output four bits of trace at a speed of 200MHz, then there is a problem. To solve this latter problem, the trace macrocell will try to compress information to reduce the bandwidth needed. However, the main method to deal with these issues is to control the trace block so that only selected trace information is gathered. For example, we might trace only execution, without recording data values, or we might trace only data accesses to a particular peripheral or during execution of a particular function.

In addition, it is common to store trace information in an on-chip memory buffer (the *Embedded Trace Buffer* (ETB)). This alleviates the problem of getting information off-chip at speed, but has an additional cost in terms of silicon area (and therefore price of the chip) and also provides a fixed limit on the amount of trace that can be captured.

The ETB stores the compressed trace information in a circular fashion, continuously capturing trace information until stopped. The size of the ETB varies between chip implementations, but a buffer of 8 or 16KB is typically enough to hold a few thousand lines of program trace.

When a program fails, if the trace buffer is enabled, you can see a portion of program history. With this program history, it is easier to walk back through your program to see what happened just before the point of failure. This is particularly useful for investigating intermittent and real-time failures, which can be difficult to identify through traditional debug methods that require stopping and starting the processor. The use of hardware tracing can significantly reduce the amount of time needed to find these failures, as the trace shows exactly what was executed, what the timing was and what data accesses occurred.

29.2.1 CoreSight

The ARM CoreSight™ technology expands on the capabilities provided by the ETM. Again its presence and capabilities in a particular system are defined by the system designer. CoreSight provides a number of extremely powerful debug facilities. It enables debug of multi-processor systems (both asymmetric and SMP) which can share debug access and trace pins, with full

control of which processors are being traced at which times. The embedded cross trigger mechanism enables tools to control multiple cores in a synchronized fashion, so that, for example when one core hits a breakpoint, all of the other cores will also be stopped.

Commercial debug tools can use trace data to provide features such as real-time views of processor registers, memory and peripherals, allowing the user to step forward and backward through the program execution. Profiling tools can use the data to show where the program is spending its time and what performance bottlenecks exist. Code coverage tools can use trace data to provide call graph exploration. Operating system aware debuggers can make use of trace (and in some cases additional code instrumentation) to provide high level system context information. Here, we list some of the available CoreSight components and give a brief description of their purpose:

Debug Access Port (DAP)

The DAP is an optional part of an ARM CoreSight system. Not every device will contain a DAP. It enables an external debugger to directly access the memory space of the system without having to put the processor into debug state. To read or write memory without a DAP might need the debugger to stop the processor and have it execute Load or Store instructions. The DAP gives an external debug tool access to all of the JTAG scan chains in a system (and therefore to all debug and trace configuration registers of the available processors).

Embedded Cross Trigger (ECT)

The ECT block is a CoreSight component which can be included within in a CoreSight system. Its purpose is to link together the debug capabilities of multiple devices in the system. For example, we can have two cores which run independently of each other. When we set a breakpoint on a program running on one core, it would be useful to be able to specify that when that core stops at the breakpoint, the other one should also be stopped (regardless of which instruction it is currently executing). The Cross Trigger Matrix and Interface within the ECT enable debug status and control information to be propagated between cores and trace macrocells.

AHB Trace Macrocell

The AMBA AHB Trace Macrocell enables the debugger to have visibility of what is happening on the system memory bus. This information is not directly obtainable from the processor ETM, as the integer core is unable to determine whether data comes from a cache or external memory.

CoreSight Serial Wire

CoreSight Serial Wire Debug gives a 2-pin connection using a Debug Access Port (DAP) which is equivalent in function to a 5-pin JTAG interface.

System Trace Macrocell (STM)

This provides a way for multiple processors (and processes) to perform printf() style debugging. Software running on any master in the system is able to access STM channels without needing to be aware of usage by others, using very simple fragments of code. This enables timestamped software instrumentation of both kernel and user space code. The timestamp information gives a delta with respect to previous events and can be extremely useful.

Trace Memory Controller (TMC)

As already described, adding additional pins to a packaged IC can significantly increase its cost. In situations where we have multiple cores (or other blocks capable of generating trace information) on a single device, it is likely that economics preclude the possibility of providing multiple trace ports. The CoreSight Trace Memory Controller can be used to combine multiple trace sources into a single bus. Controls are provided to enable prioritize and select between these multiple input sources. The trace information can be exported off-chip using a dedicated trace port, through the JTAG or serial wire interface or by re-using I/O ports of the SoC. Trace information can be stored in an ETB or in system memory.

Programmers should consult documentation specific to the device they are using to determine what trace capabilities are present and which tools are available to make use of them.

29.3 Debug monitor

We have seen how the ARM architecture provides a wide range of features accessible to an external debugger. Many of these facilities can also be used by code running on the processor – a so called debug monitor, which is resident on the target system. Monitor systems can be inexpensive, as they may not need any additional hardware. However, they take up memory space in the system and can only be used if the target system itself is actually running. They are of little value on a system which does not at least boot correctly. The breakpoint and watchpoint hardware facilities of the processor are available to a debug monitor. When Monitor mode debug is selected, breakpoint units can be programmed by code running on the ARM processor. If a BKPT instruction is executed, or a hardware breakpoint unit matches, the system behaves differently in Monitor mode. Instead of stopping the processor under control of an external hardware debugger, the processor instead takes an abort exception and this can recognize that the abort was generated by a debug event and call the Monitor code.

29.4 Debugging Linux applications

Linux is a multi-tasking operating system in which each process has its own process address space, complete with private page table mappings. This can make debug of some kinds of problems quite tricky.

We can broadly define two different debug approaches used in Linux systems.

- Linux applications are typically debugged using a GDB debug server running on the target, communicating with a host computer, usually through Ethernet. The kernel continues to operate normally while the debug session takes place. This method of debug does not provide access to the built-in hardware debug facilities. The target system is permanently in a running state. The server receives a connection request from the host debugger and then receives commands and provides data back to the host.

The host debugger sends a load request to the GDB server, which responds by starting a new process to run the application being debugged. Before execution begins, it uses the system call `ptrace()` to control the application process. All signals from this process are forwarded to the GDB server. Any signals sent to the application will go instead to the GDB server, which can deal with the signal or forward it to the application being debugged. To set a breakpoint, the GDB server inserts code which generates the **SIGTRAP** signal at the desired location in the code. When this is executed, the GDB server is called and can then perform classic debugger tasks such as examining call stack information, variables or register contents.

- For kernel debug, a JTAG-based debugger is used. The system is halted when a breakpoint is executed. This is the easiest way to examine problems such as device driver loading or incorrect operation or the kernel boot failure. Another common method is through `printk()` function calls. The `strace` tool shows information about user system calls. `Kgdb` is a source-level debugger for the Linux kernel, which works with GDB on a separate machine and enables inspection of stack traces and view of kernel state (such as PC value, timer contents, and memory). The device `/dev/kmem` enables run-time access to the kernel memory.

Of course, a Linux-aware JTAG debugger can be used to debug threads. It is usually possible only to halt all processes; one cannot halt an individual thread or process and leave others running. A breakpoint can be set either for all threads, or it can be set only on a specific thread.

As the memory map depends on which process is active, software breakpoints can usually only be set when a particular process is mapped in.

The ARM DS-5 Debugger is able to debug Linux applications via `gdbserver` and Linux kernel and Linux kernel modules via JTAG. The debug and trace features of DS-5 are described in the next section.

29.5 DS-5 debug and trace

The DS-5 Debugger takes care of downloading and connecting to the debug server. Developers need to specify the platform and the IP address. This reduces a complex task using several applications and a terminal to just a couple of steps in the IDE.

In addition, the DS-5 Debugger supports ARM CoreSight ETM, PTM and STM, to provide non-intrusive program trace that enables you to review instructions (and the associated source code) as they have occurred. It also provides the ability to debug time-sensitive issues which would otherwise not be picked up with conventional intrusive stepping techniques. The DS-5 Debugger currently uses DSTREAM to capture trace on the ETB.

The DS-5 Debugger provides a powerful tool for debugging applications on both hardware targets and models using ARM architecture-based processors. You can have complete control over the flow of the execution so that you can quickly isolate and correct errors.

The following features are provided in the DS-5 Debugger:

- Loading images and symbols.
- Running images.
- Breakpoints and watchpoints.
- Source and instruction level stepping.
- Controlling variables and register values.
- Viewing the call stack.
- Support for handling exceptions and Linux signals.
- Debug of multi-threaded Linux applications.
- Debug of Linux kernel modules, boot code and kernel porting.

The debugger supports a comprehensive set of DS-5 Debugger commands that can be executed in the Eclipse IDE, script files, or a command-line console. In addition, there is a small subset of CMM-style commands sufficient for running target initialization scripts. (CMM is a scripting language supported by some third-party debuggers.)

The DS-5 Debugger supports bare-metal debug via JTAG, Linux application debug via gdbserver, Linux kernel debug via JTAG, and Linux kernel module debug via JTAG. Debug and trace support for bare-metal SMP systems, including cross-triggering and core-dependent views and breakpoints, PTM trace, and up to 4 GB trace with DSTREAM. This support is described in the following sections.

29.5.1 Debugging Linux applications using DS-5

To debug a Linux application you can use a TCP or serial connection:

- To gdbserver running on a model that is pre-configured to boot ARM Embedded Linux.
- To gdbserver running on a hardware target.

This type of development requires gdbserver to be installed and running on the target.

29.5.2 Debugging Linux kernel modules

Linux kernel modules provide a way to extend the functionality of the kernel, and are typically used for things such as device and filesystem drivers. Modules can either be built into the kernel or can be compiled as a loadable module and then dynamically inserted and removed from a running kernel during development without the need to frequently recompile the kernel. However, some modules must be built into the kernel and are not suitable for loading dynamically. An example of a built-in module is one that is required during kernel boot and must be available prior to the root filesystem being mounted.

You can use DS-5 to set source-level breakpoints in a module provided that the debug information is loaded into the debugger. Attempts to set a breakpoint in a module before it is inserted into the kernel results in the breakpoint being pended.

When debugging a module, you must ensure that the module on your target is the same as that on your host. The code layout must be identical, but the module on your target does not need to contain debug information.

Built in module

To debug a module that has been built into the kernel using DS-5, the procedure is the same as for debugging the kernel itself:

1. Compile the kernel together with the module.
2. Load the kernel image on to the target.
3. Load the related kernel image with debug information into the debugger.
4. Debug the module as you would for any other kernel code.

Loadable module

The procedure for debugging a loadable kernel module is more complex. From a Linux terminal shell you can use the `insmod` and `rmmod` commands to insert and remove a module. Debug information for both the kernel and the loadable module must be loaded into the debugger. When you insert and remove a module the DS-5 debugger automatically resolves memory locations for debug information and existing breakpoints.

To do this, the debugger intercepts calls within the kernel to insert and remove modules. This introduces a small delay for each action while the debugger stops the kernel to interrogate various data structures.

29.5.3 Debugging Linux kernels using DS-5

To debug a Linux kernel module you can use a debug hardware agent connected to the host workstation and the running target.

29.5.4 Debugging a multi-threaded applications using DS-5

The DS-5 Debugger tracks the current thread using the debugger variable, `$thread`. You can use this variable in print commands or in expressions. Threads are displayed in the Debug Control view with a unique ID that is used by the debugger and a unique ID from the Operating System (OS). For example:

Thread 1 (OS ID 1036)

where Thread 1 is the ID used by the debugger and OS ID 1036 is the ID from the OS.

A separate call stack is maintained for each thread and the selected stack frame is shown in bold text. All the views in the DS-5 Debug perspective are associated with the selected stack frame and are updated when you select another frame.




Figure 29-1 Threading call stacks in the DS-5 Debug Control view

29.5.5 Debugging shared libraries

Shared libraries enable parts of your application to be dynamically loaded at runtime. You must ensure that the shared libraries on your target are the same as those on your host. The code layout must be identical, but the shared libraries on your target do not need to contain debug information.

You can set standard execution breakpoints in a shared library but not until it is loaded by the application and the debug information is loaded into the debugger. Pending breakpoints however, enable you to set execution breakpoints in a shared library before it is loaded by the application.

When a new shared library is loaded the DS-5 debugger re-evaluates all pending breakpoints, those with addresses that it can resolve, are set as standard execution breakpoints. Unresolved addresses remain as pending breakpoints.

The debugger automatically changes any breakpoints in a shared library to a pending breakpoint when the library is unloaded by your application.

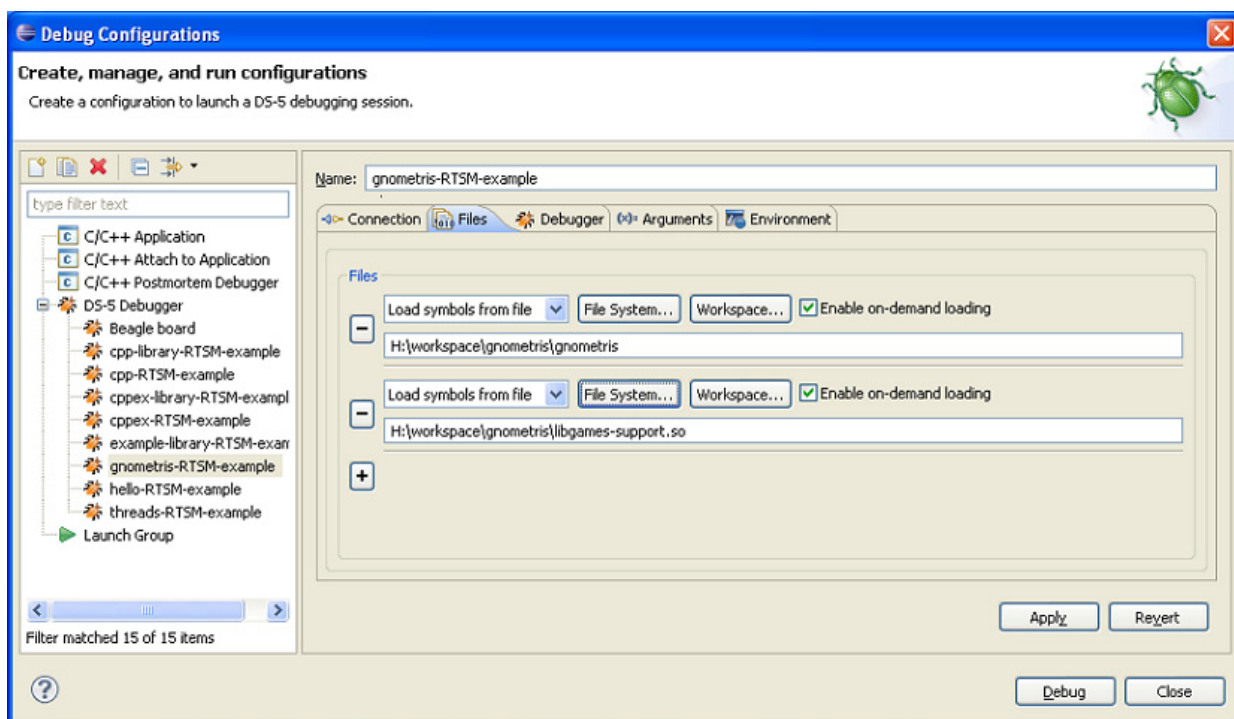


Figure 29-2 Adding shared libraries for debug using DS-5

29.5.6 Trace support in DS-5

DS-5 enables you to perform trace on your application or system. You can capture in real-time a historical, non-intrusive trace of instructions. Tracing is a powerful tool that enables you to investigate problems while the system runs at full speed. These problems can be intermittent, and are difficult to identify through traditional debugging methods that require starting and stopping the processor. Tracing is also useful when trying to identify potential bottlenecks or to improve performance-critical areas of your application.

Before the debugger can trace function executions in your application you must ensure that:

- You have a debug hardware agent, for example, an ARM DSTREAM unit with a connection to a trace stream.
- The debugger is connected to the debug hardware agent.

Trace view

When the trace has been captured the debugger extracts the information from the trace stream and decompresses it to provide a full disassembly, with symbols, of the executed code.

This view shows a graphical navigation chart that displays function executions with a navigational timeline. In addition, the disassembly trace shows function calls with associated addresses and if selected, instructions. Clicking on a specific time in the chart synchronizes the disassembly view.

In the left-hand column of the chart, percentages are shown for each function of the total trace. For example, if a total of 1000 instructions are executed and 300 of these instructions are associated with `myFunction()` then this function is displayed with 30%.

In the navigational timeline, the color coding is a “heat” map showing the executed instructions and the amount of instructions each function executes in each timeline. The darker red color showing more instructions and the lighter yellow color showing less instructions. At a scale of 1:1 however, the color scheme changes to display memory access instructions as a darker red color, branch instructions as a medium orange color, and all the other instructions as a lighter green color.




Figure 29-3 DS-5 Debugger Trace view

Trace-based profiling

Based on trace data received from a trace buffer such as the ETB, The DS-5 Debugger can generate timeline charts with information to help developers to quickly understand how their software executes on the target and which functions are using the processor the most. The timeline offers various zoom levels, and can display a heat-map based on the number of instructions per time unit or, at its highest resolution, provide per-instruction visualization color-coded by the typical latency of each group of instructions.

Appendix A

Instruction Summary

A summary of the instructions available in ARM/Thumb assembly language is given in this Appendix.

For most instructions, further explanation can be found in [Chapter 6](#). The optional condition code field (denoted by cond below) is described in Section 6.1.2 [Conditional execution on page 6-3](#). The format of the flexible operand2 used by data processing operations is described in Section 6.2.1 [Operand 2 and the barrel shifter on page 6-7](#), while the various addressing mode options for loads and stores is given in [Addressing modes on page 6-10](#).

This Appendix is intended for quick reference. If more detail about the precise operation of an instruction is required, please refer to the *ARM Architecture Reference Manual*, or to the official ARM documentation (for example the *ARM Compiler Toolchain Assembler Reference*) which can be found at <http://infocenter.arm.com/help/index.jsp>.

A.1 Instruction Summary

Instructions are listed in alphabetic order, with a description of the syntax, operands and behavior of the instruction. Not all usage restrictions are documented here, nor do we show the associated binary encoding or any detail of changes associated with older architecture versions.

A.1.1 ADC

ADC (Add with Carry) adds together the values in Rn and Operand2, with the carry flag.

Syntax

ADC{S}{cond} {Rd}, Rn, <Operand2>

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first operand.

Operand2 is a flexible second operand. See Section 6.2.1.

A.1.2 ADD

ADD adds together the values in Rn and Operand2 (or Rn and imm12).

Syntax

ADD{S}{cond} {Rd}, Rn, <Operand2>

ADD{cond} {Rd}, Rn, #imm12 (Only available in Thumb)

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first operand.

Operand2 is a flexible second operand. See Section 6.2.1.

imm12 is in the range 0-4095.

A.1.3 ADR

ADR (Address) is an instruction which loads a program or register-relative address (short range). It generates an instruction which adds or subtracts a value to the PC (in the PC-relative case). Alternatively, it can be some other register for a label defined as an offset from a base register defined with the MAP directive (see the ARM tools documentation for more detail).

Syntax

ADR{cond} Rd, label

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

label is a PC or register-relative expression.

A.1.4 ADRL

ADRL (Address) is a pseudo-instruction which loads a program or register-relative address (long range). It always generates two instructions.

Syntax

ADRL{cond} Rd, label

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

label is a PC-relative expression. The offset between label and the current location has some restrictions.

The ADRL pseudo-instruction can generate a wider range of addresses than ADR.

A.1.5 AND

AND does a bitwise AND on the values in Rn and Operand2.

Syntax

AND{S}{cond} {Rd,} Rn, <Operand2>

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first operand.

Operand2 is a flexible second operand. See Section 6.2.1.

A.1.6 ASR

ASR (Arithmetic Shift Right) shifts the value in Rn right, by the number of bit positions specified and copies the sign bit into vacated bit positions on the left. Allowed shift values are in the range 1-32. It can be considered as giving the signed value of a register divided by a power of two.

Syntax

```
ASR{S}{cond} {Rd}, Rm, Rs
ASR{S}{cond} {Rd}, Rm, imm
```

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm is the register holding the operand to be shifted.

Rs is the register which holds a shift value to apply to the value in Rm. Only the least significant byte of the register is used.

imm is a shift amount, in the range 1-32.

A.1.7 B

B (Branch) transfers program execution to the address specified by label.

Syntax

```
B{cond}{.W} label
```

where:

cond is an optional condition code. See Section 6.1.2.

label is a PC-relative expression.

.W is an optional instruction width specifier to force the use of a 32-bit instruction in Thumb.

A.1.8 BFC

BFC (Bit Field Clear) clears bits in a register. A number of bits specified by width are cleared in Rd, starting at 1sb. Other bits in Rd are unchanged.

Syntax

```
BFC{cond} Rd, #1sb, #width
```

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

1sb specifies the least significant bit to be cleared.

width is the number of bits to be cleared.

A.1.9 BFI

BFI (Bit Field Insert) copies bits into a register. A number of bits in Rd specified by width, starting at 1sb, are replaced by bits from Rn, starting at bit[0]. Other bits in Rd are unchanged.

Syntax

```
BFI{cond} Rd, Rn, #lsb, #width
```

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register which contains the bits to be copied.

lsb specifies the least significant bit in Rd to be written to.

width is the number of bits to be copied.

A.1.10 BIC

BIC (bit clear) does an AND operation on the bits in Rn, with the complements of the corresponding bits in the value of Operand2.

Syntax

```
BIC{S}{cond} {Rd}, Rn, <Operand2>
```

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first operand.

Operand2 is a flexible second operand. See Section 6.2.1.

A.1.11 BKPT

BKPT (Breakpoint) causes the processor to enter Debug state.

Syntax

```
BKPT #imm
```

where:

imm is an integer in the range 0 – 65535 (ARM) or 0 – 255 (Thumb). This integer is not used by the processor itself, but can be used by Debug tools.

A.1.12 BL

BL (Branch with Link) transfers program execution to the address specified by label and stores the address of the next instruction in the LR (R14) register.

Syntax

```
BL{cond} label
```

where:

cond is an optional condition code. See Section 6.1.2.

label is a PC-relative expression.

A.1.13 BLX

BLX (Branch with Link and eXchange) transfers program execution to the address specified by label and stores the address of the next instruction in the LR (R14) register. BLX can change the processor state from ARM to Thumb, or from Thumb to ARM. BLX label always changes the processor state from Thumb to ARM, or ARM to Thumb. BLX Rm will set the state based on bit[0] of Rm:

- Rm bit[0]=0 ARM state.
- Rm bit[0]=1 Thumb state.

Syntax

```
BLX{cond} label
BLX{cond} Rm
```

where:

cond is an optional condition code. See Section 6.1.2.

label is a PC-relative expression.

Rm is a register which holds the address to branch to.

A.1.14 BX

BX (Branch and eXchange) transfers program execution to the address specified in a register. BX can change the processor state from ARM to Thumb, or from Thumb to ARM. BX Rm will set the state based on bit[0] of Rm:

- Rm bit[0] = 0 ARM state.
- Rm bit[0] = 1 Thumb state.

Syntax

```
BX{cond} Rm
```

where:

cond is an optional condition code. See Section 6.1.2.

Rm is a register which holds the address to branch to.

A.1.15 BXJ

BXJ (Branch and eXchange Jazelle) enter Jazelle State or perform a BX branch and exchange to the address contained in Rm..

Syntax

```
BXJ{cond} Rm
```

where:

cond is an optional condition code. See Section 6.1.2.

Rm is a register which holds the address to branch to if entry to Jazelle fails.

A.1.16 CBNZ

CBNZ (Compare and Branch if Nonzero) causes a branch if the value in Rn is not zero. It does not change the PSR flags. There is no ARM or 32-bit Thumb versions of this instruction.

Syntax

CBNZ Rn, label

where:

label is a pc-relative expression.

Rn is a register which holds the operand.

A.1.17 CBZ

CBZ (Compare and Branch if Zero) causes a branch if the value in Rn is zero. It does not change the PSR flags. There is no ARM or 32-bit Thumb versions of this instruction.

Syntax

CBZ Rn, label

where:

label is a PC-relative expression.

Rn is a register which holds the operand.

A.1.18 CDP

CDP (Coprocessor Data Processing operation) performs a coprocessor operation. The purpose of this instruction is defined by the coprocessor implementer.

Syntax

CDP{cond} coproc, #opcode1, CRd, CRn, CRm{, #opcode2}

where:

cond is an optional condition code. See Section 6.1.2.

coproc is the name of the coprocessor the instruction is for. This is usually of the form pn, where n is an integer in the range 0 to 15.

opcode1 is a 4-bit coprocessor-specific opcode.

opcode2 is an optional 3-bit coprocessor-specific opcode.

CRd, CRn, CRm are coprocessor registers.

A.1.19 CDP2

CDP2 (Coprocessor Data Processing operation) performs a coprocessor operation. The purpose of this instruction is defined by the coprocessor implementer.

Syntax

CDP2{cond} coproc, #opcode1, CRd, CRn, CRm{, #opcode2}

where:

cond is an optional condition code See Section 6.1.2.

coproc is the name of the coprocessor the instruction is for. This is usually of the form pn, where n is an integer in the range 0 to 15.

opcode1 is a 4-bit coprocessor-specific opcode.

opcode2 is an optional 3-bit coprocessor-specific opcode.

CRd, CRn, CRm are coprocessor registers.

A.1.20 CHKA

CHKA (Check array) is a ThumbEE instruction. If the value in the first register is less than or equal to, the second, the IndexCheck handler is called. This instruction is only available in 16-bit ThumbEE and only when Thumb-2EE support is present.

Syntax

CHKA Rn, Rm

where:

Rn holds the size of the array.

Rm contains the array index.

A.1.21 CLREX

CLREX (Clear Exclusive) moves a local exclusive access monitor to its open-access state.

Syntax

CLREX{cond}

where:

cond is an optional condition code. See Section 6.1.2.

A.1.22 CLZ

CLZ (Count Leading Zeros) counts the number of leading zeros in the value in Rm and returns the result in Rd. The result returned is 32 if no bits are set in Rm, or 0 if bit [31] is set.

Syntax

CLZ{cond} Rd, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm is the register holding the operand.

A.1.23 CMN

CMN (Compare Negative) performs a comparison by adding the value of Operand2 to the value in Rn. The condition flags are changed, based on the result, but the result itself is discarded.

Syntax

```
CMN{cond} Rn, <Operand2>
```

where:

cond is an optional condition code. See Section 6.1.2.

Rn is the register holding the first operand.

Operand2 is a flexible second operand. See Section 6.2.1.

A.1.24 CMP

CMP (Compare) performs a comparison by subtracting the value of Operand2 from the value in Rn. The condition flags are changed, based on the result, but the result itself is discarded.

Syntax

```
CMP{cond} Rn, <Operand2>
```

where:

cond is an optional condition code. See Section 6.1.2.

Rn is the register holding the first operand.

Operand2 is a flexible second operand. See Section 6.2.1.

A.1.25 CPS

CPS (Change Processor State) can be used to change the processor mode or to enable or disable individual exception types.

Syntax

```
CPS #mode
CPSIE iflags{, #mode}
CPSID iflags{, #mode}
```

where:

mode is the number of a mode for the processor to enter.

IE Interrupt or Abort Enable.

ID Interrupt or Abort Disable.

iflags specifies one or more of:

- a = asynchronous abort.
- i = IRQ.
- f = FIQ.

A.1.26 DBG

DBG (Debug) is a hint operation, treated as a NOP by the processor, but can provide information to debug systems.

Syntax

DBG{cond} {option}

where:

cond is an optional condition code. See Section 6.1.2.

option is in the range 0-15.

A.1.27 DMB

DMB (Data Memory Barrier) requires that all explicit memory accesses in program order before the DMB instruction are observed before any explicit memory accesses in program order after the DMB instruction. See [Chapter 11](#) for a detailed description.

Syntax

DMB{cond} {option}

where:

cond is an optional condition code. See Section 6.1.2.

option is covered in depth in [Chapter 11](#).

A.1.28 DSB

DSB (Data Synchronization Barrier) requires that no further instruction executes until all explicit memory accesses, cache maintenance operations, branch prediction and TLB maintenance operations before this instruction complete. See [Chapter 11](#) for a detailed description.

Syntax

DSB{cond} {option}

where:

cond is an optional condition code. See Section 6.1.2.

option is covered in depth in [Chapter 11](#).

A.1.29 ENTERX

ENTERX causes a change from Thumb state to ThumbEE state, or has no effect in ThumbEE state. It is not available in the ARM instruction set.

Syntax

ENTERX

A.1.30 EOR

EOR performs an Exclusive OR operation on the values in Rn and Operand2.

Syntax

EOR{S}{cond} {Rd}, Rn, <Operand2>

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first operand.

Operand2 is a flexible second operand. See Section 6.2.1.

A.1.31 ERET

ERET (Exception Return) loads the PC from the ELR-hyp and loads the CPSR from SPSR-hyp when executed in Hyp mode.

When executed in a Secure or Non-Secure PLI mode, ERET behaves as:

- MOVs PC, LR in the ARM instruction set.
- SUBS PC, LR, #0 in the Thumb instruction set.

Syntax

ERET{cond} {q}

where:

cond is the optional condition code. See Section 6.1.2.

q specifies optional assembler qualifiers on the instruction. The following qualifiers are defined:

- .N(narrow), specifies that the assembler must select a 16-bit encoding for the instruction.
- .W(wide), specifies that the assembler must select a 32-bit encoding for the instruction.

If neither .W nor .N are specified, the assembler can select either 16-bit or 32-bit encodings. If both are available, it must select a 16-bit encoding.

A.1.32 HB

HB (Handler Branch) branches to a specified handler (available in ThumbEE only).

Syntax

HB{L} #HandlerID
HB{L}P #imm, #HandlerID

where:

L indicates that the instruction saves a return address in the LR.

P means that the instruction passes the value of imm to the handler in R8.

imm is an immediate value in the range 0-31 (if L is present), otherwise in the range 0-7.

HandlerID is the index number of the handler to be called, in the range 0-31 (if P is specified), otherwise in the range 0-255.

A.1.33 ISB

ISB (Instruction Synchronization Barrier) flushes the processor pipeline and ensures that context altering operations (such as ASID or other CP15 changes, branch prediction or TLB maintenance activity) *before* the ISB, are visible to the instructions fetched *after* the ISB.

See [Chapter 11](#) for a detailed description of barriers.

Syntax

ISB{cond} {option}

where:

cond is an optional condition code. See Section 6.1.2.

option can be SY (full system), which is the default and so can be omitted.

A.1.34 IT

IT (If-then) makes up to four following instructions conditional (known as the IT block). The conditions can all be the same, or some can be the logical inverse of others. IT is a pseudo-instruction in ARM state.

Syntax

IT{x{y{z}}} {cond}

where:

cond is a condition code. See Section 6.1.2 which specifies the condition for the first instruction in the IT block.

x, y and z specify the condition switch for the second, third and fourth instructions in the IT block, for example, ITTET.

The condition switch can be either:

- T (Then), which applies the condition cond to the instruction.
- E (Else), which applies the inverse condition of cond to the instruction.

A.1.35 LDC

LDC (Load Coprocessor Registers) reads a coprocessor register from memory (or multiple registers, if L is specified).

Syntax

LDC{L}{cond} coproc, CRd, [Rn]
 LDC{L}{cond} coproc, CRd, [Rn, #-}offset]{!}
 LDC{L}{cond} coproc, CRd, [Rn], #-}offset
 LDC{L}{cond} coproc, CRd, label

where:

L specifies that more than one register can be transferred (called a long transfer). The length of the transfer is determined by the coprocessor, but may not be more than 16 words.

cond is an optional condition code. See Section 6.1.2.

coproc is the name of the coprocessor the instruction is for. This is usually of the form pn, where n is an integer in the range 0 to 15.

CRd is the coprocessor register to be stored.

Rn is the register holding the base address for the memory operation.

Offset is a multiple of 4, in the range 0-1020, to be added or subtracted from Rn. If ! is present, the address including the offset is written back into Rn.

label is a word-aligned PC-relative address label.

A.1.36 LDC2

LDC2 (Load Coprocessor Registers) reads a coprocessor register from memory (or multiple registers, if L is specified).

Syntax

```
LDC2{L}{cond} coproc, CRd, [Rn]
LDC2{L}{cond} coproc, CRd, [Rn, #-offset]{}
LDC2{L}{cond} coproc, CRd, [Rn], #-offset
LDC2{L}{cond} coproc, CRd, label
```

where:

L specifies that more than one register can be transferred (called a long transfer). The length of the transfer is determined by the coprocessor, but may not be more than 16 words.

cond is an optional condition code. See Section 6.1.2.

coproc is the name of the coprocessor the instruction is for. This is usually of the form pn, where n is an integer in the range 0 to 15.

CRd is the coprocessor register to be stored.

Rn is the register holding the base address for the memory operation.

Offset is a multiple of 4, in the range 0 – 1020, to be added or subtracted from Rn. If ! is present, the address including the offset is written back into Rn.

label is a word-aligned PC-relative address label.

A.1.37 LDM

LDM (Load Multiple registers) loads one or more registers from consecutive addresses in memory at an address specified in a base register.

Syntax

```
LDM{addr_mode}{cond} Rn{!}, reglist{^}
```

where:

addr_mode is one of:

- IA – Increment address After each transfer. This is the default, and can be omitted.

- IB – Increment address Before each transfer (ARM only).
- DA – Decrement address After each transfer (ARM only).
- DB – Decrement address Before each transfer.

It is also possible to use the corresponding stack oriented addressing modes (FD, ED, EA, FA). For example LDMFD is a synonym of LDMDB.

cond is an optional condition code. See Section 6.1.2.

Rn is the base register, giving the initial address for the transfer.

! if present, specifies that the final address is written back into Rn.

Reglist is a list of one or more registers to be stored, enclosed in braces. It can contain register ranges. It must be comma separated if it contains more than one register or register range.

^ if specified (in a mode other than User or System) means one of two possible special actions will be taken:

- Data is transferred into the User mode registers instead of the current mode registers (in the case where Reglist does not contain the PC).
- If Reglist does contain the PC, the normal multiple register transfer happens and the SPSR is copied into the CPSR. This is used for returning from exception handlers.

A.1.38 LDR

LDR (Load Register) loads a value from memory to an ARM register, optionally updating the register used to give the address.

A variety of addressing options are provided. For full details of the available addressing modes, see [Addressing modes on page 6-10](#).

Syntax

```
LDR{type}{T}{cond} Rt, [Rn {, #offset}]
LDR{type}{cond} Rt, [Rn, #offset]!
LDR{type}{T}{cond} Rt, [Rn], #offset
LDR{type}{cond} Rt, [Rn, +/-Rm {, shift}]
LDR{type}{cond} Rt, [Rn, +/-Rm {, shift}]!
LDR{type}{T}{cond} Rt, [Rn], +/-Rm {, shift}
```

where:

type can be any one of:

- B – unsigned Byte. (Zero extend to 32 bits on loads.)
- SB – signed Byte. (Sign extend to 32 bits.)
- H – unsigned Halfword. (Zero extend to 32 bits on loads.)
- SH – signed Halfword. (Sign extend to 32 bits.)

or omitted, for a Word load.

T specifies that memory is accessed as if the processor was in User mode (not available in all addressing modes).

cond is an optional condition code. See Section 6.1.2.

Rn is the register holding the base address for the memory operation.

! if present, specifies that the final address is written back into Rn.

offset is a numeric value.

Rm is a register holding an offset value to be applied.

shift is either a register or immediate based shift to apply to the offset value.

A.1.39 LDR (pseudo-instruction)

LDR (Load Register) pseudo-instruction loads a register with a 32-bit immediate value or an address. It generates either a MOV or MVN instruction (if possible), or a PC-relative LDR instruction that reads the constant from the literal pool.

Syntax

```
LDR{cond}{.W} Rt, =expr
LDR{cond}{.W} Rt, label_expr
```

where:

cond is an optional condition code. See Section 6.1.2.

.W specifies that a 32-bit Thumb instruction must be used.

Rt is the register to load.

expr is a numeric value.

label_expr is a label, optionally plus or minus a numeric value.

A.1.40 LDRD

LDRD (Load Register Dual) calculates an address from a base register value and a register offset, loads two words from memory, and writes them to two registers.

Syntax

```
LDRD{cond} Rt, Rt2, [{Rn}, +/-{Rm}]{!}
LDRD{cond} Rt, Rt2, [{Rn}], +/-{Rm}
```

where:

cond is an optional condition code. See Section 6.1.2.

Rt is the first destination register. This register must be even-numbered and not R14.

Rt2 is the second destination register. This register must be <R(t+1)>.

Rn is the base register. The SP or the PC can be used.

+/- is + or omitted if the value of <Rm> is to be added to the base register value (add == TRUE), or – if it is to be subtracted (add == FALSE).

Rm contains the offset that is applied to the value of <Rn> to form the address.

A.1.41 LDREX

LDREX (Load register exclusive). Performs a load from a location and marks it for exclusive access. Byte, halfword, word and doubleword variants are provided.

Syntax

```
LDREX{cond} Rt, [Rn {, #offset}]
LDREXB{cond} Rt, [Rn]
LDREXH{cond} Rt, [Rn]
LDREXD{cond} Rt, Rt2, [Rn]
```

where:

cond is an optional condition code. See Section 6.1.2.

Rt is the register to load.

Rt2 is the second register for doubleword loads.

Rn is the register holding the address.

offset is an optional value, allowed in Thumb only.

A.1.42 LEAVEX

LEAVEX causes a change from ThumbEE state to Thumb state, or has no effect in Thumb state. It is not available in the ARM instruction set.

Syntax

```
LEAVEX
```

A.1.43 LSL

LSL (Logical Shift Left) shifts the value in Rm left by the specified number of bits, inserting zeros into the vacated bit positions.

Syntax

```
LSL{S}{cond} Rd, Rm, Rs
LSL{S}{cond} Rd, Rm, imm
```

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm is the register holding the operand to be shifted.

Rs is the register which holds a shift value to apply to the value in Rm. Only the least significant byte of the register is used.

imm is a shift amount, in the range 0-31.

A.1.44 LSR

LSR (Logical Shift Right) shifts the value in Rm right by the specified number of bits, inserting zeros into the vacated bit positions.

Syntax

```
LSR{S}{cond} Rd, Rm, Rs
LSR{S}{cond} Rd, Rm, imm
```

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm is the register holding the operand to be shifted.

Rs is the register which holds a shift value to apply to the value in Rm. Only the least significant byte of the register is used.

imm is a shift amount, in the range 1-32.

A.1.45 MCR

MCR (Move to Coprocessor from Register) writes a coprocessor register, from an ARM register. The purpose of this instruction is defined by the coprocessor implementer.

Syntax

```
MCR{cond} coproc, #opcode1, Rt, CRn, CRm{, #opcode2}
```

where:

cond is an optional condition code. See Section 6.1.2.

Rt is the ARM register to be transferred.

coproc is the name of the coprocessor the instruction is for. This is usually of the form pn, where n is an integer in the range 0 to 15.

opcode1 is a 4-bit coprocessor-specific opcode.

opcode2 is an optional 3-bit coprocessor-specific opcode.

CRn, CRm are coprocessor registers.

A.1.46 MCR2

MCR2 (Move to Coprocessor from Register) writes a coprocessor register, from an ARM register. The purpose of this instruction is defined by the coprocessor implementer.

Syntax

```
MCR2{cond} coproc, #opcode1, Rt, CRn, CRm{, #opcode2}
```

where:

cond is an optional condition code. See Section 6.1.2.

Rt is the ARM register to be transferred.

coproc is the name of the coprocessor the instruction is for. This is usually of the form pn, where n is an integer in the range 0 to 15.

opcode1 is a 4-bit coprocessor-specific opcode.

opcode2 is an optional 3-bit coprocessor-specific opcode.

CRn, CRm are coprocessor registers.

A.1.47 MCRR

MCRR (Move to Coprocessor from Registers) transfers a pair of ARM register to a coprocessor. The purpose of this instruction is defined by the coprocessor implementer.

Syntax

MCRR{cond} coproc, #opcode3, Rt, Rt2, CRm

where:

cond is an optional condition code. See Section 6.1.2.

Rt and Rt2 are the ARM registers to be transferred.

coproc is the name of the coprocessor the instruction is for. This is usually of the form pn, where n is an integer in the range 0 to 15.

CRm is a coprocessor register.

Opcode3 is an optional 4-bit coprocessor-specific opcode.

A.1.48 MCRR2

MCRR2 (Move to Coprocessor from Registers) transfers a pair of ARM register to a coprocessor. The purpose of this instruction is defined by the coprocessor implementer.

Syntax

MCRR2{cond} coproc, #opcode3, Rt, Rt2, CRm

where:

cond is an optional condition code. See Section 6.1.2.

Rt and Rt2 are the ARM registers to be transferred.

coproc is the name of the coprocessor the instruction is for. This is usually of the form pn, where n is an integer in the range 0 to 15.

CRm is a coprocessor register.

Opcode3 is an optional 4-bit coprocessor-specific opcode.

A.1.49 MLA

MLA (Multiply Accumulate) multiplies Rn and Rm, adds the value from Ra, and stores the least significant 32 bits of the result in Rd.

Syntax

MLA{S}{cond} Rd, Rn, Rm, Ra

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

Ra is the register holding the accumulate value.

A.1.50 MLS

MLS (Multiply and Subtract) multiplies Rn and Rm, subtracts the result from Ra, and stores the least significant 32 bits of the final result in Rd.

Syntax

```
MLS{S}{cond} Rd, Rn, Rm, Ra
```

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

Ra is the register holding the accumulate value.

A.1.51 MOV

MOV (Move) copies the value of Operand2 into Rd.

Syntax

```
MOV{S}{cond} Rn, <Operand2>
MOV{cond} Rd, #imm16
```

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Operand2 is a flexible second operand. See Section 6.2.1.

imm16 is an immediate value in the range 0-65535.

A.1.52 MOVT

MOVT (Move Top) writes imm16 to Rd[31:16]. It does not affect Rd[15:0].

Syntax

```
MOVT{cond} Rd, #imm16
```

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Operand2 is a flexible second operand. See Section 6.2.1.

imm16 is an immediate value in the range 0-65535.

A.1.53 MOV32

MOV32 is a pseudo-instruction which loads a register with a 32-bit immediate value or address. It generates two instructions, a MOV, MOVT pair.

Syntax

```
MOV32 Rd, expr
```

where:

Rd is the destination register.

expr is a 32-bit constant, or address label.

A.1.54 MRC

MRC (Move to Register from Coprocessor) reads a coprocessor register to an ARM register. The purpose of this instruction is defined by the coprocessor implementer.

Syntax

```
MRC{cond} coproc, #opcode1, Rt, CRn, CRm{, #opcode2}
```

where:

cond is an optional condition code. See Section 6.1.2.

Rt is the ARM register to be transferred.

coproc is the name of the coprocessor the instruction is for. This is usually of the form pn, where n is an integer in the range 0 to 15.

opcode1 is a 4-bit coprocessor-specific opcode.

opcode2 is an optional 3-bit coprocessor-specific opcode.

CRn, CRm are coprocessor registers.

A.1.55 MRC2

MRC2 (Move to Register from Coprocessor) reads a coprocessor register to an ARM register. The purpose of this instruction is defined by the coprocessor implementer.

Syntax

```
MRC2{cond} coproc, #opcode1, Rt, CRn, CRm{, #opcode2}
```


where:

cond is an optional condition code. See Section 6.1.2.

Rt is the ARM register to be transferred.

coproc is the name of the coprocessor the instruction is for. This is usually of the form pn, where n is an integer in the range 0 to 15.

opcode1 is a 4-bit coprocessor-specific opcode.

opcode2 is an optional 3-bit coprocessor-specific opcode.

CRn, CRm are coprocessor registers.

A.1.56 MRRC

MRRC (Move to Registers from Coprocessor) transfers a value from a Coprocessor to a pair of ARM registers. The purpose of this instruction is defined by the Coprocessor implementer.

Syntax

MRRC{cond} coproc, #opcode3, Rt, Rt2, CRm

where:

cond is an optional condition code, See Section 6.1.2. MRRC instructions may not specify a condition code in ARM state.

Rt and Rt2 are the ARM registers to be transferred.

coproc is the name of the coprocessor the instruction is for. This is usually of the form pn, where n is an integer in the range 0 to 15.

CRm is a coprocessor register.

Opcode3 is an optional 4-bit coprocessor-specific opcode.

A.1.57 MRRC2

MRRC2 (Move to Registers from Coprocessor) transfers a value from a Coprocessor to a pair of ARM registers. The purpose of this instruction is defined by the Coprocessor implementer.

Syntax

MRRC2{cond} coproc, #opcode3, Rt, Rt2, CRm

where:

cond is an optional condition code, See Section 6.1.2. MRRC2 instructions may not specify a condition code in ARM state.

Rt and Rt2 are the ARM registers to be transferred.

coproc is the name of the coprocessor the instruction is for. This is usually of the form pn, where n is an integer in the range 0 to 15.

CRm is a coprocessor register.

Opcode3 is an optional 4-bit coprocessor-specific opcode.

A.1.58 MRS

MRS (Move Status register or Coprocessor Register to General purpose register) can be used to read the CPSR/APSR, CP14 or CP15 Coprocessor registers.

Syntax

```
MRS{cond} Rd, psr
MRS{cond} Rn, coproc_register
MRS{cond} APSR_nzcv, DBGDSCRint
MRS{cond} APSR_nzcv, FPSCR
```

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

psr is one of: APSR, CPSR or SPSR.

coproc_register is the name of a CP14 or CP15 readable register.

DBGDSCRint is the name of a CP14 register which can be copied to the APSR.

A.1.59 MSR

MSR (Move Status Register or Coprocessor Register from General Purpose Register) can be used to write all or part of the CPSR/APSR or CP14 or CP15 registers.

Syntax

```
MSR{cond} APSR_flags, Rm
MSR{cond} coproc_register
MSR{cond} APSR_flags, #constant
MSR{cond} psr_fields, #constant
MSR{cond} psr_fields, Rm
```

where:

cond is an optional condition code. See Section 6.1.2.

Rm and Rn are the source registers.

flags can be one or more of nzcvq (ALU flags) or g (SIMD flags).

coproc_register is the name of a CP14 or CP15 readable register.

constant is an 8-bit pattern rotated by an even number of bits within a 32-bit word. (Not available in Thumb.)

psr is one of: APSR, CPSR or SPSR.

fields is one or more of:

- c control field mask byte, PSR[7:0]
- x extension field mask byte, PSR[15:8]
- s status field mask byte, PSR[23:16]
- f flags field mask byte, PSR[31:24].

A.1.60 MUL

MUL (Multiply) Multiplies Rn and Rm, and stores the least significant 32 bits of the result in Rd.

Syntax

MUL{S}{cond} {Rd}, Rn, Rm

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

A.1.61 MVN

MVN (Move Not) performs a bitwise NOT operation on the operand2 value, and places the result into Rd.

Syntax

MVN{S}{cond} Rn, <Operand2>

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Operand2 is a flexible second operand. See Section 6.2.1.

A.1.62 NOP

NOP (No Operation) does nothing.

Syntax

NOP{cond}

where:

NOP does not have to consume clock cycles. It can be removed by the processor pipeline. It is used for padding, to ensure following instructions align to a boundary.

A.1.63 ORN

ORN (OR NOT) performs an OR operation on the bits in Rn with the complement of the corresponding bits in the value of Operand2.

Syntax

ORN{S}{cond} {Rd}, Rn, <Operand2>

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first operand.

Operand2 is a flexible second operand. See Section 6.2.1.

A.1.64 ORR

Performs an OR operation on the bits in Rn with the corresponding bits in the value of Operand2.

Syntax

```
ORR{S}{cond} {Rd}, Rn, <Operand2>
```

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first operand.

Operand2 is a flexible second operand. See Section 6.2.1.

A.1.65 PKHBT

PKHBT (Pack Halfword Bottom Top) combines bits[15:0] of Rn with bits[31:16] of the shifted value from Rm.

Syntax

```
PKHBT{cond} {Rd}, Rn, Rm{, LSL #leftshift}
```

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first operand.

Rm is the register holding the second operand.

leftshift is a number in the range 0-31.

A.1.66 PKHTB

PKHTB (Pack Halfword Top Bottom) combines bits[31:16] of Rn with bits[15:0] of the shifted value from Rm.

Syntax

```
PKHTB{cond} {Rd}, Rn, Rm {, ASR #rightshift}
```

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first operand.

Rm is the register holding the second operand.

rightshift is a number in the range 1-32.

A.1.67 PLD

PLD (Preload Data) is a hint instruction which can cause data to be preloaded into the cache.

Syntax

```
PLD{cond} [Rn {, #offset}]
PLD{cond} [Rn, +/-Rm {, shift}]
PLD{cond} label
```

where:

cond is an optional condition code. See Section 6.1.2.

Rn is a base address.

offset is an immediate value, which defaults to 0 if not specified.

Rm contains an offset value and must not be PC (or SP, in Thumb state).

shift is an optional shift.

label is a PC-relative expression.

A.1.68 PLDW

PLDW (Preload data with intent to write) is a hint instruction which can cause data to be preloaded into the cache. It is available only in processors which implement multi-processing extensions.

Syntax

```
PLDW{cond} [Rn {, #offset}]
PLDW{cond} [Rn, +/-Rm {, shift}]
```

where:

cond is an optional condition code. See Section 6.1.2.

Rn is a base address.

offset is an immediate value, which defaults to 0 if not specified.

Rm contains an offset value and must not be PC (or SP, in Thumb state).

shift is an optional shift.

A.1.69 PLI

PLI (Preload instructions) is a hint instruction which can cause instructions to be preloaded into the cache.

Syntax

```
PLI{cond} [Rn {, #offset}]
PLI{cond} [Rn, +/-Rm {, shift}]
PLI{cond} label
```

where:

cond is an optional condition code. See Section 6.1.2.

Rn is a base address.

offset is an immediate value, which defaults to 0 if not specified.

Rm contains an offset value and must not be PC (or SP, in Thumb state).

shift is an optional shift.

label is a PC-relative expression.

A.1.70 POP

POP is used to pop registers off a full descending stack. POP is a synonym for LDMIA sp!, reglist.

Syntax

```
POP{cond} reglist
```

where:

cond is an optional condition code. See Section 6.1.2.

reglist is a list of one or more registers, enclosed in braces.

A.1.71 PUSH

PUSH is used to push registers on to a full descending stack. PUSH is a synonym for STMDB sp!, reglist.

Syntax

```
PUSH{cond} reglist
```

where:

cond is an optional condition code. See Section 6.1.2.

reglist is a list of one or more registers, enclosed in braces.

A.1.72 QADD

QADD (Saturating signed Add) does a signed addition and saturates the result to the signed range $-2^{31} \leq x \leq 2^{31}-1$. If saturation occurs, the Q flag is set.

Syntax

```
QADD{cond} {Rd,} Rm, Rn
```

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the register holding the operands.

A.1.73 QADD8

QADD8 (Saturating signed bitwise Add) does a signed bitwise addition (4 adds) and saturates the results to the signed range $-2^7 \leq x \leq 2^7-1$. The Q flag is not affected by this instruction.

Syntax

QADD8{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.74 QADD16

QADD16 (Saturating signed bitwise Add) does a signed halfword-wise addition (2 adds) and saturates the results to the signed range $-2^7 \leq x \leq 2^7-1$. The Q flag is not affected by this instruction.

Syntax

QADD16{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.75 QASX

QASX (Saturating signed Add Subtract eXchange) exchanges halfwords of Rm, then adds the top halfwords and subtracts the bottom halfwords and saturates the results to the signed range $-2^{15} \leq x \leq 2^{15}-1$. The Q flag is not affected by this instruction.

Syntax

QASX{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.76 QDADD

QDADD (Saturating signed Add) does a signed doubling addition and saturates the result to the signed range $-2^{31} \leq x \leq 2^{31}-1$. If saturation occurs, the Q flag is set.

Syntax

QDADD{cond} {Rd}, Rm, Rn

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

The value in Rn is multiplied by 2, saturated and then added to the value in Rm. A second saturate operation is then performed.

A.1.77 QDSUB

QDSUB (Saturating signed Doubling Subtraction) does a signed doubling subtraction and saturates the result to the signed range $-2^{31} \leq x \leq 2^{31}-1$. If saturation occurs, the Q flag is set.

Syntax

QDSUB{cond} {Rd}, Rm, Rn

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

The value in Rn is multiplied by 2, saturated and then subtracted from the value in Rm. A second saturate operation is then performed.

A.1.78 QSAX

QSAX (Saturating signed Subtract Add Exchange) exchanges the halfwords of Rm, then subtracts the top halfwords and adds the bottom halfwords and saturates the results to the signed range $-2^{15} \leq x \leq 2^{15}-1$. The Q flag is not affected by this instruction.

Syntax

QSAX{cond} {Rd}, Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.79 QSUB

QSUB (Saturating signed Subtraction) does a signed subtraction and saturates the result to the signed range $-2^{31} \leq x \leq 2^{31}-1$. If saturation occurs, the Q flag is set.

Syntax

QDSUB{cond} {Rd,} Rm, Rn

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

The value in Rn is subtracted from the value in Rm. A saturate operation is then performed.

A.1.80 QSUB8

QSUB8 (Saturating signed bitwise Subtract) does bitwise subtraction (4 subtracts), with saturation of the results to the signed range $-2^7 \leq x \leq 2^7-1$. The Q flag is not affected by this instruction.

Syntax

QSUB8{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.81 QSUB16

QSUB16 (Saturating signed halfword Subtract) does halfword-wise subtraction (2 subtracts), with saturation of the results to the signed range $-2^{15} \leq x \leq 2^{15}-1$. The Q flag is not affected by this instruction.

Syntax

QSUB16{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.82 RBIT

RBIT (Reverse bits) reverses the bit order in a 32-bit word.

Syntax

RBIT{cond} Rd, Rn

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the operand.

A.1.83 REV

REV (Reverse) converts 32-bit big-endian data into little-endian data, or 32-bit little-endian data into big-endian data.

Syntax

REV{cond} {Rd}, Rn

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the operand.

A.1.84 REV16

REV16 (Reverse byte order halfwords) converts 16-bit big-endian data into little-endian data, or 16-bit little-endian data into big-endian data.

Syntax

REV16{cond} {Rd}, Rn

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the operand.

A.1.85 REVSH

REVSH (Reverse byte order halfword, with sign extension) does a reverse byte order of the bottom halfword, and sign extends the result to 32 bits.

Syntax

REVSH{cond} Rd, Rn

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the operand.

A.1.86 RFE

RFE (Return from Exception) is used to return from an exception where the return state was saved with SRS. If ! is specified, the final address is written back into Rn.

Syntax

RFE{addr_mode}{cond} Rn{!}

where:

addr_mode is one of:

- IA – Increment address After each transfer. This is the default, and can be omitted.
- IB – Increment address Before each transfer (ARM only).
- DA – Decrement address After each transfer (ARM only).
- DB – Decrement address Before each transfer.

cond is an optional condition codes. See Section 6.1.2, and is allowed only in Thumb, using a preceding IT instruction.

Rn specifies the base register.

A.1.87 ROR

ROR (Rotate right Register) rotates a value in a register by a specified number of bits. The bits that are rotated off the right end are inserted into the vacated bit positions on the left.

Syntax

ROR{S}{cond} {Rd}, Rm, Rs
ROR{S}{cond} {Rd}, Rm, imm

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the operand.

Rm is the register holding the operand to be shifted.

Rs is the register which holds a shift value to apply to the value in Rm. Only the least significant byte of the register is used.

imm is a shift amount, in the range 1 – 31.

A.1.88 RRX

RRX (Rotate Right with extend) performs a shift right one bit on a register value. The old carry flag is shifted into bit[31]. If the S suffix is present, the old bit[0] is placed in the carry flag.

Syntax

$$\text{RRX}\{\text{S}\}\{\text{cond}\} \{ \text{Rd}, \} \text{Rm}$$

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm is the register holding the operand to be shifted.

A.1.89 RSB

RSB (Reverse Subtract) subtracts the value in Rn from the value of Operand2. This is useful because Operand2 has more options than Operand1 (which is always a register).

Syntax

$$\text{RSB}\{\text{S}\}\{\text{cond}\} \{ \text{Rd}, \} \text{Rn}, \text{<Operand2>}$$

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first operand.

Operand2 is a flexible second operand. See Section 6.2.1.

A.1.90 RSC

RSC (Reverse Subtract with Carry) subtracts Rn from Operand2. If the carry flag is clear, the result is reduced by one.

Syntax

$$\text{RSC}\{\text{S}\}\{\text{cond}\} \{ \text{Rd}, \} \text{Rn}, \text{<Operand2>}$$

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first operand.

Operand2 is a flexible second operand. See Section 6.2.1.

A.1.91 SADD8

SADD8 (Signed bytewise Add) does a signed bytewise addition (4 adds).

Syntax

SADD8{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.92 SADD16

SADD16 (Signed bitwise Add) does a signed halfword-wise addition (2 adds).

Syntax

SADD16{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.93 SASX

SASX (Signed Add Subtract Exchange) exchanges halfwords of Rm, then adds the top halfwords and subtracts the bottom halfwords.

Syntax

SASX{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.94 SBC

SBC (Subtract with Carry) subtracts the value of Operand2 from the value in Rn. If the carry flag is clear, the result is reduced by one.

Syntax

SBC{S}{cond} {Rd,} Rn, <Operand2>

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first operand.

Operand2 is a flexible second operand. See Section 6.2.1.

A.1.95 SBFX

SBFX (Signed Bit Field Extract) writes adjacent bits from one register into the least significant bits of a second register and sign extends to 32 bits.

Syntax

SBFX{cond} Rd, Rn, #lsb, #width

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register which contains the bits to be extracted.

lsb specifies the least significant bit of the bitfield.

width is the width of the bitfield.

A.1.96 SDIV

SDIV (Signed Divide). divides a 32-bit signed integer register value by a 32-bit signed integer register value, and writes the result to the destination register. This instruction is not present in all variants of the ARMv7-A architecture.

Syntax

SDIV{cond}{q} {Rd}, Rn, Rm

where:

cond is the optional condition code. See section 6.1.2.

q specifies optional assembler qualifiers on the instruction. The following qualifiers are defined:

- .N (narrow), specifies that the assembler must select a 16-bit encoding for the instruction.
- .W (wide), specifies that the assembler must select a 32-bit encoding for the instruction.

If neither .W nor .N are specified, the assembler can select either 16-bit or 32-bit encodings. If both are available, it must select a 16-bit encoding.

Rd the destination register.

Rn is the register that contains the dividend.

Rm is the register that contains the divisor.

A.1.97 SEL

SEL (Select) selects bytes from Rn or Rm, depending on the APSR GE flags.

If GE[0] is set, Rd[7:0] comes from Rn[7:0], else from Rm[7:0].

If GE[1] is set, Rd[15:8] comes from Rn[15:8], else from Rm[15:8].

If GE[2] is set, Rd[23:16] comes from Rn[23:16], else from Rm[23:16].

If GE[3] is set, Rd[31:24] comes from Rn[31:24], else from Rm[31:24].

Syntax

```
SEL{cond} {Rd,} Rn, Rm
```

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register which contains the bits to be extracted.

Rm is the register holding the second operand.

A.1.98 SETEND

SETEND (Set endianness) selects little-endian or big-endian memory access. See [Endianness on page 16-2](#) for more details.

Syntax

```
SETEND LE
SETEND BE
```

A.1.99 SEV

SEV (Send Event) causes an event to be signaled to all cores in an MPCore. See [Power and clocking on page 25-2](#) for more detail.

Syntax

```
SEV{cond}
```

where:

cond is an optional condition code. See Section 6.1.2.

A.1.100 SHADD8

SHADD8 (Signed halving bitwise Add) does a signed bitwise addition (4 adds) and halves the results.

Syntax

```
SHADD8{cond} {Rd,} Rn, Rm
```

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.101 SHADD16

SHADD16 (Signed halving bitwise Add) does a signed halfword-wise addition (2 adds) and halves the results.

Syntax

SHADD16{cond} {Rd}, Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.102 SHASX

SHASX (Signed Halving Add Subtract Exchange) exchanges halfwords of Rm, then adds the top halfwords and subtracts the bottom halfwords and halves the results.

Syntax

SHASX{cond} {Rd}, Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.103 SHSAX

SHSAX (Signed Halving Subtract Add Exchange) exchanges halfwords of Rm, then subtracts the top halfwords and adds the bottom halfwords and halves the results.

Syntax

SHSAX{cond} {Rd}, Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.104 SHSUB8

SHSUB8 (Signed halving bitwise subtraction) does a signed bitwise subtraction (4 subtracts) and halves the results.

Syntax

SHSUB8{cond} {Rd}, Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands

A.1.105 SHSUB16

SHSUB16 (Signed Halving halfword-wise Subtract) does a signed halfword-wise subtraction (2 subtracts) and halves the result.

Syntax

SHSUB16{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the register holding the operands.

A.1.106 SMC

SMC (Secure Monitor Call) is used by the ARM Security Extensions. This instruction was formerly called SMI. See [Chapter 26 Security](#) for more details.

Syntax

SMC{cond} #imm4

where:

cond is an optional condition code. See Section 6.1.2.

imm4 is an immediate value in the range 0-15, which is ignored by the processor, but can be used by the SMC exception handler.

A.1.107 SMLAxy

The SMLAxy (Signed Multiply Accumulate; $32 \leq 32 + 16 \times 16$) instruction multiplies the 16-bit signed integers from the selected halves of Rn and Rm, adds the 32-bit result to the value from Ra, and writes the result in Rd.

Syntax

SMLA<x><y>{cond} Rd, Rn, Rm, Ra

where:

<x> and <y> can be either B or T. B means use the bottom half (bits [15:0]) of a register, T means use the top half (bits [31:16]) of a register. <x> specifies which half of Rn to use, <y> does the same for Rm.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

Ra is the register which holds the accumulate value.

A.1.108 SMLAD

SMLAD (Dual Signed Multiply Accumulate; $32 \leq 32 + 16 \times 16 + 16 \times 16$) multiplies the bottom halfword of Rn with the bottom halfword of Rm, and the top halfword of Rn with the top halfword of Rm. It then adds both products to the value in Ra and writes the sum to Rd.

Syntax

SMLAD{X}{cond} Rd, Rn, Rm, Ra

where:

{X} if present, means that the most and least significant halfwords of the second operand are swapped, before the multiplication.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

Ra is the register which holds the accumulate value.

A.1.109 SMLAL

SMLAL (Signed Multiply Accumulate $64 \leq 64 + 32 \times 32$) multiplies Rn and Rm (treated as signed integers) and adds the 64-bit result to the 64-bit signed integer contained in RdHi and RdLo.

Syntax

SMLAL{S}{cond} RdLo, RdHi, Rn, Rm

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

RdLo and RdHi are the destination registers.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

A.1.110 SMLALxy

SMLALxy (Signed Multiply Accumulate; $64 \leq 64 + 16 \times 16$) multiplies the signed integer from the selected half of Rm by the signed integer from the selected half of Rn, and adds the 32-bit result to the 64-bit value in RdHi and RdLo.

Syntax

SMLAL<x><y>{cond} RdLo, RdHi, Rn, Rm

where:

<x> and <y> can be either B or T. B means use the bottom half (bits [15:0]) of a register, T means use the top half (bits [31:16]) of a register. <x> specifies which half of Rn to use, <y> does the same for Rm.

cond is an optional condition code. See Section 6.1.2.

RdLo and RdHi are the destination registers.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

A.1.111 SMLALD

SMLALD (Dual Signed Multiply Accumulate Long; $64 \leq 64 + 16 \times 16 + 16 \times 16$) multiplies the bottom halfword of Rn with the bottom halfword of Rm, and the top halfword of Rn with the top halfword of Rm and adds both products to the value in RdLo, RdHi and stores the result in RdLo and RdHi.

Syntax

SMLALD{X}{cond} RdLo, RdHi Rn, Rm

where:

{X} if present, means that the most and least significant halfwords of the second operand are swapped, before the multiplication.

cond is an optional condition code. See Section 6.1.2.

RdLo and RdHi are the destination registers.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

A.1.112 SMLAWy

SMLAW (Signed Multiply with Accumulate Wide; $32 \leq 32 \times 16 + 32$) multiplies the signed integer from the selected half of Rm by the signed integer from Rn, adds the 32-bit result to the 32-bit value in Ra, and writes the result in Rd.

Syntax

SMLAW<y>{cond} Rd, Rn, Rm, Ra

where:

<y> can be either B or T. B means use the bottom half (bits [15:0]) of Rm, T means use the top half (bits [31:16]) of Rm.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

Ra is the register which holds the accumulate value.

A.1.113 SMLS LD

SMLS LD (Dual Signed Multiply Subtract Accumulate Long; $64 \leq 64 + 16 \times 16 - 16 \times 16$) multiplies Rn[15:0] with Rm[15:0] and Rn[31:16] with Rm[31:16]. It then subtracts the second product from the first, adds the difference to the value in RdLo, RdHi, and writes the result to RdLo, RdHi.

Syntax

SMLS LD{X}{cond} RdLo, RdHi Rn, Rm

where:

{X} if present, means that the most and least significant halfwords of the second operand are swapped, before the multiplication.

cond is an optional condition code. See Section 6.1.2.

RdLo and RdHi are the destination registers and hold the value to be accumulated.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

A.1.114 SMMLA

SMMLA (Signed top word Multiply with Accumulate; $32 \leq \text{top word } (32 \times 32 + 32)$) multiplies Rn and Rm, adds Ra to the most significant 32 bits of the product, and writes the result in Rd.

Syntax

SMMLA{R}{cond} Rd, Rn, Rm, Ra

where:

R, if present means that 0×80000000 is added before extracting the most significant 32 bits. This rounds the result.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

Ra is the register holding the accumulate value.

A.1.115 SMMLS

SMMLS (Signed top word Multiply with Subtract; $32 \leq \text{top word } (32 \times 32 - 32)$) multiplies Rn and Rm, subtracts the product from the value in Ra shifted left by 32 bits, and stores the most significant 32 bits of the result in Rd.

Syntax

SMMUL{R}{cond} Rd, Rn, Rm, Ra

where:

R, if present means that 0x80000000 is added before extracting the most significant 32 bits. This rounds the result.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

Ra is the register holding the accumulate value.

A.1.116 SMMUL

SMMUL (Signed top word Multiply; 32 <= top word (32 × 32)) multiplies Rn and Rm, and writes the most significant 32 bits of the 64-bit result to Rd.

Syntax

SMMUL{R}{cond} Rd, Rn, Rm

where:

R, if present means that 0x80000000 is added before extracting the most significant 32 bits. This rounds the result.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

A.1.117 SMUAD

SMUAD (Dual Signed Multiply and Add products) multiplies Rn [15:0] with Rm [15:0] and Rn [31:16] with Rm [31:16]. It then adds the products and stores the sum to Rd.

Syntax

SMUAD{X}{cond} Rd, Rn, Rm

where:

X, if present means that the most and least significant halfwords of the second operand are exchanged before the multiplications occur.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

A.1.118 SMUSD

SMUSD (Dual Signed Multiply and Subtract products) multiplies R_n [15:0] with R_m [15:0] and R_n [31:16] with R_m [31:16]. It then subtracts the products and stores the sum to R_d .

Syntax

SMUSD{X}{cond} R_d , R_n , R_m

where:

X, if present means that the most and least significant halfwords of the second operand are exchanged before the multiplications occur.

cond is an optional condition code. See Section 6.1.2.

R_d is the destination register.

R_n is the register holding the first multiplicand.

R_m is the register holding the second multiplicand.

A.1.119 SMULxy

The SMULxy (Signed Multiply ($32 \leq 16 \times 16$)) instruction multiplies the 16-bit signed integers from the selected halves of R_n and R_m , and places the 32-bit result in R_d .

Syntax

SMUL<x><y>{cond} { R_d }, R_n , R_m

where:

<x> and <y> can be either B or T. B means use the bottom half (bits [15:0]) of a register, T means use the top half (bits [31:16]) of a register. <x> specifies which half of R_n to use, <y> does the same for R_m .

cond is an optional condition code. See Section 6.1.2.

R_d is the destination register.

R_n is the register holding the first multiplicand.

R_m is the register holding the second multiplicand.

A.1.120 SMULL

The SMULL (signed multiply long; $64 \leq 32 \times 32$) instruction multiplies R_n and R_m (treated as containing as two's complement signed integers) and places the least significant 32 bits of the result in R_dLo , and the most significant 32 bits of the result in R_dHi .

Syntax

SMULL{S}{cond} R_dLo , R_dHi , R_n , R_m

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

RdLo and RdHi are the destination registers.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

A.1.121 SMULWy

SMULWy (Signed Multiply Wide; $32 \leq 32 \times 16$) multiplies the signed integer from the chosen half of Rm with the signed integer from Rn, and places the upper 32-bits of the 48-bit result in Rd.

Syntax

SMULW<y>{cond} {Rd}, Rn, Rm

where:

<y> is either B or T. B means use the bottom half (bits [15:0]) of Rm, T means use the top half (bits [31:16]) of Rm.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

A.1.122 SRS

SRS (Store Return State) stores the LR and the SPSR of the current mode, at the address contained in the SP of the mode specified by modenum. The optional ! means that the SP value is updated. This is compatible with the normal use of the STM instruction for stack accesses.

Syntax

SRS{addr_mode}{cond} sp{!}, #modenum

where:

addr_mode is one of:

- IA – Increment address After each transfer. This is the default, and can be omitted.
- IB – Increment address Before each transfer (ARM only).
- DA – Decrement address After each transfer (ARM only).
- DB – Decrement address Before each transfer.

It is also possible to use the corresponding stack oriented addressing modes (FD, ED, EA, FA).

cond is an optional condition code. See Section 6.1.2.

modenum gives the number of the mode whose SP is used.

A.1.123 SSAT

SSAT (Signed Saturate) performs a shift and saturates the result to the signed range $-2^{\text{sat}-1} \leq x \leq 2^{\text{sat}-1}-1$. If saturation occurs, the Q flag is set.

Syntax

SSAT{cond} Rd, #sat, Rm{, shift}

where:

<y> is either B or T. B means use the bottom half (bits [15:0]) of Rm, T means use the top half (bits [31:16]) of Rm.

cond is an optional condition code. See Section 6.1.2

Rd is the destination register.

sat specifies the bit position to saturate to, in the range 1 to 32.

Rm is the register holding the second multiplicand.

shift is optional shift amount and can be either ASR #n where n is in the range (1 – 32 ARM state, 1 – 31 Thumb state) or LSL #n where n is in the range (0-31).

A.1.124 SSAT16

SSAT16 (Signed Saturate, parallel halfwords) saturates each signed halfword to the signed range $-2^{\text{sat}-1} \leq x \leq 2^{\text{sat}-1} - 1$. If saturation occurs, the Q flag is set.

Syntax

SSAT16{cond} Rd, #sat, Rn

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

sat specifies the bit position to saturate to, in the range 1 to 32.

Rn is the register holding the operand.

A.1.125 SSAX

SSAX (Signed Subtract Add Exchange) exchanges halfwords of Rm, then subtracts the top halfwords and adds the bottom halfwords.

Syntax

SSAX{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the register holding the operands.

A.1.126 SSUB8

SSUB8 (Signed halving bitwise Subtraction) does a signed bitwise subtraction (4 subtracts).

Syntax

SSUB8{cond} {Rd}, Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination registers.

Rm and Rn are the register holding the operands.

A.1.127 SSUB16

SSUB16 (Signed halfword-wise Subtract) does a signed halfword-wise subtraction (2 subtracts).

Syntax

SSUB16{cond} {Rd}, Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination registers.

Rm and Rn are the register holding the operands.

A.1.128 STC

STC (Store Coprocessor Registers) writes a coprocessor register to memory (or multiple registers, if L is specified).

Syntax

STC{L}{cond} coproc, CRd, [Rn]
 STC{L}{cond} coproc, CRd, [Rn, #{-}offset] {!}
 STC{L}{cond} coproc, CRd, [Rn], #{-}offset
 STC{L}{cond} coproc, CRd, label

where:

L specifies that more than one register can be transferred (called a long transfer). The length of the transfer is determined by the coprocessor, but may not be more than 16 words.

cond is an optional condition code. See Section 6.1.2.

coproc is the name of the coprocessor the instruction is for. This is usually of the form pn, where n is an integer in the range 0 to 15.

CRd is the coprocessor register to be stored.

Rn is the register holding the base address for the memory operation.

offset is a multiple of four, in the range 0-1020, to be added or subtracted from Rn. If ! is present, the address including the offset is written back into Rn.

label is a word-aligned PC-relative address label.

A.1.129 STC2

STC2 (Store Coprocessor registers) writes a coprocessor register to memory (or multiple registers, if L is specified).

Syntax

```
STC2{L}{cond} coproc, CRd, [Rn]
STC2{L}{cond} coproc, CRd, [Rn, #{-}offset]{!}
STC2{L}{cond} coproc, CRd, [Rn], #{-}offset
STC2{L}{cond} coproc, CRd, label
```

where:

L specifies that more than one register can be transferred (called a long transfer). The length of the transfer is determined by the coprocessor, but may not be more than 16 words.

cond is an optional condition code. See Section 6.1.2.

coproc is the name of the coprocessor the instruction is for. This is usually of the form pn, where n is an integer in the range 0 to 15.

CRd is the coprocessor register to be stored.

Rn is the register holding the base address for the memory operation.

offset is a multiple of four, in the range 0 – 1020, to be added or subtracted from Rn. If ! is present, the address including the offset is written back into Rn.

label is a word-aligned PC-relative address label.

A.1.130 STM

STM (Store Multiple registers) writes one or more registers to consecutive addresses in memory to an address specified in a base register.

Syntax

```
STM{addr_mode}{cond} Rn{!}, reglist{^}
```

where:

addr_mode is one of:

- IA – Increment address After each transfer. This is the default, and can be omitted.
- IB – Increment address Before each transfer (ARM only).
- DA – Decrement address After each transfer (ARM only).
- DB – Decrement address Before each transfer.

It is also possible to use the corresponding stack oriented addressing modes (FD, ED, EA, FA). For example STMFD is a synonym of STMDB.

cond is an optional condition code. See Section 6.1.2.

Rn is the base register, giving the initial address for the transfer.

! if present, specifies that the final address is written back into Rn.

^ if specified (in ARM state and a mode other than User or System) means that data is transferred into or out of the User mode registers instead of the current mode registers.

reglist is a list of one or more registers to be stored, enclosed in braces. It can contain register ranges. It must be comma separated if it contains more than one register or register range.

A.1.131 STR

STR (Store Register) stores a value to memory from an ARM register, optionally updating the register used to give the address.

A variety of addressing options are provided. For full details of the available addressing modes, see [Addressing modes on page 6-10](#).

Syntax

```
STR{type}{T}{cond} Rt, [Rn {, #offset}]
STR{type}{cond} Rt, [Rn, #offset]!
STR{type}{T}{cond} Rt, [Rn], #offset
STR{type}{cond} Rt, [Rn, +/-Rm {, shift}]
STR{type}{cond} Rt, [Rn, +/-Rm {, shift}]!
STR{type}{T}{cond} Rt, [Rn], +/-Rm {, shift}
```

where:

type can be any one of:

- B – unsigned Byte (Zero extend to 32 bits on loads.)
- SB – signed Byte (Sign extend to 32 bits.)
- H – unsigned Halfword (Zero extend to 32 bits on loads.)
- SH – signed Halfword (Sign extend to 32 bits.)

or omitted, for a Word load.

T specifies that memory is accessed as if the processor was in User mode (not available in all addressing modes).

cond is an optional condition code. See Section 6.1.2.

Rn is the register holding the base address for the memory operation.

! if present, specifies that the final address is written back into Rn.

offset is a numeric value.

Rm is a register holding an offset value to be applied.

shift is either a register or immediate based shift to apply to the offset value.

A.1.132 STRD

STRD (Store Register Dual) calculates an address from a base register value and a register offset, and stores two words from two registers to memory. It can use offset, post-indexed, or pre-indexed addressing.

Syntax

```
STRD{cond} Rt, Rt2, [Rn {, #+/-<imm>}]
STRD{cond} Rt, Rt2, [<Rn>, #+/-<imm>]
STRD{cond} Rt, Rt2, [<Rn>, #+/-<imm>]!
STRD{cond} Rt, Rt2, [{Rn}, +/-{Rm}]!
STRD{cond} Rt, Rt2, [{Rn}], +/-{Rm}
```

where:

cond is an optional condition code. See Section 6.1.2.

Rt is the first source register. For an ARM instruction Rt must be even-numbered and not R14.

Rt is the second source register. For an ARM instruction Rt2 must be <R(t+1)>.

Rn is the base register. The SP can be used. In the ARM instruction set for offset addressing only, the PC can be used. However, use of the PC is deprecated.

+/- is + or omitted if the value of <Rm> is to be added to the base register value (add = TRUE), or – if it is to be subtracted (add = FALSE). #0 and #-0 generate different instructions.

imm is the immediate offset used to form the address. imm can be omitted, meaning an offset of 0.

Rm contains the offset that is applied to the value of <Rn> to form the address.

A.1.133 STREX

STREX (Store register exclusive). Performs a store to a location marked for exclusive access, returning a status value if the store succeeded. Byte, halfword, word and doubleword variants are provided.

Syntax

```
STREX{cond} Rd, Rt, [Rn {, #offset}]
STREXB{cond} Rd, Rt, [Rn]
STREXH{cond} Rd, Rt, [Rn]
STREXD{cond} Rd, Rt, Rt2, [Rn]
```

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register for the return status.

Rt is the register to store.

Rt2 is the second register for doubleword stores.

Rn is the register holding the address.

offset is an optional value, allowed in Thumb only.

A.1.134 SUB

SUB (Subtract) subtracts the value Operand2 from Rn (or subtracts imm12 from Rn).

Syntax

```
SUB{S}{cond} {Rd,} Rn, <Operand2>
SUB{cond}{Rd,}, Rn, #imm12 (Only available in Thumb)
```

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first operand.

Operand2 is a flexible second operand. See Section 6.2.1.

imm12 is in the range 0-4095.

A.1.135 SVC

SVC (SuperVisor Call) causes an SVC exception (was called SWI in older documentation).

Syntax

SVC{cond} #imm

where:

cond is an optional condition code. See Section 6.1.2.

imm is an integer in the range 0 - 0xFFFF (ARM) or 0 - 0xFF (Thumb). This integer is not used by the processor itself, but can be used by exception handler code.

A.1.136 SWP

SWP (Swap registers and memory) performs the following two actions. Data from memory is loaded into Rt. Rt2 is saved to memory, at the address given by Rn. Use of this instruction is deprecated and its use is disabled by default.

Syntax

SWP{B}{cond} Rt, Rt2, [Rn]

where:

B is an optional suffix. If specified, a byte is swapped. If not present, a word is specified.

Rt is the destination register.

Rt2 is the source register and can be the same as Rt.

Rn is the register holding the address and cannot be the same as Rt or Rt2.

A.1.137 SXT

SXT (Signed Extend) extracts the specified byte and extends to 32-bit.

Syntax

SXT<extend>{cond} {Rd}, Rm {,rotation}

where:

extend must be one of:

- B16 – extends two 8-bit values to two 16-bit values.
- B – extends an 8-bit value to a 32-bit value.
- H – extends a 16-bit value to a 32-bit value.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm is the register which contains the value to be extended.

rotation can be one of ROR #8, ROR #16 or ROR #24 (or can be omitted).

A.1.138 SXTA

SXTA (Signed Extend and Add) extracts the specified byte, adds the value from Rn and extends to 32-bit.

Syntax

SXTA<extend>{cond} {Rd,} Rn, Rm {,rotation}

where:

extend must be one of:

- B16 – extends two 8-bit values to two 16-bit values.
- B – extends an 8-bit value to a 32-bit value.
- H – extends a 16-bit value to a 32-bit value.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the value to be added.

Rm is the register which contains the value to be extended.

rotation can be one of ROR #8, ROR #16 or ROR #24 (or can be omitted).

A.1.139 SYS

SYS (System coprocessor instruction) is used to execute special coprocessor instructions such as cache, branch predictor, and TLB operations. The instructions operate by writing to special write-only coprocessor registers.

Syntax

SYS{cond} instruction {,Rn}

where:

cond is an optional condition code. See Section 6.1.2.

instruction is a write-only system coprocessor register name.

Rn is the register holding the operand.

A.1.140 TBB

TBB (Table Branch Byte) causes a PC-relative forward branch using a table of single byte offsets. Rn provides a pointer to the table, and Rm supplies an index into the table. The branch length is twice the value of the byte returned from the table. The target of the branch table must be in the same execution state. There is no ARM or 16-bit Thumb version of this instruction.

Syntax

TBB [Rn, Rm]

where:

Rn is the base register which holds the address of the table of branch lengths.

Rm is a register which holds the index into the table.

A.1.141 TBH

TBH (Table Branch Halfword) causes a PC-relative forward branch using a table of halfword offsets. Rn provides a pointer to the table, and Rm supplies an index into the table. The branch length is twice the value of the halfword returned from the table. The target of the branch table must be in the same execution state.

There is no ARM or 16-bit Thumb version of this instruction.

Syntax

TBH [Rn, Rm, LSL #1]

where:

Rn is the base register which holds the address of the table of branch lengths.

Rm is a register which holds the index into the table.

A.1.142 TEQ

TEQ (Test Equivalence) does a bitwise AND operation on the value in Rn and the value of Operand2. This is the same as an ANDS instruction, except that the result is discarded.

Syntax

TEQ{cond} Rn, <Operand2>

where:

cond is an optional condition code. See Section 6.1.2.

Rn is the register holding the first operand.

Operand2 is a flexible second operand. See Section 6.2.1.

A.1.143 TST

TST (Test) does an Exclusive OR operation on the value in Rn and the value of Operand2. This is the same as an EORS instruction, except that the result is discarded.

Syntax

TST{cond} Rn, <Operand2>

where:

cond is an optional condition code. See Section 6.1.2.

Rn is the register holding the first operand.

Operand2 is a flexible second operand. See Section 6.2.1.

A.1.144 UADD8

UADD8 (Unsigned bytewise Add) does an unsigned bytewise addition (4 adds).

Syntax

UADD8{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.145 UADD16

UADD16 (Unsigned halfword-wise Add) does an unsigned halfword-wise addition (2 adds).

Syntax

UADD16{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.146 UASX

UASX (Unsigned Add Subtract Exchange) exchanges halfwords of Rm, then adds the top halfwords and subtracts the bottom halfwords.

Syntax

UASX{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.147 UBFX

UBFX (Unsigned Bit Field Extract) writes adjacent bits from one register into the least significant bits of a second register and zero extends to 32 bits.

Syntax

UBFX{cond} Rd, Rn, #lsb, #width

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register which contains the bits to be extracted.

lsb specifies the least significant bit of the bitfield.

width is the width of the bitfield.

A.1.148 UDIV

UDIV (Unsigned Divide). divides a 32-bit unsigned integer register value by a 32-bit unsigned integer register value, and writes the result to the destination register. This instruction is not present in all variants of the ARMv7-A architecture.

Syntax

UDIV{cond}{q} {Rd,} Rn, Rm

where:

cond is the optional condition code. See Section 6.1.2.

q specifies optional assembler qualifiers on the instruction. The following qualifiers are defined:

- .N (narrow), specifies that the assembler must select a 16-bit encoding for the instruction.
- .W (wide), specifies that the assembler must select a 32-bit encoding for the instruction.

If neither .W nor .N are specified, the assembler can select either 16-bit or 32-bit encodings. If both are available, it must select a 16-bit encoding.

Rd the destination register.

Rn is the register that contains the dividend.

Rm is the register that contains the divisor.

A.1.149 UHADD8

UHADD8 (Unsigned Halving bitwise Add) does an unsigned bitwise addition (4 adds) and halves the results.

Syntax

UHADD8{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.150 UHADD16

UHADD16 (Unsigned Halving halfword-wise Add) does an unsigned halfword-wise addition (2 adds) and halves the results.

Syntax

UHADD16{cond} {Rd}, Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.151 UHASX

UHASX (Unsigned Halving Add Subtract Exchange) exchanges halfwords of Rm, then adds the top halfwords and subtracts the bottom halfwords and halves the results.

Syntax

UHASX{cond} {Rd}, Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.152 UHSAX

UHSAX (Unsigned Halving Subtract Add Exchange) exchanges halfwords of Rm, then subtracts the top halfwords and adds the bottom halfwords and halves the results.

Syntax

UHSAX{cond} {Rd}, Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.153 UHSUB8

UHSUB8 (Unsigned Halving bitwise Subtraction) does an unsigned bitwise subtraction (4 subtracts) and halves the results.

Syntax

UHSUB8{cond} {Rd}, Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the register holding the operands.

A.1.154 UHSUB16

UHSUB16 (Unsigned Halving halfword-wise Subtract) does an unsigned halfword-wise subtraction (2 subtracts) and halves the result.

Syntax

UHSUB16{cond} {Rd}, Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.155 UMAAL

UMAAL (Unsigned Multiply Accumulate Long; $64 \leq 32 + 32 + 32 \times 32$) multiplies Rn and Rm (treated as unsigned integers) adds the two 32-bit values in RdHi and RdLo, and stores the 64-bit result to RdLo, RdHi.

Syntax

UMAAL{cond} RdLo, RdHi, Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

RdLo and RdHi are the destination accumulator registers.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

A.1.156 UMLAL

UMLAL (Unsigned Multiply Accumulate $64 \leq 64 + 32 \times 32$) multiplies Rn and Rm (treated as unsigned integers) and adds the 64-bit result to the 64-bit unsigned integer contained in RdHi and RdLo.

Syntax

UMLAL{S}{cond} RdLo, RdHi, Rn, Rm

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

RdLo and RdHi are the destination accumulator registers.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

A.1.157 UMULL

UMULL (Unsigned Multiply; $64 \leq 32 \times 32$) multiplies *Rn* and *Rm* (treated as unsigned integers) and stores the least significant 32 bits of the result in *RdLo*, and the most significant 32 bits of the result in *RdHi*.

Syntax

UMULL{S}{cond} *RdLo*, *RdHi*, *Rn*, *Rm*

where:

S (if specified) means that the condition code flags will be updated depending upon the result of the instruction.

cond is an optional condition code. See Section 6.1.2.

RdLo and *RdHi* are the destination registers.

Rn is the register holding the first multiplicand.

Rm is the register holding the second multiplicand.

A.1.158 UQADD8

UQADD8 (Saturating Unsigned bytewise Add) does an unsigned bytewise addition (4 adds) and saturates the results to the unsigned range $0 \leq x \leq 2^8-1$. The Q flag is not affected by this instruction.

Syntax

UQADD8{cond} {*Rd*,} *Rn*, *Rm*

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and *Rn* are the registers holding the operands.

A.1.159 UQADD16

UQADD16 (Saturating Unsigned halfword-wise Add) does an unsigned halfword-wise addition (2 adds) and saturates the results to the unsigned range $0 \leq x \leq 2^{16}-1$. The Q flag is not affected by this instruction.

Syntax

UQADD16{cond} {*Rd*,} *Rn*, *Rm*

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and *Rn* are the register holding the operands.

A.1.160 UQASX

UQASX (Saturating Unsigned Add Subtract Exchange) exchanges halfwords of R_m , then adds the top halfwords and subtracts the bottom halfwords and saturates the results to the unsigned range $0 \leq x \leq 2^{16}-1$. The Q flag is not affected by this instruction.

Syntax

UQASX{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.161 UQSAX

UQSAX (Saturating Unsigned Subtract Add Exchange) exchanges the halfwords of R_m , then subtracts the top halfwords and adds the bottom halfwords and saturates the results to the signed range $0 \leq x \leq 2^{16}-1$. The Q flag is not affected by this instruction.

Syntax

QSAX{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.162 UQSUB8

UQSUB8 (Saturating Unsigned bitwise Subtract) does bitwise subtraction (4 subtracts), with saturation of the results to the unsigned range $0 \leq x \leq 2^8-1$. The Q flag is not affected by this instruction.

Syntax

UQSUB8{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.163 UQSUB16

UQSUB16 (Saturating Unsigned halfword Subtract) does halfword-wise subtraction (two subtracts), with saturation of the results to the unsigned range $0 \leq x \leq 2^{16}-1$. The Q flag is not affected by this instruction.

Syntax

UQSUB16{cond} {Rd}, Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.164 USAD8

USAD8 (Unsigned Sum of Absolute Differences) finds the 4 differences between the unsigned values in corresponding bytes of Rn and Rm and adds the absolute values of the 4 differences, and stores the result in Rd.

Syntax

USAD8{cond} Rd, Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first operand.

Rm is the register holding the second operand.

A.1.165 USADA8

USADA8 (Unsigned Sum of Absolute Differences Accumulate) finds the 4 differences between the unsigned values in corresponding bytes of Rn and Rm and adds the absolute values of the 4 differences to the value in Ra, and stores the result in Rd.

See [Sum of absolute differences on page 6-15](#) for more information.

Syntax

USADA8{cond} Rd, Rn, Rm, Ra

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the first operand.

Rm is the register holding the second operand.

Ra is the register which holds the accumulate value.

A.1.166 USAT

USAT (Unsigned Saturate) performs a shift and saturates the result to the signed range $0 \leq x \leq 2^{\text{sat}} - 1$. If saturation occurs, the Q flag is set.

Syntax

USAT{cond} Rd, #sat, Rm{, shift}

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

sat specifies the bit position to saturate to, in the range 0 to 31.

Rm is the register holding the operand.

shift is optional shift amount and can be either ASR #n where n is in the range (1 – 32 ARM state, 1-31 Thumb state) or LSL #n where n is in the range (0 – 31).

A.1.167 USAT16

USAT16 (Unsigned Saturate, parallel halfwords) saturates each unsigned halfword to the signed range $0 \leq x \leq 2^{\text{sat}} - 1$. If saturation occurs, the Q flag is set.

Syntax

USAT16{cond} Rd, #sat, Rn

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

sat specifies the bit position to saturate to, in the range 0 to 31.

Rn is the register holding the operand.

A.1.168 USAX

USAX (Unsigned Subtract Add Exchange) exchanges halfwords of Rm, then subtracts the top halfwords and adds the bottom halfwords.

Syntax

USAX{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.169 USUB8

USUB8 (Unsigned bitwise Subtraction) does an unsigned bitwise subtraction (4 subtracts).

Syntax

USUB8{cond} {Rd,} Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.170 USUB16

USUB16 (Unsigned halfword-wise Subtract) does an unsigned halfword-wise subtraction (2 subtracts).

Syntax

USUB16{cond} {Rd}, Rn, Rm

where:

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm and Rn are the registers holding the operands.

A.1.171 UXT

UXT (Unsigned Extend) extracts the specified byte and zero extends to a 32-bit value.

Syntax

UXT<extend>{cond} {Rd}, Rm {,rotation}

where:

extend must be one of:

- B16 – extends two 8-bit values to two 16-bit values.
- B – extends an 8-bit value to a 32-bit value.
- H – extends a 16-bit value to a 32-bit value.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rm is the register which contains the value to be extended.

rotation can be one of ROR #8, ROR #16 or ROR #24 (or can be omitted).

A.1.172 UXTA

UXTA (Unsigned Extend and Add) extracts the specified byte, adds the value from Rn and zero extends to a 32-bit value.

Syntax

UXTA<extend>{cond} {Rd}, Rn, Rm {,rotation}

where:

extend must be one of:

- B16 – extends two 8-bit values to two 16-bit values.
- B – extends an 8-bit value to a 32-bit value.
- H – extends a 16-bit value to a 32-bit value.

cond is an optional condition code. See Section 6.1.2.

Rd is the destination register.

Rn is the register holding the value to be added.

Rm is the register which contains the value to be extended.

rotation can be one of ROR #8, ROR #16 or ROR #24 (or can be omitted).

A.1.173 WFE

WFE (Wait for Event). If the Event Register is not set, WFE suspends execution until one of the following events occurs:

- an IRQ interrupt (even when CPSR I-bit is set)
- an FIQ interrupt (even when CPSR F-bit is set)
- an asynchronous abort (not when masked by the CPSR A-bit)
- Debug Entry request, even when debug is disabled.
- an Event signaled by another processor using the SEV instruction.

If the Event Register is set, WFE clears it and returns immediately.

Syntax

WFE{cond}

where:

cond is an optional condition code. See Section 6.1.2.

A.1.174 WFI

WFI (Wait For Interrupt) suspends execution until one of the following events occurs:

- An IRQ interrupt (even when CPSR I-bit is set).
- An FIQ interrupt (even when CPSR F-bit is set).
- An asynchronous abort (not when masked by the CPSR A-bit).
- Debug Entry request, even when debug is disabled.

If the Event Register is set, WFE clears it and returns immediately.

Syntax

WFI{cond}

where:

cond is an optional condition code. See Section 6.1.2.

A.1.175 YIELD

YIELD indicates to the hardware that the current thread is performing a task that can be swapped out (for example, a spinlock). Hardware could use this hint to suspend and resume threads in a multithreading system.

Syntax

```
YIELD{cond}
```

where:

cond is an optional condition code. See Section 6.1.2.

Appendix B

NEON and VFP Instruction Summary

This appendix provides a summary of the assembly language instructions available to the NEON and VFP programmer. It groups the instructions into the following classifications:

For each instruction, we provide a description of the syntax, operands and behavior. Not all usage restrictions are documented here, nor do we show the associated binary encoding. For most instructions, further explanation can be found in [Chapter 8](#) and [Chapter 20](#).

This appendix is intended for quick reference. If more detail about the precise operation of an instruction is required, please refer to the *ARM Architecture Reference Manual*, or to the official ARM documentation (for example the *ARM Compiler Toolchain Assembler Reference*) which can be found at <http://infocenter.arm.com/help/index.jsp>.

The NEON and VFP instructions are divided into a number of sections:

- [NEON general data processing instructions on page B-6.](#)
- [NEON shift instructions on page B-13.](#)
- [NEON logical and compare operations on page B-17.](#)
- [NEON arithmetic instructions on page B-23.](#)
- [NEON multiply instructions on page B-32.](#)
- [NEON load and store element and structure instructions on page B-35.](#)
- [VFP instructions on page B-41.](#)
- [NEON and VFP pseudo-instructions on page B-47.](#)

Within each of these groups, instructions are listed alphabetically. [Table B-1](#) shows an alphabetic listing of all NEON and VFP instructions, showing which section of this appendix describes them and which instruction sets support the instruction.

Table B-1 NEON (Advanced SIMD) and VFP instructions

Instruction	Section	Instruction set
V{Q}{R}SHL	B2	Advanced-SIMD
V{Q}ABS	B4	Advanced-SIMD
V{Q}ADD	B4	Advanced-SIMD
V{Q}MOVN	B1	Advanced-SIMD
V{Q}SUB	B4	Advanced-SIMD
V{R}ADDHN	B4	Advanced-SIMD
V{R}HADD	B4	Advanced-SIMD
V{R}SHR{N}	B2	Advanced-SIMD
V{R}SRA	B2	Advanced-SIMD
V{R}SUBHN	B4	Advanced-SIMD
VABA{L}	B4	Advanced-SIMD
VABD{L}	B4	Advanced-SIMD
VABS	B7	VFP
VACGE	B3	Advanced-SIMD
VACGT	B3	Advanced-SIMD
VACLE	B8	Advanced-SIMD
VACLT	B8	Advanced-SIMD
VADD	B7	VFP
VADDL	B4	Advanced-SIMD
VADDW	B4	Advanced-SIMD
VAND	B3	Advanced-SIMD
	B8	Advanced-SIMD
VBIC	B3	Advanced-SIMD
VBIF	B3	Advanced-SIMD
VBIT	B3	Advanced-SIMD
VBSL	B3	Advanced-SIMD
VCEQ	B3	Advanced-SIMD
VCGE	B3	Advanced-SIMD
VCGT	B3	Advanced-SIMD

Table B-1 NEON (Advanced SIMD) and VFP instructions (continued)

Instruction	Section	Instruction set
VCLE	B3	Advanced-SIMD
	B8	Advanced-SIMD
VCLS	B4	Advanced-SIMD
VCLT	B3	Advanced-SIMD
	B8	Advanced-SIMD
VCLZ	B4	Advanced-SIMD
VCMP	B7	VFP
VCNT	B4	Advanced-SIMD
VCVT	B1	Advanced-SIMD
	B7	VFP
VCVTB	B7	VFP
VCVTT	B7	VFP
VDIV	B7	VFP
VDUP	B1	Advanced-SIMD
VEOR	B3	Advanced-SIMD
VEXT	B1	Advanced-SIMD
VFMA	B5	Advanced-SIMD
	B7	VFP
VFMS	B5	Advanced-SIMD
	B7	VFP
VFNMA	B7	VFP
VFNMS	B7	VFP
VHSUB	B4	Advanced-SIMD
VLD	B8	Advanced-SIMD, VFP
VLDM	B6	Advanced-SIMD, VFP
VLDR	B6	Advanced-SIMD, VFP
VMAX	B4	Advanced-SIMD
VMIN	B4	Advanced-SIMD
VMLA	B7	VFP
VMLA{L}	B5	Advanced-SIMD
VMLS	B7	VFP
VMLS{L}	B5	Advanced-SIMD

Table B-1 NEON (Advanced SIMD) and VFP instructions (continued)

Instruction	Section	Instruction set
VMOV	B1	Advanced-SIMD
	B3	Advanced-SIMD
	B7	VFP
	B6	Advanced-SIMD, VFP
VMOV2	B8	Advanced-SIMD
VMOVL	B1	Advanced-SIMD
VMRS	B6	Advanced-SIMD, VFP
	B6	Advanced-SIMD, VFP
VMUL	B7	VFP
VMUL{L}	B5	Advanced-SIMD
VMVN	B1	Advanced-SIMD
	B3	Advanced-SIMD
VNEG	B7	VFP
VNMLA	B7	VFP
VNMLS	B7	VFP
VNMUL	B7	VFP
VORN	B3	Advanced-SIMD
	B8	Advanced-SIMD
VORR	B3	Advanced-SIMD
VPADAL	B4	Advanced-SIMD
VPADD{L}	B4	Advanced-SIMD
VPMAX	B4	Advanced-SIMD
VPMIN	B4	Advanced-SIMD
VPOP	B6	Advanced-SIMD, VFP
VPUSH	B6	Advanced-SIMD, VFP
VQ{R}DMULH	B5	Advanced-SIMD
VQ{R}SHR{U}N	B2	Advanced-SIMD
VQDMLAL	B5	Advanced-SIMD
VQDMLSL	B5	Advanced-SIMD
VQDMULL	B5	Advanced-SIMD
VQMOVUN	B1	Advanced-SIMD
VQSHL	B2	Advanced-SIMD
VQSHLU	B2	Advanced-SIMD

Table B-1 NEON (Advanced SIMD) and VFP instructions (continued)

Instruction	Section	Instruction set
VRECPE	B4	Advanced-SIMD
VRECPS	B4	Advanced-SIMD
VREV	B1	Advanced-SIMD
VRSQRTE	B4	Advanced-SIMD
VRSQRTS	B4	Advanced-SIMD
VSHL	B2	Advanced-SIMD
VSHLL	B2	Advanced-SIMD
VSLI	B2	Advanced-SIMD
VSQRT	B7	VFP
VSRI	B2	Advanced-SIMD
VST	B8	Advanced-SIMD, VFP
VSTM	B6	Advanced-SIMD, VFP
VSTR	B6	Advanced-SIMD, VFP
VSUB	B7	VFP
VSUBL	B4	Advanced-SIMD
VSUBW	B4	Advanced-SIMD
VSWP	B1	Advanced-SIMD
VTBL	B1	Advanced-SIMD
VTBX	B1	Advanced-SIMD
VTRN	B1	Advanced-SIMD
VTST	B3	Advanced-SIMD
VUZP	B1	Advanced-SIMD
VZIP	B1	Advanced-SIMD

B.1 NEON general data processing instructions

This section covers NEON data processing instructions, including those which extract or manipulate values in registers, perform type conversion or other operations.

B.1.1 VCVT (fixed-point or integer to floating-point)

VCVT (Vector Convert) converts each element in a vector and places the results in the destination vector. The possible conversions are:

- Floating-point to integer.
- Integer to floating-point.
- Floating-point to fixed-point.
- Fixed-point to floating-point.

Integer or fixed-point to floating-point conversions use round to nearest.

Floating-point to integer or fixed-point conversions use round towards zero.

Syntax

```
VCVT{cond}.type Qd, Qm {, #fbits}
VCVT{cond}.type Dd, Dm {, #fbits}
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

type specifies the data types for the elements of the vectors and can be:

- S32.F32 floating-point to signed integer or fixed-point.
- U32.F32 floating-point to unsigned integer or fixed-point.
- F32.S32 signed integer or fixed-point to floating-point.
- F32.U32 unsigned integer or fixed-point to floating-point.

Qd and Qm specify the destination and operand vectors for a quadword operation.

Dd and Dm specify the destination and operand vectors for a doubleword operation.

fbits specifies the number of fraction bits in the fixed point number, in the range 0-32. If fbits is not present, the conversion is between floating-point and integer.

B.1.2 VCVT (between half-precision and single-precision floating-point)

VCVT (Vector Convert), with half-precision extension, converts each element in a vector and places the results in the destination vector. The conversion can be:

- From half-precision floating-point to single-precision floating-point (F32.F16).
- Single-precision floating-point to half-precision floating-point (F16.F32).

This instruction is present only in NEON systems with the half-precision extension.

Syntax

```
VCVT{cond}.F32.F16 Qd, Dm
VCVT{cond}.F16.F32 Dd, Qm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Qd and Dm specify the destination vector for the single-precision results and the half-precision operand vector.

Dd and Qm specify the destination vector for half-precision results and the single-precision operand vector.

B.1.3 VDUP

VDUP (Vector Duplicate) duplicates a scalar into every element of the destination vector. The source can be either a NEON scalar or an ARM register.

Syntax

```
VDUP{cond}.size Qd, Dm[x]
VDUP{cond}.size Dd, Dm[x]
VDUP{cond}.size Qd, Rm
VDUP{cond}.size Dd, Rm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

size is 8, 16, or 32.

Qd specifies the destination register for a quadword operation.

Dd specifies the destination register for a doubleword operation.

Dm[x] specifies the NEON scalar or Rm the ARM register.

B.1.4 VEXT

VEXT (Vector Extract) extracts 8-bit elements from the bottom end of the second operand vector and the top end of the first, concatenates them, and stores the result in the destination vector.

Syntax

```
VEXT{cond}.8 {Qd,} Qn, Qm, #imm
VEXT{cond}.8 {Dd,} Dn, Dm, #imm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#)) An ARM VEXT instruction must be unconditional.

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation.

imm gives the number of 8-bit elements to extract from the bottom of the second operand vector, between 0-7 for doubleword operations, or 0-15 for quadword operations.

VEXT can also be written as a pseudo-instruction. In this case, a datatype of 16, 32, or 64 can be used and #imm refers to halfwords, words, or doublewords, with a corresponding reduction in the permitted range.

If the value is 16, 32, or 64, the syntax is a pseudo-instruction for a VEXT instruction specifying the equivalent number of bytes. The assembler treats values greater than 8 in the following way:

VEXT.16 D0, D1, #x is treated as VEXT.8 D0, D1, #(x*2)

VEXT.32 D0, D1, #x is treated as VEXT.8 D0, D1, #(x*4)

VEXT.64 Q0, Q1, #x is treated as VEXT.8 Q0, Q1, #(x*8).

B.1.5 VMOV

VMOV (Vector Bitwise Move) (immediate), places an immediate value into every element of the destination register.

Syntax

```
VMOV{cond}.datatype Qd, #imm
VMOV{cond}.datatype Dd, #im
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is I8, I16, I32, I64, or F32.

Qd or Dd specify the NEON register for the result.

imm is an immediate value of the type specified by datatype, which is replicated to fill the destination register.

B.1.6 VMVN

VMVN (Vector Bitwise NOT) (immediate), places the bitwise inverse of an immediate integer value into every element of the destination register.

Syntax

```
VMVN{cond}.datatype Qd, #imm
VMVN{cond}.datatype Dd, #imm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is one of I8, I16, I32, I64, or F32.

Qd or Dd specify the NEON register for the result.

imm is an immediate value of the type specified by datatype, which is replicated to fill the destination register.

B.1.7 VMOVL, V{Q}MOVN, VQMOVUN

VMOVL (Vector Move Long) takes each element in a doubleword vector and sign or zero-extends them to twice their original length. The results are stored in a quadword vector.

VMOVN (Vector Move and Narrow) copies the least significant half of each element of a quadword vector into the corresponding element of a doubleword vector.

VQMOVN (Vector Saturating Move and Narrow) copies each element of the operand vector to the corresponding element of the destination. The result element is half the width of the operand element, and values are saturated to the result width.

VQMOVUN (Vector Saturating Move and Narrow, signed operand with Unsigned result) copies each element of the operand vector to the corresponding element of the destination. The result element is half the width of the operand element and values are saturated to the result width.

Syntax

```
VMOVL{cond}.datatype Qd, Dm
V{Q}MOVN{cond}.datatype Dd, Qm
VQMOVUN{cond}.datatype Dd, Qm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Q specifies that the results are saturated.

datatype is one of:

- S8, S16, S32 for VMOVL.
- U8, U16, U32 for VMOVL.
- I16, I32, I64 for VMOVN.
- F32.U32 unsigned integer or fixed-point to floating-point.
- U16, U32, U64 for VQMOVN.

Qd and Dm specify the destination vector and the operand vector for VMOVL.

Dd and Qm specify the destination vector and the operand vector for V{Q}MOV{U}N.

B.1.8 VREV

VREV16 (Vector Reverse halfwords) reverses the order of 8-bit elements within each halfword of the vector and stores the result in the corresponding destination vector.

VREV32 (Vector Reverse words) reverses the order of 8-bit or 16-bit elements within each word of the vector, and stores the result in the corresponding destination vector.

VREV64 (Vector Reverse doublewords) reverses the order of 8-bit, 16-bit, or 32-bit elements within each doubleword of the vector, and stores the result in the destination vector.

Syntax

```
VREVn{cond}.size Qd, Qm
VREVn{cond}.size Dd, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

n is one of 16, 32, or 64.

size is one of 8, 16, or 32, and must be less than n.

Qd and Qm specify the destination and operand registers for a quadword operation.

Dd and Dm specify the destination and operand registers for a doubleword operation.

B.1.9 VSWP

VSWP (Vector Swap) exchanges the contents of two vectors, which can be either doubleword or quadword.

Syntax

```
VSWP{cond}{.datatype} Qd, Qm
VSWP{cond}{.datatype} Dd, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Qd and Qm specify the vectors for a quadword operation.

Dd and Dm specify the vectors for a doubleword operation.

B.1.10 VTBL

VTBL (Vector Table Lookup) uses byte indexes in a control vector to look up byte values in a table and generate a new vector. Indexes which are out of range return 0.

Syntax

```
VTBL{cond}.8 Dd, list, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Dd specifies the destination vector.

Dm specifies the index vector.

list specifies the vectors containing the table. It is one of:

- {Dn}
- {Dn, D(n+1)}
- {Dn, D(n+1), D(n+2)}
- {Dn, D(n+1), D(n+2), D(n+3)}
- {Qn, Q(n+1)}

All of the registers in list must be in the range D0 – D31 or Q0 – Q15 and are not permitted to wrap around the end of the register bank.

B.1.11 VTBX

VTBX (Vector Table Extension) uses byte indexes in a control vector to look up byte values in a table and generate a new vector, but leaves the destination element unchanged when an out of range index is used.

Syntax

```
VTBX{cond}.8 Dd, list, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Dd specifies the destination vector.

Dm specifies the index vector.

list specifies the vectors containing the table. It is one of:

- {Dn}

- {Dn,D(n+1)}
- {Dn,D(n+1),D(n+2)}
- {Dn,D(n+1),D(n+2),D(n+3)}
- {Qn,Q(n+1)}

All of the registers in list must be in the range D0 – D31 or Q0 – Q15 and are not permitted to wrap around the end of the register bank.

B.1.12 VTRN

VTRN (Vector Transpose) treats the elements of its operand vectors as elements of 2×2 matrices, and transposes them.

Syntax

```
VTRN{cond}.size Qd, Qm
VTRN{cond}.size Dd, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

size is one of 8, 16, or 32.

Qd and Qm specify the vectors for a quadword operation.

Dd and Dm specify the vectors, for a doubleword operation.

B.1.13 VUZP

VUZP (Vector Unzip) de-interleaves the elements of two vectors.

Syntax

```
VUZP{cond}.size Qd, Qm
VUZP{cond}.size Dd, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

size is one of 8, 16, or 32.

Qd and Qm specify the vectors for a quadword operation.

Dd and Dm specify the vectors, for a doubleword operation.

B.1.14 VZIP

VZIP (Vector Zip) interleaves the elements of two vectors.

Syntax

```
VZIP{cond}.size Qd, Qm
VZIP{cond}.size Dd, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

size is one of 8, 16, or 32.

Qd and Qm specify the vectors for a quadword operation.

Dd and Dm specify the vectors, for a doubleword operation.

B.2 NEON shift instructions

This section covers NEON instructions which perform logical shift or insert operations.

B.2.1 VSHL, VQSHL, VQSHLU, and VSHLL (by immediate)

Vector Shift Left (by immediate) instructions take each element in a vector of integers, left shift them by an immediate value, and write the result to the destination vector. The instruction variants are:

- VSHL (Vector Shift Left) discards bits shifted out of the left of each element.
- VQSHL (Vector Saturating Shift Left) sets the sticky QC flag (FPSCR bit[27]) if saturation occurs.
- VQSHLU (Vector Saturating Shift Left Unsigned) sets the sticky QC flag (FPSCR bit[27]) if saturation occurs.
- VSHLL (Vector Shift Left Long) takes each element in a doubleword vector, left shifts them by an immediate value, and places the results in a quadword vector.

Syntax

```
V{Q}SHL{U}{cond}.datatype {Qd,} Qm, #imm
V{Q}SHL{U}{cond}.datatype {Dd,} Dm, #imm
VSHLL{cond}.datatype Qd, Dm, #imm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Q indicates that results are saturated if they overflow.

U is only permitted if Q is also specified and indicates that the results are unsigned, even though the operands are signed.

datatype is one of:

- I8, I16, I32, I64 for VSHL.
- S8, S16, S32 for VSHLL, VQSHL, or VQSHLU.
- U8, U16, U32 for VSHLL or VQSHL.
- S64 for VQSHL or VQSHLU.
- U64 for VQSHL.

Qd and Qm are the destination and operand vectors for a quadword operation.

Dd and Dm are the destination and operand vectors for a doubleword operation.

Qd and Dm are the destination and operand vectors, for a long operation.

imm is the immediate value that sets the size of the shift and must be in the range:

- 1 to size(datatype) for VSHLL.
- 1 to (size(datatype) – 1) for VSHL, VQSHL, or VQSHLU.

A shift of 0 is permitted, but the code disassembles to VMOV or VMOVL.

B.2.2 V{Q}{R}SHL

VSHL (Vector Shift Left) (by signed variable) shifts each element in a vector by a value from the least significant byte of the corresponding element of a second vector, and stores the results in the destination vector. If the shift value is positive, the operation is a left shift. Otherwise, it is a right shift. The results can be optionally saturated, rounded, or both. The sticky QC flag is set if saturation occurs.

Syntax

```
V{Q}{R}SHL{cond}.datatype {Qd,} Qm, Qn.  
V{Q}{R}SHL{cond}.datatype {Dd,} Dm, Dn.
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Q indicates that results are saturated if they overflow.

R indicates that each result is rounded. If R is not specified, each result is truncated.

datatype is one of S8, S16, S32, S64, U8, U16, U32, or U64.

Qd, Qm, and Qn specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dm, and Dn specify the destination, first operand and second operand registers for a doubleword operation.

B.2.3 V{R}SHR{N}, V{R}SRA

V{R}SHR{N} (Vector Shift Right) (by immediate value) right shifts each element in a vector by an immediate value and stores the results in the destination vector. The results can be optionally rounded, or narrowed, or both.

V{R}SRA (Vector Shift Right (by immediate value) and Accumulate) right shifts each element in a vector by an immediate value, and accumulates the results into the destination vector. The results can be optionally rounded.

Syntax

```
V{R}SHR{cond}.datatype {Qd,} Qm, #imm  
V{R}SHR{cond}.datatype {Dd,} Dm, #imm  
V{R}SRA{cond}.datatype {Qd,} Qm, #imm  
V{R}SRA{cond}.datatype {Dd,} Dm, #imm  
V{R}SHRN{cond}.datatype Dd, Qm, #imm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

R indicates that the results are rounded. If R is not present, the results are truncated.

datatype is one of:

- S8, S16, S32, S64 for V{R}SHR or V{R}SRA.
- U8, U16, U32, U64 for V{R}SHR or V{R}SRA.
- I16, I32, I64 for V{R}SHRN.

Qd and Qm are the destination and operand vectors, for a quadword operation.

Dd and Dm are the destination and operand vectors for a doubleword operation.

Dd and Qm are the destination vector and the operand vector, for a narrow operation.

imm is the immediate value specifying the size of the shift, in the range 0 to $(\text{size}(\text{datatype}) - 1)$.

B.2.4 VQ{R}SHR{U}N

VQ{R}SHR{U}N (Vector Saturating Shift Right, Narrow) (by immediate value) with optional Rounding) right shifts each element in a quadword vector of integers by an immediate value, and stores the results in a doubleword vector. The sticky QC flag (FPSCR bit[27]) is set if saturation occurs.

Syntax

VQ{R}SHR{U}N{cond}.datatype Dd , Qm , #imm

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

R indicates that the results are rounded. If R is not present, the results are truncated.

U indicates that the results are unsigned, although the operands are signed.

datatype is one of:

- S16, S32, S64 for VQ{R}SHRN or VQ{R}SHRUN.
- U16, U32, U64 for VQ{R}SHRN only.

Dd and Qm are the destination vector and the operand vector.

imm is the immediate value specifying the size of the shift, in the range 0 to $(\text{size}(\text{datatype}) - 1)$.

B.2.5 VSLI

VSLI (Vector Shift Left and Insert) left shifts each element in a vector by an immediate value, and inserts the results in the destination vector. Bits shifted out of the left of each element are lost.

Syntax

VSLI{cond}.size { Qd ,} Qm , #imm
VSLI{cond}.size { Dd ,} Dm , #imm

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

size is one of 8, 16, 32, or 64.

Qd and Qm are the destination and operand vectors for a quadword operation.

Dd and Dm are the destination and operand vectors for a doubleword operation.

Dd and Qm are the destination vector and the operand vector.

imm is the immediate value that sets the size of the shift, in the range 0 to $(\text{size} - 1)$

B.2.6 VSRI

VSRI (Vector Shift Right and Insert) right shifts each element in a vector by an immediate value, and inserts the results in the destination vector. Bits shifted out of the right of each element are lost.

Syntax

```
VSRI{cond}.size {Qd,} Qm, #imm  
VSRI{cond}.size {Dd,} Dm, #imm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

size is one of 8, 16, 32, or 64.

Qd and Qm are the destination and operand vectors for a quadword operation.

Dd and Dm are the destination and operand vectors for a doubleword operation.

imm is the immediate value that sets the size of the shift, in the range 1 to size.

B.3 NEON logical and compare operations

This section covers NEON instructions which perform bitwise boolean operations and comparisons.

B.3.1 VACGE and VACGT

Vector Absolute Compare takes the absolute value of each element in a vector, and compares it with the absolute value of the corresponding element in a second vector. If the condition is true, the corresponding element in the destination vector is set to all ones. If the condition is not true, it is set to all zeros.

Syntax

```
VACop{cond}.F32 {Qd,} Qn, Qm
VACop{cond}.F32 {Dd,} Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

op is either GE (Absolute Greater than or Equal) or GT (Absolute Greater Than).

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation. The result datatype is I32.

B.3.2 VAND

VAND performs a bitwise logical AND operation between values in two registers, and places the results in the destination register.

Syntax

```
VAND{cond}{.datatype} {Qd,} Qn, Qm
VAND{cond}{.datatype} {Dd,} Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is an optional data type, which is ignored by the assembler.

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation.

B.3.3 VBIC (immediate)

VBIC (Vector Bitwise Clear) (immediate) takes each element of the destination vector, does a bitwise AND Complement with an immediate value, and stores the result in the destination vector.

Syntax

```
VBIC{cond}.datatype Qd, #imm
VBIC{cond}.datatype Dd, #imm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is one of I8, I16, I32, or I64.

Qd or Dd is the NEON register for the source and result.

imm is the immediate value.

Immediate values can be specified as a pattern which the assembler repeats to fill the destination register. Alternatively, you can specify the entire value, provided it meets certain requirements. The permitted patterns for immediate value are 0x00XY or 0xXY00 for I16, or 0x000000XY, 0x0000XY00, 0x00XY0000 or 0xXY000000 for I32. If you choose I8 or I64 datatypes, the assembler produces the I16 or I32 equivalent.

B.3.4 VBIC (register)

VBIC (Vector Bitwise Clear) performs a bitwise logical AND complement operation between values in two registers, and places the results in the destination register.

Syntax

```
VBIC{cond}{.datatype} {Qd,} Qn, Qm
VBIC{cond}{.datatype} {Dd,} Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is an optional data type, which is ignored by the assembler.

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation.

B.3.5 VBIF

VBIF (Vector Bitwise Insert if False) inserts each bit from the first operand into the destination, if the corresponding bit of the second operand is 0, otherwise it does not change the destination bit.

Syntax

```
VBIF{cond}{.datatype} {Qd,} Qn, Qm
VBIF{cond}{.datatype} {Dd,} Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is an optional data type, which is ignored by the assembler.

Qd , Qn , and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd , Dn , and Dm specify the destination, first operand and second operand registers for a doubleword operation.

B.3.6 VBIT

VBIT (Vector Bitwise Insert if True) inserts each bit from the first operand into the destination if the corresponding bit of the second operand is 1, otherwise it does not change the destination bit.

Syntax

```
VBIT{cond}{.datatype} {Qd}, Qn, Qm
VBIT{cond}{.datatype} {Dd}, Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is an optional data type, which is ignored by the assembler.

Qd , Qn , and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd , Dn , and Dm specify the destination, first operand and second operand registers for a doubleword operation.

B.3.7 VBSL

VBSL (Vector Bitwise Select) selects each destination bit from the first operand if the corresponding destination bit was 1 or from the second operand if the corresponding destination bit was 0.

Syntax

```
VBSL{cond}{.datatype} {Qd}, Qn, Qm
VBSL{cond}{.datatype} {Dd}, Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is an optional data type, which is ignored by the assembler.

Qd , Qn , and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd , Dn , and Dm specify the destination, first operand and second operand registers for a doubleword operation.

B.3.8 VCEQ, VCGE, VCGT, VCLE, and VCLT

Vector Compare takes each element in a vector and compares it with the corresponding element of a second vector (or zero). If the condition is true, the corresponding element in the destination vector is set to all ones. Otherwise, it is set to all zeros.

Syntax

```
VCop{cond}.datatype {Qd}, Qn, Qm
```

VCop{cond}.datatype {Dd,} Dn, Dm
 VCop{cond}.datatype {Qd,} Qn, #0
 VCop{cond}.datatype {Dd,} Dn, #0

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

op is one of:

- EQ – Equal.
- GE – Greater than or Equal.
- GT – Greater Than.
- LE – Less than or Equal (only if the second operand is #0).
- LT – Less Than (only if the second operand is #0).

datatype is one of:

- I8, I16, I32, or F32 for VCEQ.
- S8, S16, S32, U8, U16, U32, or F32 for VCGE, VCGT, VCLE, or VCLT (except #0 form).
- S8, S16, S32, or F32 for VCGE, VCGT, VCLE, or VCLT (#0 form).

The result datatype is:

- I32 for operand datatypes I32, S32, U32, or F32.
- I16 for operand datatypes I16, S16, or U16.
- I8 for operand datatypes I8, S8, or U8.

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation. #0 replaces Qm or Dm for comparisons with zero.

B.3.9 VEOR

VEOR performs a bitwise logical exclusive OR operation between values in two registers, and places the results in the destination register.

Syntax

VEOR{cond}{.datatype} {Qd,} Qn, Qm
 VEOR{cond}{.datatype} {Dd,} Dn, Dm

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation.

B.3.10 VMOV

VMOV, Vector Move (register) copies a value from a source register to a destination register.

Syntax

VMOV{cond}{.datatype} Qd, Qm

`VMOV{cond}{.datatype} Dd, Dm`

where:

`cond` is an optional conditional code. (See [Conditional execution on page 6-3](#))

`Qd` and `Qm` specify the destination vector and the source vector, for a quadword operation.

`Dd` and `Dm` specify the destination and source vector, for a doubleword operation.

B.3.11 VMVN

VMVN, Vector Bitwise NOT (register) inverts the value of each bit from the source register and stores the result to the destination register.

Syntax

`VMVN{cond}{.datatype} Qd, Qm`
`VMVN{cond}{.datatype} Dd, Dm`

where:

`cond` is an optional conditional code. (See [Conditional execution on page 6-3](#))

`Qd` and `Qm` specify the destination vector and the source vector, for a quadword operation.

`Dd` and `Dm` specify the destination and source vector, for a doubleword operation.

B.3.12 VORN

VORN performs a bitwise logical OR NOT operation between values in two registers, and places the results in the destination register.

Syntax

`VORN{cond}{.datatype} {Qd,} Qn, Qm`
`VORN{cond}{.datatype} {Dd,} Dn, Dm`

where:

`cond` is an optional conditional code. (See [Conditional execution on page 6-3](#))

`Qd`, `Qn`, and `Qm` specify the destination, first operand and second operand registers for a quadword operation.

`Dd`, `Dn`, and `Dm` specify the destination, first operand and second operand registers for a doubleword operation.

B.3.13 VORR (immediate)

VORR (Bitwise OR) (immediate) takes each element of the destination vector, does a bitwise OR with an immediate value, and stores the result into the destination vector.

Syntax

`VORR{cond}.datatype Qd, #imm`
`VORR{cond}.datatype Dd, #imm`

where:

`cond` is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is one of I8, I16, I32, or I64.

Qd or Dd is the NEON register for the source and result.

imm is the immediate value.

Immediate values can be specified as a pattern which the assembler repeats to fill the destination register. Alternatively, you can specify the entire value, provided it meets certain requirements. The permitted patterns for immediate value are 0x00XY or 0xXY00 for I16, or 0x000000XY, 0x0000XY00, 0x00XY0000 or 0xXY000000 for I32. If you choose I8 or I64 datatypes, the assembler produces the I16 or I32 equivalent.

B.3.14 VORR (register)

VORR (Bitwise OR) (register) performs a bitwise logical OR operation between values in two registers, and places the results in the destination register.

Syntax

```
VORR{cond}{.datatype} {Qd,} Qn, Qm
VORR{cond}{.datatype} {Dd,} Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation.

B.3.15 VTST

VTST (Vector Test Bits) takes each element in a vector and does a bitwise AND with the corresponding element in a second vector. If the result is zero, the corresponding element in the destination vector is set to all zeros. Otherwise, it is set to all ones.

Syntax

```
VTST{cond}.size {Qd,} Qn, Qm
VTST{cond}.size {Dd,} Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

size is one of 8, 16, or 32.

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation.

B.4 NEON arithmetic instructions

This section describe NEON instructions which perform arithmetic operations such as addition, subtraction or reciprocal calculation.

B.4.1 VABA{L}

VABA (Vector Absolute Difference and Accumulate) subtracts the elements of one vector from the corresponding elements of another and accumulates the absolute values of the results into the elements of the destination vector. A long version of this instructions is available.

Syntax

```
VABA{cond}.datatype {Qd,} Qn, Qm
VABA{cond}.datatype {Dd,} Dn, Dm
VABAL{cond}.datatype Qd, Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is one of S8, S16, S32, U8, U16, or U32.

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation.

Qd, Dn and Dm specify the destination, first operand and second operand vectors for a long operation.

B.4.2 VABD{L}

VABD (Vector Absolute Difference) subtracts the elements of one vector from the corresponding elements of another and places the absolute values of the results into the elements of the destination vector. A long version of the instruction is available.

Syntax

```
VABD{cond}.datatype {Qd,} Qn, Qm
VABD{cond}.datatype {Dd,} Dn, Dm
VABDL{cond}.datatype Qd, Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is one of:

- S8, S16, S32, U8, U16, U32 or F32 for VABD.
- S8, S16, S32, U8, U16, or U32 for VABDL.

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation.

Qd, Dn, and Dm specify the destination, first operand and second operand vectors for a long operation.

B.4.3 V{Q}ABS

VABS (Vector Absolute) takes the absolute value of each element in a vector, and stores the results in the destination. The floating-point version only clears the sign bit. A saturating version of the instructions is available. The sticky QC flag is set if saturation occurs.

Syntax

```
V{Q}ABS{cond}.datatype Qd, Qm
V{Q}ABS{cond}.datatype Dd, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Q indicates that saturation is performed if any of the results overflow.

datatype is one of:

- S8, S16, S32 for VABS or VQABS
- F32 for VABS only.

Qd and Qm specify the destination and operand vectors for a quadword operation.

Dd and Dm specify the destination and operand vectors for a doubleword operation.

B.4.4 V{Q}ADD, VADDL, VADDW

VADD (Vector Add) adds corresponding elements in two vectors, and stores the results in the destination vector. VADD has long, wide and saturating variants.

Syntax

```
V{Q}ADD{cond}.datatype {Qd,} Qn, Qm ;
V{Q}ADD{cond}.datatype {Dd,} Dn, Dm ;
VADDL{cond}.datatype Qd, Dn, Dm ;
VADDW{cond}.datatype {Qd,} Qn, Dm ;
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Q indicates that saturation is performed if any of the results overflow.

datatype is one of:

- I8, I16, I32, I64, F32 for VADD
- S8, S16, S32 for VQADD, VADDL or VADDW
- U8, U16, U32 for VQADD, VADDL or VADDW
- S64, U64 for VQADD.

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation.

Qd, Dn, and Dm specify the destination, first operand and second operand vectors for a long operation.

Qd, Qn and Dm specify the destination, first operand and second operand vectors for a wideoperation.

B.4.5 V{R}ADDHN

V{R}ADDHN (Vector Add and Narrow, selecting High half) adds corresponding elements in two vectors, selects the most significant halves of the results, and stores them in the destination vector.

Syntax

V{R}ADDHN{cond}.datatype Dd, Qn, Qm

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

R specifies that each result is rounded. If R is not specified, each result is truncated.

datatype is one of I16, I32, or I64.

Dd, Qn, and Qm are the destination vector, the first operand vector, and the second operand vector.

B.4.6 VCLS

VCLS (Vector Count Leading Sign Bits) counts the number of consecutive bits following the topmost bit that are the same as that bit, in each element in a vector, and stores the results in a destination vector.

Syntax

VCLS{cond}.datatype Qd, Qm
VCLS{cond}.datatype Dd, Dm

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is one of S8, S16, or S32.

Qd and Qm specify the destination and operand vectors for a quadword operation.

Dd and Dm specify the destination and operand vectors for a doubleword operation.

B.4.7 VCLZ

VCLZ (Vector Count Leading Zeros) counts the number of consecutive zeros, starting from the top bit, in each element in a vector, and stores the results in a destination vector.

Syntax

VCLZ{cond}.datatype Qd, Qm
VCLZ{cond}.datatype Dd, Dm

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is one of S8, S16, or S32.

Qd and Qm specify the destination and operand vectors for a quadword operation.

Dd and Dm specify the destination and operand vectors for a doubleword operation.

B.4.8 VCNT

VCNT (Vector Count Set Bits) counts the number of bits that are one in each element in a vector, and stores the results in a destination vector.

Syntax

```
VCNT{cond}.datatype Qd, Qm
VCNT{cond}.datatype Dd, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype must be I8.

Qd and Qm specify the destination and operand vectors for a quadword operation.

Dd and Dm specify the destination and operand vectors for a doubleword operation.

B.4.9 V{R}HADD

VHADD (Vector Halving Add) adds corresponding elements in two vectors, shifts each result right one bit and stores the results in the destination vector. Results can be either rounded or truncated.

Syntax

```
V{R}HADD{cond}.datatype {Qd,} Qn, Qm
V{R}HADD{cond}.datatype {Dd,} Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

R specifies that each result is rounded. If R is not specified, each result is truncated.

datatype is one of S8, S16, S32, U8, U16, or U32.

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation.

B.4.10 VHSUB

VHSUB (Vector Halving Subtract) subtracts the elements of one vector from the corresponding elements of another vector, shifts each result right one bit, and stores the result in the destination vector. Results are always truncated.

Syntax

```
VHSUB{cond}.datatype {Qd,} Qn, Qm
VHSUB{cond}.datatype {Dd,} Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is one of S8, S16, S32, U8, U16, or U32.

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation.

B.4.11 VMAX and VMIN

VMAX (Vector Maximum) compares corresponding elements in two vectors, and writes the larger of them into the corresponding element in the destination vector.

VMIN (Vector Minimum) compares corresponding elements in two vectors, and writes the smaller value into the corresponding element in the destination vector.

Syntax

```
VMAX{cond}.datatype Qd, Qn, Qm
VMAX{cond}.datatype Dd, Dn, Dm
VMIN{cond}.datatype Qd, Qn, Qm
VMIN{cond}.datatype Dd, Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is one of S8, S16, S32, U8, U16, U32, or F32.

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation.

B.4.12 V{Q}NEG

VNEG (Vector Negate) negates each element in a vector, and places the results in a second vector. The floating-point version only inverts the sign bit. A saturating version of the instruction is available. The sticky QC flag is set if saturation occurs.

Syntax

```
V{Q}NEG{cond}.datatype Qd, Qm
V{Q}NEG{cond}.datatype Dd, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Q indicates that saturation is performed if any of the results overflow.

datatype is one of:

- S8, S16, S32 for VNEG, or VQNEG
- F32 or F64 for VNEG.

Qd and Qm specify the destination and operand vectors for a quadword operation.

Dd and Dm specify the destination and operand vectors for a doubleword operation.

B.4.13 VPADD{L}, VPADAL

VPADD (Vector Pairwise Add) adds adjacent pairs of elements of two vectors, and stores the results in the destination vector.

VPADDL (Vector Pairwise Add Long) adds adjacent pairs of elements of a vector, sign or zero extends the results to twice their original width and stores the results in the destination vector.

VPADAL (Vector Pairwise Add and Accumulate Long) adds adjacent pairs of elements of a vector, and accumulates the results into the elements of the destination vector.

Syntax

```
VPADD{cond}.datatype {Dd,} Dn, Dm
VPopL{cond}.datatype Qd, Qm
VPopL{cond}.datatype Dd, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

op is either ADD or ADA.

datatype is one of:

- I8, I16, I32, F32 for VPADD.
- S8, S16, S32 for VPADDL or VPADAL.
- U8, U16, U32 for VPADDL or VPADAL.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a VPADD.

Qd and Qm specify the destination and operand vectors for a quadword operation.

Dd and Dm specify the destination and operand vectors for a doubleword operation.

B.4.14 VPMAX and VPMIN

VPMAX (Vector Pairwise Maximum) compares adjacent pairs of elements in two vectors and writes the larger of each pair into the corresponding element in the destination vector.

VPMIN (Vector Pairwise Minimum) compares adjacent pairs of elements in two vectors, and writes the smaller of each pair into the corresponding element in the destination vector. Operands and results must be doubleword vectors for the latter two pairwise operations.

Syntax

```
VPMAX{cond}.datatype Dd, Dn, Dm
VPMIN{cond}.datatype Dd, Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is one of S8, S16, S32, U8, U16, U32, or F32.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation.

B.4.15 VRECPE

VRECPE (Vector Reciprocal Estimate) finds an approximate reciprocal of each element in a vector, and stores the results in a destination vector.

Syntax

```
VRECPE{cond}.datatype Qd, Qm
VRECPE{cond}.datatype Dd, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is either U32 or F32.

Qd and Qm specify the destination and operand vectors for a quadword operation.

Dd and Dm specify the destination and operand vectors for a doubleword operation.

B.4.16 VRECPS

VRECPS (Vector Reciprocal Step) multiplies the elements of one vector by the corresponding elements of another, subtracts each of the results from 2.0, and stores the final results into the elements of the destination. This instruction is used as part of the Newton-Raphson iteration algorithm.

Syntax

```
VRECPS{cond}.F32 {Qd,} Qn, Qm
VRECPS{cond}.F32 {Dd,} Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is either U32 or F32.

Qd and Qm specify the destination and operand vectors for a quadword operation.

Dd and Dm specify the destination and operand vectors for a doubleword operation.

B.4.17 VRSQRTE

VRSQRTE (Vector Reciprocal Square Root Estimate) finds an approximate reciprocal square root of each element in a vector and stores the results in a destination vector.

Syntax

```
VRSQRTE{cond}.datatype Qd, Qm
VRSQRTE{cond}.datatype Dd, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is either U32 or F32.

Qd and Qm specify the destination and operand vectors for a quadword operation.

Dd and Dm specify the destination and operand vectors for a doubleword operation.

B.4.18 VRSQRTS

VRSQRTS (Vector Reciprocal Square Root Step) multiplies the elements of one vector by the corresponding elements of another, subtracts each of the results from 3.0, divides these results by 2.0, and places the final results into the elements of the destination. This instruction is used as part of the Newton-Raphson iteration algorithm.

Syntax

```
VRSQRTS{cond}.F32 {Qd,} Qn, Qm
VRSQRTS{cond}.F32 {Dd,} Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Qd and Qm specify the destination and operand vectors for a quadword operation.

Dd and Dm specify the destination and operand vectors for a doubleword operation.

B.4.19 V{Q}SUB, VSUBL and VSUBW

VSUB (Vector Subtract) subtracts the elements of one vector from the corresponding elements of another vector and stores the results in the destination vector. VSUB has Long, Wide and Saturating variants.

Syntax

```
V{Q}SUB{cond}.datatype {Qd,} Qn, Qm
V{Q}SUB{cond}.datatype {Dd,} Dn, Dm
VSUBL{cond}.datatype Qd, Dn, Dm
VSUBW{cond}.datatype {Qd,} Qn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Q specifies that saturation should be performed if any of the results overflow.

datatype is one of:

- I8, I16, I32, I64, F32 for VSUB
- S8, S16, S32 for VQSUB, VSUBL, or VSUBW
- U8, U16, U32 for VQSUB, VSUBL, or VSUBW
- S64, U64 for VQSUB.

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation.

Qd, Dn, and Dm specify the destination, first operand and second operand vectors for a long operation.

Qd, Qn and Dm specify the destination, first operand and second operand vectors for a wideoperation.

B.4.20 V{R}SUBHN

V{R}SUBHN (Vector Subtract and Narrow, selecting High Half) subtracts the elements of one vector from the corresponding elements of another vector, selects the most significant halves of the results and stores them in the destination vector. Results can be either rounded or truncated for both operations.

Syntax

V{R}SUBHN{cond}.datatype Dd, Qn, Qm

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

R specifies that each result is rounded. If R is not specified, each result is truncated.

datatype is one of I16, I32, or I64.

Dd, Qn, and Qm are the destination vector, the first operand vector, and the second operand vector.

B.5 NEON multiply instructions

This section describe NEON instructions which perform multiplication or multiply-accumulate.

B.5.1 VFMA, VFMS

VFMA (Vector Fused Multiply Accumulate) multiplies corresponding elements in two vectors, and accumulates the results into the destination vector.

VFMS (Vector Fused Multiply Subtract) multiplies corresponding elements in two operand vectors, subtracts the products from the corresponding elements of the destination vector, and stores the final results in the destination vector. The result of the multiply is not rounded before performing the accumulate or subtract operation.

Syntax

```
Vop{cond}.F32 {Qd,} Qn, Qm
Vop{cond}.F32 {Dd,} Dn, Dm
Vop{cond}.F64 {Dd,} Dn, Dm
Vop{cond}.F32 {Sd,} Sn, Sm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

op is one of FMA or FMS.

Sd, Sn, and Sm are the destination and operand vectors for word operation.

Dd, Dn, and Dm are the destination and operand vectors for doubleword operation.

Qd, Qn, and Qm are the destination and operand vectors for quadword operation.

B.5.2 VMUL{L}, VMLA{L}, and VMLS{L}

VMUL (Vector Multiply) multiplies corresponding elements in two vectors and stores the results in the destination vector.

VMLA (Vector Multiply Accumulate) multiplies corresponding elements in two vectors and adds the results to the corresponding element of the destination vector.

VMLS (Vector Multiply Subtract) multiplies elements in two vectors, subtracts the results from corresponding elements of the destination vector, and stores the results in the destination vector.

Syntax

```
Vop{cond}.datatype {Qd,} Qn, Qm
Vop{cond}.datatype {Dd,} Dn, Dm
VopL{cond}.datatype Qd, Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

op is one of:

- MUL, Multiply
- MLA, Multiply Accumulate
- MLS, Multiply Subtract.

datatype is one of:

- I8, I16, I32, F32 for VMUL, VMLA, or VMLS
- S8, S16, S32 for VMULL, VMLAL, or VMLSL
- U8, U16, U32 for VMULL, VMLAL, or VMLSL
- P8 for VMUL or VMULL.

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation.

Qd, Dn, and Dm specify the destination, first operand and second operand vectors for a long operation.

B.5.3 VMUL{L}, VMLA{L}, and VMLS{L} (by scalar)

VMUL (Vector Multiply by scalar) multiplies each element in a vector by a scalar and stores the results in the destination vector.

VMLA (Vector Multiply Accumulate) multiplies each element in a vector by a scalar and accumulates the results into the corresponding elements of the destination vector.

VMLS (Vector Multiply Subtract) multiplies each element in a vector by a scalar and subtracts the results from the corresponding elements of the destination vector and stores the final results in the destination vector.

Syntax

```
Vop{cond}.datatype {Qd,} Qn, Dm[x]
Vop{cond}.datatype {Dd,} Dn, Dm[x]
VopL{cond}.datatype Qd, Dn, Dm[x]
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

op is one of:

- MUL – Multiply.
- MLA – Multiply Accumulate.
- MLS – Multiply Subtract.

datatype is one of:

- I16, I32, F32 for VMUL, VMLA, or VMLS.
- S16, S32 for VMULL, VMLAL, or VMLSL.
- U16, U32 for VMULL, VMLAL, or VMLSL.

Qd, Qn, and Qm specify the destination, first operand and second operand registers for a quadword operation.

Dd, Dn, and Dm specify the destination, first operand and second operand registers for a doubleword operation.

Qd, Dn, and Dm specify the destination, first operand and second operand vectors for a long operation.

Dm[x] is the scalar holding the second operand.

B.5.4 VQ{R}DMULH (by vector or by scalar)

Vector Saturating Doubling Multiply Returning High Half instructions multiply their operands and double the result. They return only the high half of the results. If any of the results overflow, they are saturated and the sticky QC flag is set.

Syntax

```
VQ{R}DMULH{cond}.datatype {Qd,} Qn, Qm
VQ{R}DMULH{cond}.datatype {Dd,} Dn, Dm
VQ{R}DMULH{cond}.datatype {Qd,} Qn, Dm[x]
VQ{R}DMULH{cond}.datatype {Dd,} Dn, Dm[x]
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

R specifies that each result is rounded. If R is not specified, each result is truncated.

datatype is either S16 or S32.

Qd and Qn are the destination and first operand vector, for a quadword operation.

Dd and Dn are the destination and first operand vector for a doubleword operation.

Qm or Dm specifies the vector holding the second operand, for a by vector operation.

Dm[x] is the scalar holding the second operand for a by scalar operation.

B.5.5 VQDMULL, VQDMLAL, and VQDMLSL (by vector or by scalar)

Vector Saturating Doubling Multiply Long instructions multiply their operands and double the results. The instruction variants are:

- VQDMULL stores the results in the destination register.
- VQDMLAL adds the results to the values in the destination register.
- VQDMLSL subtracts the results from the values in the destination register.

If any of the results overflow, they are saturated and the sticky QC flag is set.

Syntax

```
VQDopL{cond}.datatype Qd, Dn, Dm
VQDopL{cond}.datatype Qd, Dn, Dm[x]
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

op is one of:

- MUL, Multiply.
- MLA, Multiply Accumulate.
- MLS, Multiply Subtract.

datatype is either S16 or S32.

Qd and Dn are the destination vector and the first operand vector.

Dm is the vector holding the second operand in the case of a by vector operation.

Dm[x] is the scalar holding the second operand for a by scalar operation.

B.6 NEON load and store element and structure instructions

This section describe NEON instructions which load or store data to or from memory or ARM integer registers.

B.6.1 VLDn and VSTn (single n-element structure to one lane)

VLDn (Vector Load single n-element structure to one lane) loads an n-element structure from memory into one or more NEON registers. Elements of the register that are not loaded remain unchanged.

VSTn (Vector Store single n-element structure to one lane) stores an n-element structure to memory from one or more NEON registers.

Syntax

```
Vopn{cond}.datatype list, [Rn{:align}]{!}  
Vopn{cond}.datatype list, [Rn{:align}], Rm
```

where:

op is either LD or ST.

n is one of 1, 2, 3, or 4.

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is one of 8, 16 or 32.

The following table shows the permitted options.

Table B-2 Permitted combinations of parameters

n	datatype	list	align	Alignment
1	8	{Dd[x]}	-	Standard only
	16	{Dd[x]}	16	2-byte
	32	{Dd[x]}	32	4-byte
2	8	{Dd[x], D(d+1)[x]}	16	2-byte
		{Dd[x], D(d+1)[x]}	32	4-byte
	16	{Dd[x], D(d+2)[x]}	32	4-byte
		{Dd[x], D(d+1)[x]}	64	8-byte
		{Dd[x], D(d+2)[x]}	64	8-byte
3	8	{Dd[x], D(d+1)[x], D(d+2)[x]}	-	Standard only
	16 or 32	{Dd[x], D(d+1)[x], D(d+2)[x]}	-	Standard only
		{Dd[x], D(d+2)[x], D(d+4)[x]}	-	Standard only
4	8	{Dd[x], D(d+1)[x], D(d+2)[x], D(d+3)[x]}	32	4-byte
	16	{Dd[x], D(d+1)[x], D(d+2)[x], D(d+3)[x]}	64	8-byte

Table B-2 Permitted combinations of parameters (continued)

n	datatype	list	align	Alignment
		{Dd[x], D(d+2)[x], D(d+4)[x], D(d+6)[x]}	64	8-byte
32		{Dd[x], D(d+1)[x], D(d+2)[x], D(d+3)[x]}	64 or 128	8-byte or 16-byte
		{Dd[x], D(d+2)[x], D(d+4)[x], D(d+6)[x]}	64 or 128	8-byte or 16-byte

list is a list of NEON registers in the range D0-D31, subject to the limitations given in the table.

Rn is the ARM register containing the base address (cannot be PC). If ! is present, Rn is updated to (Rn + the number of bytes transferred by the instruction). The update occurs after all the loads or stores have been performed.

Rm is the ARM register containing an offset from the base address. If Rm is present, then Rn is updated to (Rn + Rm) after the memory accesses have been performed. Rm cannot be SP or PC.

B.6.2 VLDn (single n-element structure to all lanes)

VLDn (Vector Load single n-element structure to all lanes) loads multiple copies of an n-element structure from memory into one or more NEON registers.

Syntax

```
VLDn{cond}.datatype list, [Rn{:align}]{!}
VLDn{cond}.datatype list, [Rn{:align}], Rm
```

where:

n is one of 1, 2, 3, or 4.

cond is an optional conditional code (see [Conditional execution on page 6-3](#)).

datatype is one of 8, 16 or 32.

The following table shows the permitted options:

Table B-3 Permitted combinations of parameters

n	datatype	list	align	Alignment
1	8	{Dd[]}	-	Standard only
		{Dd[], D(d+1)[]}	-	Standard only
	16	{Dd[]}	16	2 byte
		{Dd[], D(d+1)[]}	16	2 byte
	32	{Dd[]}	32	4 byte
		{Dd[], D(d+1)[]}	32	4 byte
2	8	{Dd[], D(d+1)[]}	8	byte
		{Dd[], D(d+2)[]}	8	byte
	16	{Dd[], D(d+1)[]}	16	2 byte
		{Dd[], D(d+2)[]}	16	2 byte

Table B-3 Permitted combinations of parameters (continued)

n	datatype	list	align	Alignment
	32	{Dd[], D(d+1)[]}	32	4 byte
		{Dd[], D(d+2)[]}	32	4 byte
3	8, 16, or 32	{Dd[], D(d+1)[], D(d+2)[]}	-	Standard only
		{Dd[], D(d+2)[], D(d+4)[]}	-	Standard only
4	8	{Dd[], D(d+1)[], D(d+2)[], D(d+3)[]}	32	4 byte
		{Dd[], D(d+2)[], D(d+4)[], D(d+6)[]}	32	4 byte
	16	{Dd[], D(d+1)[], D(d+2)[], D(d+3)[]}	64	8 byte
		{Dd[], D(d+2)[], D(d+4)[], D(d+6)[]}	64	8 byte
	32	{Dd[], D(d+1)[], D(d+2)[], D(d+3)[]}	64 or 128	8 byte or 16 byte
		{Dd[], D(d+2)[], D(d+4)[], D(d+6)[]}	64 or 128	8 byte or 16 byte

list is a list of NEON registers in the range D0-D31, subject to the limitations given in the table.

Rn is an ARM register containing the base address. Rn cannot be PC. If ! is specified, Rn is updated to (Rn + the number of bytes transferred by the instruction). The update occurs after the memory accesses are performed.

Rm is an ARM register containing an offset from the base address. If Rm is present, Rn is updated to (Rn + Rm) after the address is used to access memory. Rm cannot be SP or PC.

B.6.3 VLDn and VSTn (multiple n-element structures)

VLDn (Vector Load multiple n-element structures) loads multiple n-element structures from memory into one or more NEON registers, with de-interleaving (unless n == 1). Every element of each register is loaded.

VSTn (Vector Store multiple n-element structures) writes multiple n-element structures to memory from one or more NEON registers, with interleaving (unless n == 1). Every element of each register is stored.

Syntax

```
Vopn{cond}.datatype list, [Rn{:align}][]{!}
Vopn{cond}.datatype list, [Rn{:align}], Rm
```

where:

op is either LD or ST.

n is one of 1, 2, 3, or 4.

cond is an optional conditional code (see [Conditional execution on page 6-3](#)).

datatype is one of 8, 16 or 32.

The following table shows the permitted options:

Table B-4 Permitted combinations of parameters

n	datatype	list	align	Alignment
1	8, 16, 32, or 64	{Dd}	64	8 byte
		{Dd, D(d+1)}	64 or 128	8 byte or 16 byte
		{Dd, D(d+1), D(d+2)}	64	8 byte
		{Dd, D(d+1), D(d+2), D(d+3)}	64, 128, or 256	8 byte, 16 byte, or 32 byte
2	8, 16, or 32	{Dd, D(d+1)}	64, 128	8 byte or 16 byte
		{Dd, D(d+2)}	64, 128	8 byte or 16 byte
		{Dd, D(d+1), D(d+2), D(d+3)}	64, 128, or 256	8 byte, 16 byte, or 32 byte
3	8, 16, or 32	{Dd, D(d+1), D(d+2)}	64	8 byte
		{Dd, D(d+2), D(d+4)}	64	8 byte
4	8, 16, or 32	{Dd, D(d+1), D(d+2), D(d+3)}	64, 128, or 256	8 byte, 16 byte, or 32 byte
		{Dd, D(d+2), D(d+4), D(d+6)}	64, 128, or 256	8 byte, 16 byte, or 32 byte

list is a list of NEON registers in the range D0-D31, subject to the limitations given in the table.

Rn is the ARM register containing the base address. Rn cannot be PC. If ! is specified, Rn is updated to (Rn + the number of bytes transferred by the instruction). The update occurs after the memory accesses are performed.

Rm is an ARM register containing an offset from the base address. If Rm is present, Rn is updated to (Rn + Rm) after the address is used to access memory. Rm cannot be SP or PC.

B.6.4 VLDR and VSTR

VLDR loads a single extension register from memory, using an address from an ARM core register, with an optional offset.

VSTR saves the contents of a NEON or VFP register to memory.

One word is transferred if Fd is a VFP single-precision register, otherwise two words are transferred.

This instruction is present in both NEON and VFP instruction sets.

Syntax

```
VLDR{cond}{.size} Fd, [Rn{, #offset}]
VSTR{cond}{.size} Fd, [Rn{, #offset}]
VLDR{cond}{.size} Fd, label
VSTR{cond}{.size} Fd, label
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

size is an optional data size specifier, which is 32 if Fd is an S register, or 64 otherwise.

Fd is the extension register to be loaded or saved. For a NEON instruction, it must be a D register. For a VFP instruction, it can be either a D or S register.

Rn is the ARM register holding the base address for the transfer.

offset is an optional numeric expression. It must be a multiple of 4, within the range –1020 to +1020. The value is added to the base address to form the address used for the transfer. Label is a PC-relative expression and must align to a word boundary within ± 1 KB of the current instruction.

B.6.5 VLDM, VSTM, VPOP, and VPUSH

NEON and VFP register load multiple (VLDM), store multiple (VSTM), pop from stack (VPOP), push onto stack (VPUSH).

These instructions are present in both NEON and VFP instruction sets.

Syntax

```
VLDMmode{cond} Rn{!}, Registers
VSTMmode{cond} Rn{!}, Registers
VPOP{cond} Registers
VPUSH{cond} Registers
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

mode is one of:

- IA – Increment address after each transfer. This is the default, and can be omitted.
- DB – Decrement address before each transfer.
- EA – Empty Ascending stack operation. This is the same as DB for loads and IA for saves.
- FD – Full Descending stack operation. This is the same as IA for loads, and DB for saves.

Rn is the ARM register holding the base address for the transfer. If ! is specified, the updated base address must be written back to Rn. If ! is not specified, the mode must be IA.

Registers is a list of one or more consecutive NEON or VFP registers enclosed in braces, { }. The list can be comma-separated, or in range format. S, D, or Q registers can be specified, but they must not be mixed. The number of registers must not exceed 16 D registers, or 8 Q registers.

VPOP is equivalent to VLDM sp! and VPUSH is equivalent to VSTMDB sp!.

B.6.6 VMOV (between two ARM registers and an extension register)

Transfer contents between two ARM registers and a 64-bit NEON or VFP register, or two consecutive 32-bit VFP registers.

This instruction is present in both NEON and VFP instruction sets.

Syntax

```
VMOV{cond} Dm, Rd, Rn
VMOV{cond} Rd, Rn, Dm
VMOV{cond} Sm, Sm1, Rd, Rn
VMOV{cond} Rd, Rn, Sm, Sm1
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Dm is a 64-bit extension register.

Sm is a VFP 32-bit register and Sm1 is the next consecutive VFP 32-bit register after Sm.

Rd and Rn are the ARM registers.

B.6.7 VMOV (between an ARM register and a NEON scalar)

Transfer contents between an ARM register and a NEON scalar.

This instruction is present in both NEON and VFP instruction sets.

Syntax

```
VMOV{cond}{.size} Dn[x], Rd
VMOV{cond}{.datatype} Rd, Dn[x]
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

size can be 8, 16, or 32(default) for NEON, or 32 for VFP.

datatype can be U8, S8, U16, S16, or 32 (default). For VFP instructions, datatype must be 32 or omitted.

Dn[x] is the NEON scalar.

Rd is the ARM register.

B.6.8 VMRS and VMSR

VMRS transfers the contents of NEON or VFP system register FPSCR into Rd.

VMSR transfers the contents of Rd into a NEON or VFP system register, FPSCR.

These instructions are present in both NEON and VFP instruction sets.

Syntax

```
VMRS{cond} Rd, extsysreg
VMSR{cond} extsysreg, Rd
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

extsysreg specifies a NEON and VFP system register, one of:

- FPSCR
- FPSID
- FPEXC.

Rd is the ARM register. Rd can be APSR_nzcv, if extsysreg is FPSCR. Here, the floating-point status flags are transferred into the corresponding flags in the ARM APSR.

These instructions stall the ARM until all current NEON or VFP operations complete.

B.7 VFP instructions

This section describes VFP floating-point instructions.

B.7.1 VABS

Floating-point absolute value (VABS). This instruction can be scalar, vector, or mixed. VABS takes the contents of the specified register, clears the sign bit and stores the result.

Syntax

```
VABS{cond}.F32 Sd, Sm
VABS{cond}.F64 Dd, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Sd and Sm are the single-precision registers for the result and operand.

Dd and Dm are the double-precision registers for the result and operand.

B.7.2 VADD

VADD adds the values in the operand registers and places the result in the destination register.

Syntax

```
VADD{cond}.F32 {Sd,} Sn, Sm
VADD{cond}.F64 {Dd,} Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Sd, Sn, and Sm are the single-precision registers for the result and operands.

Dd, Dn, and Dm are the double-precision registers for the result and operands.

B.7.3 VCMPI

VCMPI (Floating-Point Compare) subtracts the value in the second operand register (or 0 if the second operand is #0) from the value in the first operand register and sets the VFP condition flags depending on the result. VCMPI is always scalar.

Syntax

```
VCMPI{cond}.F32 Sd, Sm
VCMPI{cond}.F32 Sd, #0
VCMPI{cond}.F64 Dd, Dm
VCMPI{cond}.F64 Dd, #0
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Sd and Sm are the single-precision registers holding the operands.

Dd and Dm are the double-precision registers holding the operands.

B.7.4 VCVT (between single-precision and double-precision)

VCVT converts the single-precision value in *Sm* to double-precision and stores the result in *Dd*, or converts the double-precision value in *Dm* to single-precision, storing the result in *Sd*. VCVT is always scalar.

Syntax

```
VCVT{cond}.F64.F32 Dd, Sm
VCVT{cond}.F32.F64 Sd, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Dd is a double-precision register for the result, with *Sm* a single-precision register which holds the operand.

Sd is a single-precision register for the result with *Dm* a double-precision register holding the operand.

B.7.5 VCVT (between floating-point and integer)

VCVT forms which convert from floating-point to integer or from integer to floating-point. VCVT is always scalar.

Syntax

```
VCVT{R}{cond}.type.F64 Sd, Dm
VCVT{R}{cond}.type.F32 Sd, Sm
VCVT{cond}.F64.type Dd, Sm
VCVT{cond}.F32.type Sd, Sm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

R makes the operation use the rounding mode specified by the FPSCR. If *R* is not specified, VCVT rounds towards zero.

type is either U32 (unsigned 32-bit integer) or S32 (signed 32-bit integer).

Sd is a single-precision register for the result.

Dd is a double-precision register for the result.

Sm is a single-precision register holding the operand.

Dm is a double-precision register holding the operand.

B.7.6 VCVT (between floating-point and fixed-point)

Convert between floating-point and fixed-point numbers. In all cases the fixed-point number is contained in the least significant 16 or 32 bits of the register. VCVT is always scalar.

Syntax

```
VCVT{cond}.type.F64 Dd, Dd, #fbits
VCVT{cond}.type.F32 Sd, Sd, #fbits
VCVT{cond}.F64.type Dd, Dd, #fbits
VCVT{cond}.F32.type Sd, Sd, #fbits
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

type can be any one of:

- S16, 16-bit signed fixed-point number.
- U16, 16-bit unsigned fixed-point number.
- S32, 32-bit signed fixed-point number.
- U32, 32-bit unsigned fixed-point number.

Sd is a single-precision register for the operand and result.

Dd is a double-precision register for the operand and result.

fbits is the number of fraction bits in the fixed-point number, in the range 0-16 (if type is S16 or U16), or 1-32 (if type is S32 or U32).

B.7.7 VCVTB, VCVTT (half-precision extension)

These instructions convert between half-precision and single-precision floating-point numbers:

- VCVTB uses the lower half (bits [15:0]) of the single word register to obtain or store the half-precision value.
- VCVTT uses the upper half (bits [31:16]) of the single word register to obtain or store the half-precision value.

VCVTB and VCVTT are always scalar.

Syntax

```
VCVTB{cond}.type Sd, Sm
VCVTT{cond}.type Sd, Sm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

type is either F32.F16 (convert from half-precision to single-precision) or F16.F32 (convert from single-precision to half-precision).

Sd is a single word register for the result.

Sm is a single word register for the operand.

B.7.8 VDIV

VDIV divides the value in the first operand register by the value in the second operand register, and places the result in the destination register. The instructions can be scalar, vector, or mixed.

Syntax

```
VDIV{cond}.F32 {Sd,} Sn, Sm
VDIV{cond}.F64 {Dd,} Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Sd, Sn, and Sm are the single-precision registers for the result and operands.

Dd, Dn, and Dm are the double-precision registers for the result and operands.

B.7.9 VFMA, VFNMA, VFMS, VFNMS

VFMA (Fused Floating-point Multiply Accumulate (with optional Negation)) multiplies the operand registers, adds the value from the destination register and stores the final result in the destination register. The result of the multiply is not rounded before the accumulation.

VFMS (Fused Floating-point Multiply Subtract (with optional Negation)) multiplies the values in the operand registers, subtracts the product from the destination register value and places the final result in the destination register. The result of the multiply is not rounded before the subtraction.

These instructions are always scalar.

Syntax

```
VF{N}op{cond}.F64 {Dd,} Dn, Dm
VF{N}op{cond}.F32 {Sd,} Sn, Sm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

op is either MA or MS.

N negates the final result.

Sd, Sn, and Sm are the single-precision registers for the result and operands.

Dd, Dn, and Dm are the double-precision registers for the result and operands.

Qd, Qn, and Qm are the double-precision registers for the result and operands.

B.7.10 VMOV

VMOV puts a floating-point immediate value into a single-precision or double-precision register, or copies one register into another register. This instruction is always scalar.

Syntax

```
VMOV{cond}.F32 Sd, #imm
VMOV{cond}.F64 Dd, #imm
VMOV{cond}.F32 Sd, Sm
VMOV{cond}.F64 Dd, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Sd is the single-precision destination register.

Dd is the double-precision destination register.

imm is the floating-point immediate value.

Sm is the single-precision source register.

Dm is the double-precision source register.

B.7.11 VMOV

Transfer contents between a single-precision floating-point register and an ARM register.

Syntax

```
VMOV{cond} Rd, Sn
VMOV{cond} Sn, Rd
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Sn is the VFP single-precision register.

Rd is the ARM register.

B.7.12 VMUL, VMLA, VMLS, VNMUL, VNMLA, and VNMLS

VMUL (Floating-point Multiply (with optional Negation)) multiplies the values in the operand registers and stores the result in the destination register.

VMLA (Floating-point Multiply Accumulate (with optional Negation)) multiplies the values in the operand registers, adds the value from the destination register, and stores the final result in the destination register.

VMLS (Floating-point Multiply and Multiply Subtract (with optional Negation)) multiplies the values in the operand registers, subtracts the result from the value in the destination register, and stores the final result in the destination register.

The final result is negated if the N option is used.

These instructions can be scalar, vector, or mixed.

Syntax

```
V{N}MUL{cond}.F32 {Sd,} Sn, Sm
V{N}MUL{cond}.F64 {Dd,} Dn, Dm
V{N}MLA{cond}.F32 Sd, Sn, Sm
V{N}MLA{cond}.F64 Dd, Dn, Dm
V{N}MLS{cond}.F32 Sd, Sn, Sm
V{N}MLS{cond}.F64 Dd, Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

N negates the final result.

Sd, Sn, and Sm are the single-precision registers for the result and operands.

Dd, Dn, and Dm are the double-precision registers for the result and operands.

B.7.13 VNEG

VNEG (Floating-point negate). This instruction can be scalar, vector, or mixed. VNEG takes the contents of the specified register, inverts the sign bit and stores the result.

Syntax

```
VNEG{cond}.F32 Sd, Sm
```

VNEG{cond}.F64 Dd, Dm

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Sd and Sm are the single-precision registers for the result and operand.

Dd and Dm are the double-precision registers for the result and operand.

B.7.14 VSQRT

VSQRT (Floating-point square root) takes the square root of the contents of Sm or Dm, and places the result in Sd or Dd. It can be scalar, vector, or mixed.

Syntax

VSQRT{cond}.F32 Sd, Sm
VSQRT{cond}.F64 Dd, Dm

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Sd and Sm are the single-precision registers for the result and operand.

Dd and Dm are the double-precision registers for the result and operand.

B.7.15 VSUB

VSUB subtracts the value in the second operand register from the value in the first operand register, and places the result in the destination register. The instructions can be scalar, vector, or mixed.

Syntax

VSUB{cond}.F32 {Sd,} Sn, Sm
VSUB{cond}.F64 {Dd,} Dn, Dm

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

Sd, Sn, and Sm are the single-precision registers for the result and operands.

Dd, Dn, and Dm are the double-precision registers for the result and operands.

B.8 NEON and VFP pseudo-instructions

This section describes pseudo-instructions which will be translated by the assembler to a real machine instruction.

B.8.1 VACLE and VACLT

Vector Absolute Compare takes the absolute value of each element in a vector, and compares with the absolute value of the corresponding element of a second vector. If the condition is true, the corresponding element in the destination vector is set to all ones. If false, it is set to all zeros. This produces the corresponding VACGE and VACGT instructions, with the operands reversed.

Syntax

```
VACop{cond}.datatype {Qd,} Qn, Qm
VACop{cond}.datatype {Dd,} Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

op is either LE (Absolute Less than or Equal) or LT (Absolute Less Than).

datatype must be F32.

Qd or Dd is the NEON register for the result. The result datatype is always I32.

Qn or Dn is the NEON register holding the first operand,

Qm or Dm is the NEON register which holds the second operand.

B.8.2 VAND (immediate)

VAND (bitwise AND immediate) takes each element of the destination vector, does a bitwise AND with an immediate value, and stores the result into the destination.

Syntax

```
VAND{cond}.datatype Qd, #imm
VAND{cond}.datatype Dd, #imm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is one of I8, I16, I32, or I64.

Qd or Dd is the NEON register for the result.

imm is the immediate value.

If datatype is I16, the immediate value must have one of the following forms:

- 0xFFXY
- 0XYFF.

If datatype is I32, the immediate value must have one of the following forms:

- 0xFFFFFXY
- 0FFFFXYFF
- 0FFXYFFFF
- 0XYFFFFFFF.

B.8.3 VCLE and VCLT

Vector Compare takes the value of each element in a vector, and compares it with the value of the corresponding element of a second vector. If the condition is true, the corresponding element in the destination vector is set to all ones. If false, it is set to all zeros. These pseudo-instructions produce the corresponding VCGE and VCGT instructions, with the operands reversed.

Syntax

```
VCop{cond}.datatype {Qd,} Qn, Qm
VCop{cond}.datatype {Dd,} Dn, Dm
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

op is either LE (Less than or Equal) or LT (Less Than).

datatype is one of S8, S16, S32, U8, U16, U32, or F32.

Qd or Dd is the NEON register for the result. The result datatype is as follows:

- I32 for operand datatypes I32, S32, U32, or F32.
- I16 for operand datatypes I16, S16, or U16.
- I8 for operand datatypes I8, S8, or U8.
- U32 32-bit unsigned fixed-point number.

Qn or Dn is the NEON register holding the first operand.

Qm or Dm is the NEON register which holds the second operand.

B.8.4 VLDR pseudo-instruction

The VLDR pseudo-instruction loads a constant value into every element of a 64-bit NEON vector (or a VFP single-precision or double-precision register).

Syntax

```
VLDR{cond}.datatype Dd,=constant
VLDR{cond}.datatype Sd,=constant
```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is one of In, Sn, Un (where *n* is one of 8, 16, 32, or 64) or F32 for NEON. For VFP code, use either F32 or F64.

Dd or Sd is the register to be loaded.

constant is an immediate value of the appropriate type for datatype. If an instruction is available that can generate the constant directly, the assembler uses it. Otherwise, it generates a doubleword literal pool entry containing the constant and use a VLDR instruction to load the constant.

B.8.5 VLDR and VSTR (post-increment and pre-decrement)

The VLDR and VSTR pseudo-instructions load or store extension registers with post-increment and pre-decrement.

Syntax

```

op{cond}{.size} Fd, [Rn], #offset ; post-increment
op{cond}{.size} Fd, [Rn, #-offset]! ; pre-decrement

```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

op is either VLDR or VSTR.

size is 32 if Fd is an S register, or 64 if Fd is a D register. For a NEON instruction, it must be a doubleword (Dd) register. For a VFP instruction, it can be a double-precision (Dd) or single precision (Sd) register.

Rn is the ARM register that sets the base address for the transfer.

The post-increment instruction increments the base address in the register by the offset value, after the transfer. The pre-decrement instruction decrements the base address in the register by the offset value, and then does the transfer using the new address in the register. These pseudo-instructions assemble to VLDM or VSTM instructions.

B.8.6 VMOV2

The VMOV2 pseudo-instruction generates an immediate value and places it in every element of a NEON vector, without a load from a literal pool. It always generates two instructions typically a VMOV or VMVN followed by a VBIC or VORR. VMOV2 can generate any 16-bit immediate value, and a restricted range of 32-bit and 64-bit immediate values.

Syntax

```

VMOV2{cond}.datatype Qd, #constant
VMOV2{cond}.datatype Dd, #constant

```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is one of:

- I8, I16, I32, or I64.
- S8, S16, S32, or S64.
- U8, U16, U32, or U64.
- F32.

Qd or Dd is the extension register to be loaded.

constant is an immediate value of the appropriate type for datatype.

B.8.7 VORN

VORN (Bitwise OR NOT) (immediate) takes each element of the destination vector, does a bitwise OR complement with an immediate value and stores the result into the destination vector.

Syntax

```

VORN{cond}.datatype Qd, #imm
VORN{cond}.datatype Dd, #imm

```

where:

cond is an optional conditional code. (See [Conditional execution on page 6-3](#))

datatype is one of I8, I16, I32, or I64.

Qd or Dd is the NEON register for the result.

imm is the immediate value.

If datatype is I16, the immediate value must have one of the following forms:

- 0xFFXY
- 0xXYFF.

If datatype is I32, the immediate value must have one of the following forms:

- 0xFFFFFFXY.
- 0xFFFFXYFF.
- 0xFFXYFFFF.
- 0XYFFFFFFF.

Appendix C

Building Linux for ARM Systems

The objective of this appendix is to enable the reader to build a Linux system for an ARM platform. Though the process for doing this will vary slightly for different processors and platforms, most of the steps will be very similar to those outlined here.

A working Linux system has two primary components, namely the kernel and the root filesystem. The kernel is the core of the operating system and acts as a resource manager. It is loaded into RAM by the bootloader and then executed during the boot process, as described in [Chapter 15](#). The root filesystem contains system libraries, applications, tools and utilities, for example, command line interfaces or shells, graphical interfaces (such as the X window system), text editors (such as vi, emacs, gedit) and advanced applications like web browsers or office suites.

The root filesystem is often located in persistent storage, such as a hard disk or a flash device. However, the root filesystem can also reside in primary memory (RAM). This book will not cover RAM-based filesystems in great detail. We will first cover the steps for building the kernel and then the root filesystem, before analyzing how these fit together to get the system running.

C.1 Building the Linux kernel

This section explains the steps required to build the kernel for an ARM processor based platform, either natively on the target platform or cross-compiled on, for example, an x86 PC. The platform we build the kernel for is called the *target* platform. The platform on which we build is called the *build* platform. We assume that the reader has some experience in using Linux, C compilers and has a machine (either x86 or ARM processor based) running a Linux distribution like Ubuntu – for other distributions some of the steps of the build procedure might differ slightly. We also assume that the build machine has a working internet connection. If this machine is ARM processor based, it is not necessary for the build platform to be exactly the same as the target platform.

The Linux kernel can be viewed as consisting two parts. One part is architecture independent and consists of components like process schedulers, system call interfaces architecture-independent device drivers and high-level network subsystems. The other part is closely related to the hardware platform for which the kernel is being built. This consists of board initialization code and drivers corresponding to a specific hardware platform. While building a kernel one has to be sure of having the correct set of initialization code and drivers for the platform at hand.

The Linux kernel sources can be found at , <http://www.kernel.org/>.

There are two ways to get the kernel source code. The first is to download a compressed tar file. The exact name of this file will naturally depend upon the version of the kernel selected, for example 2.6.34. When downloaded, this file must be uncompressed using a command similar to the following.

```
tar xjvf linux-2.6.34.tar.bz2
```

The other option for obtaining the kernel source is to obtain the source tree using GIT (git) commands. Linux is developed using the GIT version control system, which can be installed using a command similar to the one below for Ubuntu.

```
sudo apt-get install git-core
```

Obviously, a working internet connection is required for the above steps and those which follow.

The Linux source tree can be cloned with a command similar to the following:

```
git clone git://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux-2.6.git
```

This command will copy the official Linux kernel repository to the current directory, which can take some time. You can then check out a specific tag, which is a local operation and fairly quick.

```
git checkout -b v2.6.34
```

Follow the instructions in [Prebuilt versions of GNU toolchains on page 3-9](#) in order to obtain and install a suitable toolchain. Also ensure that the toolchain is accessible on your path.

When the source tree is in place, the kernel needs to be configured to match the hardware platform and desired kernel features. The standard method is to use a command such as:

```
make ARCH=arm realview_defconfig
```

which generates a default configuration file for the RealView platform file and stores it as `.config`.

There are several methods available for configuring the kernel. The most commonly used provides a text based interface based on the ncurses library. It can be invoked using the command:

```
make ARCH=arm menuconfig
```

This gives a configuration screen for selecting or omitting features of the kernel, as shown in [Figure C-1](#).




Figure C-1 Kernel Configuration screen using make menuconfig

If this command fails, and the configuration menu does not appear, this could be because the ncurses header files are not installed on your build host. You can install them by executing:

```
sudo apt-get install libncurses5-dev
```

The other alternative is to use a graphical Xconfig tool; this uses the Qt GUI library and can be invoked using the command:

```
make ARCH=arm xconfig
```

[Figure C-2 on page C-4](#) shows the Kernel Configuration screen from this tool.




Figure C-2 Kernel Configuration screen using make xconfig

When the kernel is configured correctly then it can be built using a simple make command as below – exact details can differ slightly depending on the bootloader used:

```
make ARCH=arm CROSS_COMPILE=arm-none-linux-gnueabi-
```

The CROSS_COMPILE value should be set to the toolchain cross-compilation prefix (*Prebuilt versions of GNU toolchains on page 3-9*) or should be completely left out for native compilation.

The output of the compilation would be in the form of a compressed kernel zImage. This can usually be found in the path <source root>/arch/arm/boot as a file named zImage.

When compiling natively on an ARM processor based system, the CROSS_COMPILE... parameter should be left out. If a different cross compilation toolchain than the codesourcery Linux EABI toolchain is used, arm-none-linux-eabi- might need to be modified to reflect the name of the toolchain executables. The above also assumes that the cross compilation toolchain executable directory is listed in your PATH environment variable.

This operation requires the mkimage utility to be installed. The Ubuntu package name for this tool is uboot-mkimage.

If the hardware platform uses the U-Boot bootloader, the kernel image will need to be converted into a form accepted by U-boot. For example, in Ubuntu the `uboot-mkimage` package can be installed followed by the command below, instead of the simple `make`, to create an “ubootified” image.

```
make ARCH=arm CROSS_COMPILE=arm-none-linux-gnueabi- uImage
```

The built kernel image would be found as the file named `uImage`, again in the path `<source root>/arch/arm/boot`.

Now that the kernel is built, the filesystem needs to be created in order to have a working ARM Linux system. The next section deals with creating the filesystem.

C.2 Creating the Linux filesystem

Creating a filesystem for an ARM processor based platform is relatively straightforward, using a Ubuntu distribution. The procedure described here does not actually compile the filesystem, rather it downloads pre-built packages to create the filesystem. A full compilation of the filesystem would take many hours of compilation time and can be technically challenging.

For carrying out the process of putting together the filesystem the rootstock package needs to be installed, which can be done using the command below:

```
sudo apt-get install rootstock
```

When the package is installed, navigate to the directory to be used for creating and saving the filesystem. The filesystem can be created in the form of a compressed tarball using the command below.

```
sudo rootstock --fqdn ubuntu --login ubuntu --password ubuntu --imagesize 3G --seed ubuntu-desktop
```

In this case, the system would be based on the ubuntu-desktop seed which includes a desktop windowing system. Other *seeds* could be used instead, for example xubuntu-desktop (for a smaller, more lightweight desktop system) or build-essential (for a text based interface).

Note

This step might take a significant amount of time, depending on the speed of the internet connection.

The following steps describe the process of preparing a disk drive and transferring the filesystem to it. This disk would then be used as the root filesystem on the ARM processor based platform. The term disk drive is used here in a loose sense and refers to a variety of secondary storage devices, for example, hard disks, compact-flash drives, or USB drives. After plugging in the disk drive to the computer which was used to create the filesystem, the following command can be used to check the connected drives

```
sudo fdisk -l
```

Note

The disk drive being used for storing the root filesystem of the ARM processor based platform will be formatted by the following steps. The process of formatting the disk will destroy any data that might exist on the disk. Therefore the user needs to ensure that the disk does not contain any useful data prior to formatting, as this data will be permanently erased. It is also easy to lose the data elsewhere on the system if the wrong device is specified here.

The disk drive to be used needs to be identified correctly in the list given by the command above (for example, /dev/sdb). Assuming that /dev/sdb is the correct device or disk drive, the following command is needed to partition the drive correctly.

```
sudo fdisk /dev/sdb
```

On entering the above command, the fdisk prompt will be displayed. Now a sequence of fdisk command needs to be entered which are in the form of single characters.

Type *m* to display the help for the possible set of fdisk commands. Use *d* to delete any existing partitions on the disk. You can then create a new partition using *n*. The character *w* is used to write the changes to the disk and exit. You can check that the partition has been written correctly by starting fdisk again using the above command, followed by a *p* for printing the partition table.

When the partition table has been written correctly, the disk partition that has been created can be formatted. Usually the first partition created on the disk identified by (for example) `/dev/sdb` is denoted as `/dev/sdb1`. The command for formatting this partition as an ext-3 filesystem is as below.

```
sudo mkfs.ext3 /dev/sdb1
```

Now, the previously created tarball for the Linux filesystem needs to be decompressed into this disk partition. To do this, the disk partition needs to be *mounted* which either can be done manually with the `mount` command, or automatically by unplugging and re-plugging in the drive. In the automatic mounting case, the disk is usually mounted at the location `/media/<disk_directory>`. Navigate to the directory at which the disk is mounted and uncompress the file-stem tarball into it using the command below:

```
sudo tar zxvf <path_to_tarred_file_system/file.tgz>
```

Now that both the kernel and filesystem have been created, the two need to be brought together to have a working Linux system.

Note

It is also possible to store the root filesystem in primary memory, using a RAMDISK. However, this alternative will not be described here and is left to advanced users.

C.3 Putting it together

The steps for creating the kernel and the filesystem are generic and common to many different boards or platforms. However, the steps for having the kernel programmed onto the platform can vary between boards. For most boards, the kernel needs to be transferred to some form of secondary memory that either exists on the board or is connected to it. For example, on ARM Versatile boards, the kernel needs to be copied on to a flash device on the board. There can be two different kinds of procedures for doing this depending on the board.

Use the documentation for your development board to find the appropriate method for getting your kernel installed, and follow the instructions. You should also ensure that the root filesystem which was created earlier is connected correctly or copied onto the target board.

In order to obtain a list of possible U-Boot commands type `help` at the U-Boot prompt.

The U-Boot bootloader is popular on many ARM processor based platforms. It is free and open-source. Other than U-Boot, there are also other bootloaders which might be proprietary for certain ARM processor based platforms. In order to interact with the bootloader, it might be necessary to connect a serial interface between the board or platform and a personal computer running a serial terminal. Again, refer to the board's user manual for more details.

The bootloader can pass a set of parameters to the kernel during the boot process. Among these is the kernel command line, known in U-Boot as *bootargs* (short for boot arguments). These parameters are used by the kernel for some initialization configurations. An example command for setting up the bootargs on ARM Versatile boards is shown below. This needs to be entered at the U-Boot prompt.

```
setenv bootargs root=/dev/sda1 mem=512M@0x20000000 ip=dhcp console=ttyAMA0 clcd=xvga
rootwait
```

The bootargs in this case specifies the following:

- The root filesystem is in `/dev/sda1`.
- The memory region to be used by the operating system in the form of a size (512MB) and a starting address location. This isn't required if the bootloader passes the correct `ATAG_MEM` atag.
- The IP address (for example, 192.168.0.7) or the mechanism for obtaining it (in this case DHCP).
- The display interface details.
- The delay required before trying to mount the root filesystem which can be needed for devices to be recognized (for example, USB disks).

The command above can be followed by a `saveenv` command in U-Boot, to save the changes into flash and make them permanent. In order to obtain a list of possible U-Boot commands type `help` at the U-Boot prompt. Type `help <command>` at the U-Boot prompt, to get more details regarding a particular command.

The bootloader also needs a boot command to start the boot process automatically. In the simplest case the `bootcmd` can be set as follows, where `0x41000000` is the location where the kernel is stored for example in flash memory.

```
setenv bootcmd bootm 0x41000000
```

If the kernel image exists at a different location, or on the network, a `cp` or `tftp` command can precede the `bootm` command. U-Boot commands on the same line need to be separated by a “;”, for example:

```
setenv bootcmd cp 0x65000000 0x41000000 0x8000000; bootm 0x41000000
```

Again, a `saveenv` command would be required, to ensure that changes are saved to secondary storage. Now that the `bootcmd` is in place, the system needs to be restarted such that the `bootcmd` is used automatically and the system should start running. The sequence of boot commands can also be entered manually at the U-Boot prompt, to try different options without saving them permanently.

Index

A

AAPCS xiii, 17-1
 ABI xiii, 17-1
 Aborts
 Abort handlers 14-2
 Accelerator Coherency Port xiii, 22-3
 Access alignment 16-5
 ACP xiii, 22-3
 ACTLR 6-20
 Additional multipliers 6-9
 Address Space ID xiii, 10-16
 Advanced Extensible Interface xiii
 Advanced High-performance Bus xiii
 Advanced Microcontroller Bus Architecture xiii
 Advanced Peripheral Bus xiii
 Advanced SIMD technology 2-7
 AHB xiii
 AHB Trace Macrocell 29-5
 AHB trace macrocell 29-5
 AMBA xiii
 Amdahl's Law 24-1
 AMP xiii, 21-7
 APB xiii
 Application Binary Interface xiii, 17-1
 Application profile 2-3
 Application Program Status Register xiii, 4-4, 4-5

Application Specific Integrated Circuit xiii, 1-3
 APSR xiii, 4-4, 4-5
 Architectural configurations 22-1
 Architecture 2-1, 2-8
 Architecture history and extensions 2-4, 2-5
 Architecture Reference Manual (ARM ARM) xiii
 Architecture versions 2-3, 2-5
 Application profile 2-3
 ARMv4T 2-4
 ARMv5TE 2-4
 ARMv6 2-4
 ARMv7A 2-3, 2-4
 Microcontroller profile 2-3
 Real-time profile 2-3
 ARM Architecture 2-1
 ARM Architecture Procedure Call Standard xiii, 17-1
 ARM Assembler 5-9
 Comparison with other forms 5-2
 Identifying 5-12
 Interworking 5-11
 Mixing C and assembler 17-7
 Porting 16-10
 Power instructions 25-3
 Syntax 5-9
 Directives 5-10
 Label 5-9

ARM Compiler toolchain 3-9, 5-1
 ARM DS-5 3-11
 Streamline 3-12
 ARM History 1-2
 ARM Instruction Set 4-1, 4-2, 5-4, 6-9, 6-22
 Basics 6-2
 Bit manipulation instructions 6-22
 Branches 6-13
 Byte reversal 6-23
 Cache preload 6-22
 Condition Codes 6-5
 Conditional execution 6-3
 Constant values 6-2
 Coprorocessor instructions 6-19
 Data processing 6-6
 Integer register SIMD instructions 6-14
 Integer SIMD instructions 6-14
 Memory instructions 6-10
 Miscellaneous instructions 6-19
 Operand2 6-7
 PRS modification 6-22
 Saturating arithmetic 6-18
 Status Flags 6-5
 ARM Processors 1-1, 2-1
 Software toolchains 3-8
 ARM Registers 4-1, 4-4
 ARM Thumb Procedure Call Standard xiii, 17-1
 ARM tools assembly language 5-9

armar 3-10
 armasm 3-9
 armcc 3-9
 armcc optimization options 19-7
 armlink 3-10
 ASIC xiii, 1-3
 ASID xiii, 10-16
 Assembly language 5-1
 Assigning interrupts 13-2
 Asymmetric multi-processing xiii, 21-7
 ATPCS xiii, 17-1
 Auxiliary Control Register 6-20
 AXI xiii

B

Bare-metal booting 15-2
 BE8 xiii, 16-4
 Big-endian 16-2
 big.LITTLE 2-7, 28-1
 Cluster migration 28-6
 Configuration 28-2
 CPU migration 28-7
 Execution modes 28-4
 Full MP operation 28-9
 Managing a system 28-2
 Structure of s system 28-3
 Bit manipulation instructions 6-22
 BKPT (Instruction) 6-23
 Board Support Package xiii, 3-4
 Architecture-specific code 3-4
 Generic device drivers 3-4
 Platform-specific code 3-4
 Processor-specific code 3-4
 Boot code 15-1
 Configuration 15-6
 Boot process
 Linux 14-5
 Booting a bare-metal system 15-2
 Booting Linux 15-7
 Booting SMP systems 22-10
 Bootloaders 3-6, 15-7
 Branch prediction 4-10
 Branch prediction logic 4-10
 Branch Target Address Cache xiii, 4-10
 Branch Target Buffer xiii, 4-10
 Branches and interworking 16-11
 BSP xiii, 3-4
 BTAC xiii, 4-10
 BTB xiii, 4-10
 Building Linux on ARM C-1
 BusyBox 3-6
 Byte Invariant Big-Endian Mode xiii

C

Cache coherency 22-2
 MESI protocol 22-2
 MOESI protocol 22-3
 Cache thrashing 24-4
 Cachegrind 18-7

Caches 9-1
 Allocation policies
 Read 9-12
 Write 9-12
 Architecture 9-6
 Basics of operation 9-3
 Cache controller 9-6
 Direct mapped 9-7
 Drawbacks 9-4
 Invalidating 9-16
 Memory hierarchy 9-5
 parity and ECC 9-19
 Performance and hit rate 9-15
 Policies 9-12
 Replacement policy 9-12
 Set associative 9-8
 Terminology 9-9
 Index 9-9
 Line 9-9
 Set 9-9
 Tag 9-10
 Way 9-9
 Write and fetch buffers 9-14
 Write policy 9-12
 CBAR 6-21
 CISC xiii, 5-2
 Clobber list 17-8
 CLREX (Instruction) 22-7
 Cluster migration 28-6
 Compiler optimization 19-3
 Compiler packing of structures 16-8
 Complex Instruction Set Computer xiii
 Complex Instruction Set computer 5-2
 Condition codes 6-5
 Configuration Base Address Register 6-21
 Context ID Register 6-21
 CONTEXTIDR 6-21
 Coprocessor Access Control Register 6-20
 Coprocessor 15 xiii, 6-19
 Instruction syntax 6-21
 Register summary 6-20
 Core 2-2
 CoreSight 29-4
 CoreSight Serial Wire 29-5
 Cortex-A Series architecture 2-8
 Pipelines 2-9
 Processors 2-9
 Cortex-A series architecture
 Cortex-A15 2-14, 27-1
 Cortex-A5 2-11
 Cortex-A7 2-11
 Cortex-A8 2-13
 Cortex-A9 2-14
 Cortex-A15 processor 2-14, 27-1
 Cortex-A5 processor 2-11
 Cortex-A7 processor 2-11
 Cortex-A8 processor 2-13
 Cortex-A9 processor 2-14
 MPCore 21-4
 Register renaming 4-9
 CPACR 6-20
 CPSR xiii, 4-4, 4-6
 CPSR E bit 16-3, 22-12

CPU migration 28-7
 CP15 xiii, 6-19
 Instruction syntax 6-21
 register summary 6-20
 Current Program Status Register xiii, 4-4, 4-6

D

DAP xiii, 29-5
 Data Fault Address Register 6-20
 Data Fault Status Register 6-20
 Data Memory Barrier xiii
 Data packing/unpacking 6-16
 Byte selection 6-17
 Data processing instructions 6-6
 Data Synchronization Barrier xiii
 DBX xiii
 DDR xiii
 Debug 29-1, 29-5
 ARM trace hardware 29-4
 CoreSight 29-4
 Debug Access Port (DAP) 29-5
 Debug and Trace 29-9
 Debugging Linux applications 29-8
 Embedded cross trigger 29-5
 Hardware 29-2
 Monitor 29-7
 System trace macrocell 29-5
 Trace memory controller 29-6
 Debug Access Port xiii, 29-5
 Debug events 29-2
 Debugging Linux Applications 29-8
 Device memory 11-4
 DFAR 6-20
 DFSR 6-20
 Digital Signal Processing xiii
 Direct Bytecode Execution xiii
 Direct mapped caches 9-7
 Direct Memory Access xiii, 9-4
 DMA xiii, 9-4
 DMB xiii
 Dormant mode 25-3
 Double Data Rate (SDRAM) xiii
 DSB xiii
 DSP xiii, 2-6
 DSTREAM xiii
 DS-5 xiii, 3-11, 18-4
 Trace support in 29-12
 DVFS xiv, 25-4
 Dynamic Voltage/Frequency Scaling xiv, 25-4

E

EABI xiv
 ECC xiv, 9-19
 ECT xiv, 29-5
 Embedded ABI xiv
 Embedded Cross Trigger xiv, 29-5
 Embedded Linux 3-3

- Embedded systems 1-4
 - Real-time behavior 1-4
 - Time to market 1-4
- Embedded Trace Buffer xiv
- Embedded Trace Macrocell xiv
- Endianness 16-2
- Error Correcting Code xiv
- ETB xiv
- ETM xiv
- Example platforms
 - BeagleBoard 3-13
 - Gumstix 3-13
 - PandaBoard 3-13
 - Pandora 3-13
- Exception handling 12-1
 - SVC 14-4
- Exception handling 14-1
- Exception program flow
 - Linux 14-5
- Exceptions
 - Entering an exception handler 12-7
 - Exception mode summary 12-5
 - Exception priorities 12-6
 - Exit from an exception handler 12-7, 12-8
 - Return instruction 12-10
 - Types 12-3
 - Aborts 12-3
 - Exceptional instructions 12-4
 - Interrupts 12-3
 - Reset 12-3, 12-4
 - Vector table 12-9
- External interrupt request 13-2

F

- False sharing 24-4
- Fast Context Switch Extension 10-18
- FCSE 10-18
- FIQ xiv, 12-3, 25-3
- Floating point 7-1, 7-4
 - Basics 7-2
 - Flag meanings 7-5
 - Interpreting the flags 7-6
 - Optimization 7-11
 - Rounding algorithm 7-4
- Floating Point Exception Register 7-5
- Floating Point Status and Control Register xiv
- Floating Point System and Control Register 7-4
- Floating Point System ID Register 7-4
- FPEXC 7-5
- FPSCR xiv, 7-4
- FPSID 7-4
- fromef 3-10
- Ftrace 18-7
- Function inlining 19-3

G

- GCC xiv
- GCC optimization options 19-5
- General Purpose Input/Output xiv, 1-4
- Generic Interrupt Controller xiv, 13-1, 13-4
- GIC xiv, 13-1, 13-4
- GIF xiv
- GNU 3-8
- GNU Assembler 5-10
 - Directives 5-6
 - Expressions 5-8
 - Introduction 5-5
 - Invoking 5-5
 - Sections 5-6
 - Syntax 5-5
- GNU Compiler Collection xiv
- GNU tools 5-8
- GPIO xiv, 1-4
- Gprof xiv, 18-3
- Graphics Interchange Format xiv

H

- Halt debug mode 29-3
- Handling interrupts 22-6
- Handling interrupts 13-1
- Harvard architecture xiv
- Heterogeneous multiprocessing 21-3, 28-1
- History of ARM 1-2
- HVC (Instruction) 27-5
- Hypervisor 27-1
- Hypervisor exception model 27-5
- Hypervisor mode 27-3

I

- ICCIAR 13-6
- IDE xiv
- Identifying assembler 5-12
- IEEE-754 7-2
- IFAR 6-20
- IFSR 6-20
- Indirect branches 4-10
- Initializing the memory system 15-8
- Inlining 19-3
- Instruction Fault Address Register 6-20
- Instruction Fault Status Register 6-20
- Instruction pipelines 4-7
- Instruction Set Architecture xiv, 5-2
- Instruction sets 5-4
 - ARM 5-4
 - Thumb 5-4
- Instruction Synchronization Barrier xiv
- Instructions
 - CLREX 22-7
 - HVC 27-5
 - LDREX 22-7
 - STREX 22-7
 - SVC 6-22, 12-4
 - SWI 12-4

- WFI 25-3

- Integer register SIMD instructions 6-14
- Integer SIMD instructions 6-14
- Integrated Development Environment xiv
- Intermediate Physical Address 10-19
- Interrupt Acknowledge Register (ICCIAR) 13-6
- Interrupt dispatch
 - Linux 14-5
- Interrupt handling 13-1, 22-6
- Interrupt Request xiv
- Interrupt Service Routine xiv
- Interrupts 13-2
 - Configuration 13-5
 - Handling in SMP systems 22-6
 - Nested handling 13-3
 - Non-nested handling 13-2
 - Simplistic handling 13-2
- Invalidating and cleaning cache memory 9-16
- IPA 10-19
- IRQ xiv, 12-3, 25-3
- ISA xiv, 5-2
- ISB xiv
- ISR xiv

J

- Jazelle xiv, 2-6

K

- Kernel image 15-8

L

- Large 10-19
- Large Physical Address Extension 2-7, 10-3, 10-14, 10-19
- LDREX (Instruction) 22-7
- Level 1 page tables 10-4
- Level 2 page tables 10-7
- Linaro 3-4
- Link register(R14) 4-5
- Linux
 - Boot process 14-5
 - Booting 15-7
 - Building on ARM C-1
 - Debugging 29-8
 - Distributions 3-2
 - Exception program flow 14-5
 - Interrupt dispatch 14-5
 - Kernel entry 15-8
 - Kernel parameters 15-8
 - Kernel start-up code 15-9
 - perf events 18-7
 - Platform specific actions 15-9
 - Reset handler 15-7
 - Virtual memory view 15-10
- Linux on ARM processors 3-2

- Linux terminology 3-2
 - Files 3-3
 - Libraries 3-3
 - Process 3-2
 - Scheduler 3-3
 - System calls 3-3
 - Thread 3-2
- Linux tools
 - BusyBox 3-6
 - QEMU 3-6
 - Scratchbox 3-6
 - Tianocore 3-7
 - U-Boot 3-6
 - UEFI 3-7
- Little-endian 16-2
- Load-Exclusive operations 22-7
- Loop unrolling 19-4
- LPAAE xiv, 2-7, 2-14, 10-3, 10-14, 10-19, 27-2
- L1 Caches 9-5
- L2 Cache 9-5
 - Controller 9-18
 - Maintenance 9-18

M

- Main ID Register 6-20
- Memory access ordering 16-10
- Memory access permissions 10-13
- Memory attributes 10-13
 - page table entries 10-13
- Memory barriers 16-10
- Memory instructions 6-10
 - Addressing modes 6-10
 - Multiple transfers 6-11
- Memory Management Unit xiv, 10-1
 - Domains 10-15
- Memory Protection Unit xiv
- Memory translation 27-4
- Memory types
 - Device memory 11-4
 - Strongly-ordered 11-4
- MESI protocol xiv
- Micro-architecture 2-1
- Micro-architecture optimization
 - Complex addressing modes 19-16
- Microcontroller profile 2-3
- MMU xiv, 10-1
- Monitor debug mode 29-3
- Monitor Vector Base Address Register 6-21
- MPIDR 6-20
- MPU xiv
- Multi-issue pipelines 4-8
- Multiplication instructions 6-9
- Multiplication operations 6-9
- Multiply-accumulate 2-6
- Multi-processing 21-1
 - ARM systems 21-3
 - Asymmetric 21-7
 - Symmetric 21-5
- Multiprocessor Affinity Register 6-20
- Multi-tasking 10-16
- Mutex 23-6
- MVBAR 6-21
- MVFR0/1
 - Media and VFP Feature Registers 7-5

N

- NEON 2-7
 - ABI restriction 20-5
 - Alignment 20-8
 - Architecture 8-4
 - Assembler 20-5
 - Data types 8-4
 - Detecting 20-5
 - Instruction set 8-7
 - Intrinsics 20-2
 - Introduction 8-1
 - Libraries 20-2
 - Long instructions 8-6
 - Memory access optimization 20-7
 - Narrow instructions 8-7
 - NEON C compiler and assembler 20-2
 - Normal instructions 8-6
 - Optimizing code 20-7
 - Power saving 20-9
 - Registers 8-5
 - Scheduling 20-8
 - SIMD 8-2
 - Types in C 20-3
 - Variables and constants 20-3
 - Vectorization 20-2
 - VFP Commonality 8-4
 - Wide instructions 8-6
 - Writing NEON code 20-1
- NEON Instructions B-1
 - Arithmetic B-23
 - General data processing B-6
 - Load/Store and Structure B-35
 - Logical and comparison B-17
 - Multiply B-32
 - Pseudo instructions B-47
 - Shift B-13
- Nested interrupt handling 13-3
- NMI xiv
- Non-Maskable Interrupt xiv
- Non-nested interrupt handling 13-2
- Non-shareable memory 11-5
- NOP (Instruction) 6-23

O

- Operand combinations 16-13
- Oprofile xiv, 18-4
- Optimization 18-1
- Optimizing 19-1, 19-4
 - armcc options 19-7
 - Associativity effects 19-10
 - Data abort 19-12
 - Data cache 19-8

- Eliminating common sub-expressions 19-3
- Floating point 7-11
- GCC options 19-5
- Instruction cache usage 19-10
- Loop interchange 19-9
- Loop tiling 19-8
- L2/Outer Cache Usage 19-11
- Memory system 19-8
- NEON 20-7
- Prefetching memory block access 19-12
- Source code modification 19-13
- Structure alignment 19-10
- TLB Usage 19-11
- OS use of page tables 10-16

P

- Page sizes 10-12
- Page table Entries
 - Memory attributes 10-13
- Page Table Entry xiv
- Page Table Walking 10-3
- Page tables
 - Level 1 10-4
 - Level 2 10-7
- Parallelization 23-1, 24-1
 - Bandwidth 24-3
 - Deadlock 24-4
 - Decomposition methods 23-2
 - False sharing 24-4
 - Inter-thread communication 23-6
 - Issues 24-1
 - Livelock 24-4
 - Performance issues 24-3
 - Bandwidth 24-3
 - Cache contention 24-3
 - False sharing 24-3
 - Reentrancy 24-2
 - Synchronization mechanisms 23-8
 - Thread affinity 23-7
 - Thread dependencies 24-3
 - Thread safety 24-2
 - Threaded performance 23-6
 - Threading libraries 23-5
 - Threading models 23-4
- PCI xiv
- PCS xiv
- Performance monitor 18-5
- Performance Monitor Unit xiv
- Peripheral Component Interconnect xiv
- Physically Indexed Physically Tagged xiv, 9-11
- Pipelines 2-9, 4-7
 - Multi-issue 4-8
 - Pipeline artifacts 4-8
- PIPT xiv, 9-11
- PLE xiv
- PMU xiv
- PoC xiv, 9-17
- Point of Coherency xiv, 9-17
- Point of Unification xiv, 9-17

Porting 16-1, 16-13
 ARM assembler to ARMv7 16-10
 ARM code to Thumb 16-11
 Branches and interworking 16-11
 Compiler packing 16-8
 Unsigned and signed char 16-7
 Using PC as an operand 16-11
 POSIX 23-5
 PoU xv, 9-17
 Power and Clocking 25-2
 Power management 25-1
 PPI xv
 Preload Engine xv
 Private Peripheral Interrupt xv
 Privilege 4-3
 Privilege model 27-3
 Procedure Call Standard xv, 17-4
 Procedure call standard 17-2
 Processor 2-1
 Processor instruction pipelines 4-7
 Processor Modes 4-1, 4-3
 Processors 2-9
 Profiler output 18-3
 Profiles
 Application 2-3
 Microcontroller 2-3
 Real-time 2-3
 Profiling 18-1
 Event based sampling 18-2
 Time based sampling 18-2
 Profiling in SMP systems 24-5
 Program counter (R15) 4-4
 Program Status Register xv
 PSR xv
 PTE xv

Q

QEMU xv
 Qualcomm Scorpion 2-15

R

Race conditions 24-3
 RCT xv
 Real-time profile 2-3
 RealView Compilation Tools xv
 Reduced Instruction Set Computer xv, 5-2
 Reentrant 13-3, 22-7
 Register renaming 4-9
 Registers 4-1, 6-20
 ACTLR 6-20
 CBAR 6-21
 CONTEXTIDR 6-21
 CPACR 6-20
 DFAR 6-20
 DFSR 6-20
 IFAR 6-20
 IFSR 6-20
 MVBAR 6-21
 NEON 8-5

NEON usage 17-4
 SCR 6-20
 SCTLR 6-20
 TTBCR 6-20
 TTBR 6-20
 VBAR 6-21
 VFP 7-4
 VFP usage 17-4
 Renaming registers 4-9
 Reset handler
 Linux 15-7
 Return stack 4-10
 RISC xv, 5-2
 Runtime Compiler Target xv
 RVCT xv
 R13 4-5
 R14 4-5
 R15 4-4

S

Saturated arithmetic instructions 2-6
 Saturating arithmetic 6-18
 Saved Program Status Register xv
 SCR 6-20, 26-3
 SCTLR 6-20
 SCU xv, 21-3, 22-2
 Secure Configuration Register 6-20, 26-3
 Secure systems 26-1
 Security 26-1
 Multi-processor system extensions 26-4
 Normal world 26-4
 Secure world 26-4
 Security Extensions 2-6
 Security Extensions and Virtualization 27-6
 Semaphores 23-8
 Set associative caches 9-8
 SGI xv
 Shareable memory 11-5
 Shared Peripheral Interrupt xv
 Short-Descriptor Format 10-3
 SIMD xv, 2-7, 8-2
 SIMD instructions
 Data packing/unpacking 6-16
 Sum of absolute differences 6-15
 Simplistic interrupt handling 13-2
 Single Instruction Multiple Data xv, 8-2
 SiP xv, 1-3
 SMP 22-1
 SMP systems
 booting 22-10
 Snoop Control Unit xv, 21-3, 22-2
 SoC xv, 1-3
 Software Generated Interrupt xv
 Software Interrupt xvi
 Source code optimization
 Division/modulo 19-15
 Inline assembler 19-16
 Linker 19-17
 Loop fusion 19-13
 Loop termination 19-13

Pointer aliasing 19-15
 Reducing stack/heap use 19-13, 19-14
 Variable selection 19-14
 SPI xv
 Spinlock code 3-4, 22-9, 23-8
 Spinlocks 22-7
 SPSR xv
 Stack and heap 17-6
 Stack pointer (R13) 4-5
 Standby mode 25-2
 Status flags 6-5
 STM 29-5
 Store-Exclusive operations 22-7
 Streamline xv, 3-12, 18-4
 STREX (Instruction) 22-7
 Strongly-ordered memory 11-4, 12-4
 Supervisor Call xvi
 SVC exception handling 14-4
 SVC (Instruction) xvi, 6-22, 12-4
 SWI (Instruction) xvi, 12-4
 Symmetric multi-processing 21-5, 22-1, 22-6
 Booting 22-10
 Linux 22-10
 Processor ID 22-10
 Cache coherency 22-2
 Cache maintenance broadcast 22-5
 Exclusive accesses 22-7
 Private memory region 22-12
 Timers 22-12
 TLB 22-5
 Watchdogs 22-12
 Synchronization mechanisms
 Completion 23-8
 Lock-free 23-8
 Semaphores 23-8
 Spinlocks 23-8
 System Control Processor xiii
 System Control Register 6-20
 System in Package 1-3
 System in Package (SiP) xv
 System on Chip xv, 1-3
 System Trace Macrocell 29-5
 System-on-Chip 1-3

T

TCM xvi
 Thread dependencies 24-3
 Priority inversion 24-3
 Thread Local Storage xvi
 Thumb xvi, 2-6
 Porting ARM code to 16-11
 Thumb Execution Environment 2-6
 Thumb instruction set 4-2
 ThumbEE 2-6
 Thumb-2 technology xvi, 2-6, 4-2
 Tightly Coupled Memory xvi
 TLB xvi, 10-10
 TLB Coherency 10-11
 TLS xvi
 TMC 29-6

- Toolchains 3-8
 - ARM Compiler 3-9
 - GNU 3-8
- Trace memory Controller 29-6
- Translation Lookaside Buffer xvi, 10-10
- Translation Table Base xvi
 - Register 0 and 1 10-17
- Translation Table Base Control Register 6-20
- Translation Table Base Register xvi, 6-20, 10-19
- TrustZone xvi, 2-6, 26-2
 - Architecture 26-2
- TTB xvi, 10-17
- TTBCR 6-20
- TTBR xvi, 6-20, 10-19
- TTBR0 xvi, 6-20, 10-19
- TTBR1 xvi, 6-20, 10-19

U

- UAL xvi, 5-9
- UART xvi
- U-Boot xvi, 3-6
- UEFI xvi, 3-7
- Unaligned access 19-17
- Undefined instruction handling 14-3
- Unified Assembly Language xvi, 5-9
- Unified Extensible Firmware Interface xvi, 3-7
- Universal Asynchronous
 - Receiver/Transmitter xvi
- Unprivileged mode 4-3
- User Mode xvi
- USR xvi

V

- Valgrind 18-7
- VBAR 6-21
- Vector Base Address Register 6-21
- Vector Floating Point xvi, 2-6
- Vectored Interrupt Controller xvi
- VFP xvi, 2-6, 7-4
 - Support in ARM Compiler 7-9
 - Support in GCC 7-8
 - Support in Linux 7-10
- VFP Instructions 7-7, B-41
- VFP Registers 7-4
- VIC xvi
- VIPT xvi, 9-11
- Virtual Machine ID xvi
- Virtual machine monitor 27-1
- Virtual memory 10-3
- Virtual Memory Systems Architecture xvi
- Virtualization 2-7, 27-1
- Virtualization and Security Extensions 27-6
- Virtualization Extensions 4-5, 27-1, 27-3
 - Privilege model 27-3

- Virtually Indexed Physically Tagged xvi, 9-11
- VMID xvi
- VMSA xvi

W

- Weakly ordered model 9-2
- WFI (Instruction) 6-23, 25-3
- Write and fetch buffers 9-14

Z

- zImage 15-8